

50.003: Elements of Software Construction

Week 5

Basics of Concurrency: Requirements

Course Plan

Week	Cohort Class 1	Cohort Class 2	Cohort Class 3	Remarks
1 (Jan 28)	Software Development Process			
2 (Feb 4)	Software Design and UML		Project Meeting I	Problem Set 1
3 (Feb 11)	Software Design and UML			
4 (Feb 18)	Software Testing			Problem Set 2, Quiz 1
5 (Feb 25)	Basics of Concurrency		Project Meeting II	
6 (Mar 4)	Thread Safe Java Class			Problem Set 3
7 (Mar 11)	Recess			Problem Set 4
8 (Mar 18)	Testing Concurrent Programs		Project Meeting III	Quiz 2
9 (Mar 25)	Performance			Problem Set 5
10 (Apr 1)	Design Patterns			Problem Set 6; Quiz 3
11 (Apr 8)	Advanced Software Testing		Project Meeting IV	Problem Set 7
12 (Apr 15)	Software Debugging, Code smells and Maintenance			Quiz 4
13 (Apr 22)	Final Project Presentation (15 minutes for each group)			Project Report/Code Due
14 (May 2)	Final Exam			

Example

NSA intercepted a RSA-encrypted secret message which tells the location of a terrorist act, we believe that the act is going to happen one week from now, we need your help in decrypting the message.

Requirement/Analysis

“Given a semi-prime, your program outputs its prime factors within 6 days.”

green: pre-condition

red: post-condition

purple: non-functional requirement

Total Correctness: If the precondition is satisfied, it is guaranteed that the method terminates and satisfies the postcondition.

Design/Implementation

Design

- Use the trial division method
- [http://en.wikipedia.org/wiki/Integer factorization](http://en.wikipedia.org/wiki/Integer_factorization)

Implementation

- FactorPrime.java

Testing

Correctness Testing

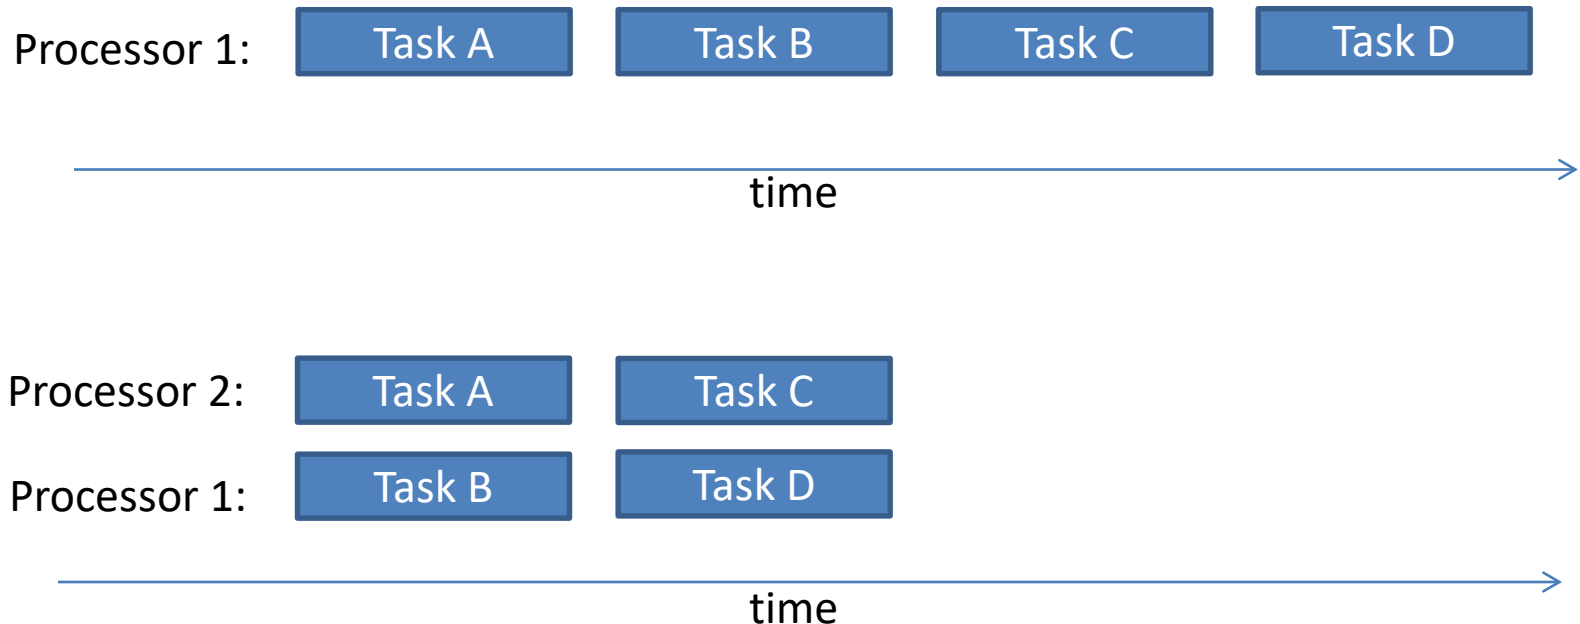
- 4294967297 (famous Fermat Number)
- 1127451830576035879
- 160731047637009729259688920385507056726966793490579598495689711866432421212774967029895340327197901756096014299132623454583177072050452755510701340673282385647899694083881316194642417451570483466327782135730575564856185546487053034404560063433614723836456790266457438831626375556854133866958349817172727462462516466898479574402841071703909138062456567624565784254101568378407242273207660892036869708190688033351601539401621576507964841597205952722487750670904522932328731530640706457382162644738538813247139315456213401586618820517823576427094125197001270350087878270889717445401145792231674098948416888868250143592026973853973785120217077951766546939577520897245392186547279572494177680291506578508962707934879124914880885500726439625033021936728949277390185399024276547035995915648938170415663757378637207011391538009596833354107737156273037494727858302028663366296943925008647348769272035532265048049709827275179381252898675965528510619258376779171030556482884535728812916216625430187039533668677528079544176897647303445153643525354817413650848544778690688201005274443717680593899

Performance Testing

FactorPrimeTest.java

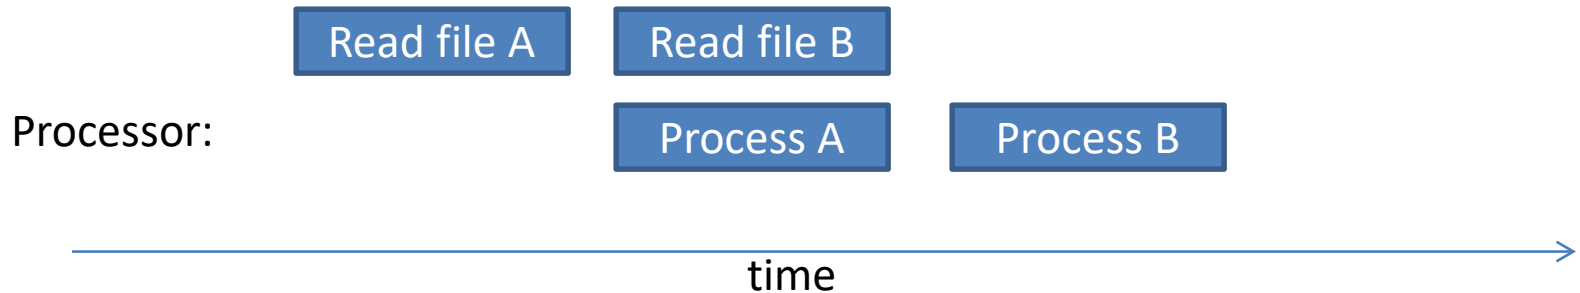
Concurrency: Benefit

Better resource utilization: with K processors, **ideally** we can be K times faster*



Concurrency: Benefit

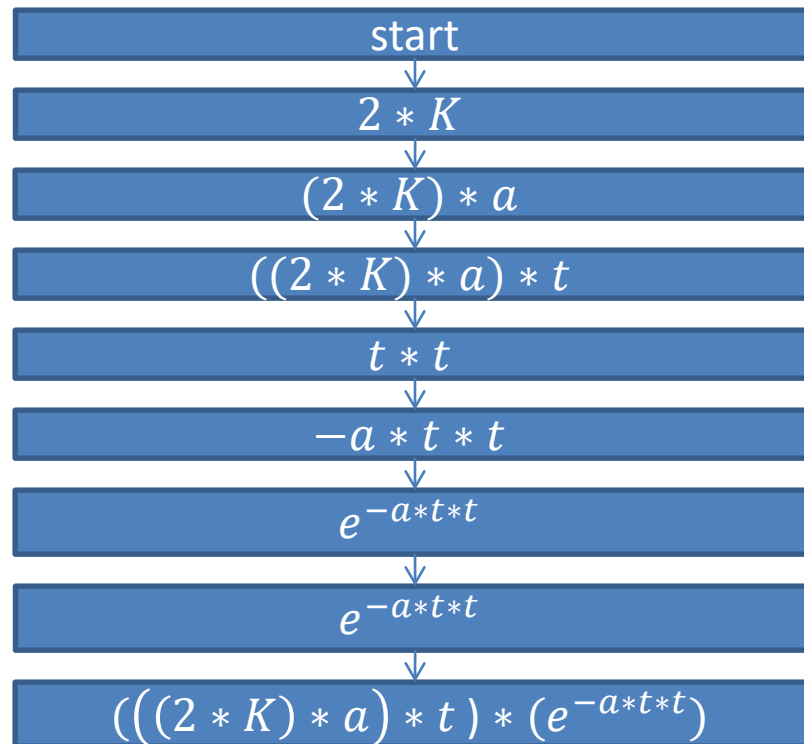
Can we get better performance with 1 processor only?



Threads

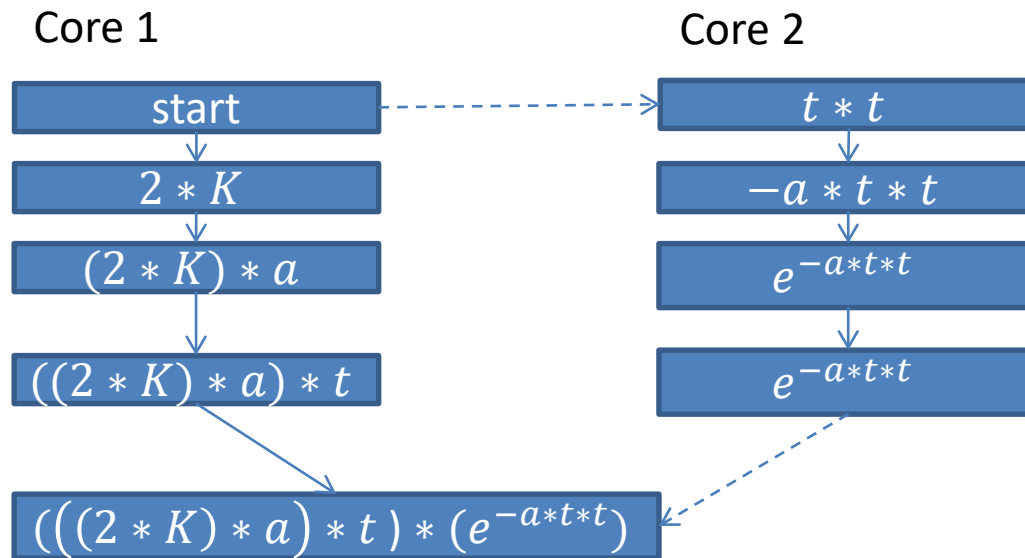
Example: evaluate $2 * K * a * t * e^{-a*t*t}$

Core 1



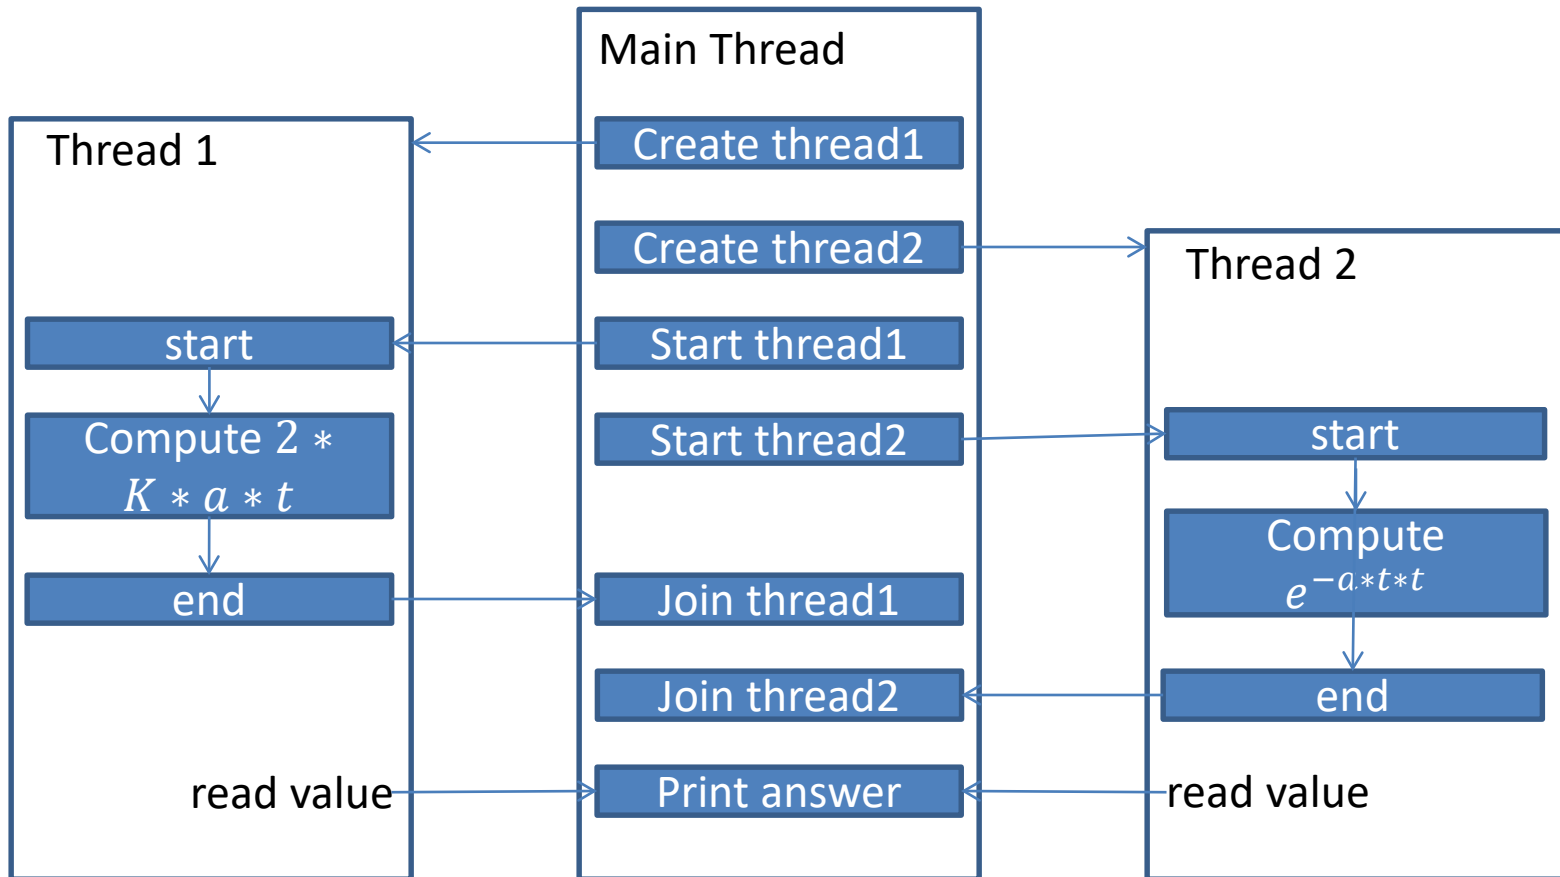
Threads

Example: evaluate $2 * K * a * t * e^{-a * t * t}$

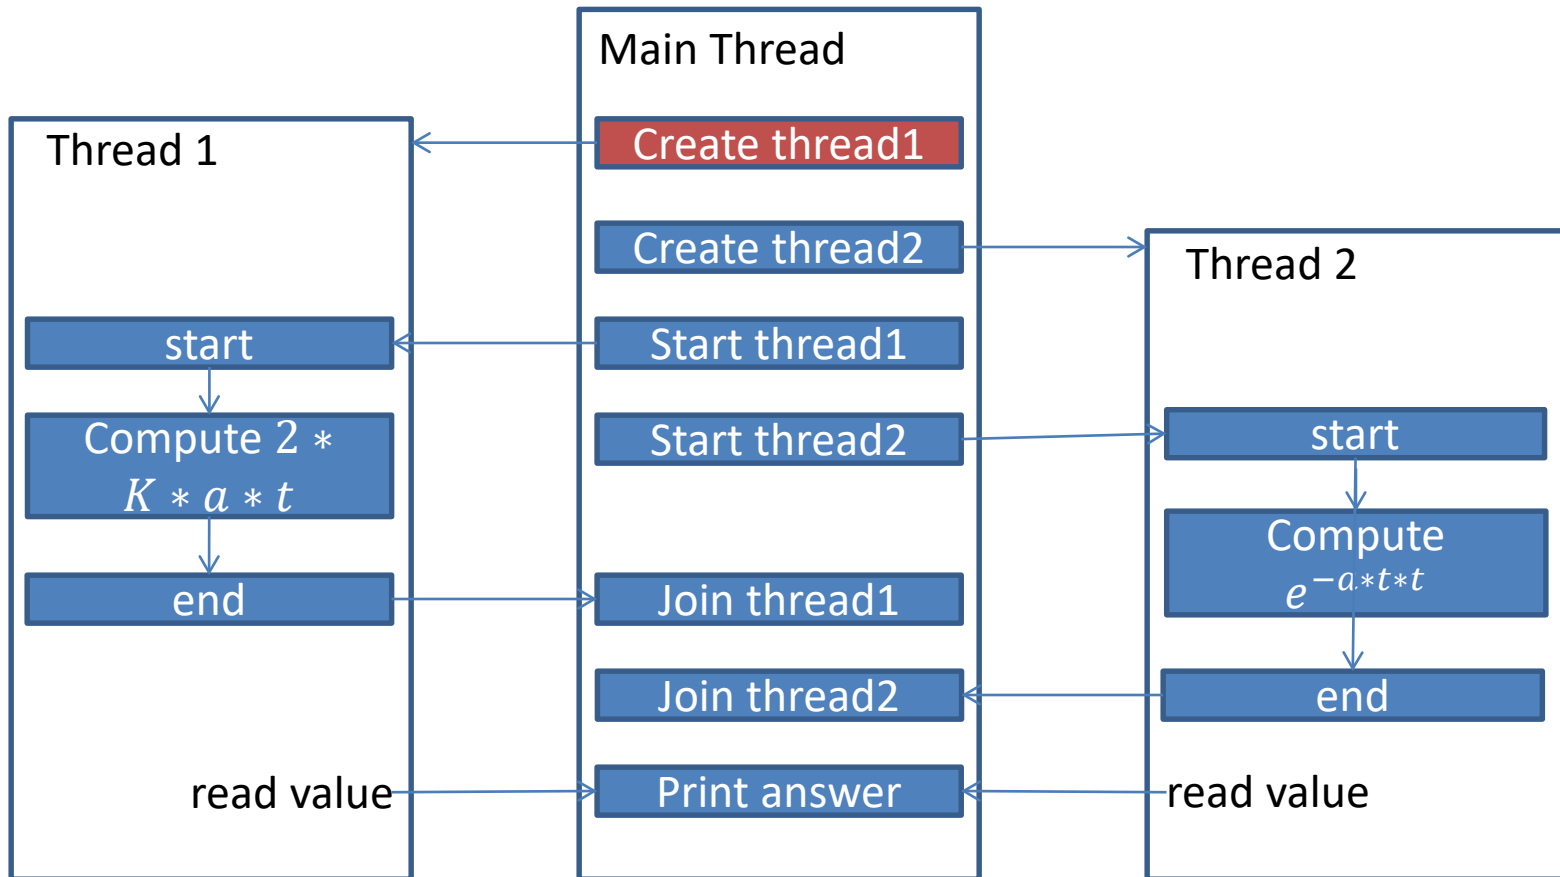


HelloWorldThread.java

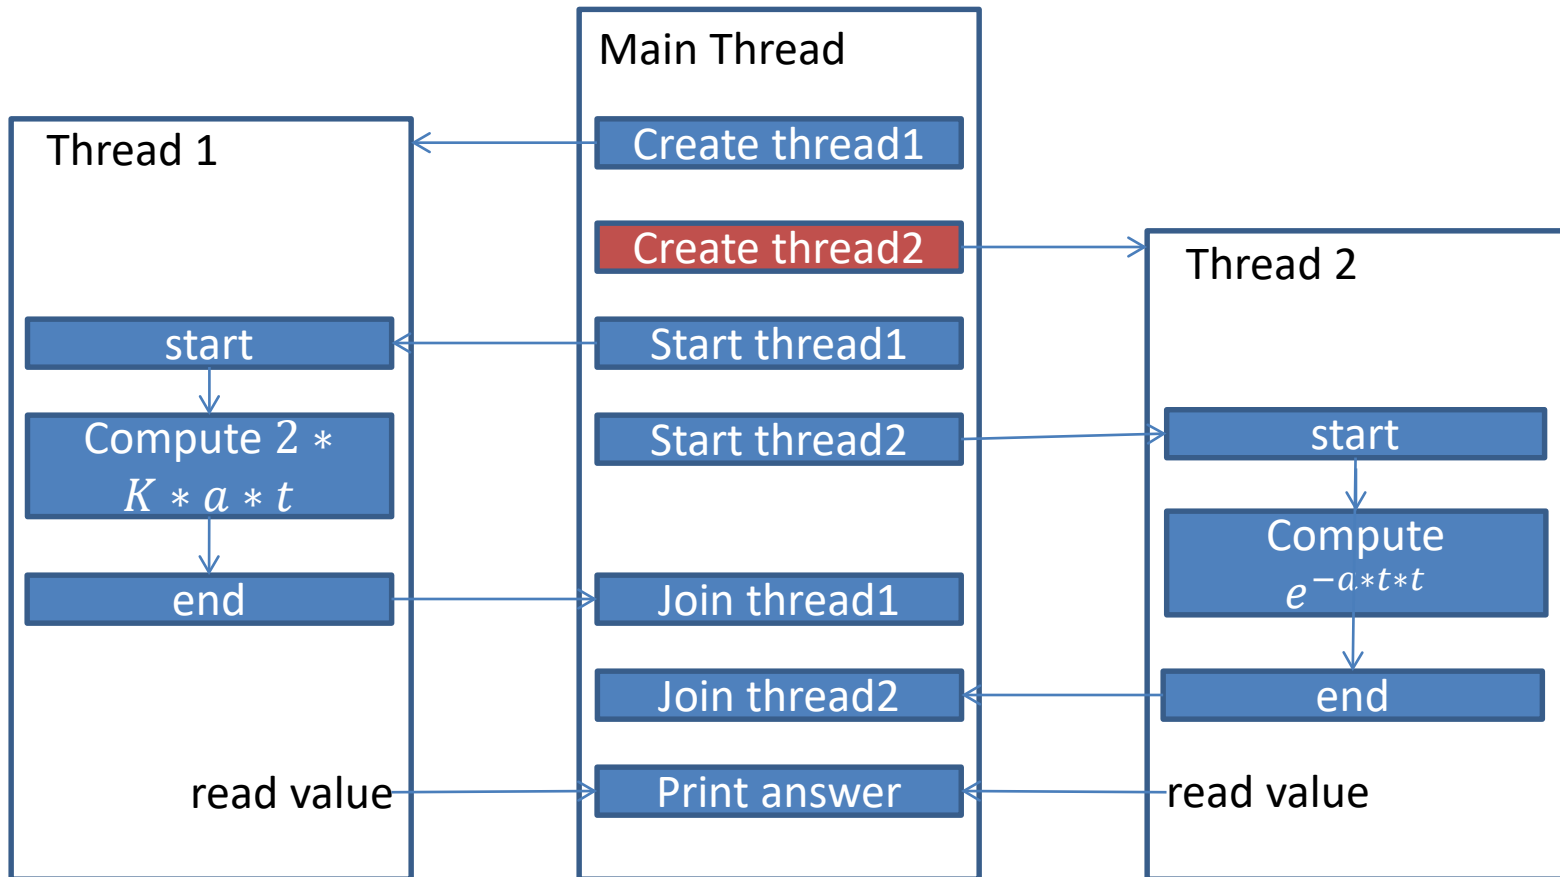
Hello World



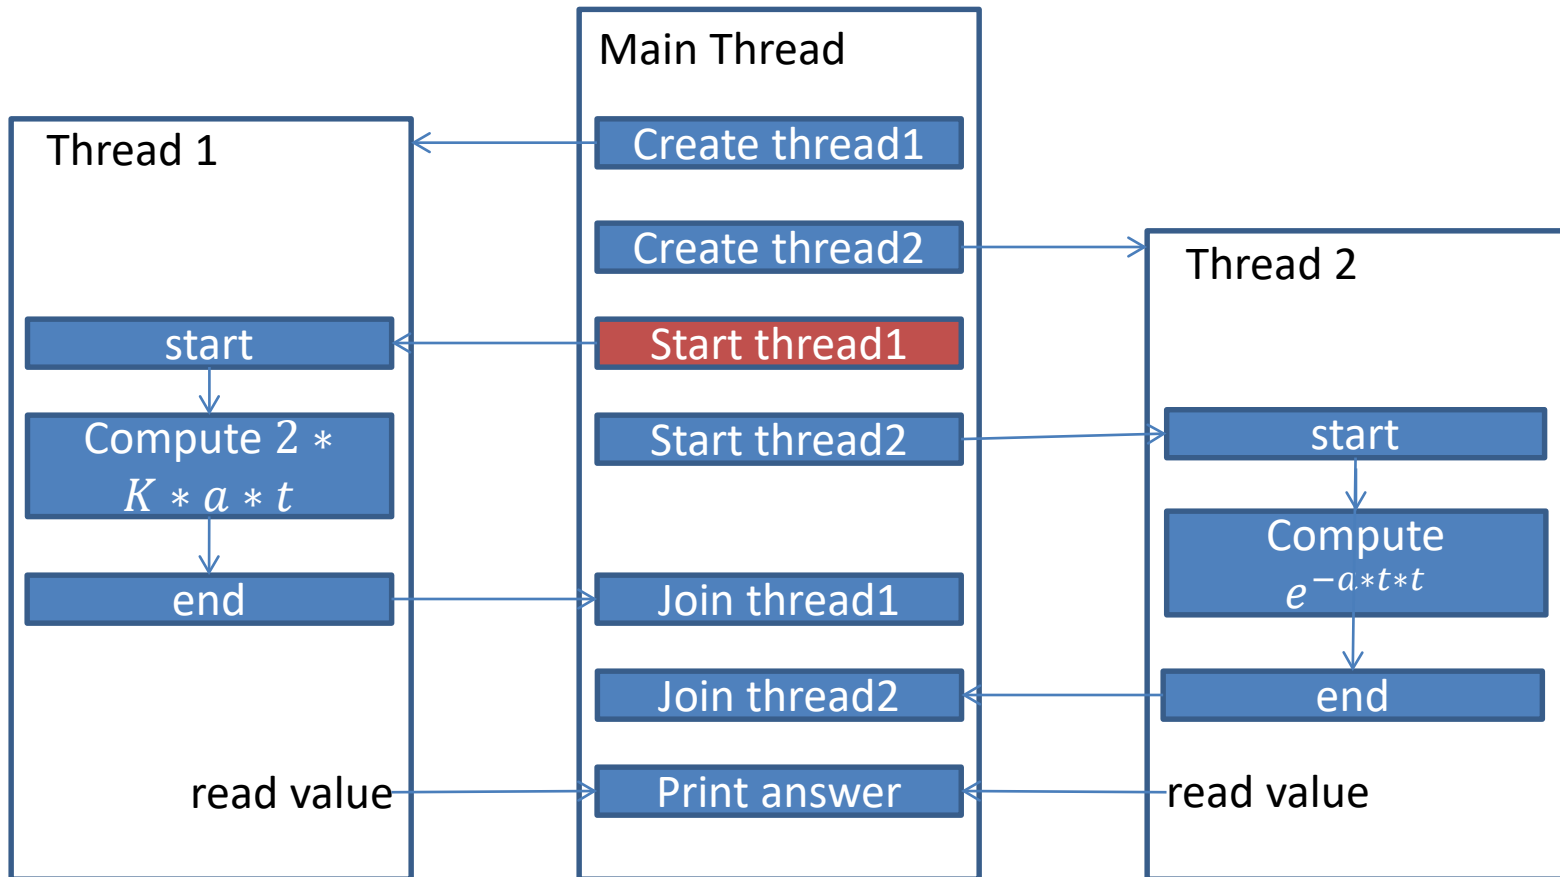
Hello World



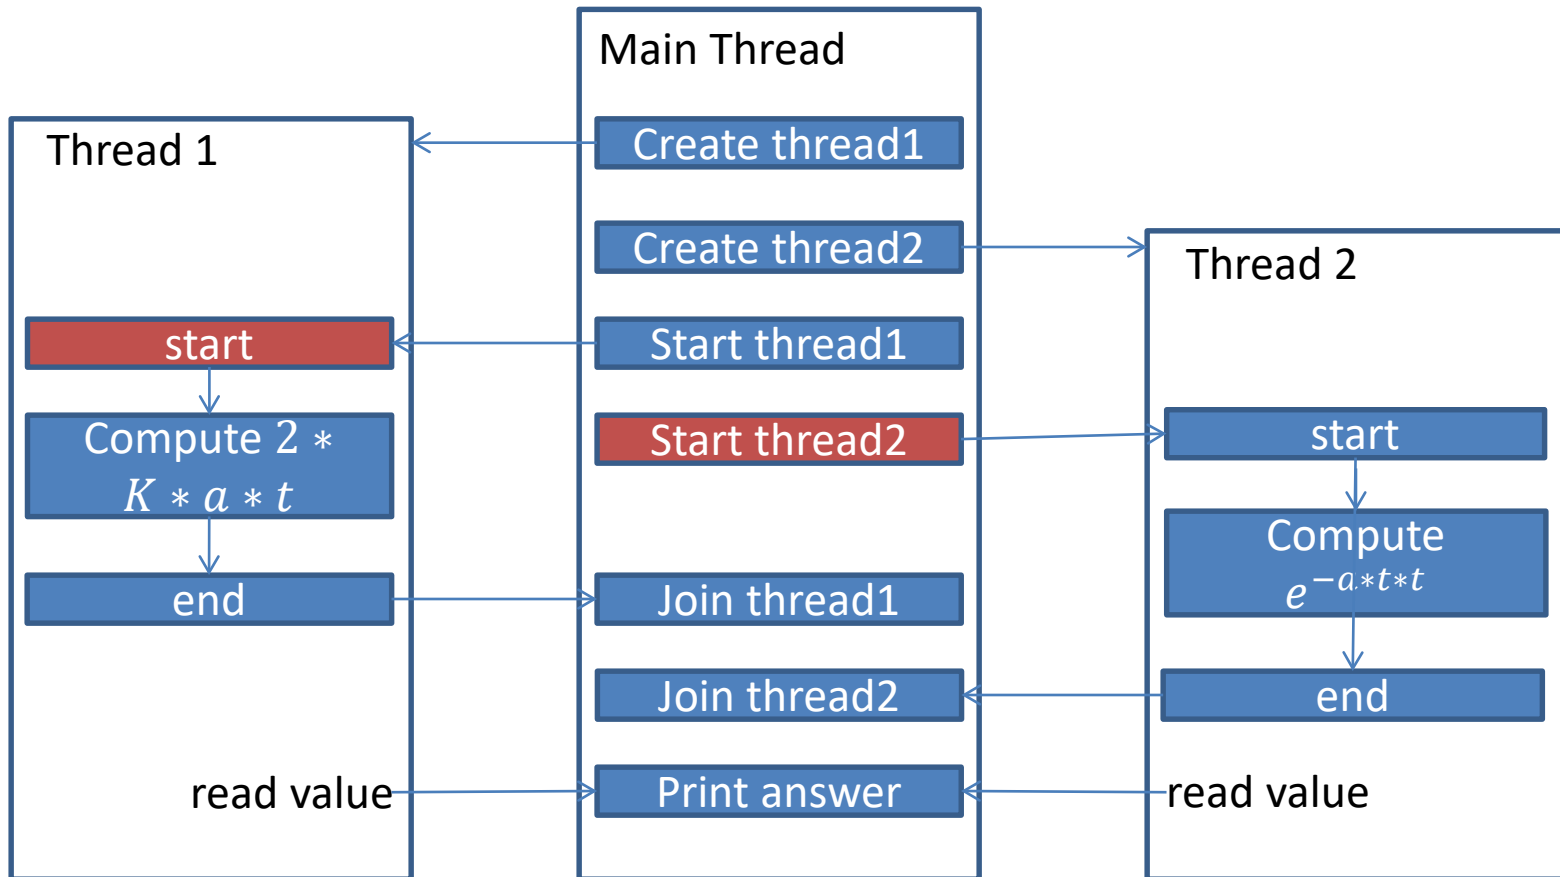
Hello World



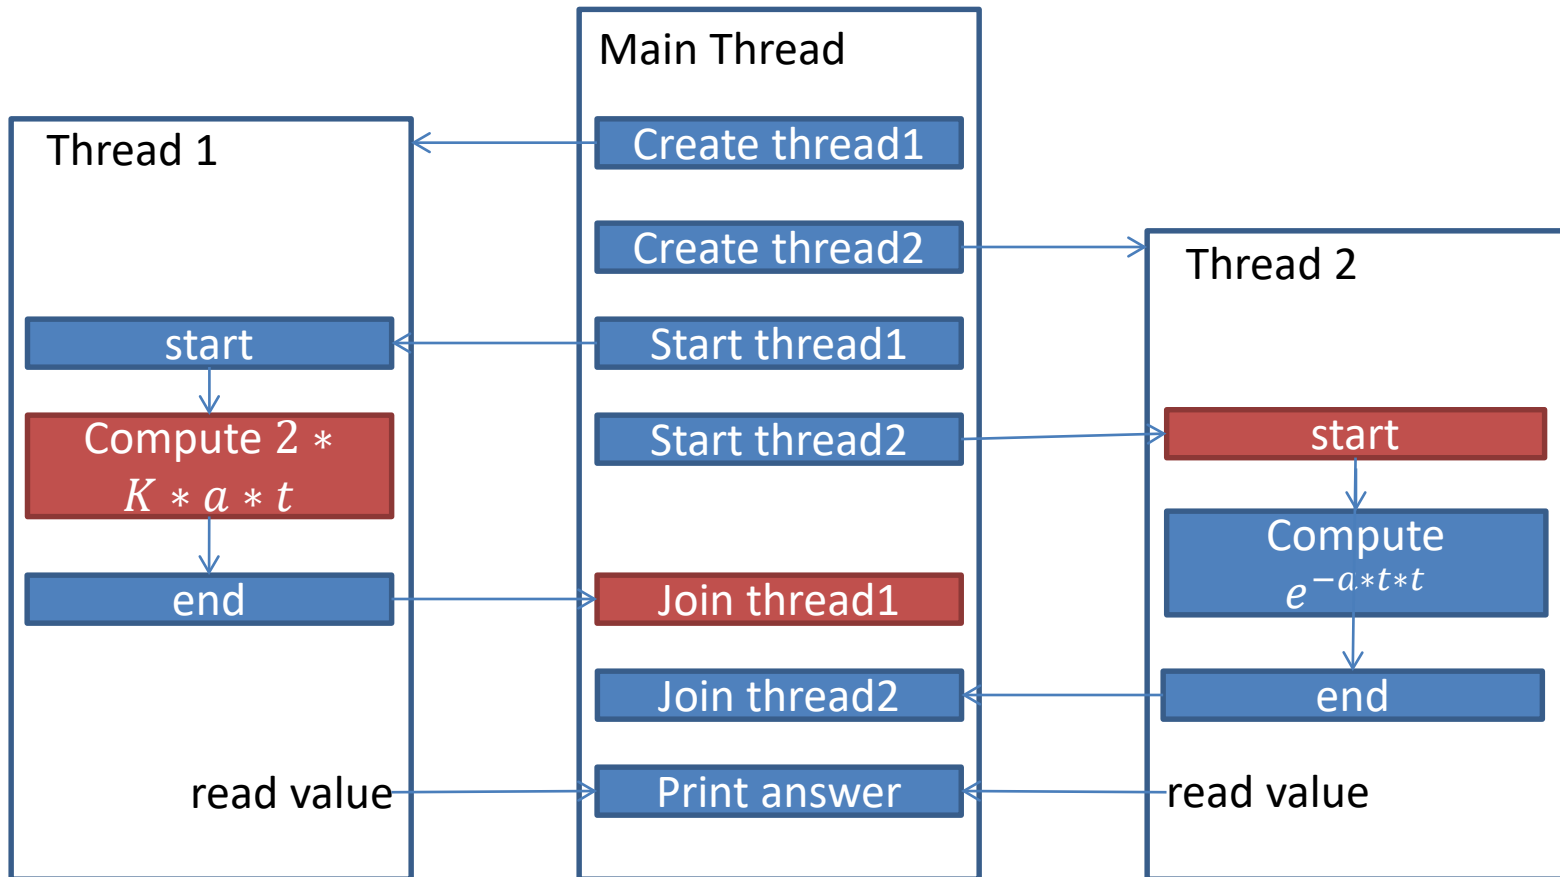
Hello World



Hello World

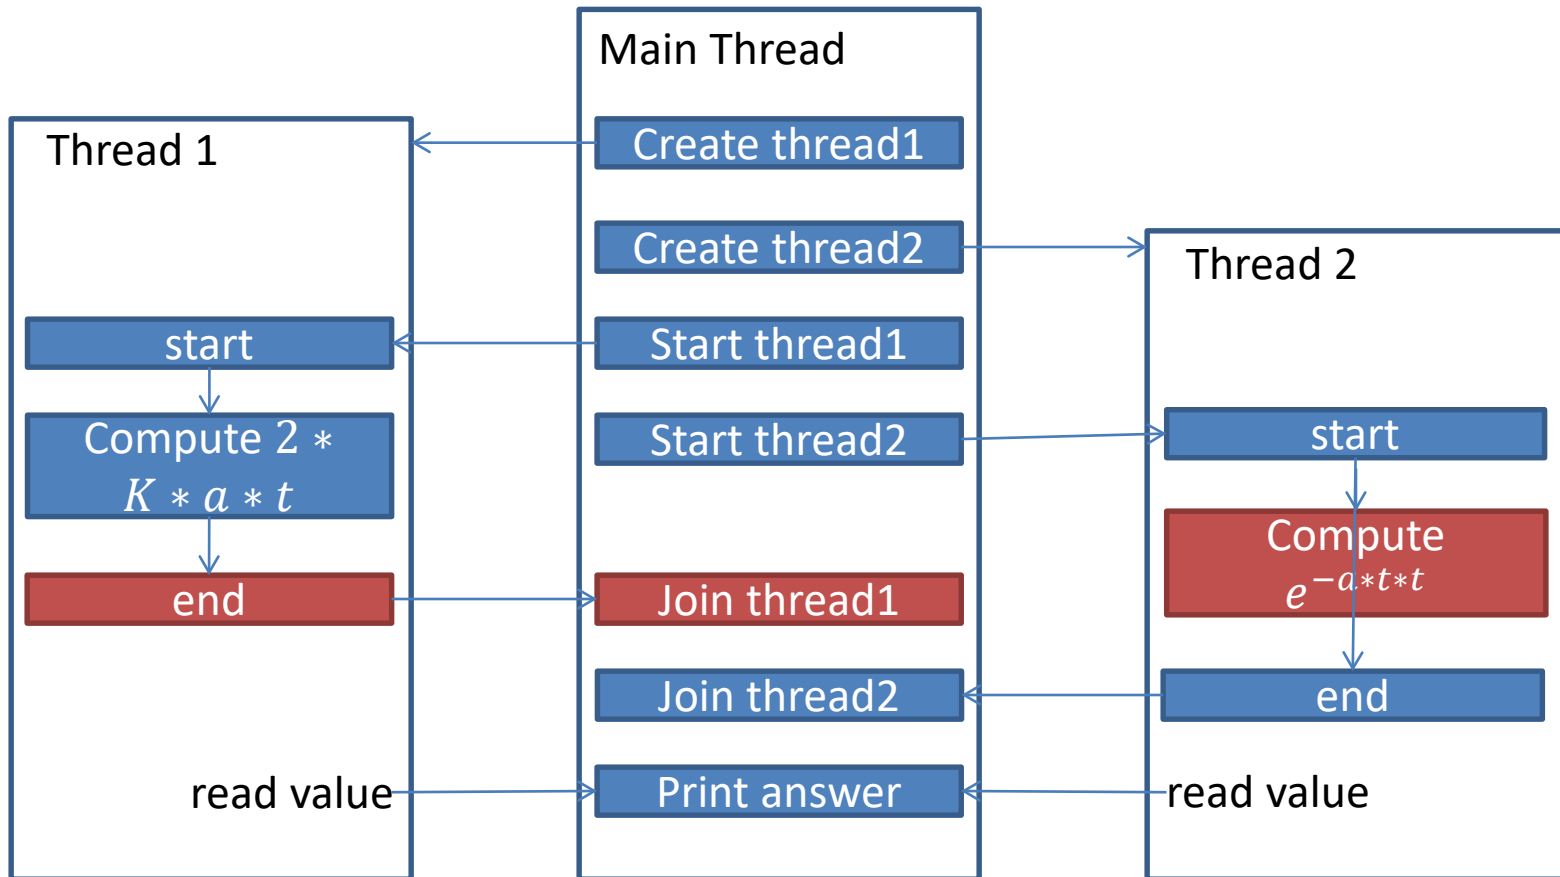


Hello World

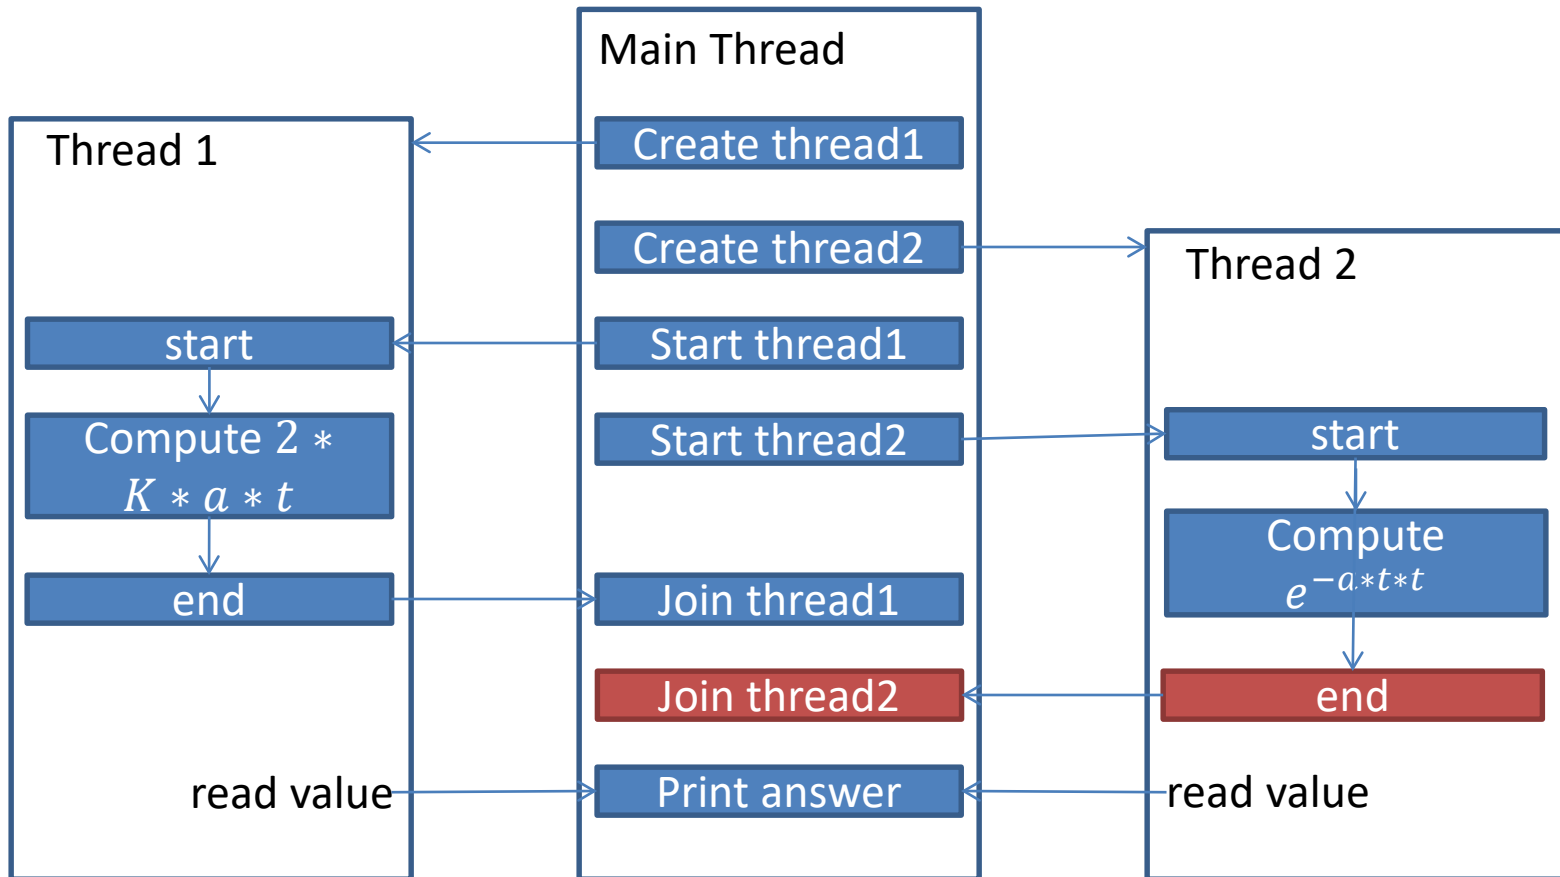


Which thread moves first is non-deterministic

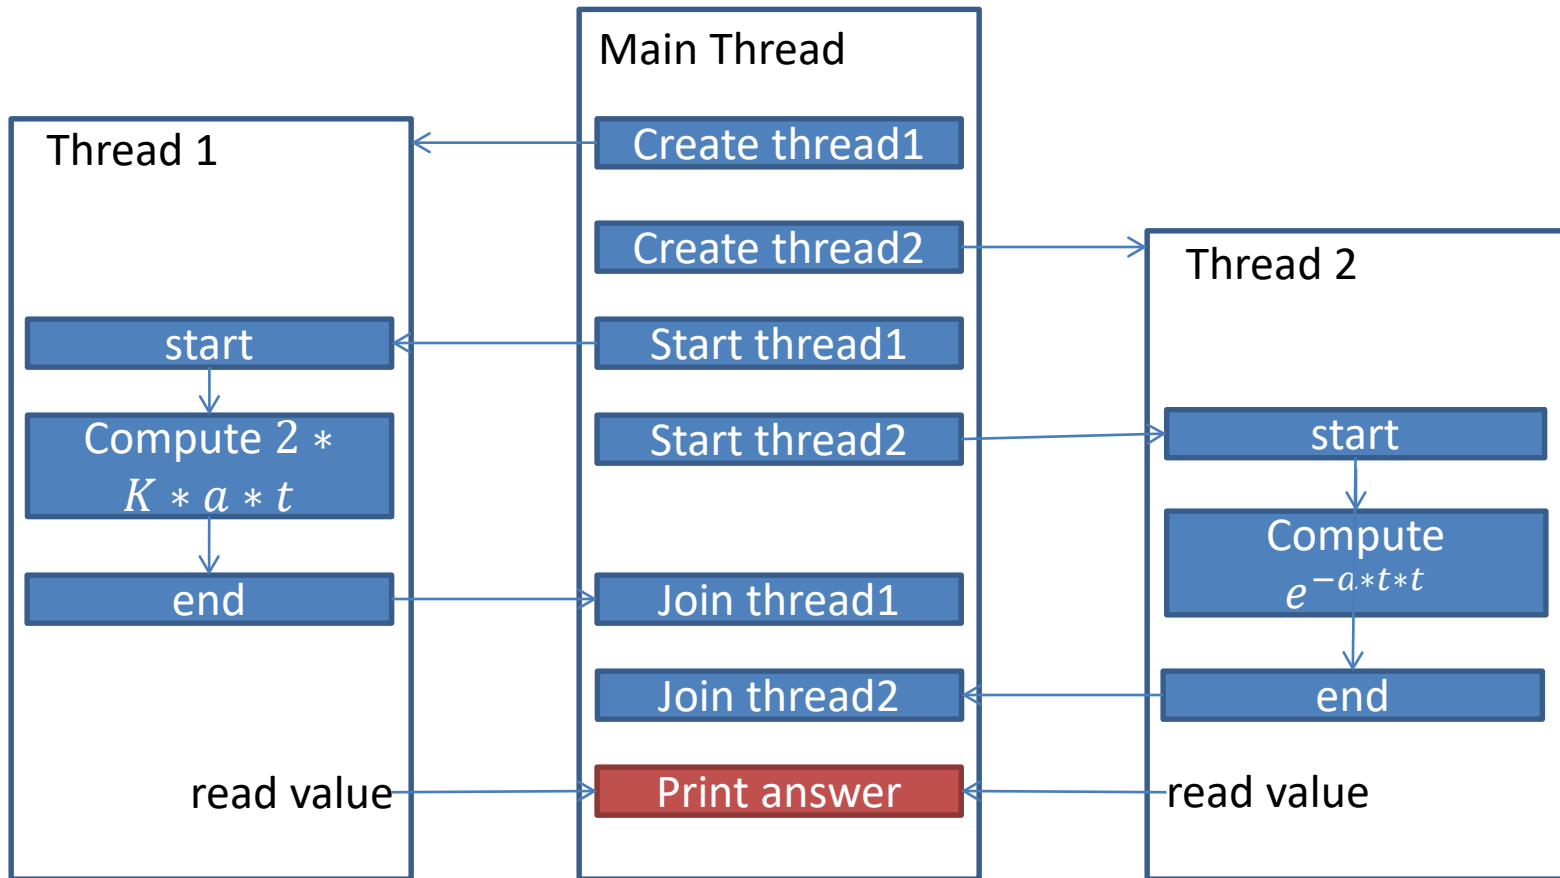
Hello World



Hello World



Hello World



Runnable interface

- “implements Runnable” vs. “extends Thread”
 - Use Runnable that you can inherit other classes
 - Runnable allows better separation of the actual computation and thread control.

RunnableExample.java

Cohort Exercise 1

Given FactorPrime.java, write a program so that multi-threads factor the number at the same time. Print the factor as soon as it is found.

How to Stop a Thread

Given FactorPrime.java, write a program so that multi-threads factor the number at the same time. Print the factor as soon as it is found.

How do we stop a thread as soon as a factor has been found?

How to Stop a Thread

- No
 - `destroy()`
 - or `stop()`
 - or `stop(Throwable obj)`
 - or `suspend()`
- Yes
 - `Interrupt()`

interrupt()

- Example 1: interrupt() results in Exception
- Example 2: handle interrupt() explicitly in run() by checking Thread.interrupted()

InterruptExample1.java and InterruptExample2.java

Cohort Exercise 2

Continue with Exercise 1, modify your program such that once the factors are found, stop all threads as soon as possible by interrupting the threads. (Use a static Boolean variable found for communication between threads.)

Test your program with FactorPrimeTest.java.

It's Not That Simple

- Multi-threaded programs could have race conditions, visibility issues, deadlocks, etc.

Problem 1: Race Condition

Example

Program `FirstError.java` has 10000 threads which concurrently increment a static variable (initially 0) by 1.

What is the Problem?

- A sequential program consisted of a sequence of instructions (and a memory), where each instruction executed one after the other (to modify the memory, etc.).
- The sequential paradigm has the following two characteristics: the textual order of statements specifies their order of execution; successive statements must be executed without any overlap (in time) with one another.

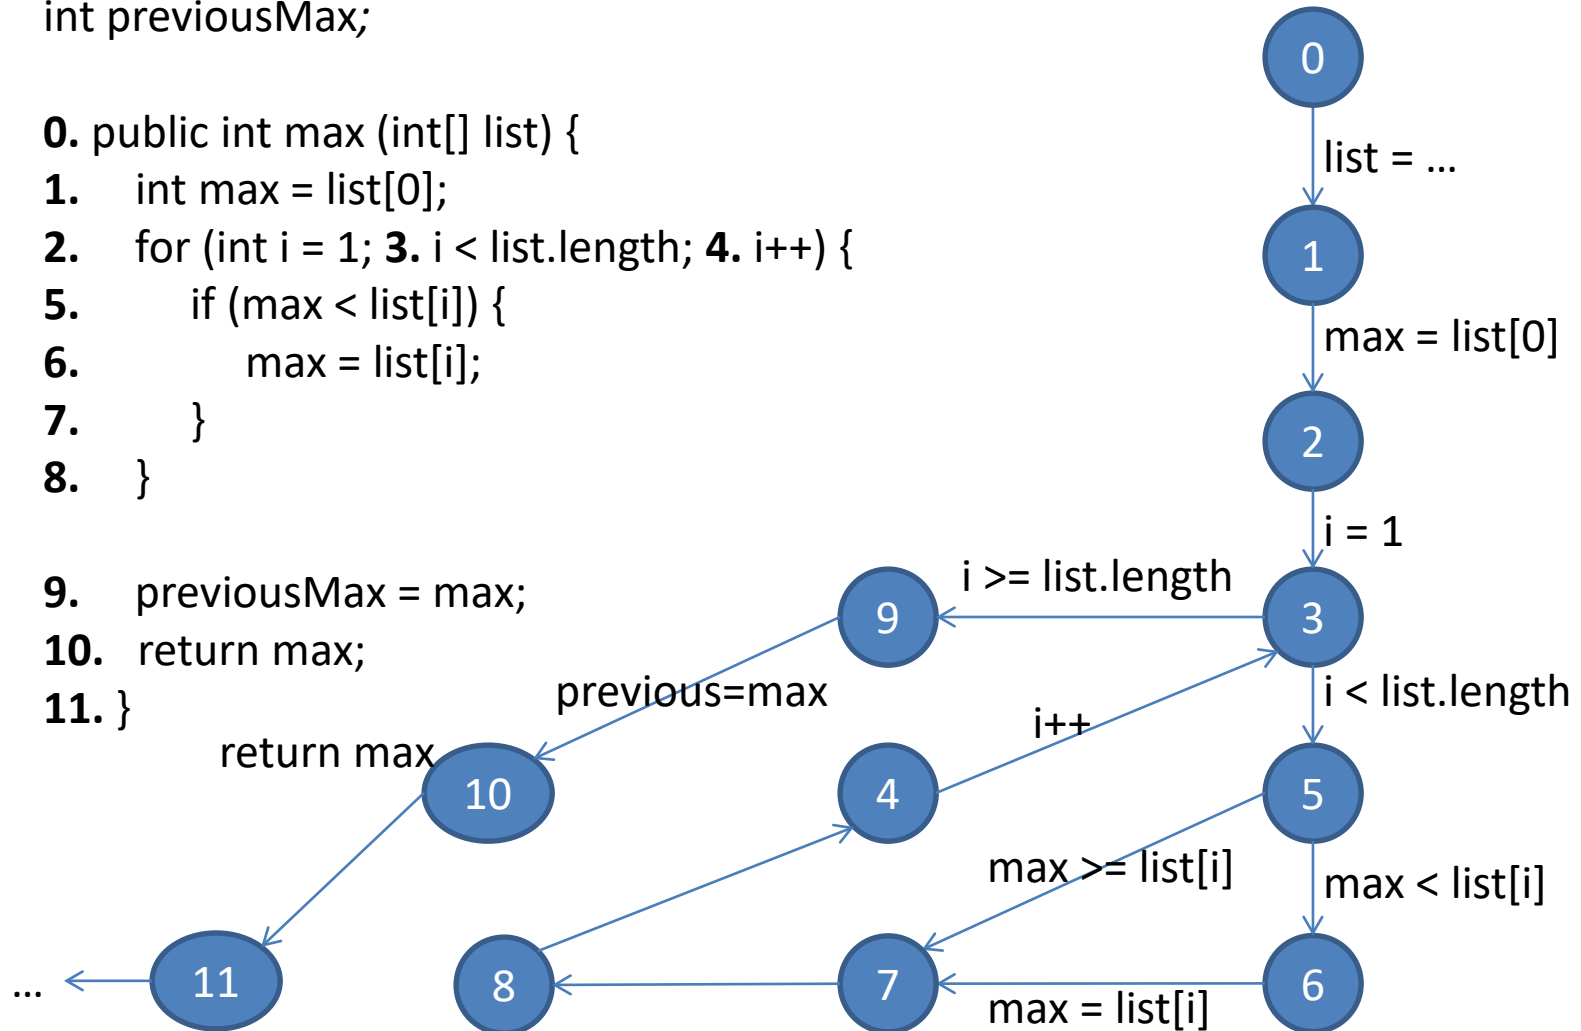
Both are not true in concurrent programs.

The Illusion

int previousMax;

```
0. public int max (int[] list) {  
1.   int max = list[0];  
2.   for (int i = 1; 3. i < list.length; 4. i++) {  
5.     if (max < list[i]) {  
6.       max = list[i];  
7.     }  
8.   }
```

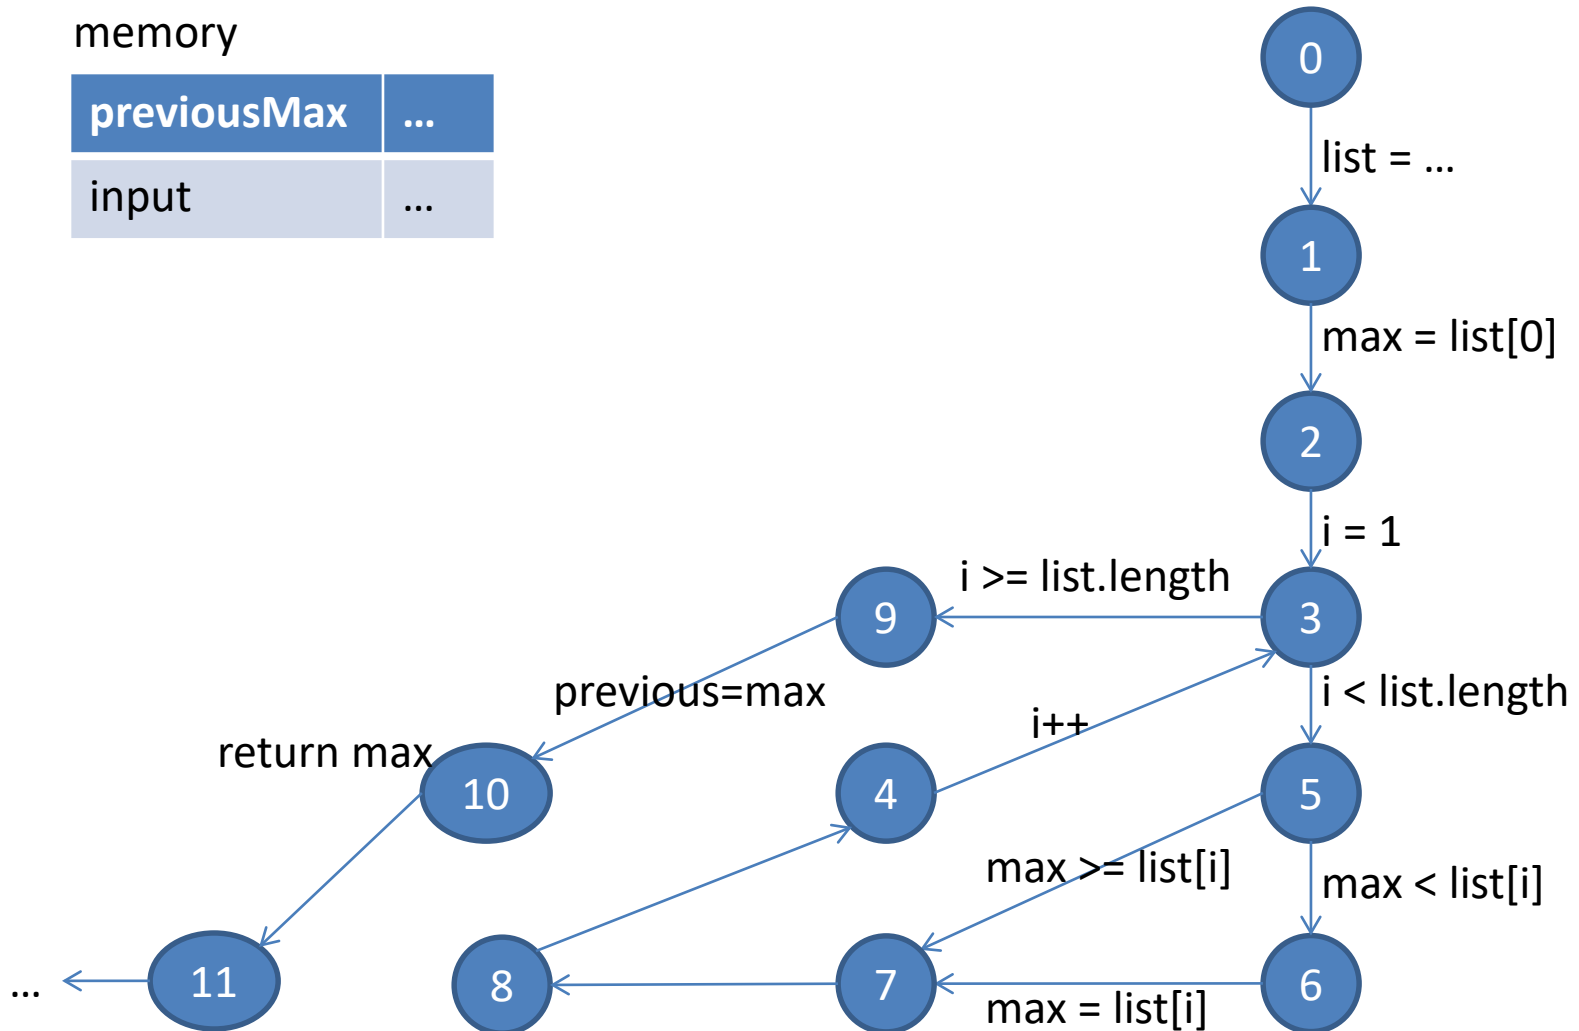
```
9.   previousMax = max;  
10.  return max;  
11. }
```



Control Flow Graph

memory

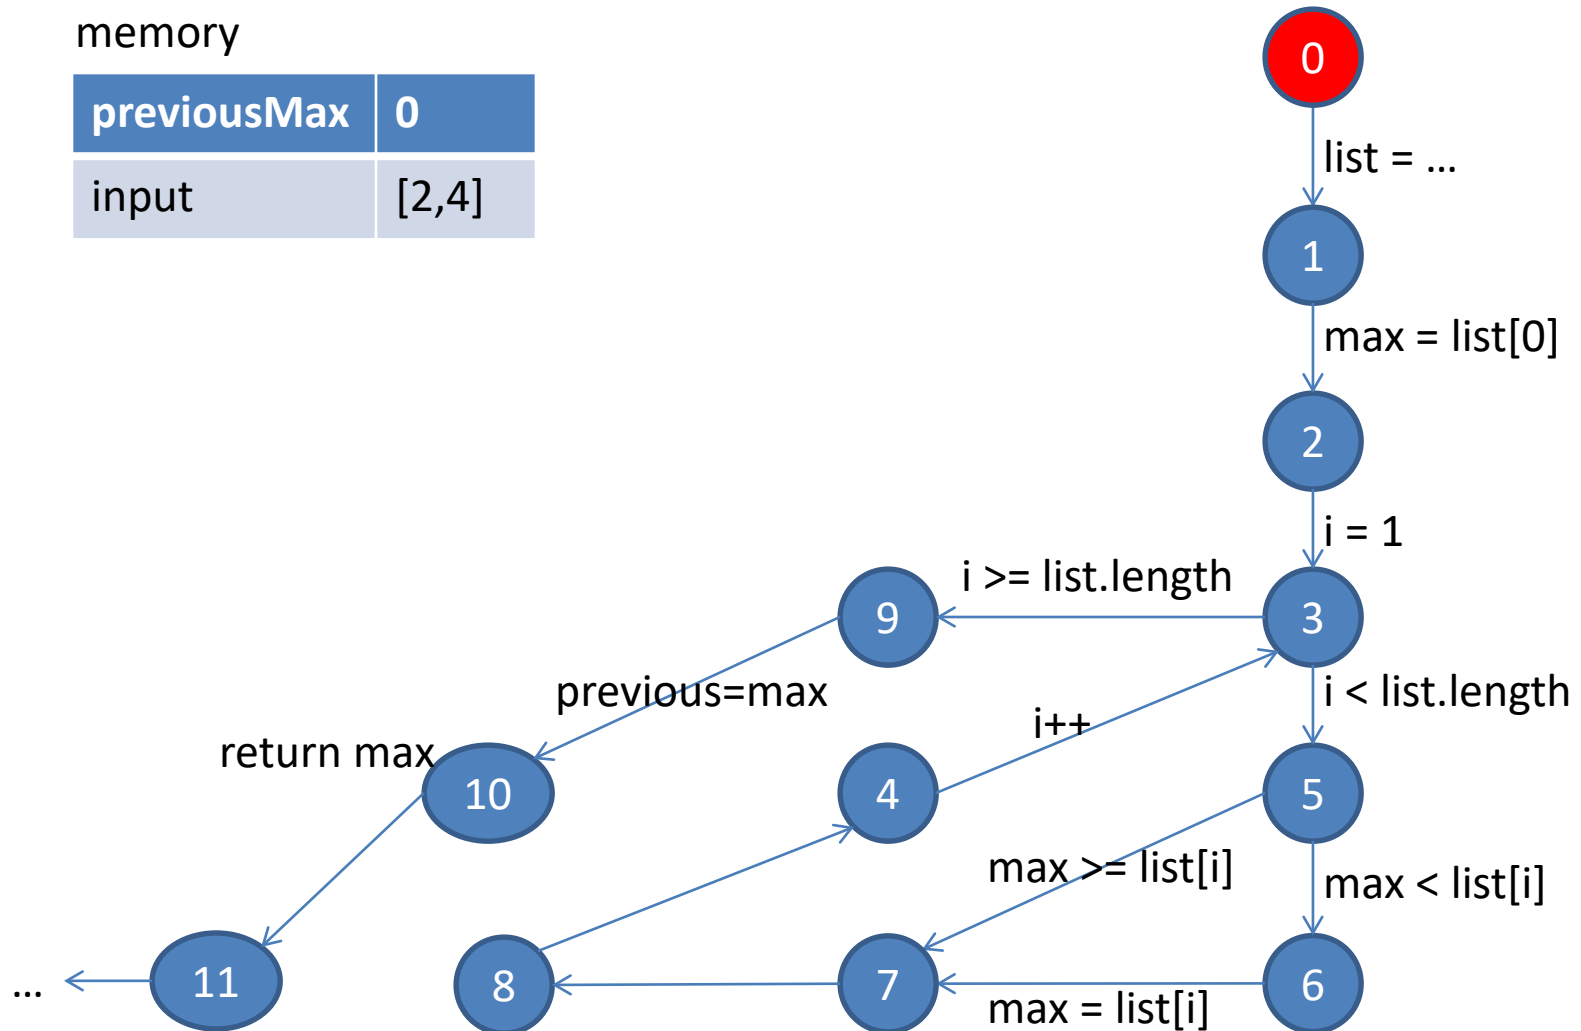
previousMax	...
input	...



System Execution

memory

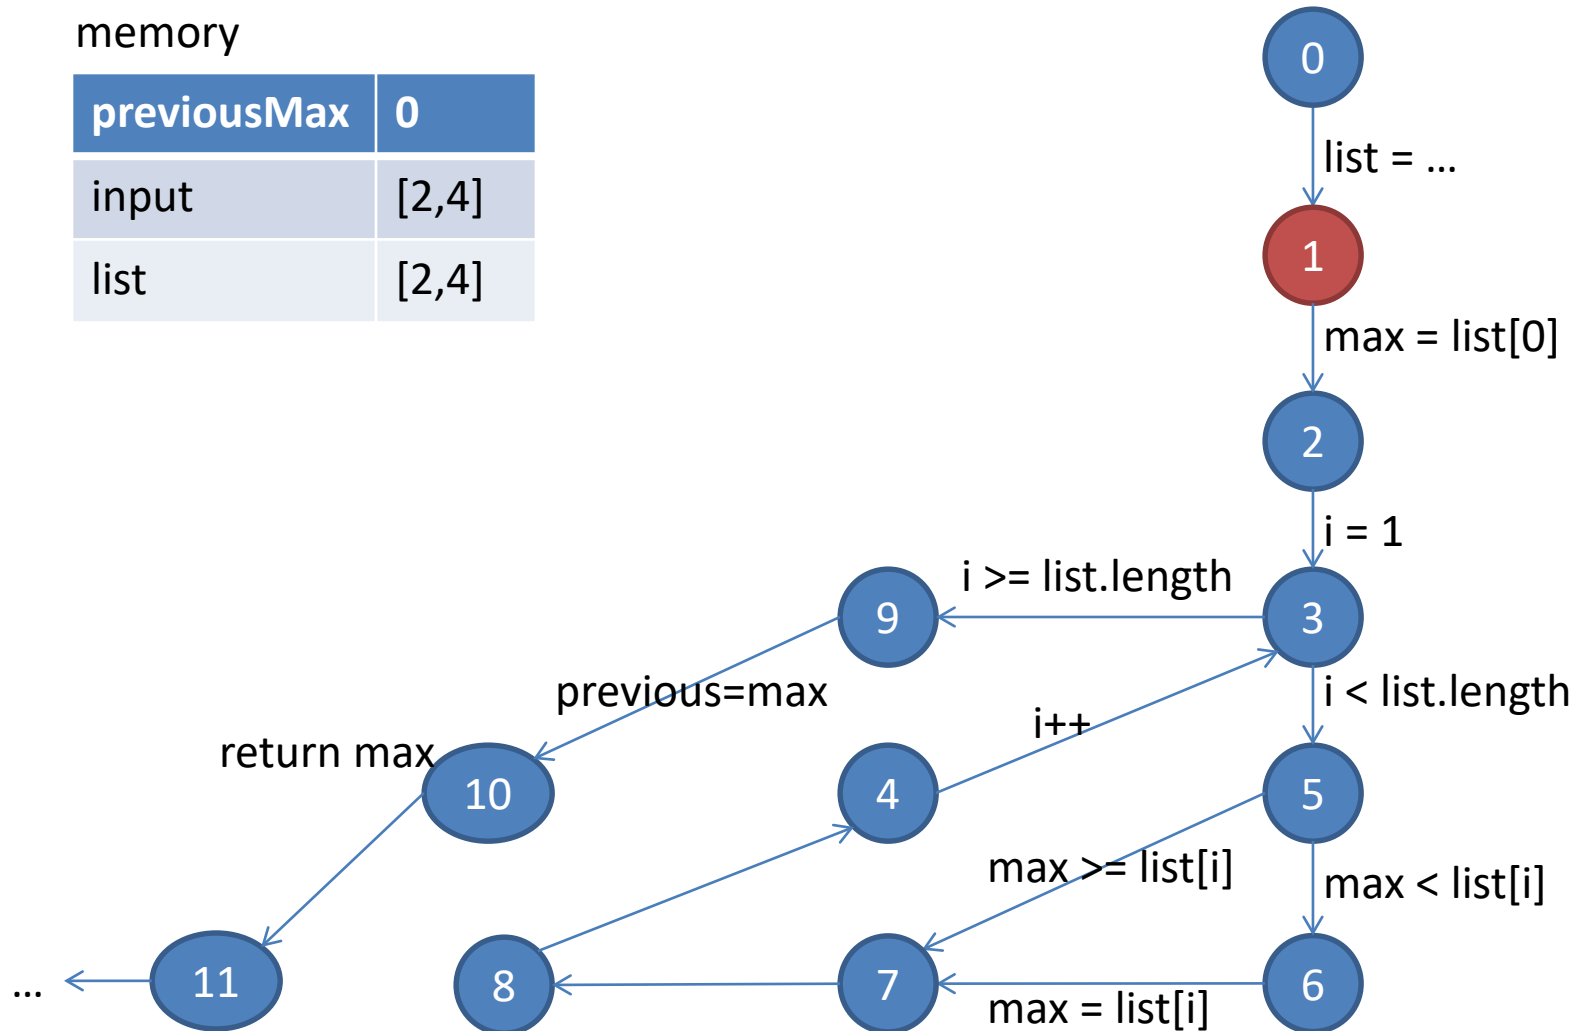
previousMax	0
input	[2,4]



System Execution

memory

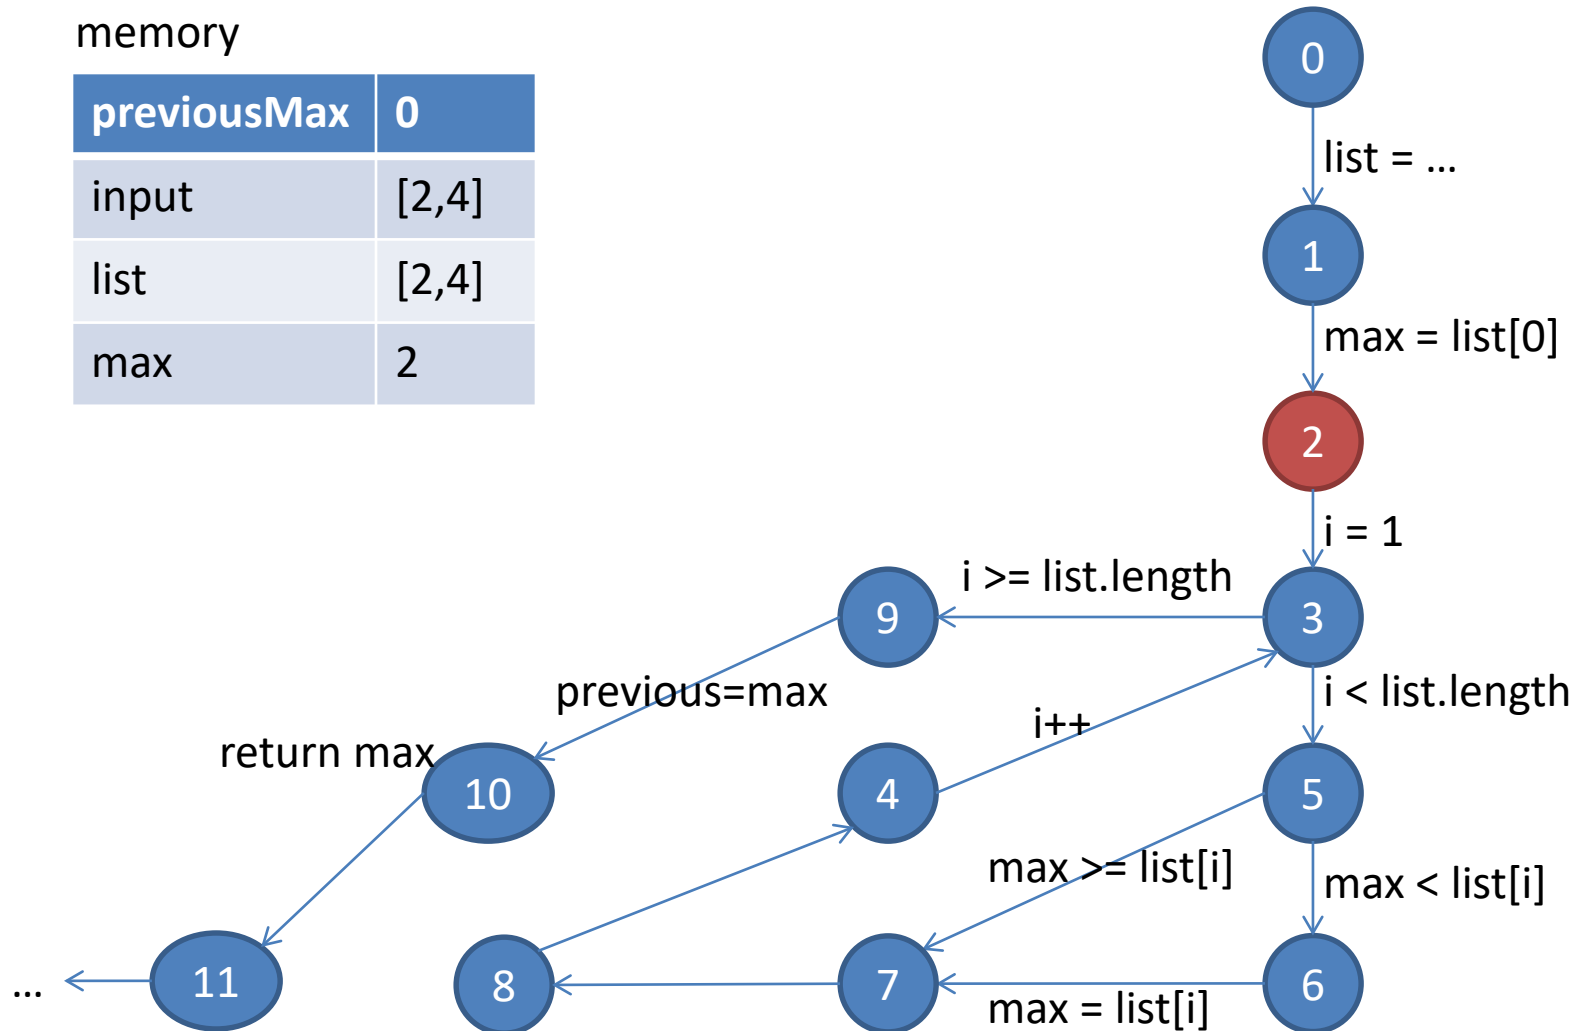
previousMax	0
input	[2,4]
list	[2,4]



System Execution

memory

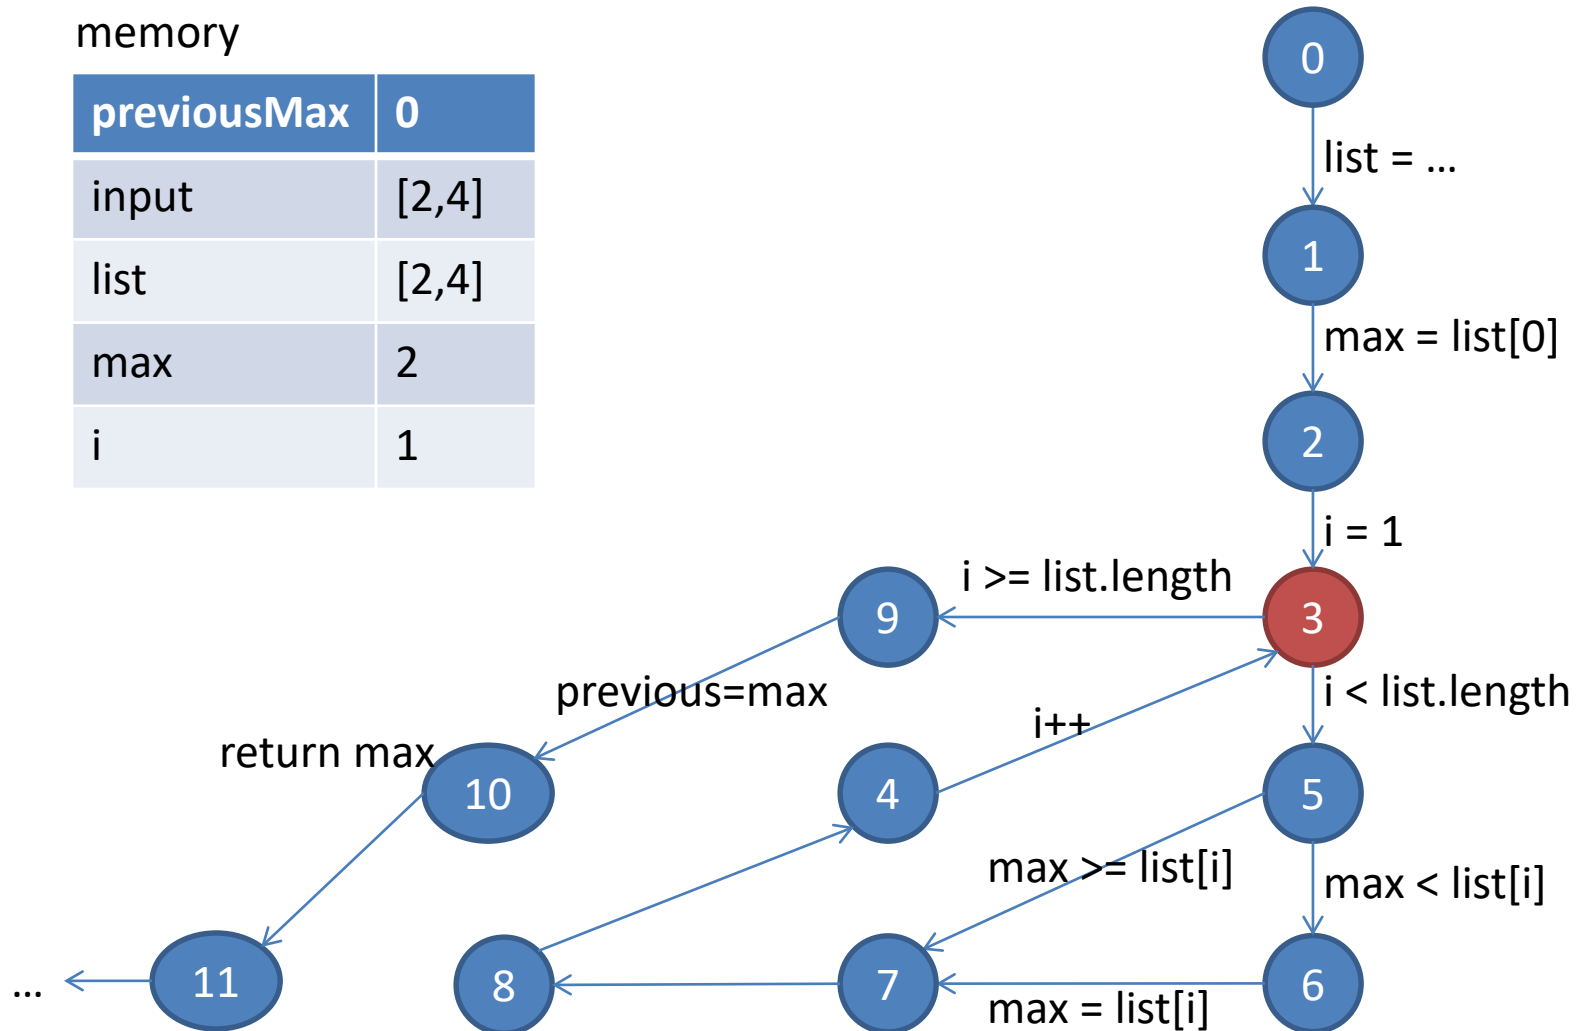
previousMax	0
input	[2,4]
list	[2,4]
max	2



System Execution

memory

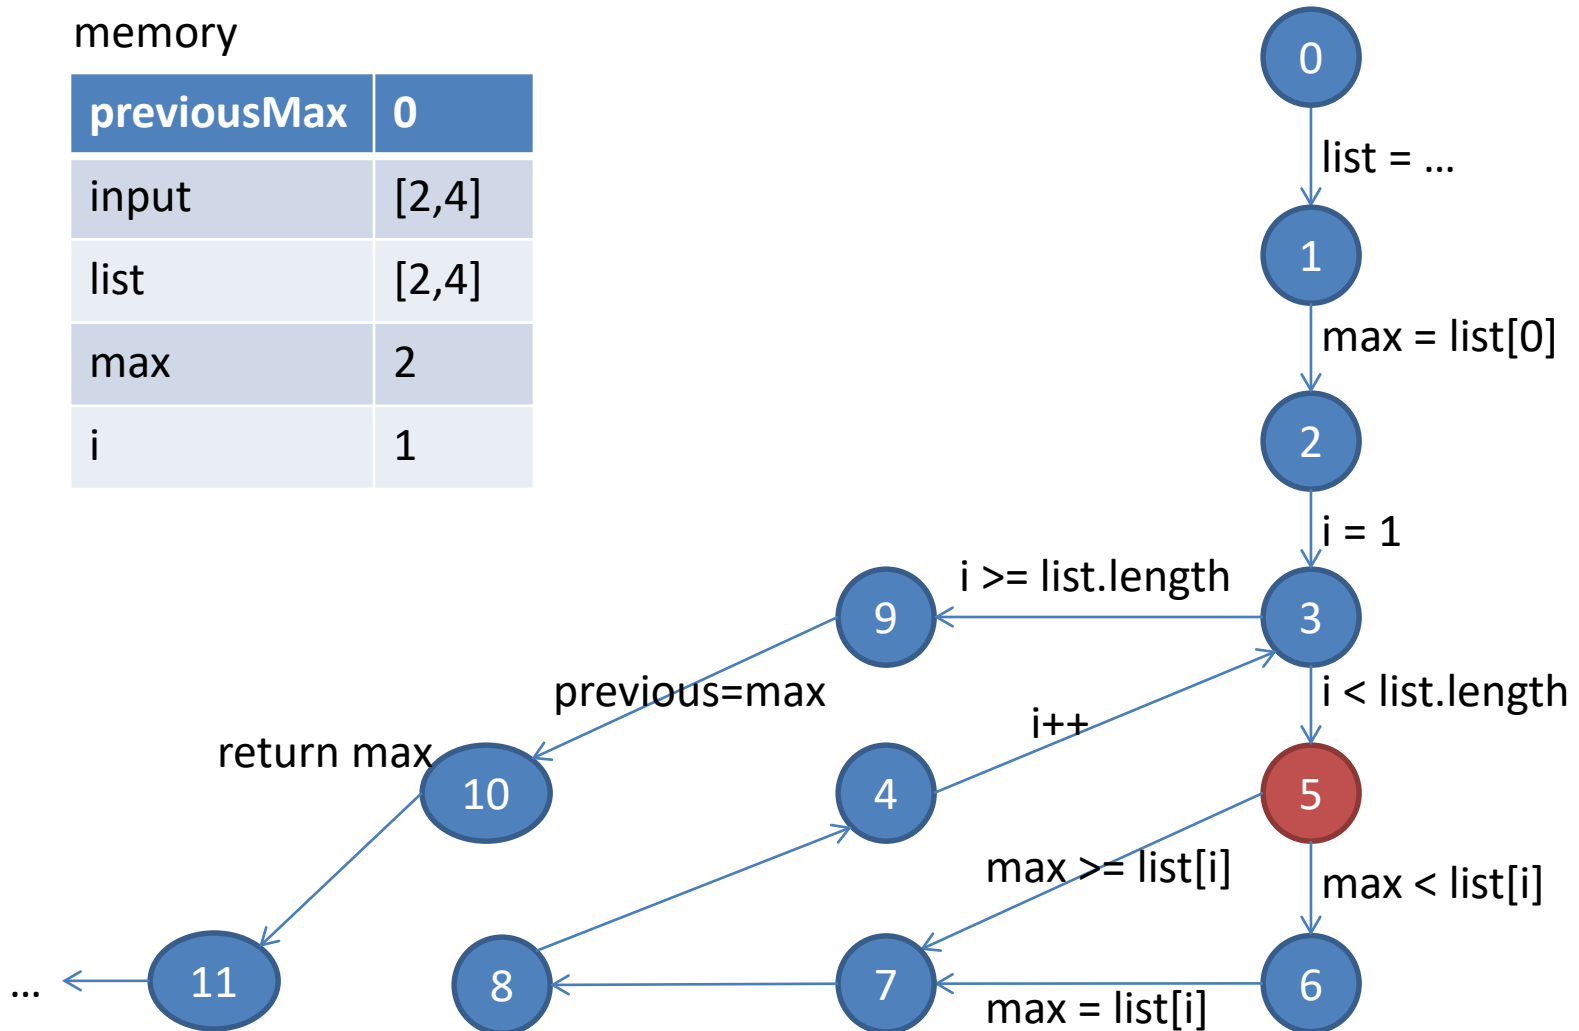
previousMax	0
input	[2,4]
list	[2,4]
max	2
i	1



System Execution

memory

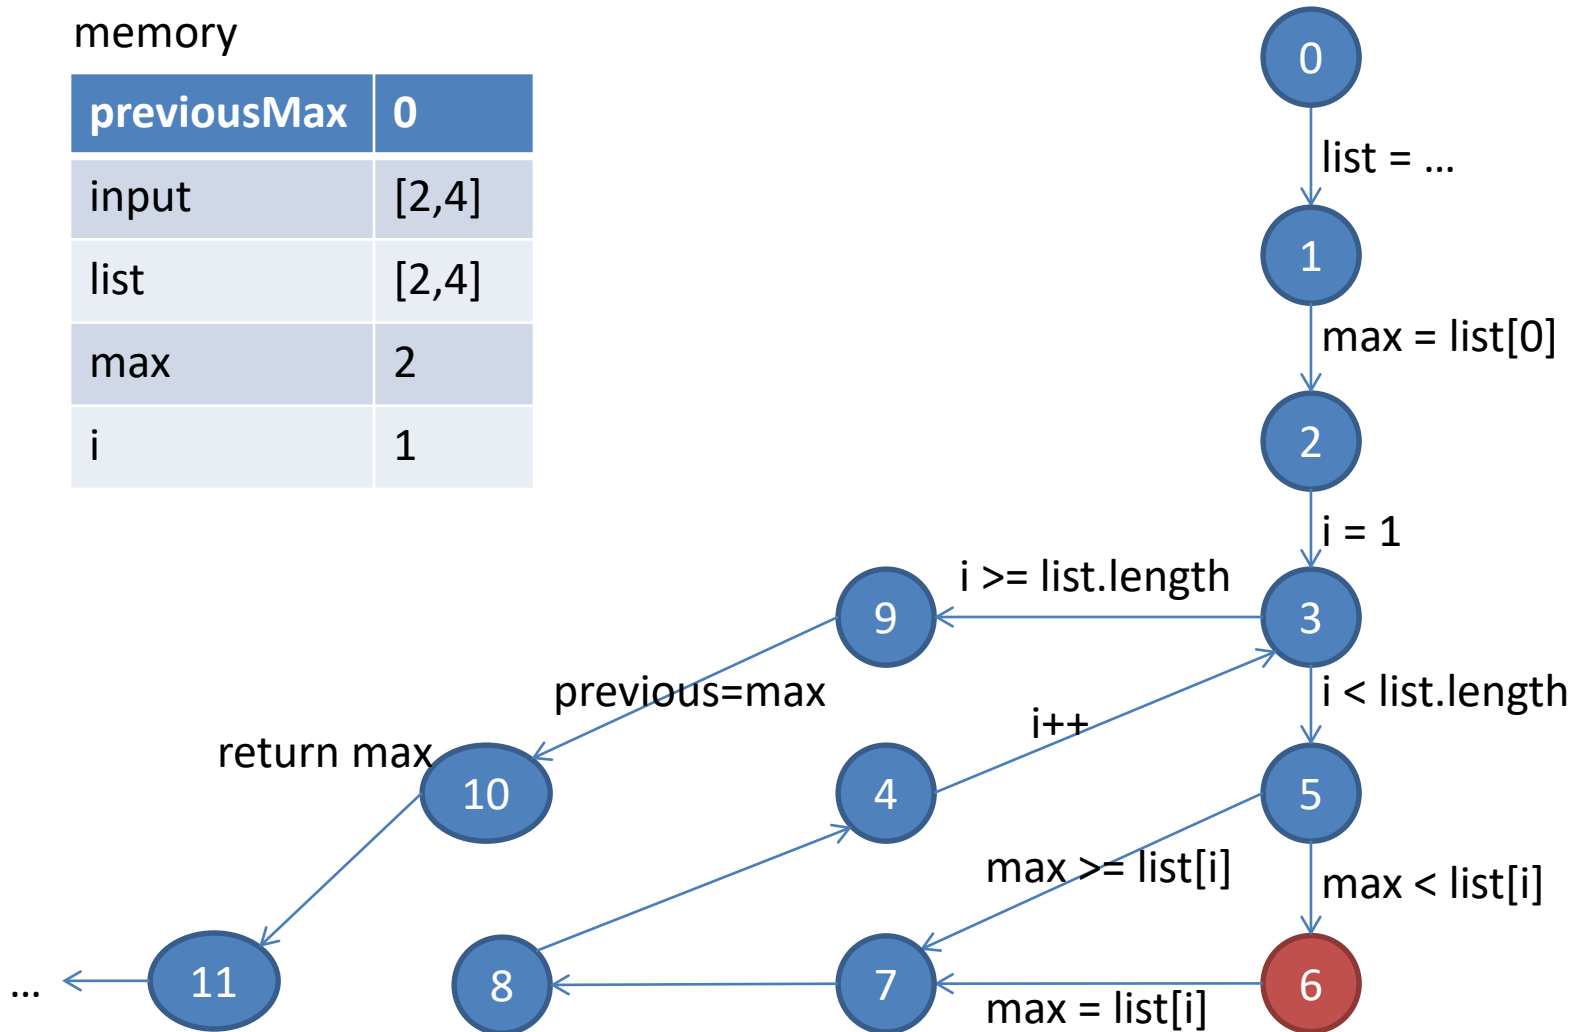
previousMax	0
input	[2,4]
list	[2,4]
max	2
i	1



System Execution

memory

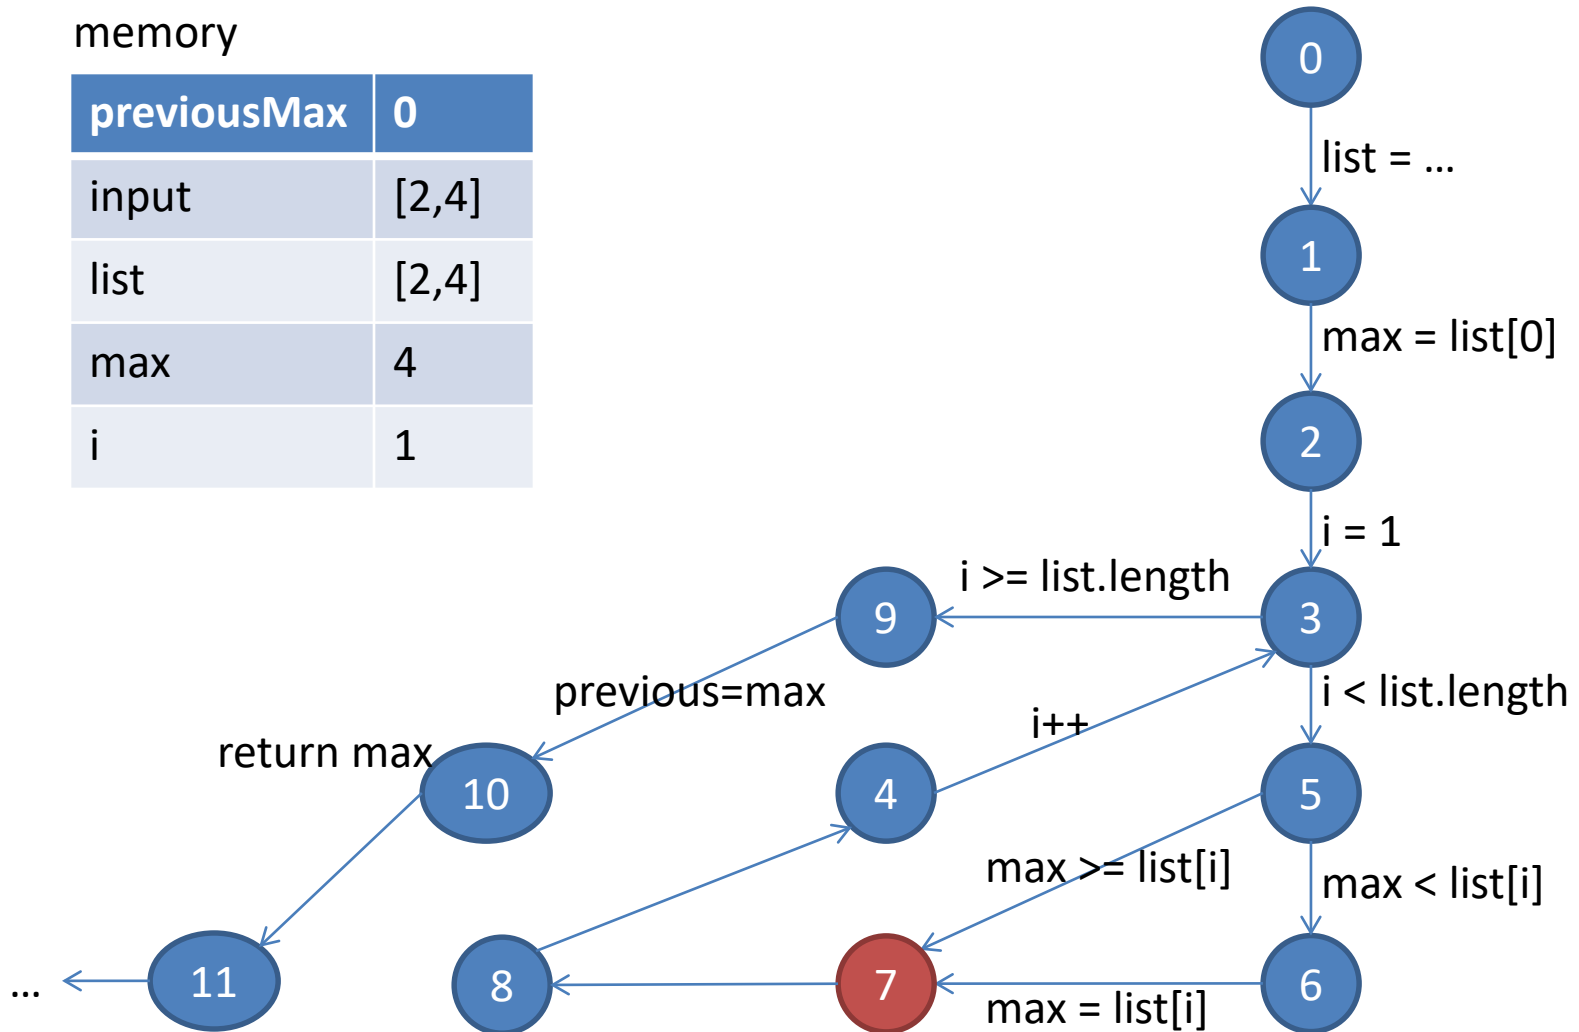
previousMax	0
input	[2,4]
list	[2,4]
max	2
i	1



System Execution

memory

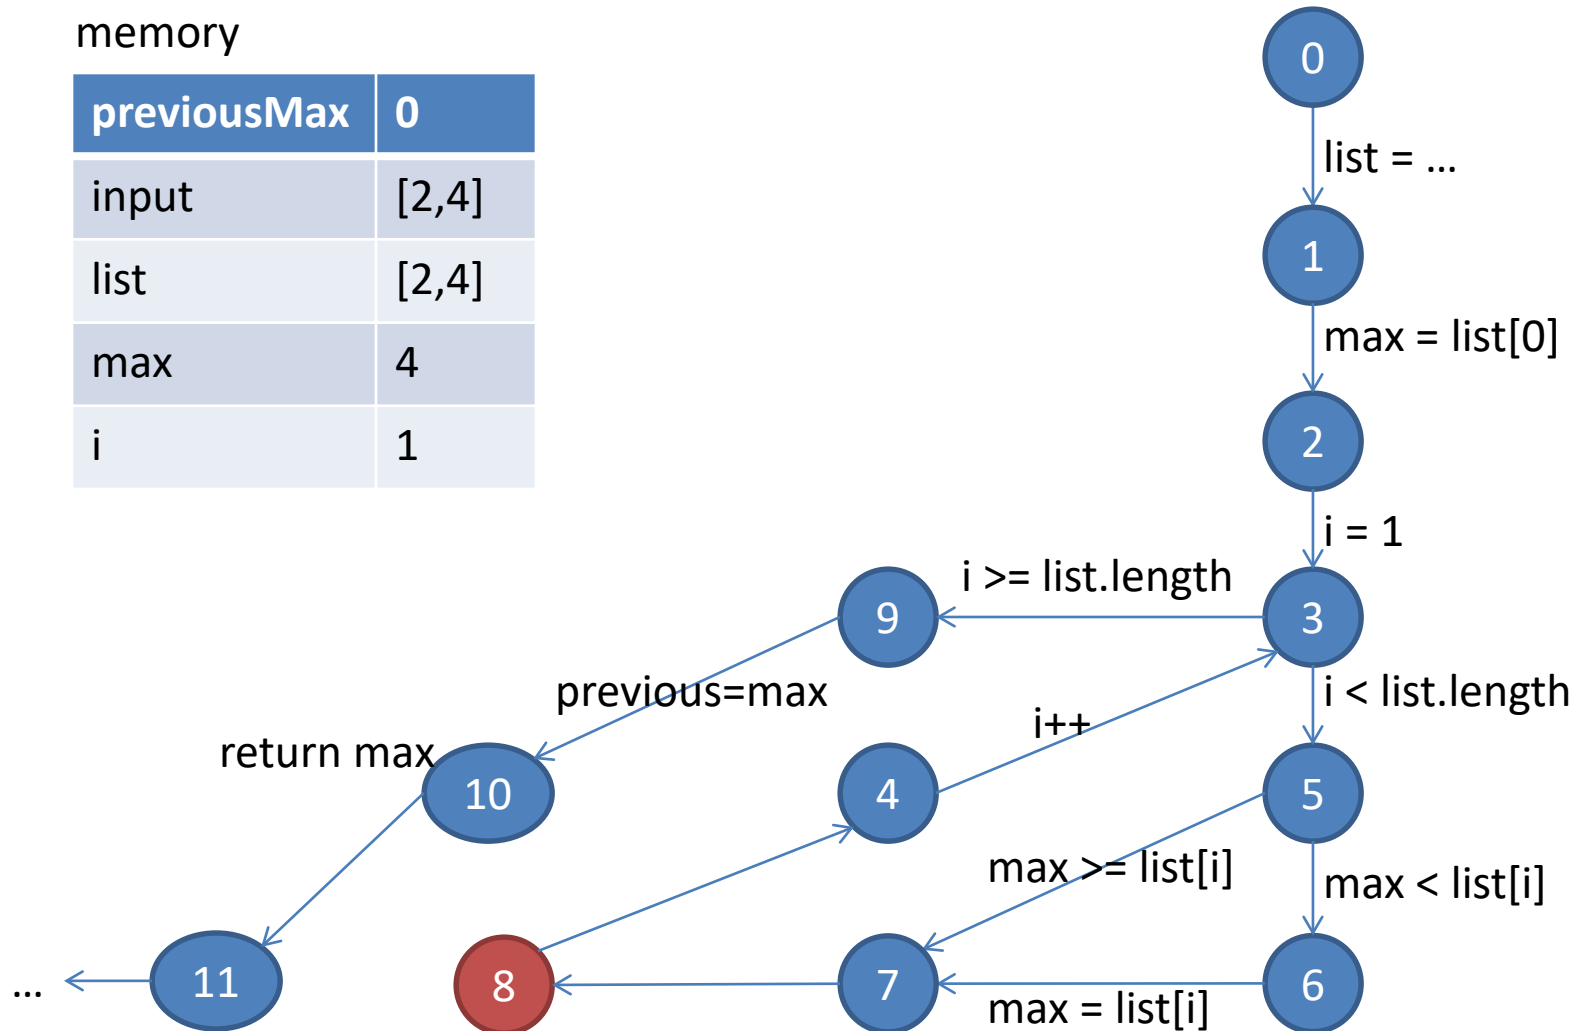
previousMax	0
input	[2,4]
list	[2,4]
max	4
i	1



System Execution

memory

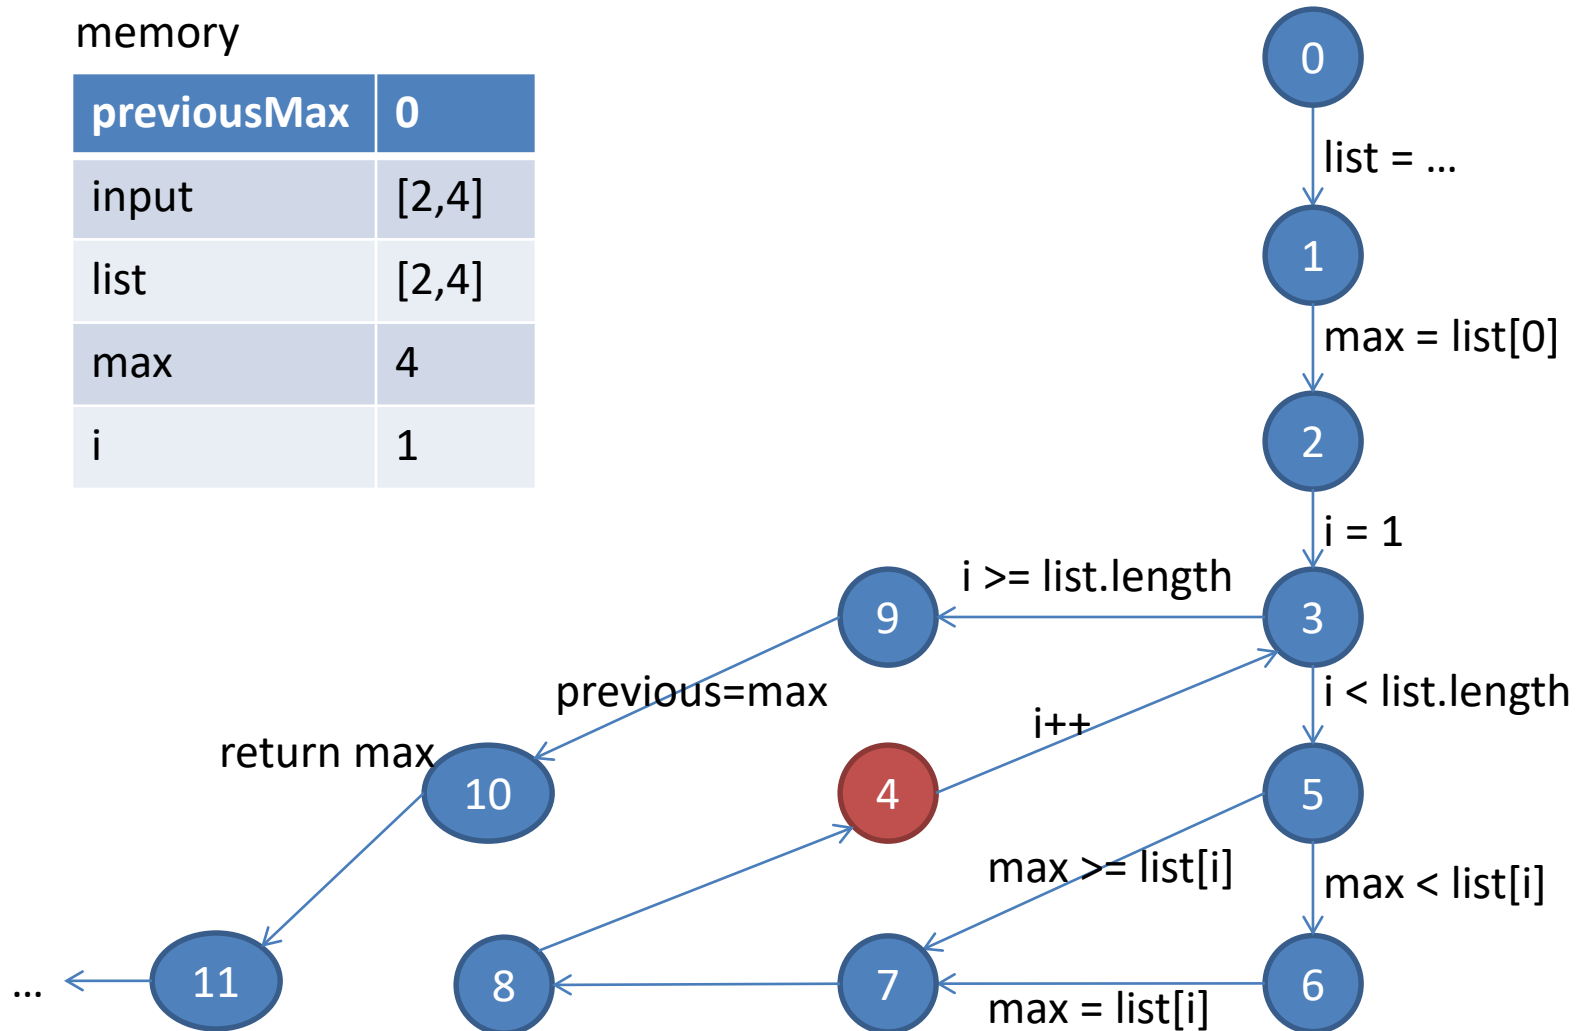
previousMax	0
input	[2,4]
list	[2,4]
max	4
i	1



System Execution

memory

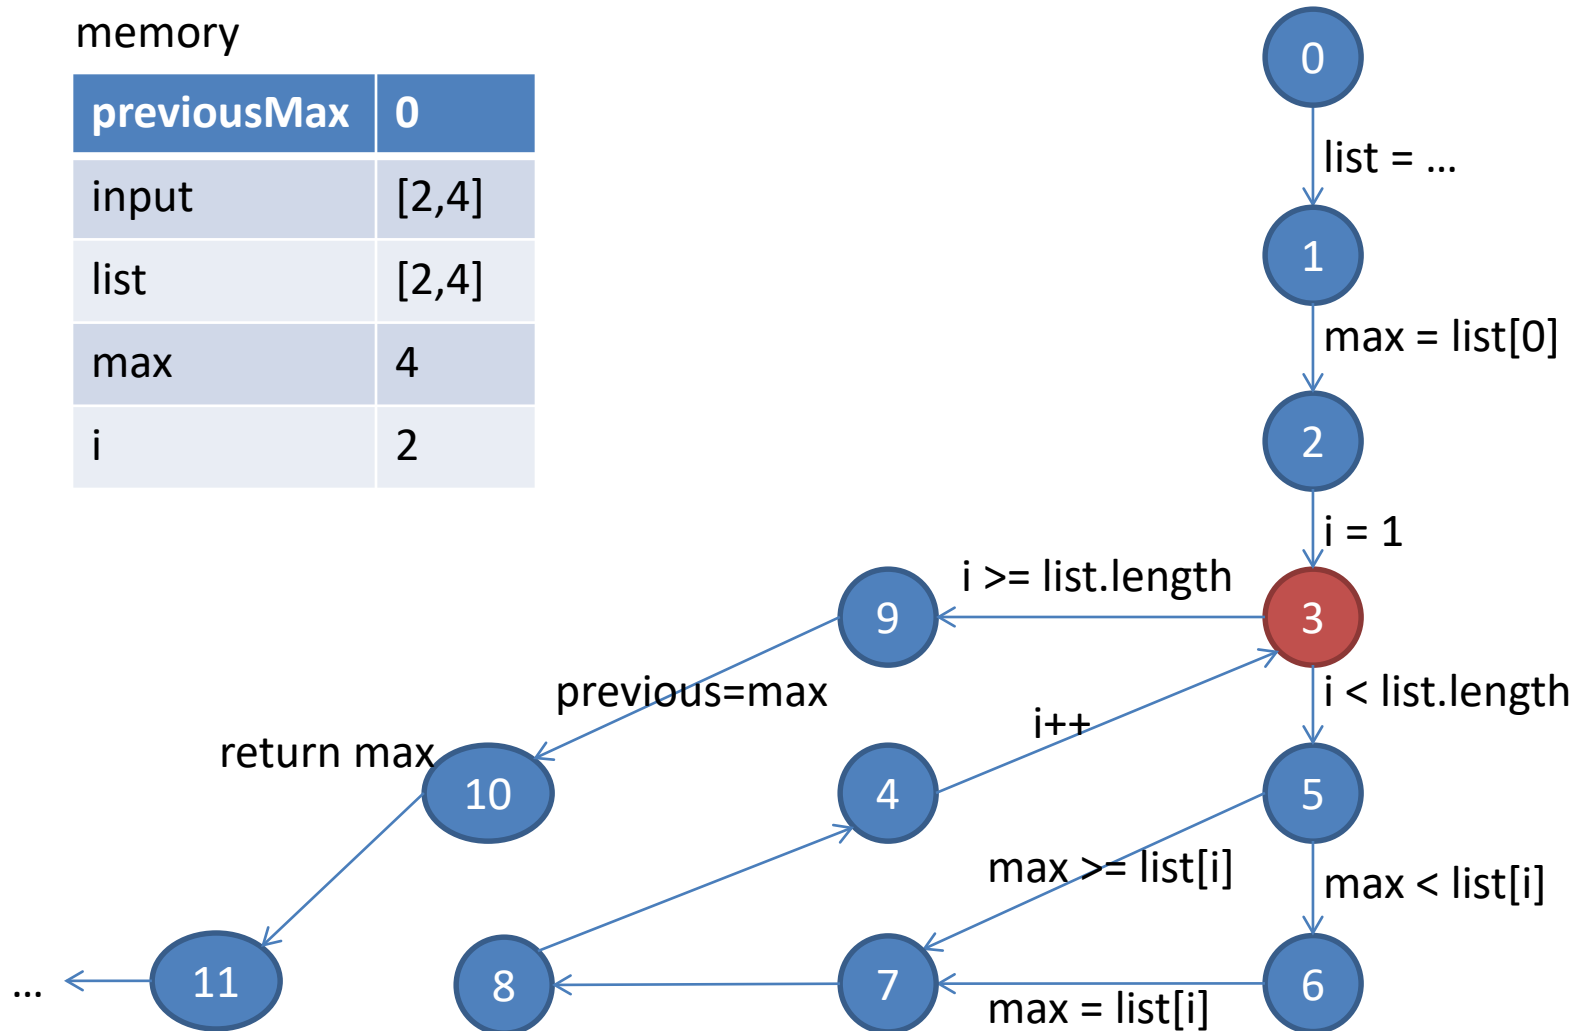
previousMax	0
input	[2,4]
list	[2,4]
max	4
i	1



System Execution

memory

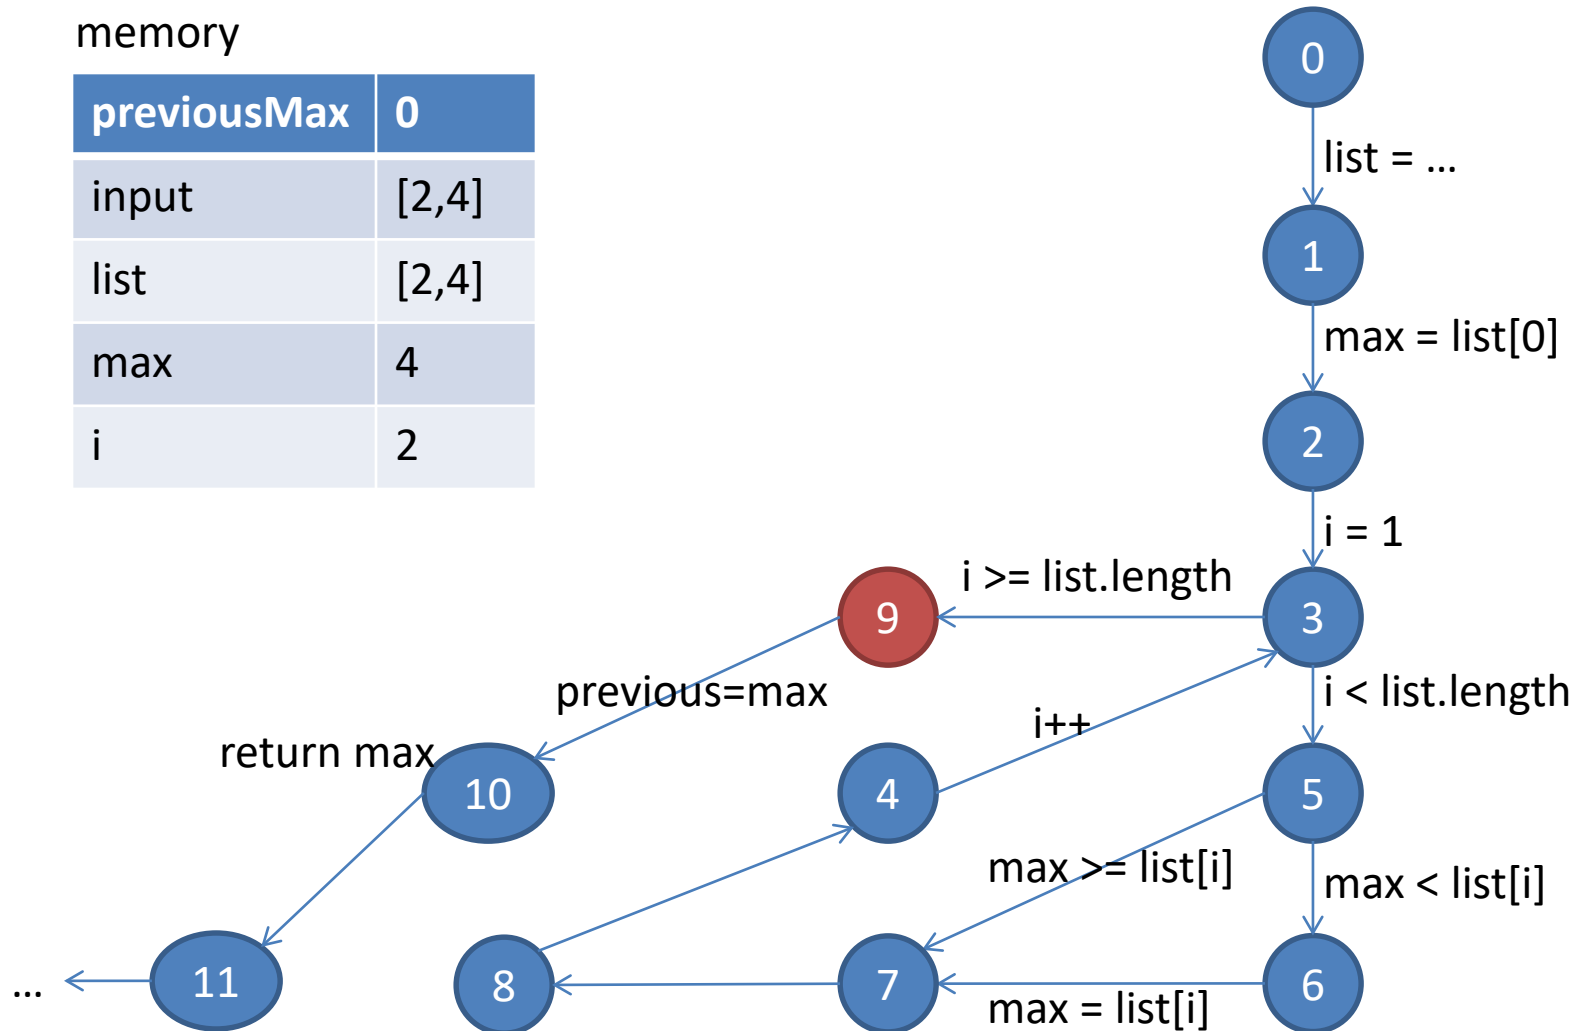
previousMax	0
input	[2,4]
list	[2,4]
max	4
i	2



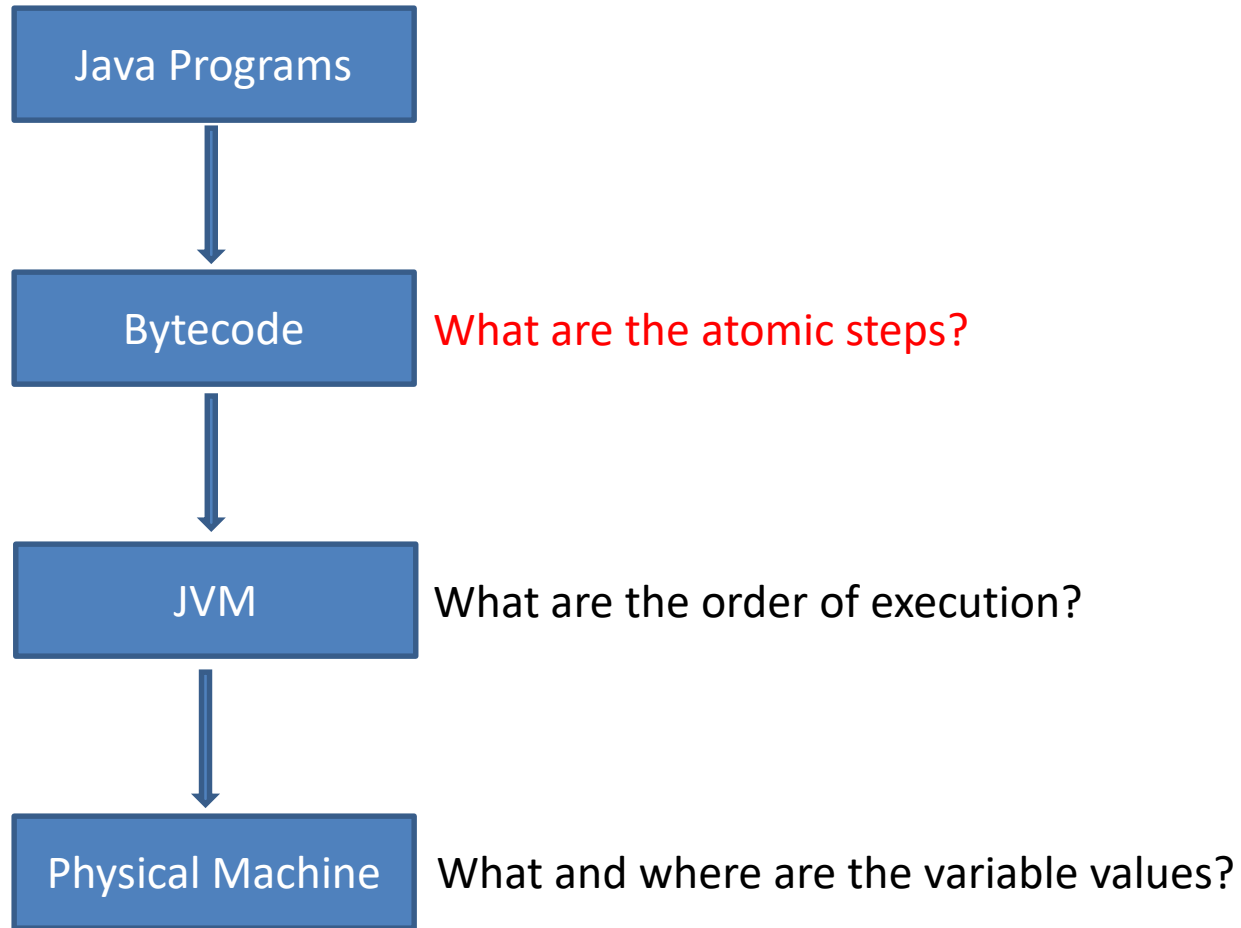
System Execution

memory



previousMax	0
input	[2,4]
list	[2,4]
max	4
i	2



Reality is Messy




Scheduling



Scheduler

The scheduler is 'un-predictable'

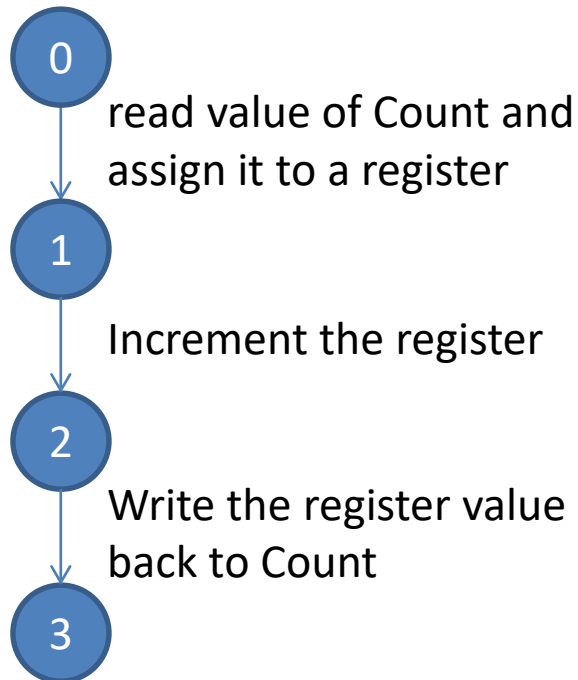
Is This Real?



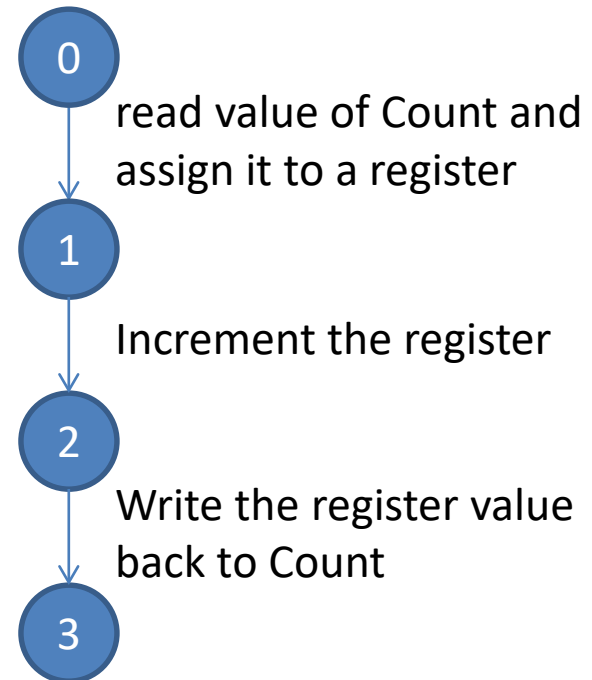
This is assuming that count++ is one step. Or is it?

What Really Happened?

Thread1





Thread2



For double type, even read/write is not atomic!




What Really Happened?

Thread1 Thread2



What Really Happened?

Thread1 Thread2



Is this correct?

A Joke

Some people, when confronted with a problem, think, 'I know, I'll use threads' - and then two they hav erpoblesms.

AtomicXXX

If the **shared variable** is of a type Boolean, int, int array, and etc., use classes in package `java.util.concurrent.atomic`

Example:

```
AtomicInteger x = new AtomicInteger(0)  
x.incrementAndGet() //increments x by 1 atomically
```

Example: `FirstErrorFixed.java`

Compound Actions

Sometime it is not sufficient to simply use an AtomicXXX object.

```
//withdraw from a bank account
```

```
//check and update
```

```
if (amount >= 1000) {
```

```
    amount = amount - 1000;
```

```
}
```

Refer to example: SecondError.java


Intrinsic Locks

- Every Java object can implicitly act as a lock for purposes of synchronization.

```
synchronized (lock) {  
    //Access shared state guarded by lock  
}
```

- Intrinsic locks acts as mutexes (mutual exclusion locks), i.e., at most one thread may own the lock.
- Since only one thread at a time can execute a block of code guarded by a given lock, the synchronized blocks guarded by the same lock execute atomically with respect to one another.

How Lock Works



SecondErrorFixed.java

Cohort Exercise 3

Assuming that we would like to maintain that “saving + cash = 5000” always, fix the following class: LockStaticVariables.java.



Hint: Think about what is the lock?

*if synchronise (this)
it will not work for the amount thing
cause there will be different locks, you are only
sync-ing your thread with ur thread.*

wait() and notify()

- Busy waiting is not efficient
 - Consider a voting system with two threads. One collects votes and the other is waiting to count the votes when the voting is completed.


```
while (true) {  
    synchronized(this) {  
        if (votingComplete) {  
            break;  
        }  
    }  
}
```

- Use wait()/notifyAll() to avoid busy waiting

Voting.java

Cohort Exercise 4


- Producer/Consumer Pattern



- Exercise: fix the Buffer class in `BufferExample.java` so that it is thread-safe and efficient



Problem 2: Visibility



How could we know where?

Example: FactorThread.java

Problem 3: Execution Ordering

Thread 1

1: r2 = A;

2: B = 1;

Thread 2

3: r1 = B;

4: A = 2;

- Initially A, B, r1 and r2 are all 0.
- What are the values of the variables after both threads complete?
- Is it possible to have B = 1 and r2 = 2 and A = 2 and r1 = 1?



What are the order of execution?

- Java compiler might switch the order of sequential statements (e.g., for efficiency)
- Example: line 2 and line 3 might be switched

1. `x++;`




2. `y++;`

3. `x++;`

How could we know the order of execution?
Self-read: Java Memory Model


Problem 4: Deadlock



- Each philosopher needs two forks to eat.
- Each philosopher picks the one on the left first.

Deadlock

- Deadlock is the situation when two or more threads are both waiting for the others to complete, forever.




Cohort Exercise 5

Given DiningPhil.java, modify it so as to demonstrate the deadlock.

Lock-Ordering Deadlock

```
public class LeftRightDeadlock {  
    private final Object left = new Object ();  
    private final Object right = new Object ();  
  
    public void leftRight () {  
        synchronized (left) {  
            synchronized (right) {  
                doSomething();  
            }  
        }  
    }  
  
    public void rightLeft () {  
        synchronized (right) {  
            synchronized (left) {  
                doSomethingElse();  
            }  
        }  
    }  
}
```



Example

```
public void transferMoney (Account from, Account to, int amount) {  
    synchronized (from) {  
        synchronized (to) {  
            if (from.getBalance() < amount) {  
                //raiseException  
            }  
            else {  
                from.debit(amount);  
                to.credit(amount)  
            }  
        }  
    }  
}
```

*deadlock when you try
to transfer from A to B
and B to A.*

*A->B
lock A
wait for B*

*B->A
lock B
wait for A*

Is it deadlocking?



Example

```
public void transferMoney (Account from, Account to, int amount) {  
    synchronized (from) {  
        synchronized (to) {  
            if (from.getBalance() < amount) {  
                //raiseException  
            }  
            else {  
                from.debit(amount);  
                to.credit(amount)  
            }  
        }  
    }  
}
```

How can transferMoney deadlock?

Thread A: transferMoney(myAccount, yourAccount, 1)

Thread B: transferMoney(yourAccount, myAccount, 1)

Check out: DemonstrateDeadlock.java

Cohort Exercise 6

- Given DLExample.java, explain whether it is possibly deadlocking.
- Write a test case which potentially demos the deadlock.



*to explain deadlock, need 2 threads and 2 locks
each thread, which is calling which methods, lock they hold*

if only 1 lock, worst case is you having to wait till the other thread is done.

Other Liveness Hazards

- Deadlock is the most widely encountered liveness hazard.
- Starvation occurs when a thread is denied access to resources it needs in order to make progress.
 - Often caused by use of thread priority or executing infinite loops with a lock held.

Avoid using thread priority, since they increase platform dependence and can cause liveness problems.

Other Liveness Hazards (cont'd)

- Poor responsiveness
 - may be caused by poor lock management.
- Livelock: a thread, while not blocked, still cannot make progress because it keeps retrying an operation that will always fail.
 - e.g., when two overly polite people are walking in the opposite direction in a hallway.

Requirements

Multi-threaded programs have at least the following additional requirements.

- No race condition
- No visibility issue
- No execution ordering problem
- No deadlocks
- Efficiency!

How do we make sure our multi-threaded program satisfies these requirements?