

50.005 (Spring 2019)

Networking Basics

Objectives: What is a network or internetwork? What are the key challenges in designing a large-scale global internetwork? What are some basic design principles that can help overcome the challenges?

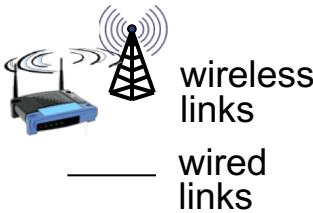
NS1: March 2019

Textbook (**K&R**): Sections 1.1, 1.3, 1.5

What is a network (Internet)?

- We all know ...
 - We use the network everyday
 - We're using it *now*
 - All our favorite gadgets are *online!*
 - The Internet means: my iPhone, my Powerbook, the web, my email, WhatsApp, FB, Twitter, ...
- But no authoritative formal answer ...
 - Depends on who you ask, what's the purpose

What's the Internet: “nuts and bolts” view

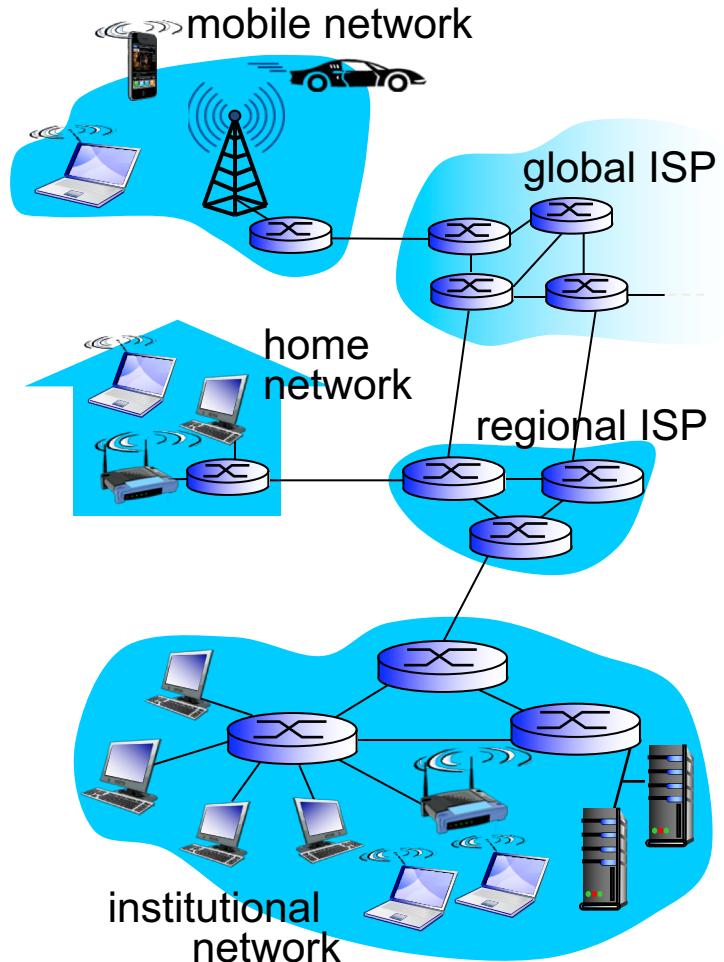


- millions of connected computing devices:
 - **hosts** = *end systems*
 - running *network apps*

❖ **communication links**

- fiber, copper, radio, satellite
- transmission rate: *bandwidth*

- ❖ **Packet switches:** forward packets (chunks of data)
- *routers* and *switches*



“Fun” internet appliances



IP picture frame
<http://www.ceiva.com/>



Internet refrigerator



Web-enabled toaster +
weather forecaster



Tweet-a-watt:
monitor energy use



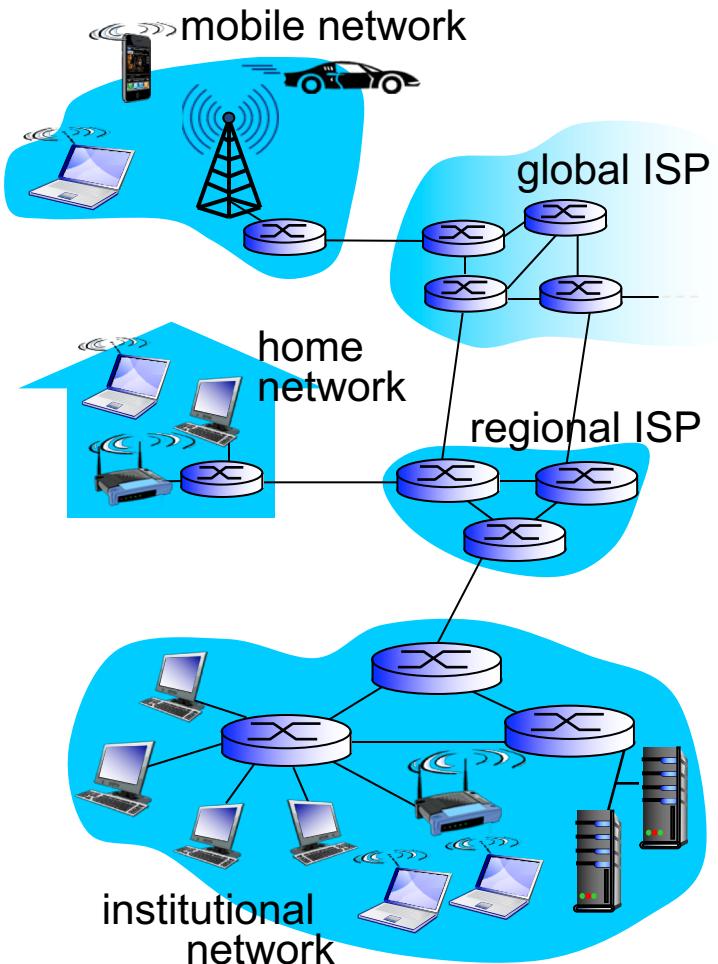
Slingbox: watch,
control cable TV remotely



Internet phones

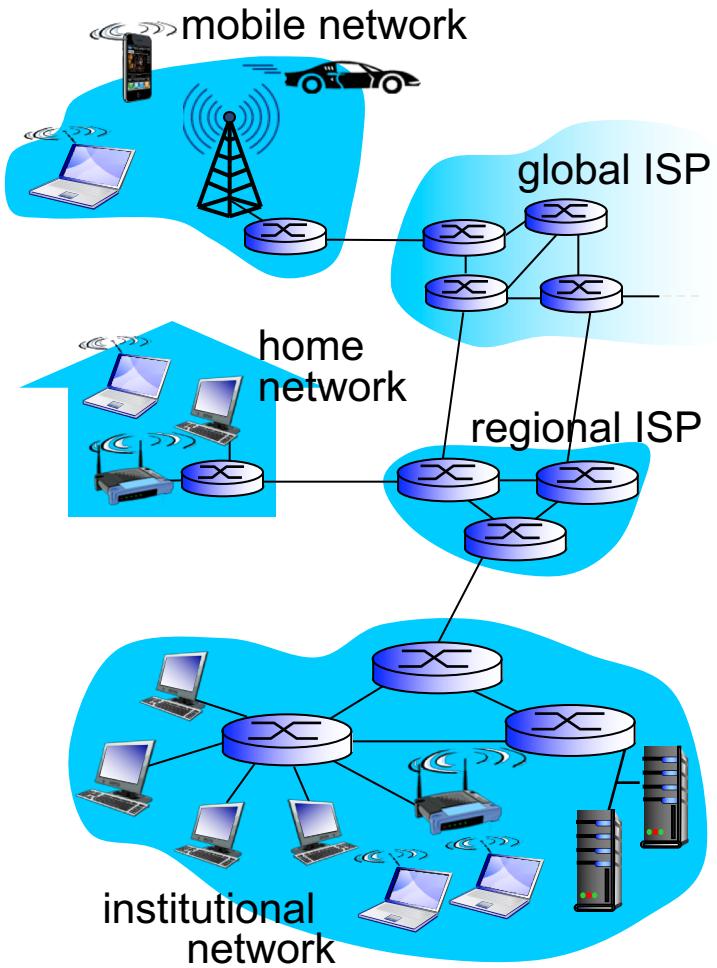
What's the Internet: “nuts and bolts” view

- *Internet: “network of networks”*
 - Interconnected ISPs
- *protocols* control sending, receiving of messages
 - e.g., TCP, IP, HTTP, Skype, 802.11
- *Internet standards (open, free)*
 - RFC: Request for comments
 - IETF: Internet Engineering Task Force



What's the Internet: a service view

- *Infrastructure that provides services to applications:*
 - Web, VoIP, email, games, e-commerce, social nets, ...
- *provides programming interface to apps*
 - hooks that allow sending and receiving app programs to “connect” to Internet
 - provides service options, analogous to postal service
 - *Socket API*



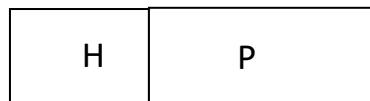
Let's build an Internet!

- One that connects billions of people, and many times more devices?
- Wow, it's so difficult! Well ...
- Know your enemy
 - What are the basic *challenges*?
- Know your weapons
 - What are the design *principles* to overcome the challenges?

Challenge 0: Interoperability

- How to make everything talk? They're all different!
- How do we (you and I) do it?
 - Fix the language (e.g., English)
 - Fix some conventions, etiquette
- Now make it more formal & precise
 - Standardization (IETF/RFC) into communication *protocols*
 - HTTP, TCP, UDP, IP, ARP, whatever-P
- Each “P” deals with unit of data called *packet*

H - header (interpreted)
P – payload (uninterpreted)



Principle of
encapsulation
(info hiding)

What's a protocol?

human protocols:

- “what’s the time?”
- “I have a question”
- introductions

... specific msgs sent

... specific actions taken
when msgs received, or
other events

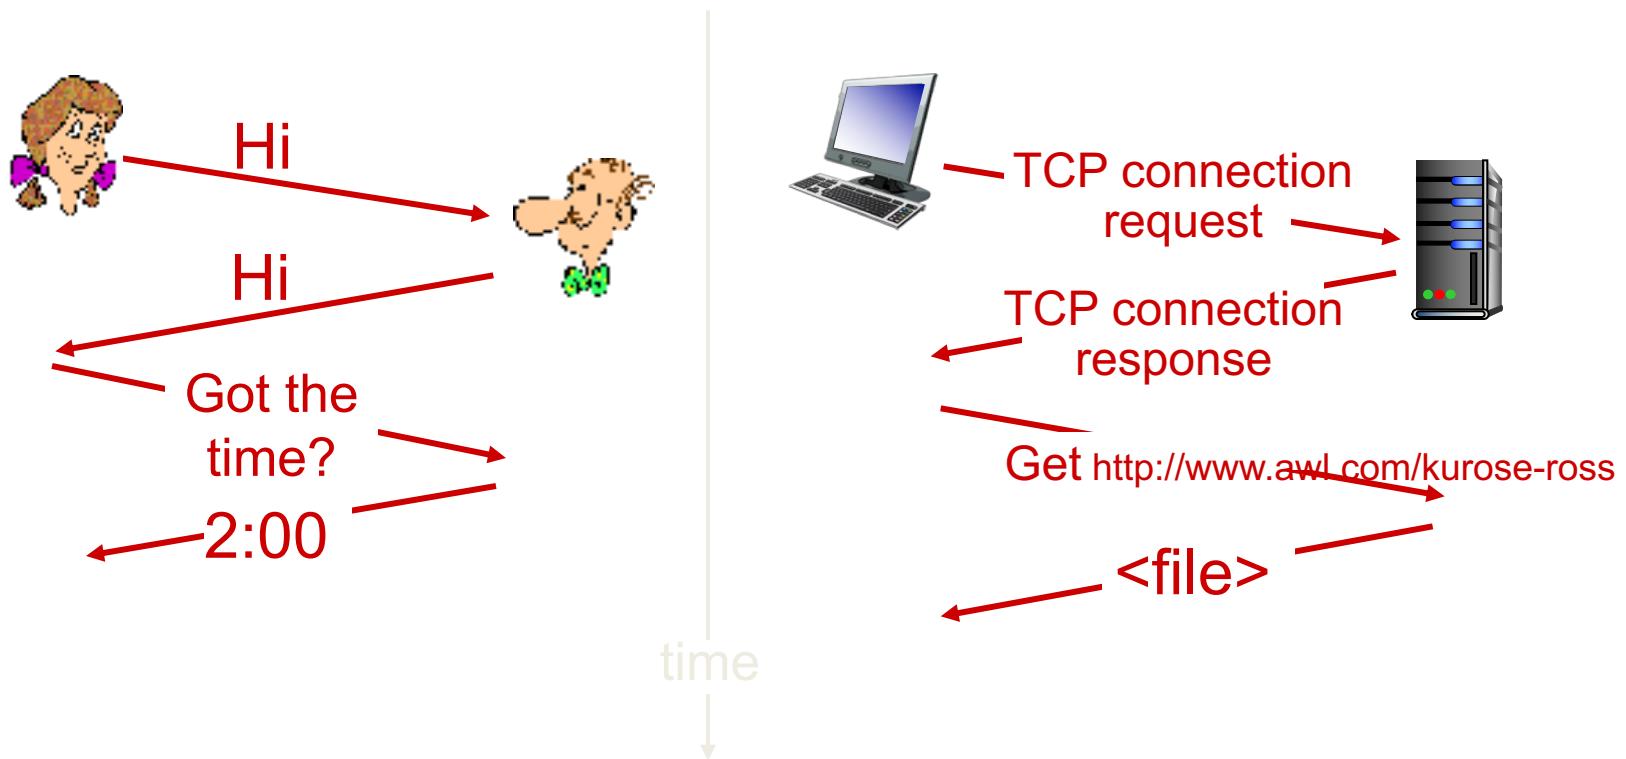
network protocols:

- machines rather than humans
- all communication activity in Internet governed by protocols

*protocols define format, order
of msgs sent and received
among network entities,
and actions taken on msg
transmission, receipt*

What's a protocol?

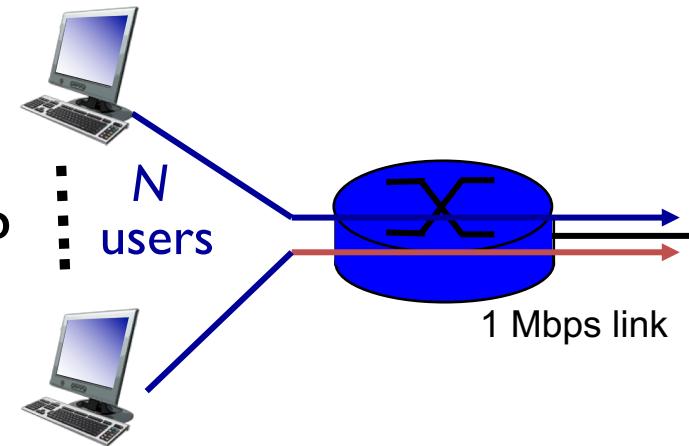
a human protocol and a computer network protocol:



Q: other human protocols?

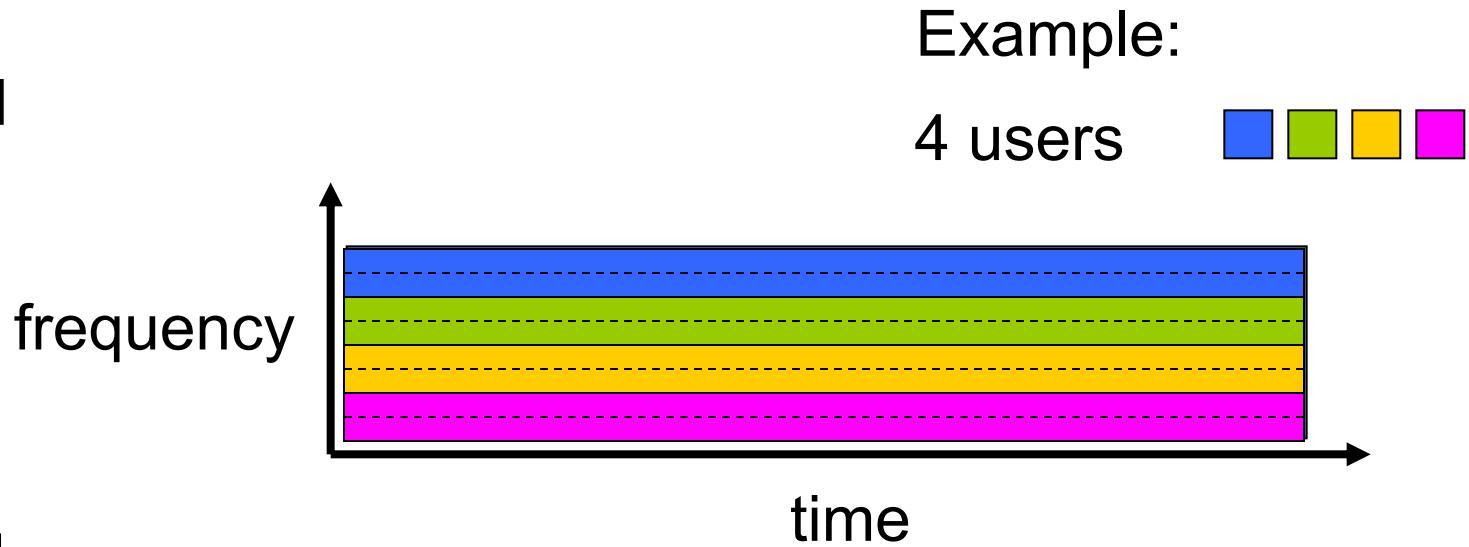
Challenge 1: Sharing among many

- Internet is a *shared* resource
- How do you share a function room?
 - Book a time and own that time slot
 - TDM! (time division multiplexing)
- TDM works well for some applications
 - Phone calls: smooth and same bit rate all the time
 - Simple: Allocate fixed rate to each established call!
- But, data traffic is *bursty*
 - How do you surf? Look + think -> click a link -> look + think ...
(net silent) (net busy)
- Packet magic: *packet switching* & *statistical multiplexing*

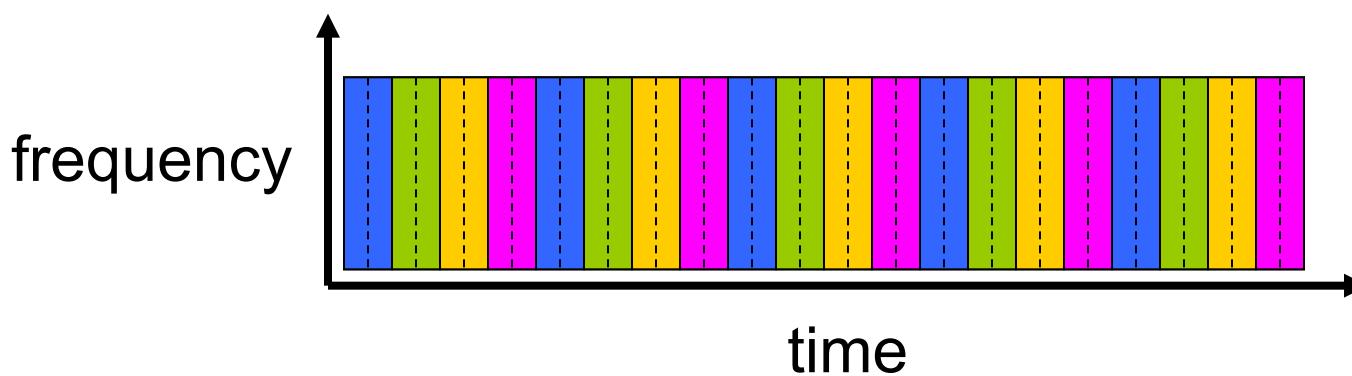


Circuit switching: FDM and TDM

FDM



TDM



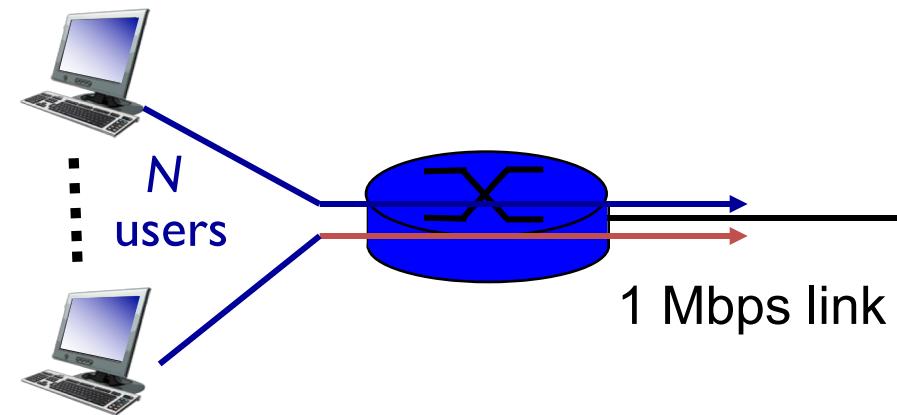
Circuit switching (vs. packet switching) isn't flexible: blue can't use what green doesn't use

Statistical multiplexing w/ packet switching

Why? Bursty data

Example:

- 1 Mb/s link
- each user:
 - 100 kb/s when “active”
 - active 10% of time
- *circuit-switching (i.e., TDM or FDM)*: Fixed, dedicated fraction of link
 - 10 users flat
- *packet switching*: Packet occupies link *on demand*
 - How many?



Activity 1.1

- Can I get away with 35 users?
- Prob. only first 11 of them active at the same time?
- Prob. any (exactly) 11 of them active at the same time?
- Prob. > 10 of them active at the same time?
 - Formula: CDF of binomial distribution
 - Exact answer: write a program (homework), or
<http://www.danielsoper.com/statcalc3/calc.aspx?id=71>

Challenge 2:

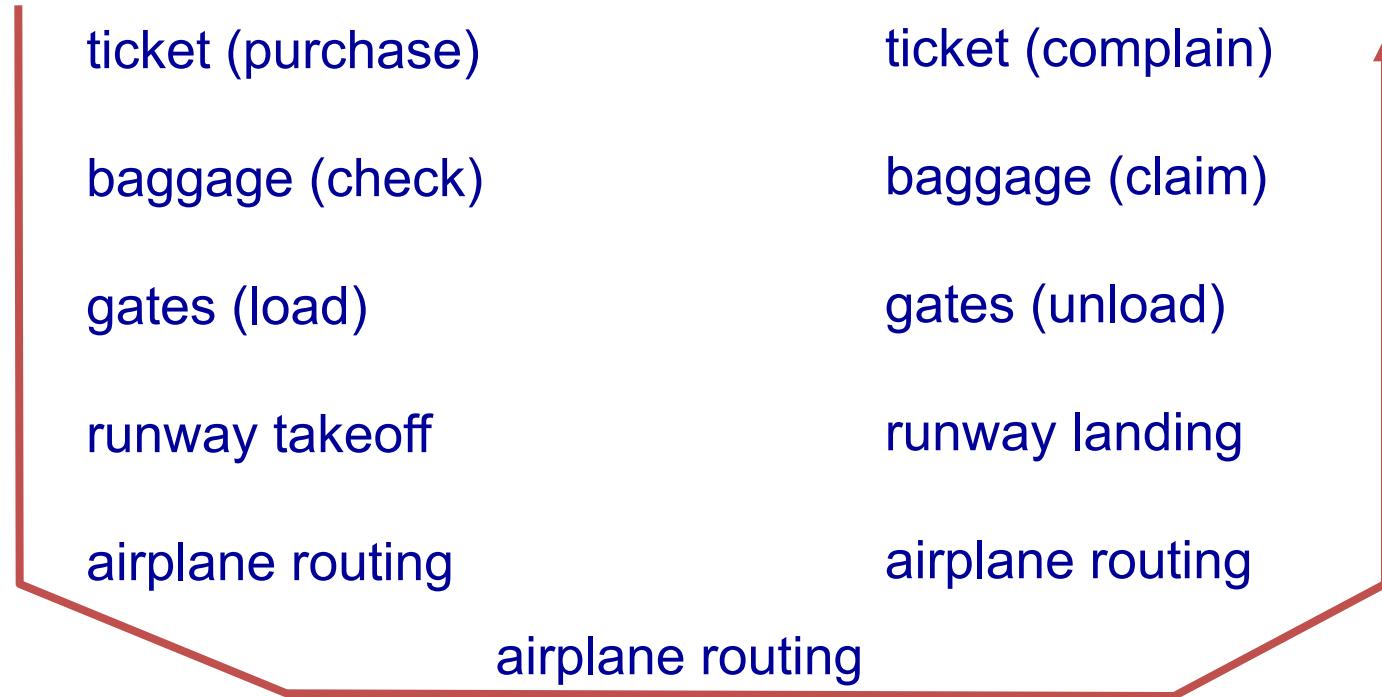
Complex interacting components

- Network needs many “pieces”: naming, routing, reliability, (lots of) applications, ...
- Modularity: divide-&-conquer complexity
 - Think functions, processes (sep. of interface from implementation; easy to maintain/update)
- But not enough
 - Also interactions between modules
 - *Emergent behavior* gives (not-so-nice) surprises
 - i.e., behavior isn’t observed in modules individually, but it arises only when you put them together (cf. chemistry: add water to HCL -> generates heat and causes water to spill out)
- Example of emergent behavior (OS revisited)
 - Priority scheduling + locking = *priority inversion*

Activity #1.2: Priority Inversion

- 3 processes: A of priority 1, B of 10, C of 100
(static priority: higher runs first)
- A and C access same critical section S guarded by lock L; B doesn't synchronize w/ A or C
- Consider A and C only: If A runs first; then C runs and now wants to enter S, how long C has to wait at most?
- Now consider A, B, C: If A runs first, followed by B, followed by C, and C now wants to enter S, how long might C have to wait (worst-case scenario)?

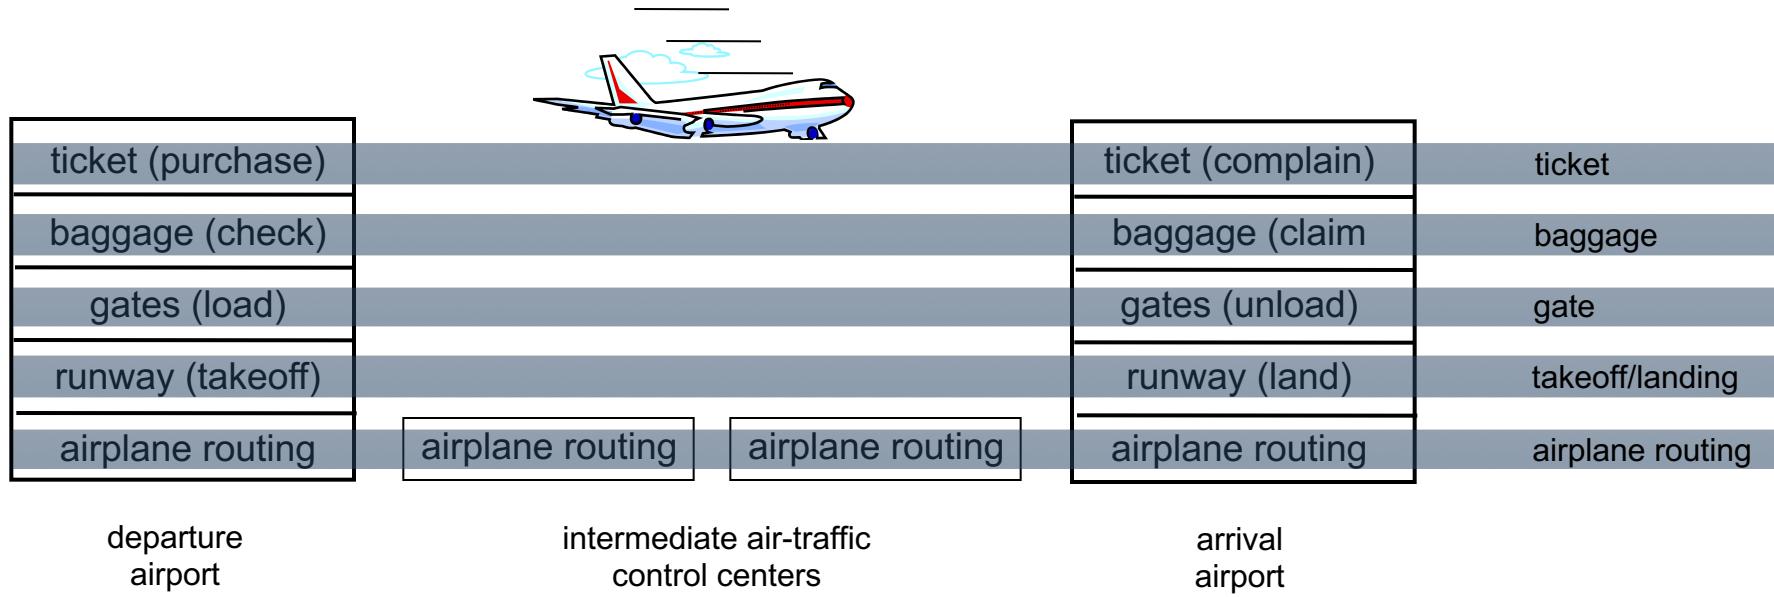
Layering: reducing interactions



Organization of air travel: a series of steps

To reduce emergent behavior, we should minimize interactions between modules; *layering* is an approach to achieving this goal

Layering of airline functionality



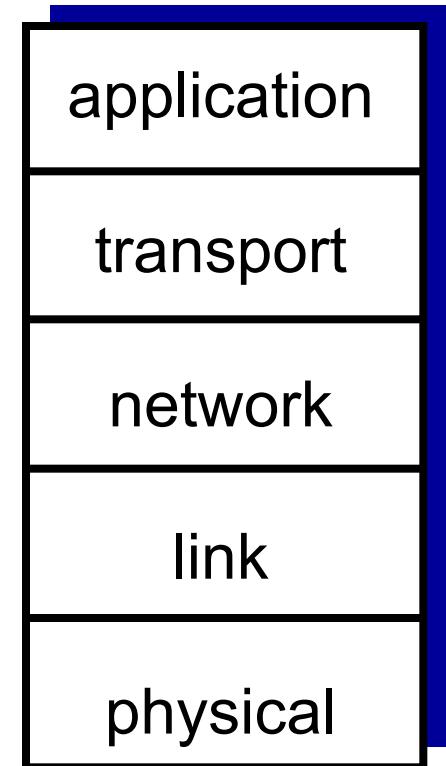
Layers: each layer is like a module, it implements a service

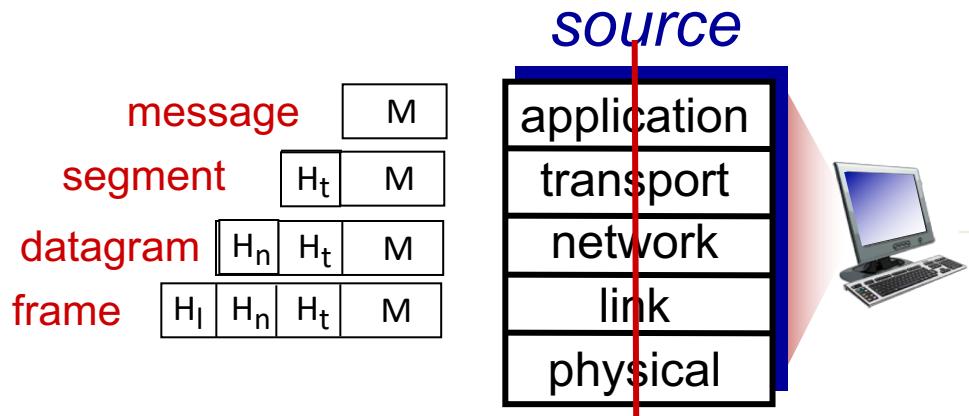
- via its own internal-layer actions
- relying on services provided by layer below
- in turn provides services to layer above
- Layer i uses service of layer $i - 1$ to provide service for $i + 1$

Question: N modules; how many possible interactions if they interact freely? How about with layering?

Internet protocol stack

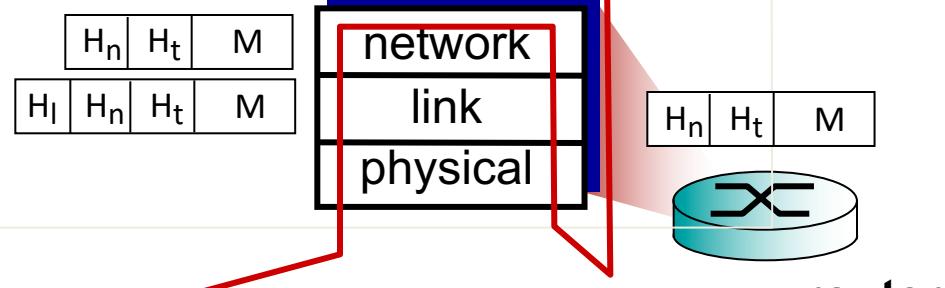
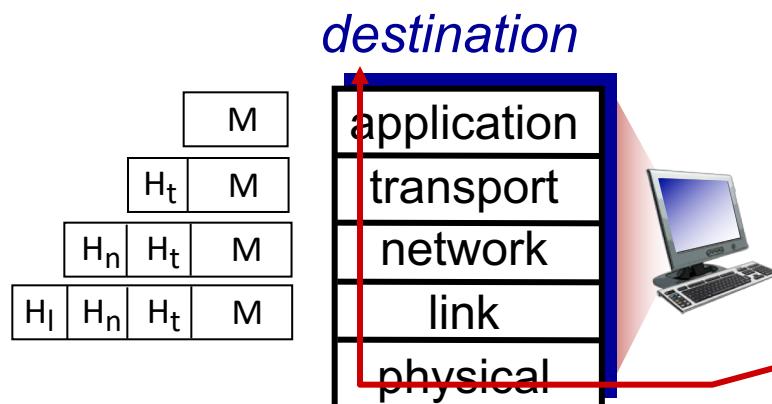
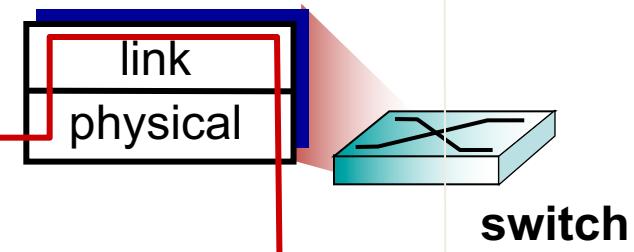
- *application*: supporting network applications
 - e.g., FTP, SMTP (email), HTTP (web)
- *transport*: process-process data transfer
 - e.g., TCP (reliable), UDP (not reliable)
- *network*: routing of datagrams from source to destination
 - e.g., IP, routing protocols
- *link*: data transfer between neighboring network elements
 - Ethernet, 802.11 (WiFi), PPP
- *physical*: bits “on the wire”





Cascaded headers

Note: As packet goes *down* protocol stack, header is *added* by each layer (protocol); as it goes *up* protocol stack, each layer *strips off* its header



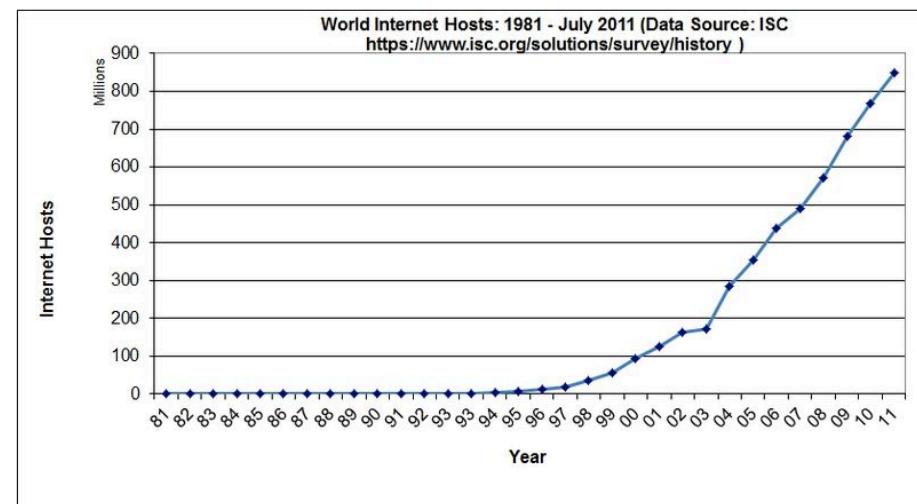
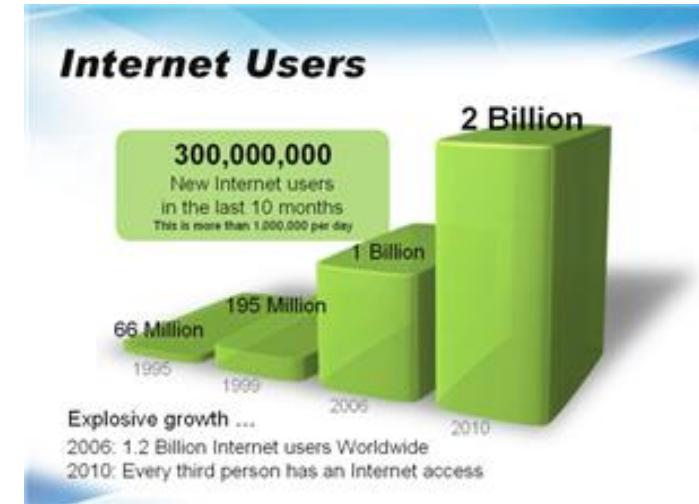
NB: With layering, header vs. payload is now *relative* to the layer; e.g., transport layer's header is network layer's payload

Challenge 3: Scalability

Scalability means how to manage system as the system size grows

Bigger & bigger size of Internet

- Sustained fast growth
- Increasing structural complexity
- *Incommensurate scaling!*
 - *I.e., different aspects of network grow at different speeds*



Can you build bigger and bigger pyramids?



Assume dimension D.
Weight of pyramid
grows like volume, i.e.,
 $O(D^3)$. Support is by
contact area between
lower and immediately
higher bricks – grows
like $O(D^2)$. So we have
**incommensurate
scaling**.

Designing for scalability

In internet, when the number of nodes N grows, the (pairwise) interconnect between them grows like N^2 : incommensurate scaling

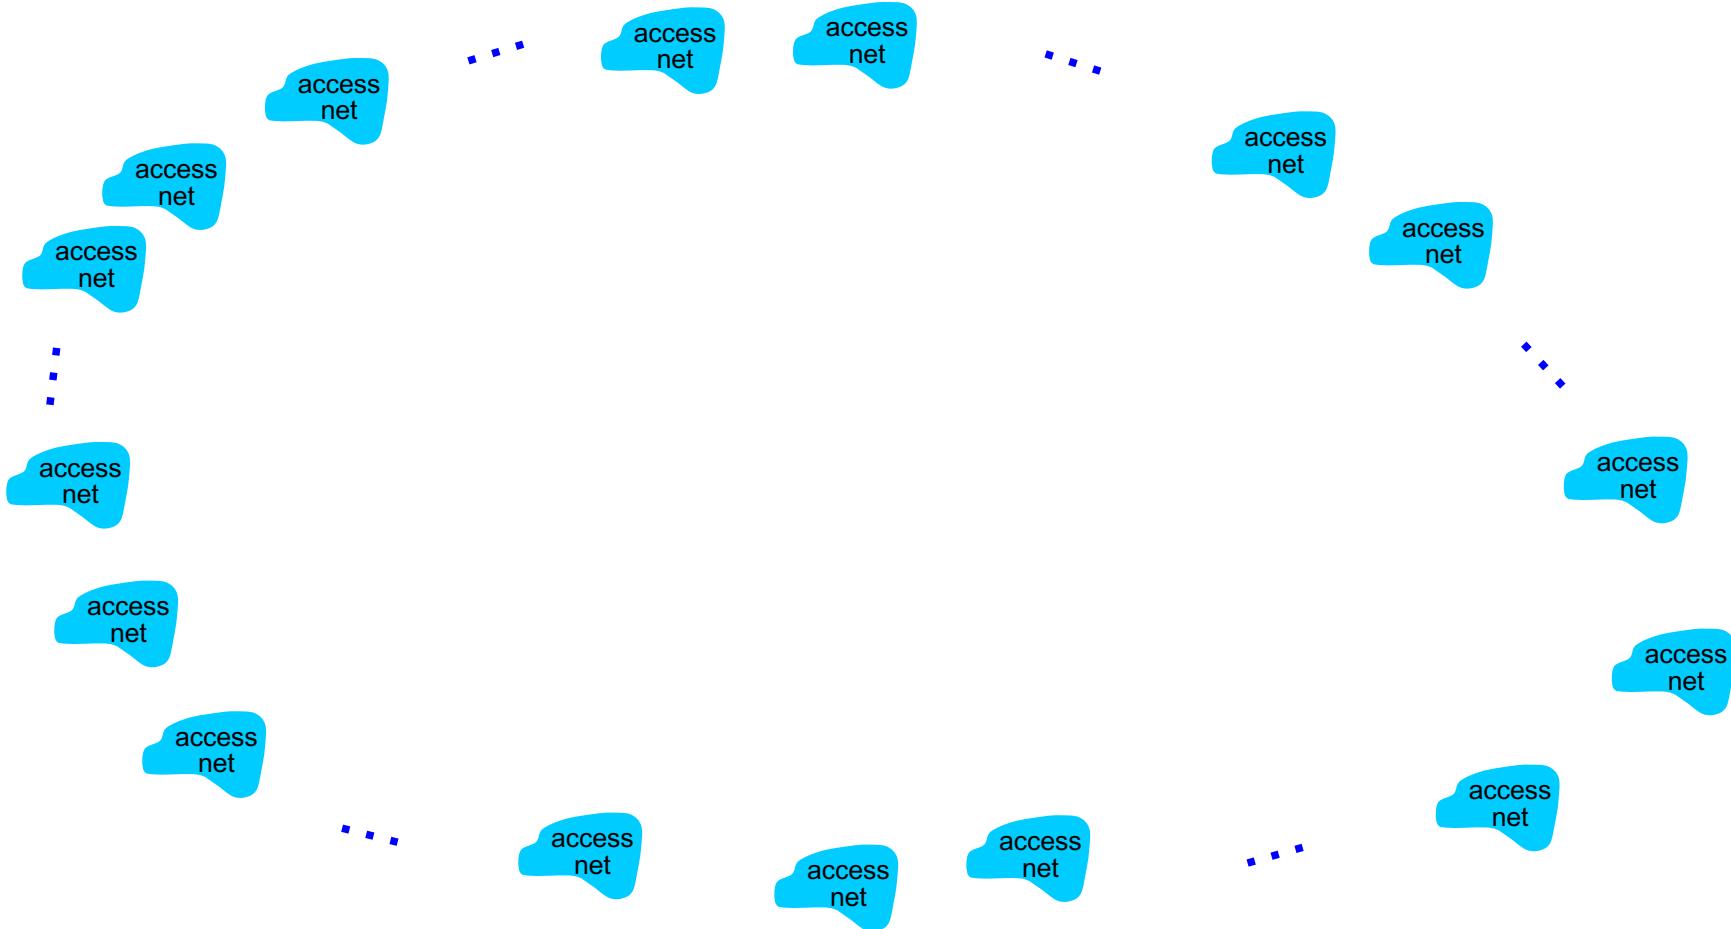
Use design principle of *hierarchy* to manage the large size and structure

- Network of networks (recursive interconnect)
- Access, regional, national, global ISPs



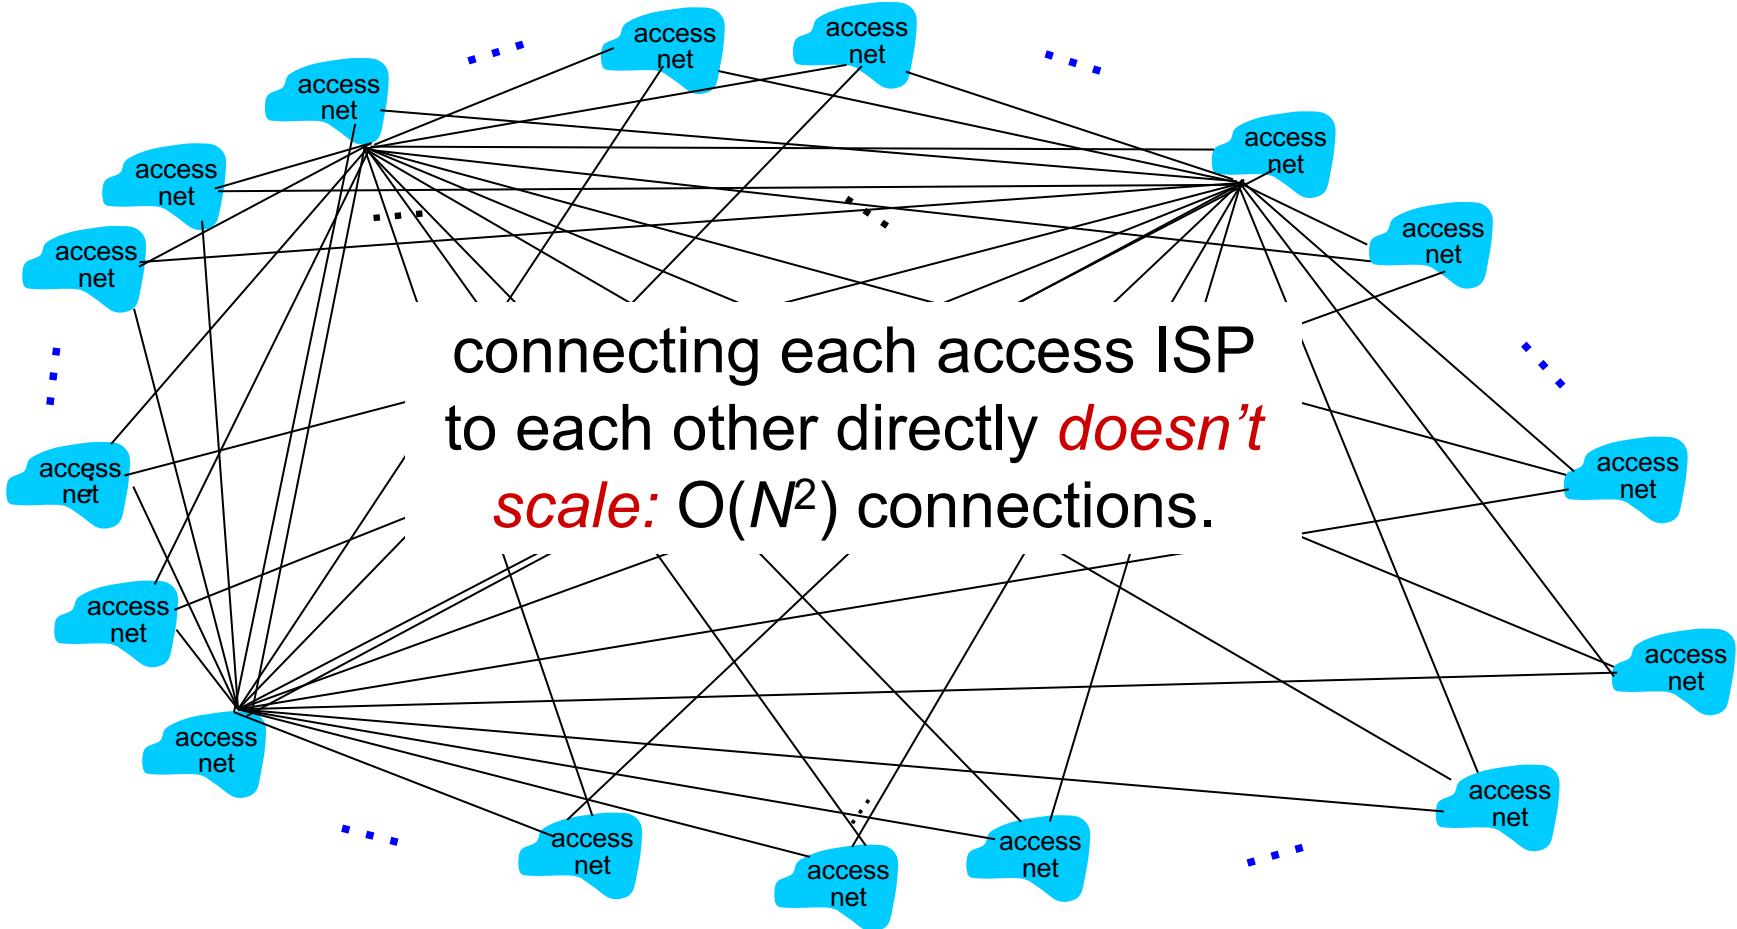
Internet structure: network of networks

Question: given *millions* of access ISPs, how to connect them together?



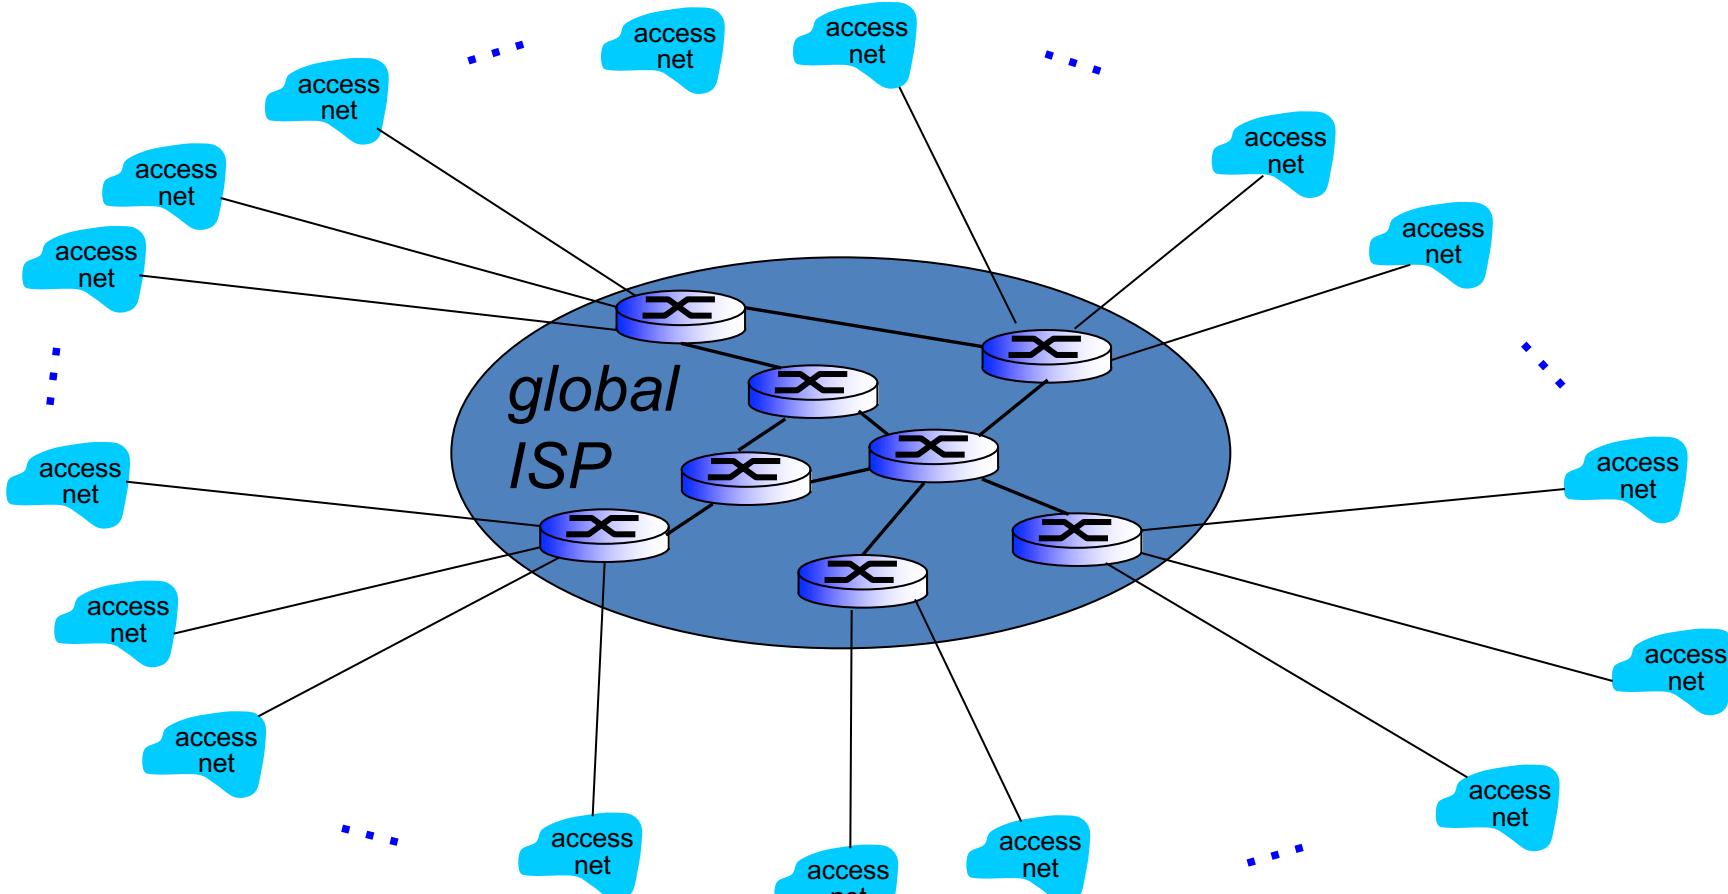
Internet structure: network of networks

Option: connect each access ISP to every other access ISP?



Internet structure: network of networks

Option: connect each access ISP to a global transit ISP? Customer and provider ISPs have economic agreement.

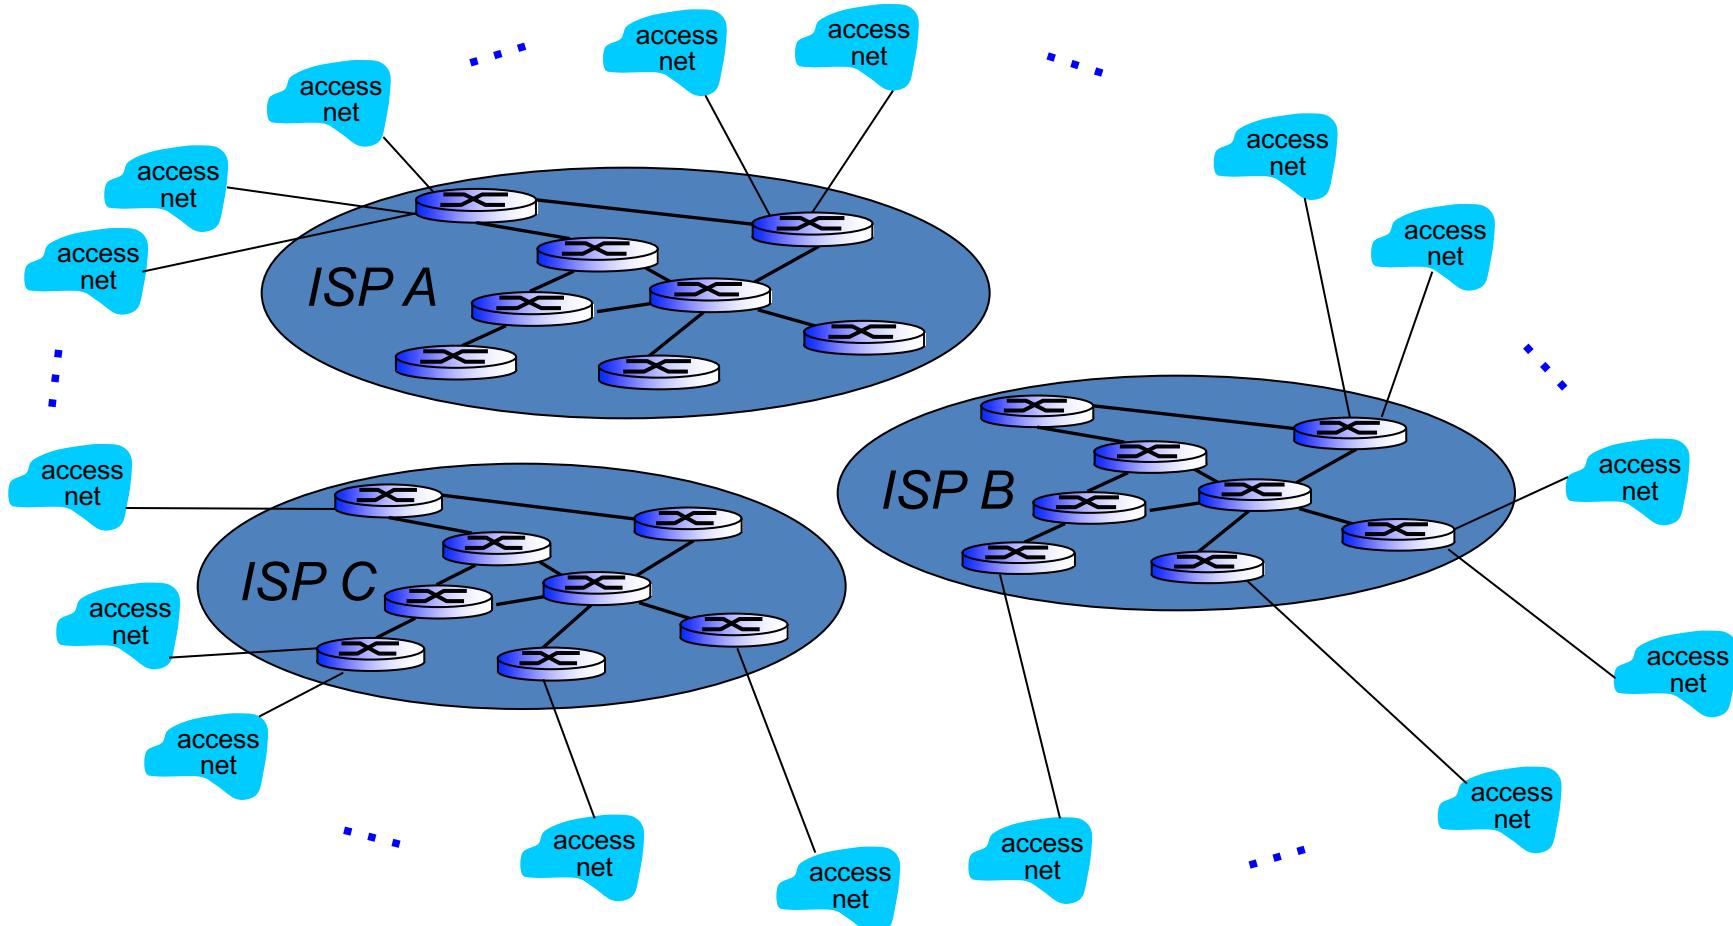


*NB: Use backbone links in network core that are *shared* by many pairs of access networks (many fewer links needed than N^2) – like highways (PIE, AYE) in Singapore that connect many districts*

Internet structure: network of networks

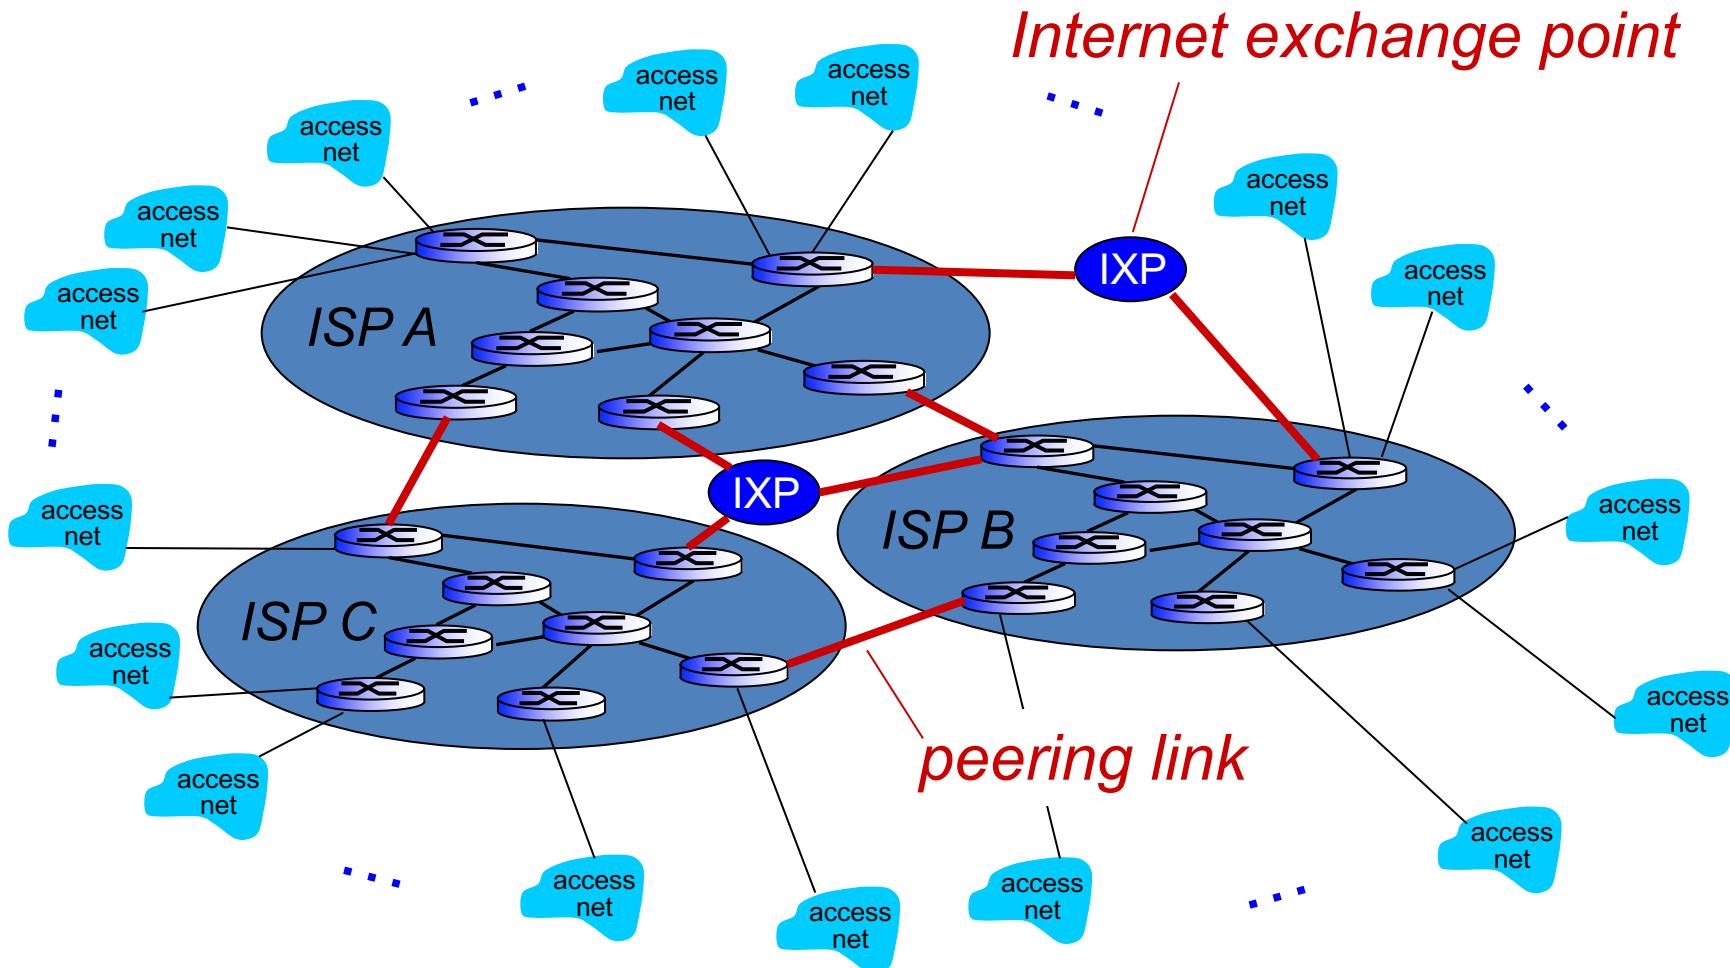
But if one global ISP is viable business, there will be competitors

....



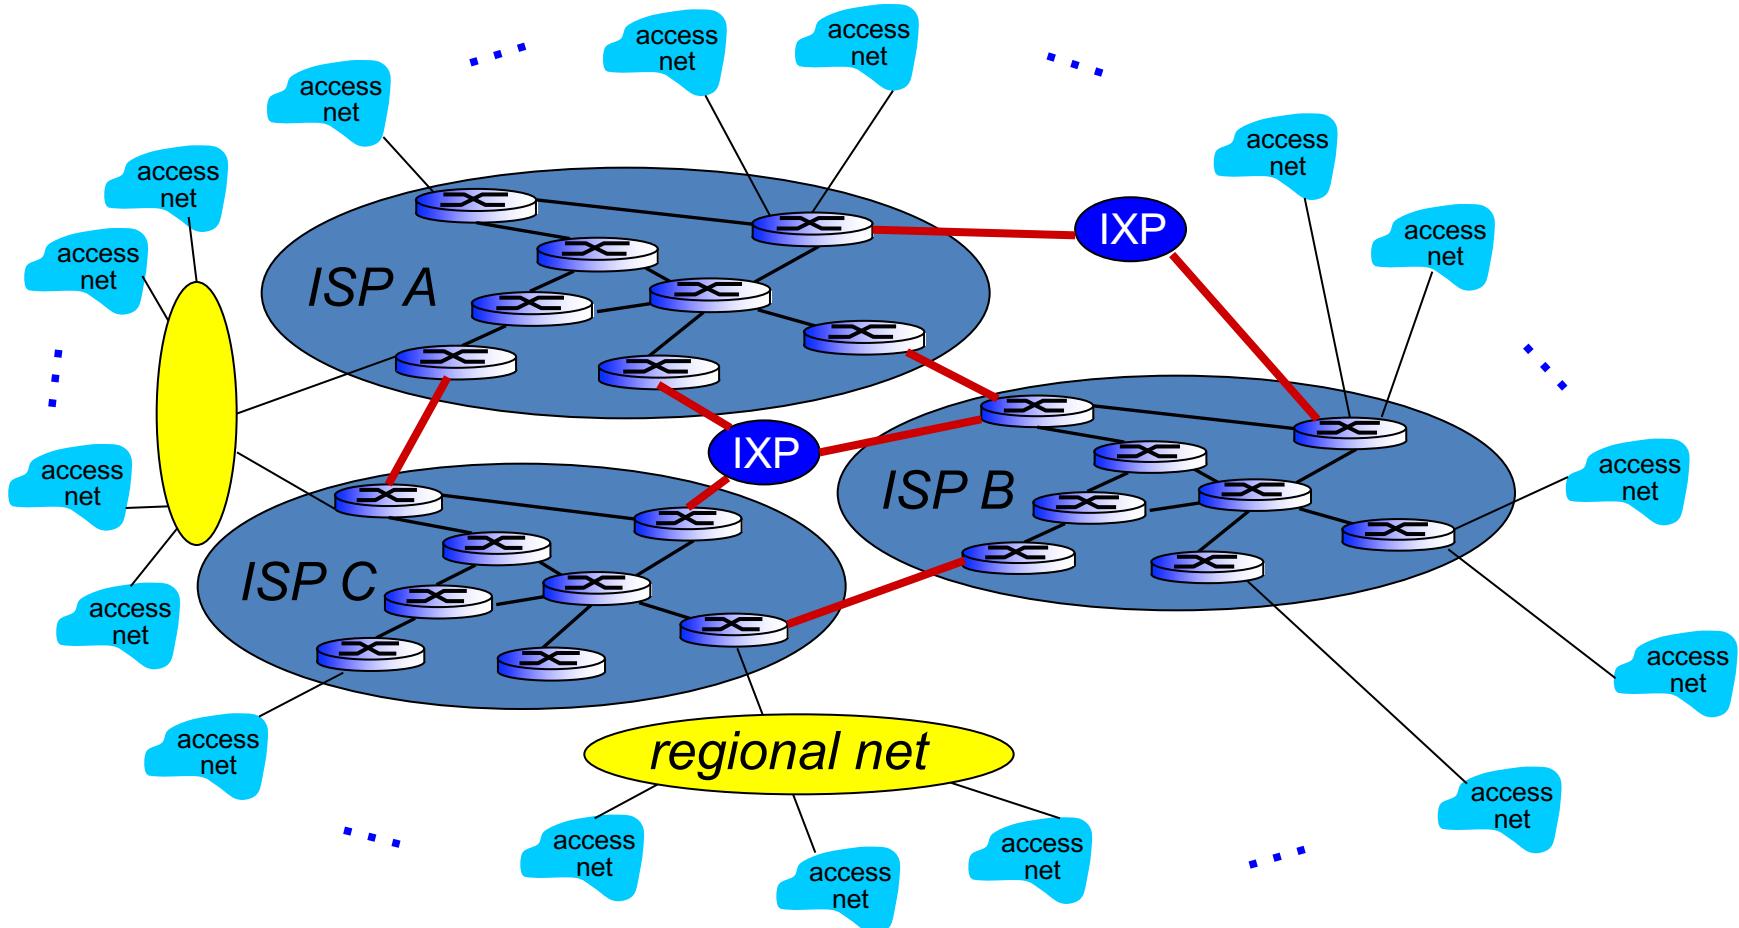
Internet structure: network of networks

But if one global ISP is viable business, there will be competitors
.... which must be interconnected



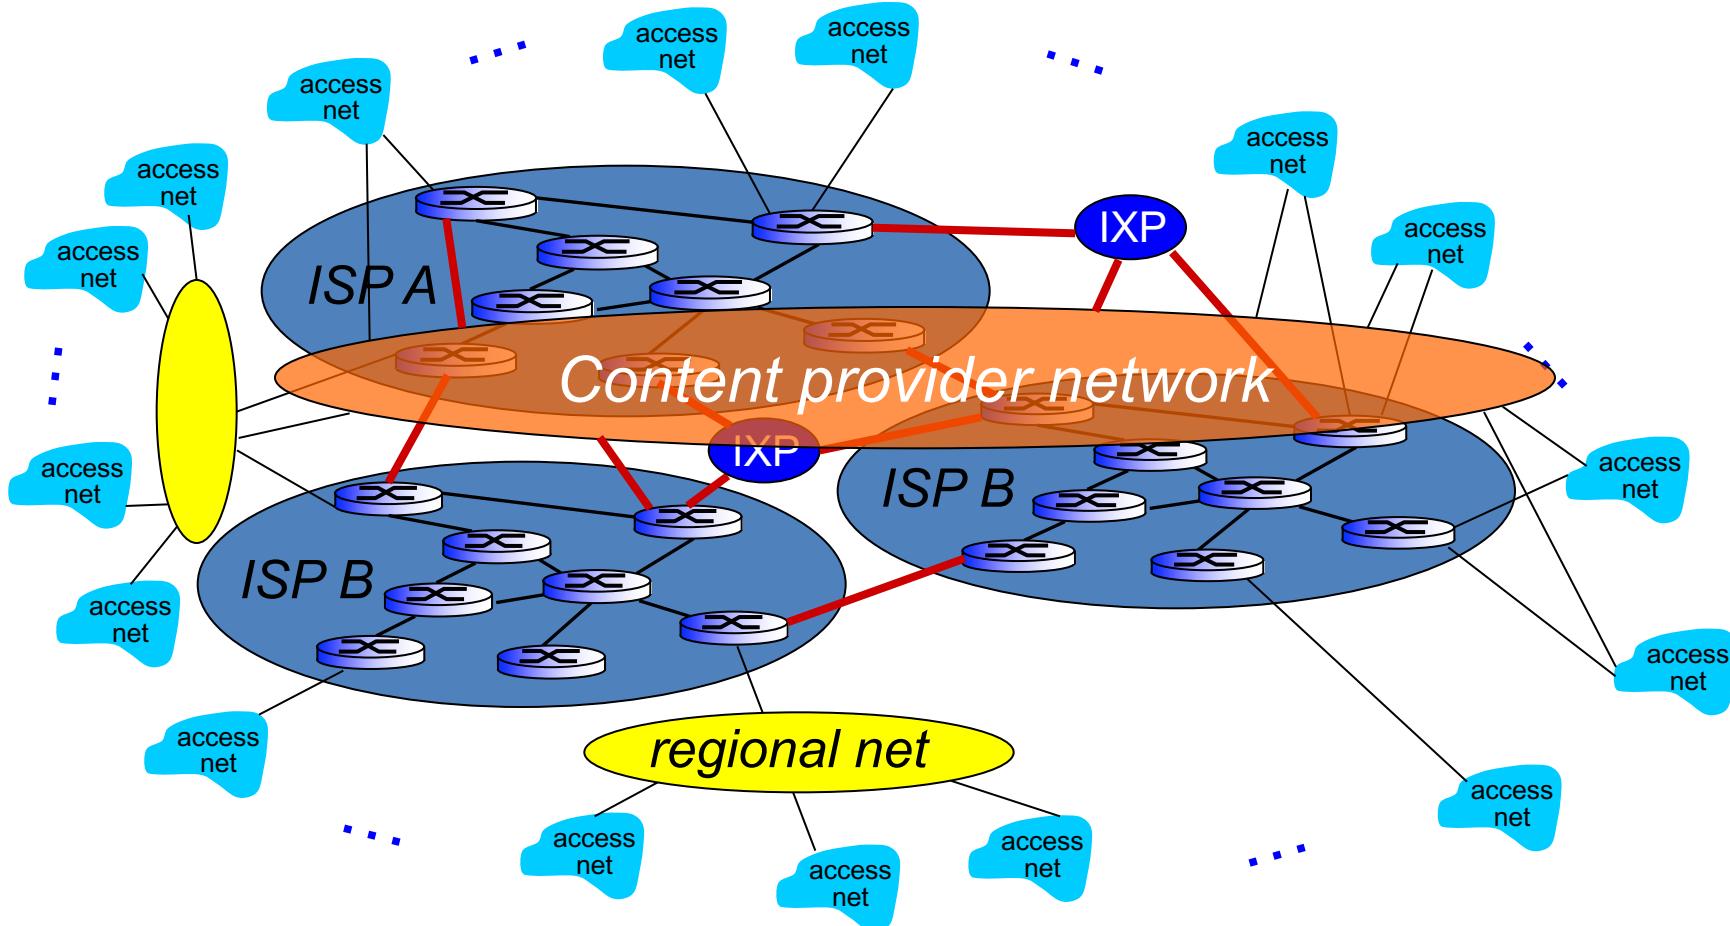
Internet structure: network of networks

... and regional networks may arise to connect access nets to ISPs

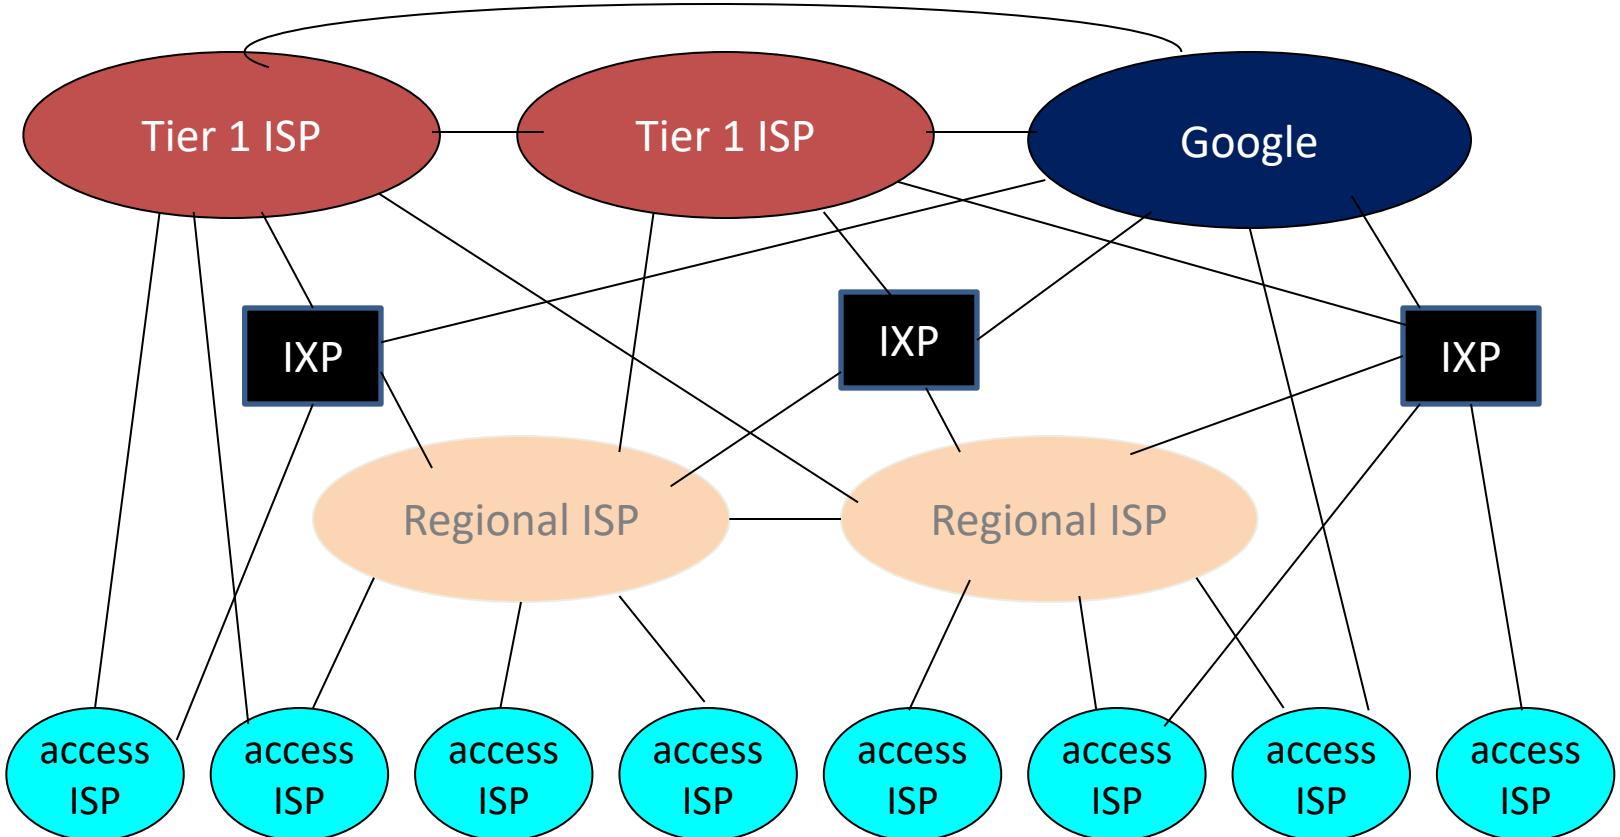


Internet structure: network of networks

... and content provider networks (e.g., Google, Microsoft, Akamai) may run their own network, to bring services, content close to end users



Internet structure: network of networks



- at center: small # of well-connected large networks
 - “tier-1” commercial ISPs (e.g., Level 3, Sprint, AT&T, NTT), national & international coverage (*tier 1* represents the *largest global ISPs*)
 - **content provider network** (e.g, Google): private network that connects its data centers to Internet, often bypassing tier-1, regional ISPs

Tier-I ISP: e.g., Sprint

