



•
5 0 . 0 0 5 C S E

Natalie Agus
Information Systems Technology and Design
SUTD

MAKING PROGRAM

```
1 //  
2 // cKnowledge.c  
3 // CSE_CCodes  
4 //  
5 // Created by Natalie Agus on 14/1/19.  
6 // Copyright © 2019 Natalie Agus. All rights reserved.  
7 //  
8  
9 #include "cKnowledge.h"  
10  
11  
12 void func1(void) { printf("Function 1 Called\n"); }  
13 void func2(void) { printf("Function 2 Called\n"); }  
14 void func3(void) { printf("Function 3 Called\n"); }  
15  
16 int testFunctionPointers(void)  
17 {  
18     static void (*ptr[3])(void) = {func1, func2, func3};  
19     int k=0;  
20     for(k=0;k<3;k++)  
21         ptr[k]();  
22     return 0;  
23 }
```

Your code
in plaintext format, commonly



Compiled

A program, it is executable

A set of instructions, that can tell the kernel
what to do to create a **process image**

→ The start (standard) state when
you open an app (at last saved
state (if any))

Process image: everything about a process at a
point of time, its heap, stack, registers value,
data, instructions, as well as its PCB data
structure

MAKING PROCESS



A program, executable

A set of instructions, that can tell the kernel what to do to create a **process image**

A program doesn't change over time.
It is passive in nature.

PC →



Executed

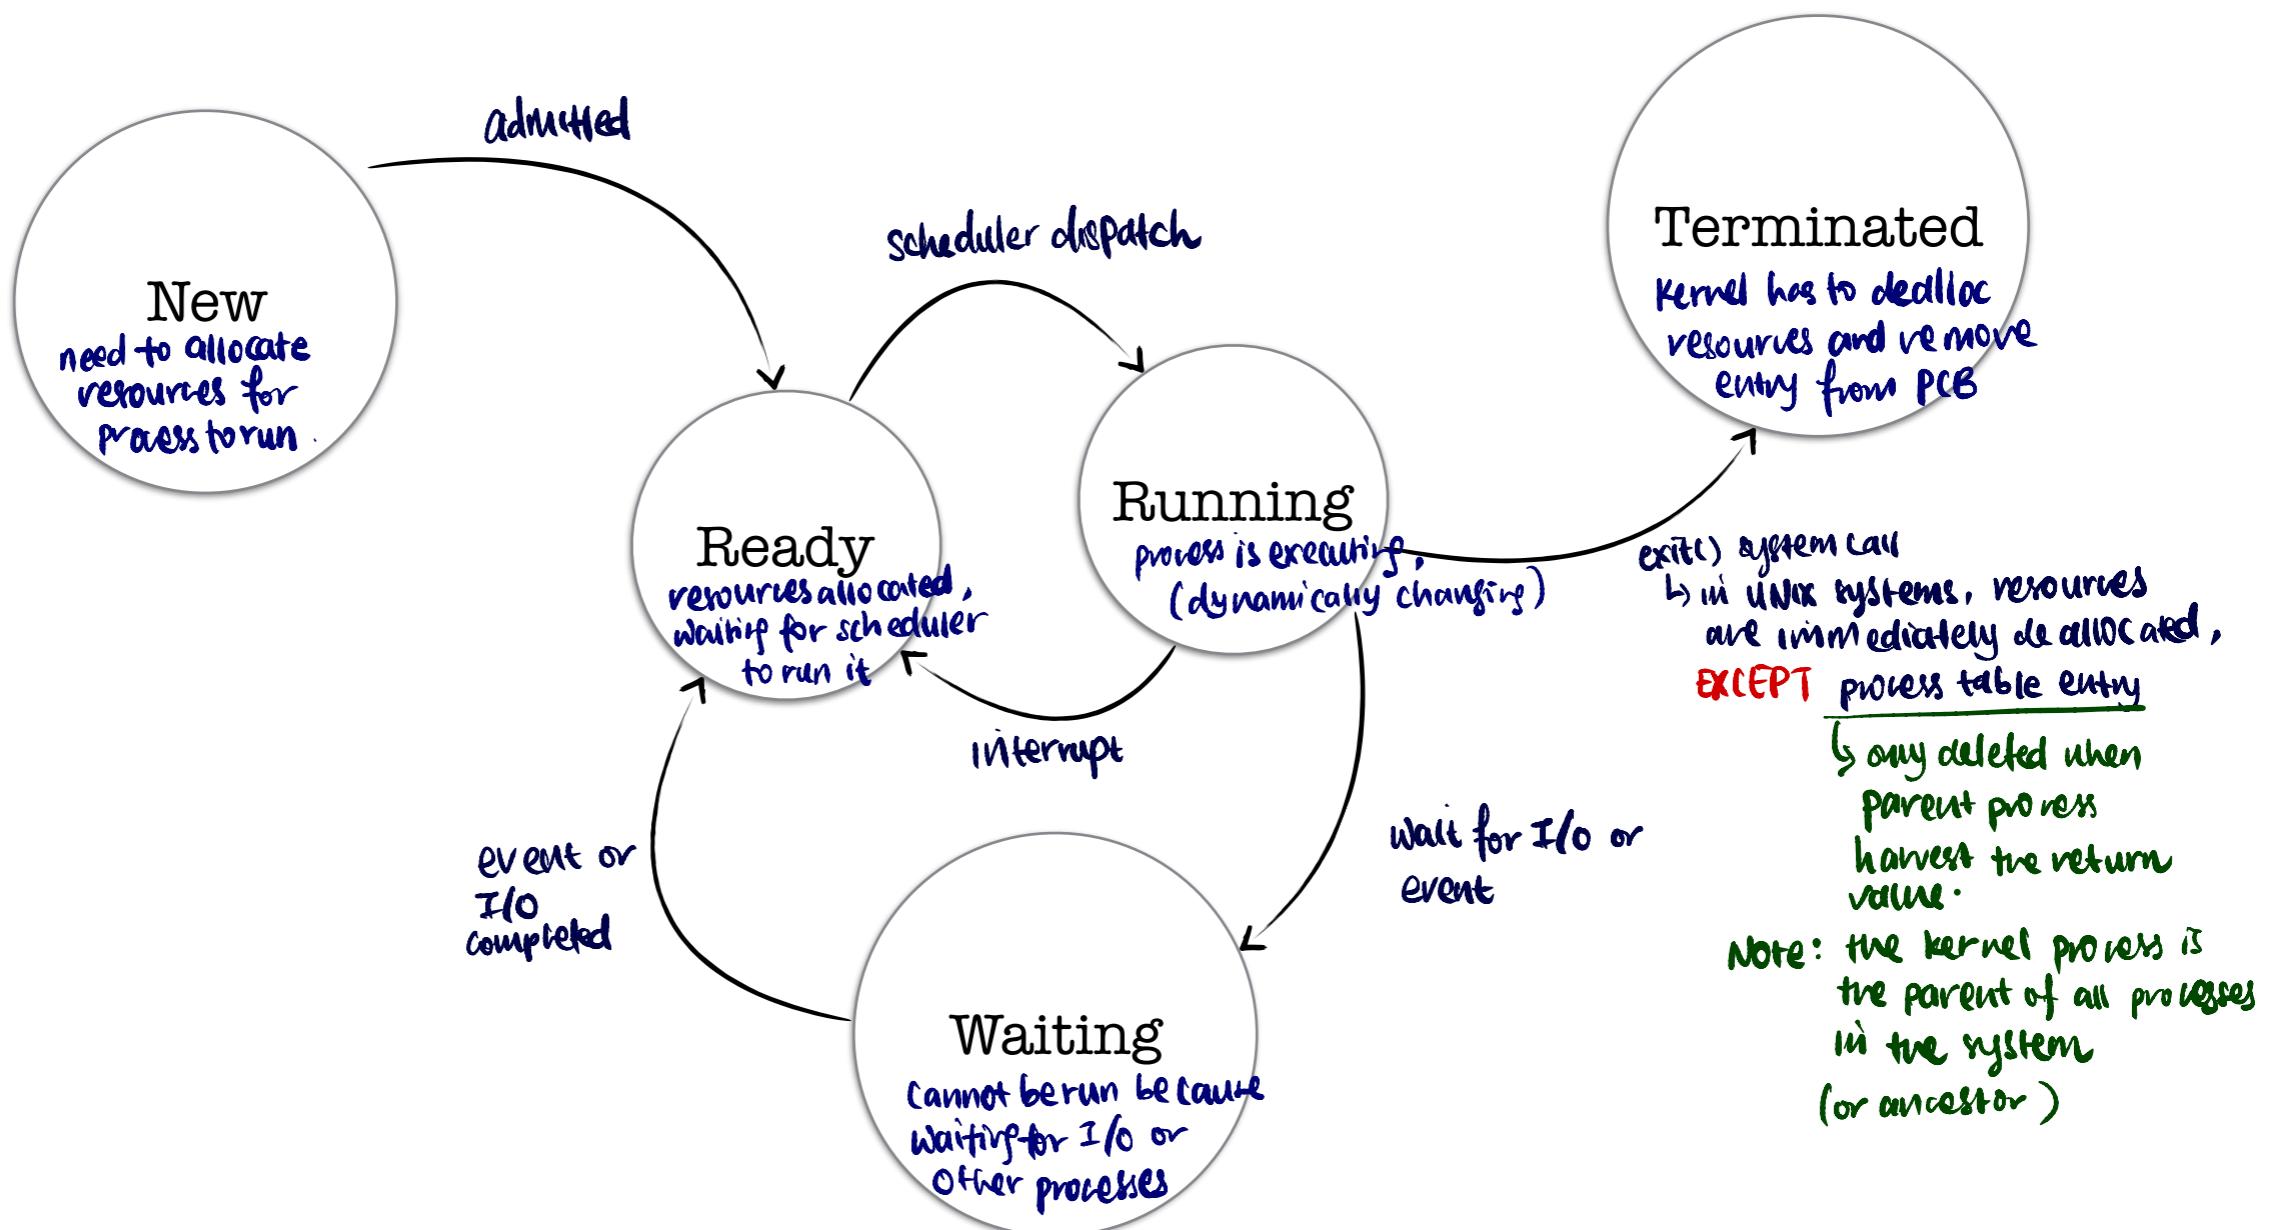
A **process image** from the executable is loaded into the **private** VM space of the process (can be on RAM or disk). This gives the initial state of a process.

Also, an entry in the process table is created to indicate that a new process is going to be run in the system.

Then as the PC moves and execute the instructions, a **process is happening**.

A process changes over time.
It is active in nature.

SCHEDULING STATE TRANSITION



PROCESS CONTROL BLOCK

A *data structure* that stores information about each process **for scheduling purposes**.

The OS scheduler manage the scheduling of processes, and keeps track of all the processes' PCB in the process table

The name of each process in the system

now, ready, running, waiting, terminated

```
pid_t pid; /* process identifier */  
long state; /* state of the process */  
unsigned int time_slice /* scheduling information */  
struct task_struct *parent; /* this process's parent */  
struct list_head children; /* this process's children */  
struct files_struct *files; /* list of open files */  
struct mm_struct **mm; /* address space of this process */
```

struct task_struct

Other typical process information:

- Think about why these infos are important for the scheduler and are important for context switch
1. Process state
 2. Program counters
 3. CPU registers
 4. CPU scheduling information (priority, scheduling protocol)
 5. Memory management: pagetable, value of base and limit regs
 6. I/O: list of open files, connected devices
 7. Accounting: time running, resources used

THE PROCESS TABLE

↳ each entry in the process table is called the process control block

↳ a process table (proc table) contains the info of all processes in the system

	0	2	1	0	0	1
P1	P2	P3	P4	P5	P6	P7

A curved arrow originates from the row of process P1 and points to the row of process P7, illustrating a context switch between these two processes.

Context switch

- Step ① : save states of P_i
- Step ② : load states of P_j
- Step ③ : run P_j

some terminologies :

PROCESS SCHEDULING QUEUES

Three main types of queue:

1. Job queue : **set of all processes in the system**
2. Ready queue : **resides in RAM, ready to exec**
3. Device queue : **wait for I/O devices**

→ 1 queue **per device**

the PCB contains all process info.
these queues serve as a **fast lookup**
data structure for the scheduler to select a process.

↓
after selecting, will
refer and update
relevant PCB entries
for context switch.

Processes may migrate among different types of queue

CONTEXT SWITCH

Interleave executions between processes in the system

Is an **overhead** because switching does no useful work apart from giving the "**concurrency**" illusion for single core CPUs

SCHEDULING

INTERLEAVED EXECUTION

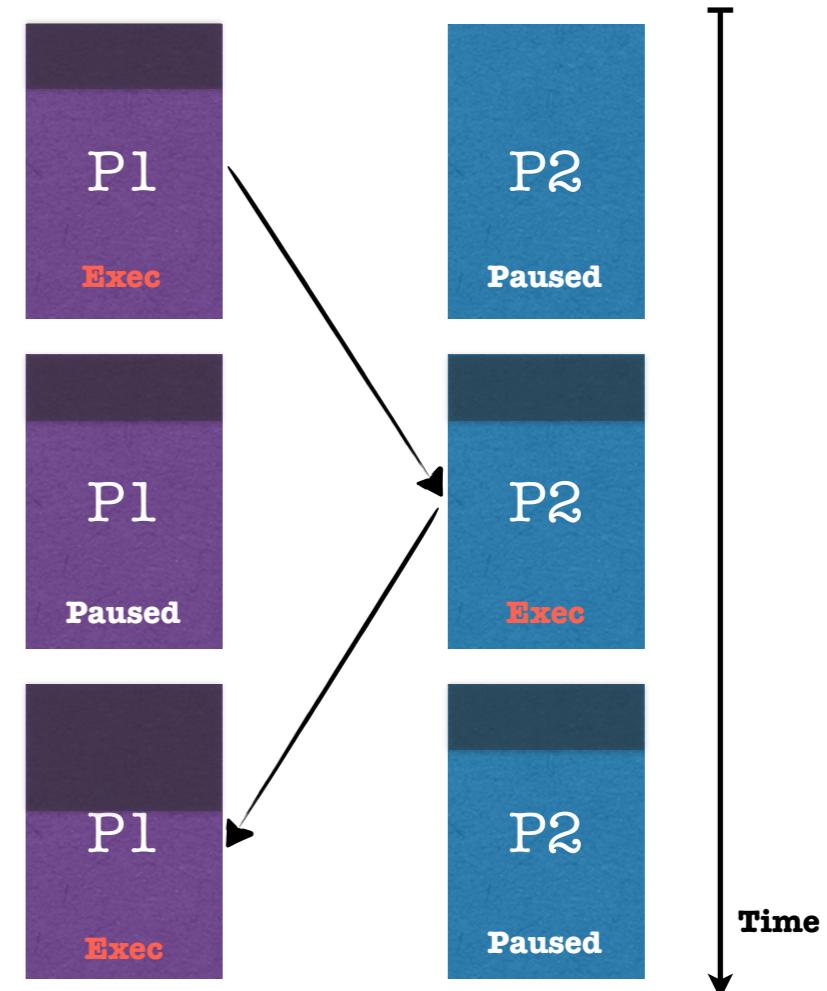
Processes take turns to be executed bit by bit

CONCURRENCY

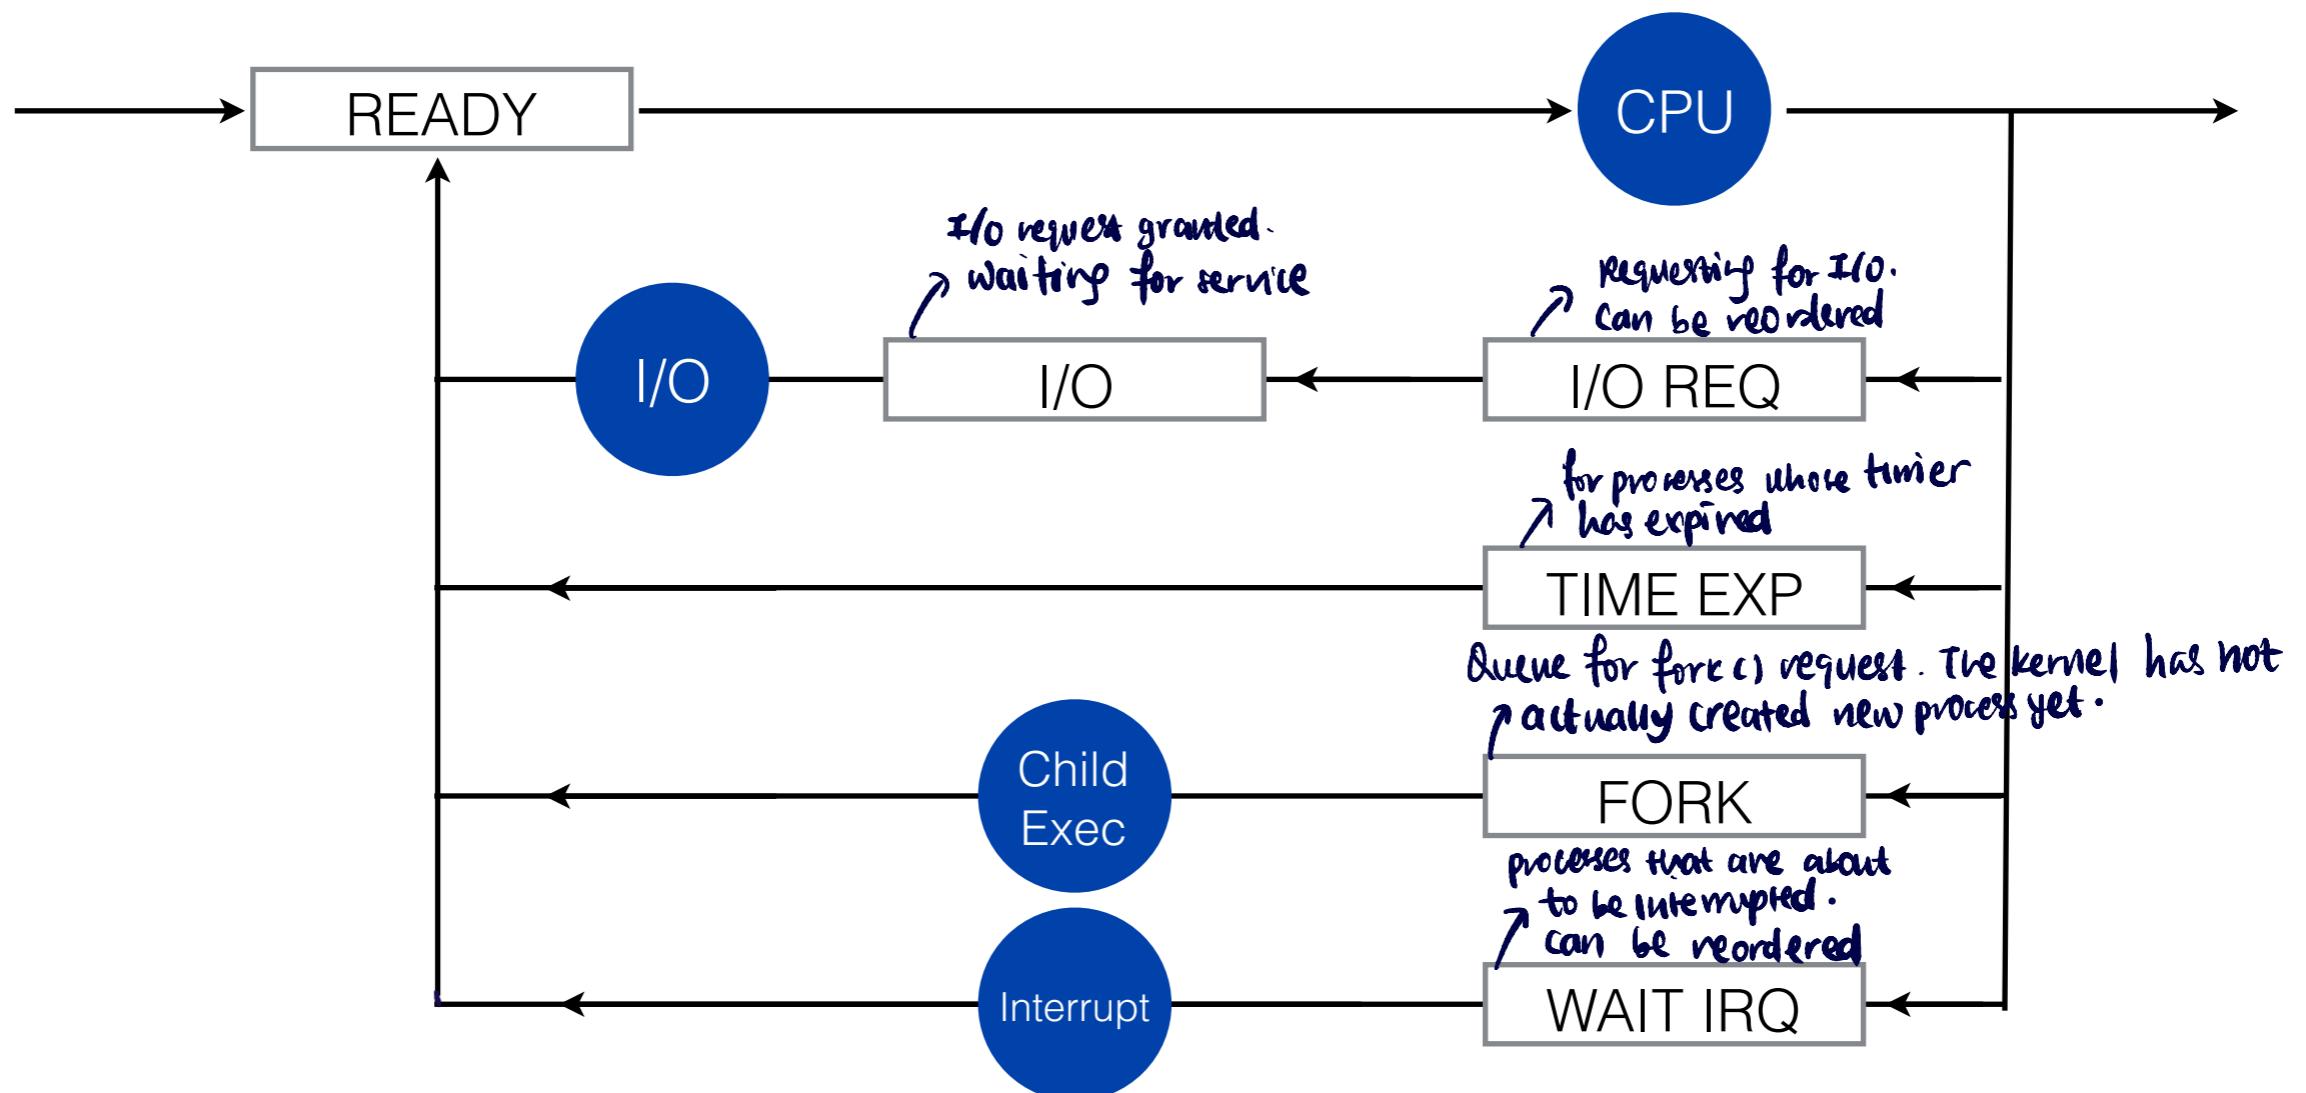
Interleaving process execution gives the illusion of concurrency, where all the programs seem to "run" at the same time

PARALLELISM

Only multi-core systems (multiple CPU) achieves true parallelism

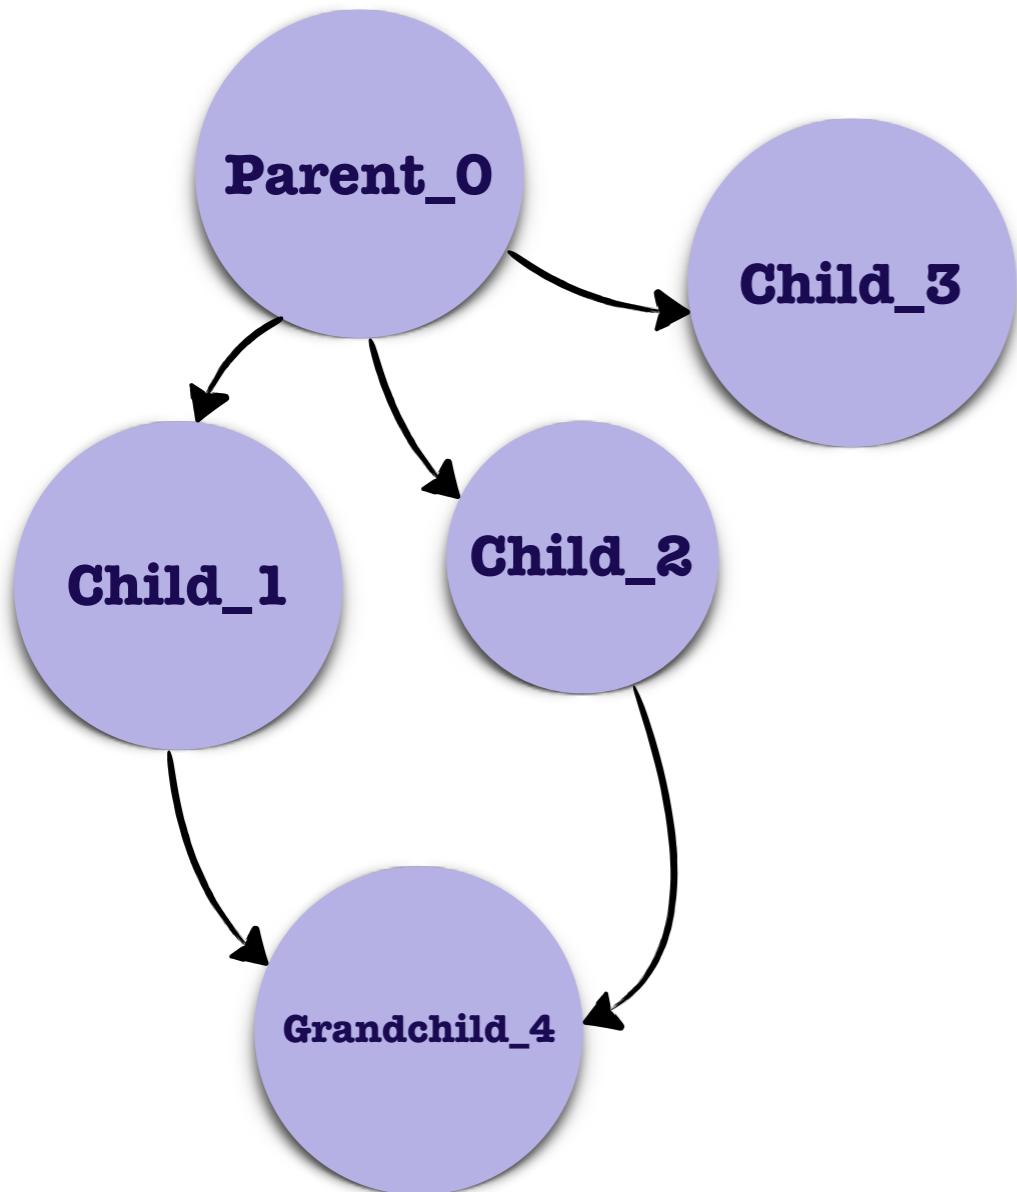


PROCESS QUEUING DIAGRAM



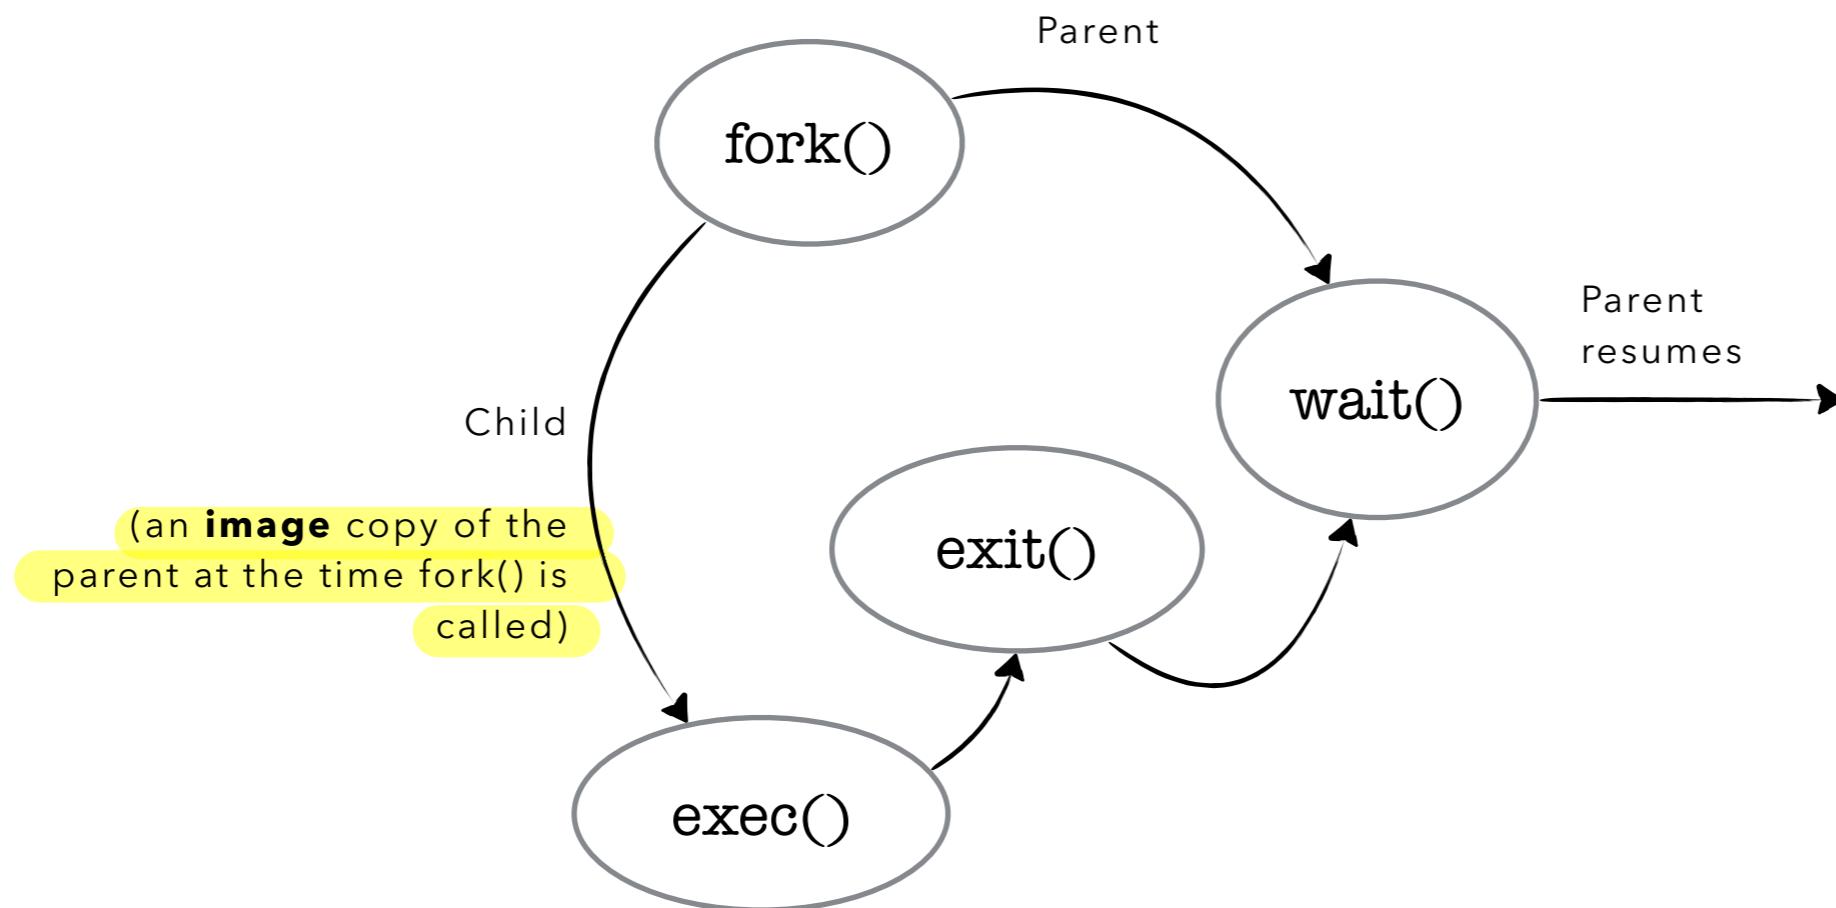
PROCESS CREATION

New processes are created by **fork()** system call in UNIX-based machines



- Each process can create another process, forming process tree.
- Processes can execute **concurrently**
- Parents and children belong to **different virtual address space**, but they can share resources and communicate as well using **shared memory**
- Children is a **duplicate** of parent upon creation
- Parents can wait for children using `wait()`, but not the other way around

FORK() SYSTEM CALL



1. `exec()` replaces the current process image with a new process image.
2. Loads the program into the new process space and runs it from the **entry** point

Therefore, `fork()` system call is necessary before `exec(program)` so that the new executable program does not replace the parent's process space .

FORK SYSTEM CALL IN C

```
void fork_example_3(){

    int returnValue = fork();
    // child process because return value zero
    if (returnValue == 0){
        printf("Hello from Child!\n");
        exit(0);
    }
    // The value of returnValue in parent process is the child's pid
    // parent process because return value non-zero.
    else{
        printf("Hello from Parent!\n");
        wait();
    }
    //parent or child can run concurrently
    //the order of output is unclear, because
    //we don't know how OS execute them
}
```

Output possibility 1:

Hello from Child!
Hello from Parent!

Output possibility 2:

Hello from Parent!
Hello from Child!

There are 2 possible outputs because
we do not know which process will
the scheduler decide to run first.

```
void fork_example_3(){

    int returnValue = fork();
    // child process because return value zero
    if (returnValue == 0){
        printf("Hello from Child!\n");
        exit(0);
    }

    // parent process because return value non-zero.
    else{
        printf("Hello from Parent!\n");
        wait();
    }
    //parent or child can run concurrently
    //the order of output is unclear, because
    //we don't know how OS execute them
}
```

————— Parent

Note: see more `fork()` examples
in the google drive

Child —————
(exact copy of parent's
code, continued from
`fork()` onwards)

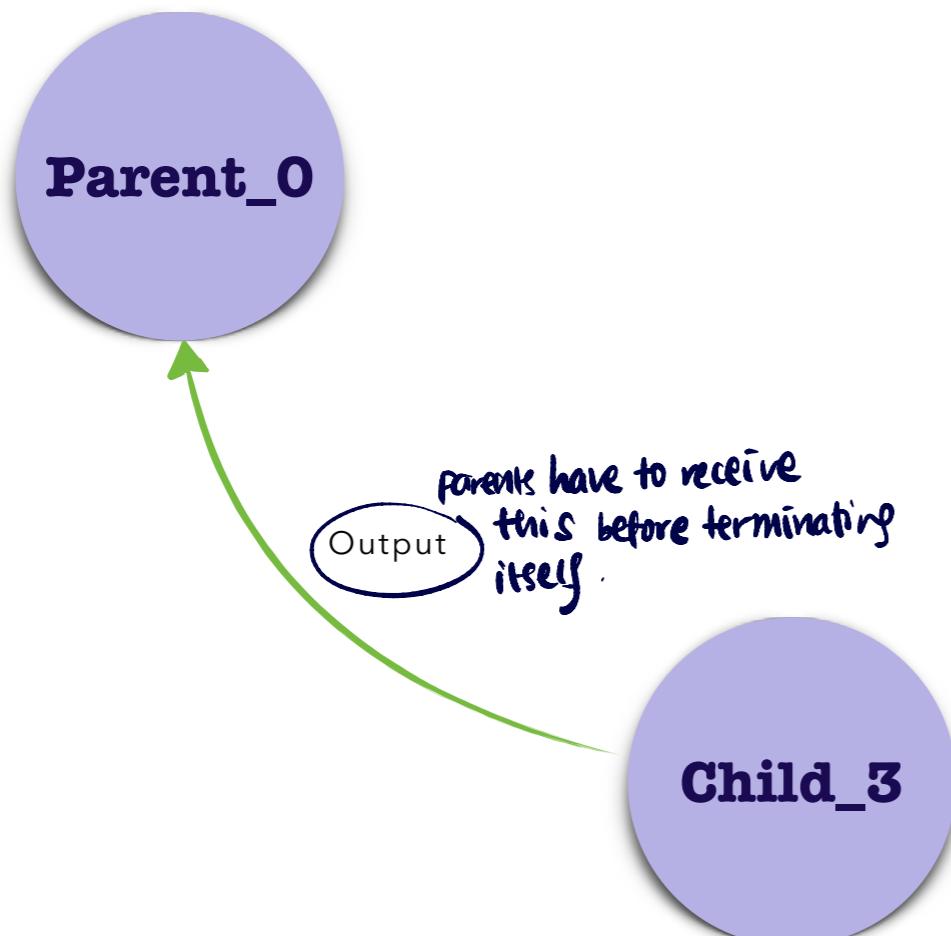
```
void fork_example_3(){

    int returnValue = fork(); -Entry point
    // child process because return value zero
    if (returnValue == 0){
        printf("Hello from Child!\n");
        exit(0);
    }

    // parent process because return value non-zero.
    else{
        printf("Hello from Parent!\n");
        wait();
    }
    //parent or child can run concurrently
    //the order of output is unclear, because
    //we don't know how OS execute them
}
```

PROCESS TERMINATION

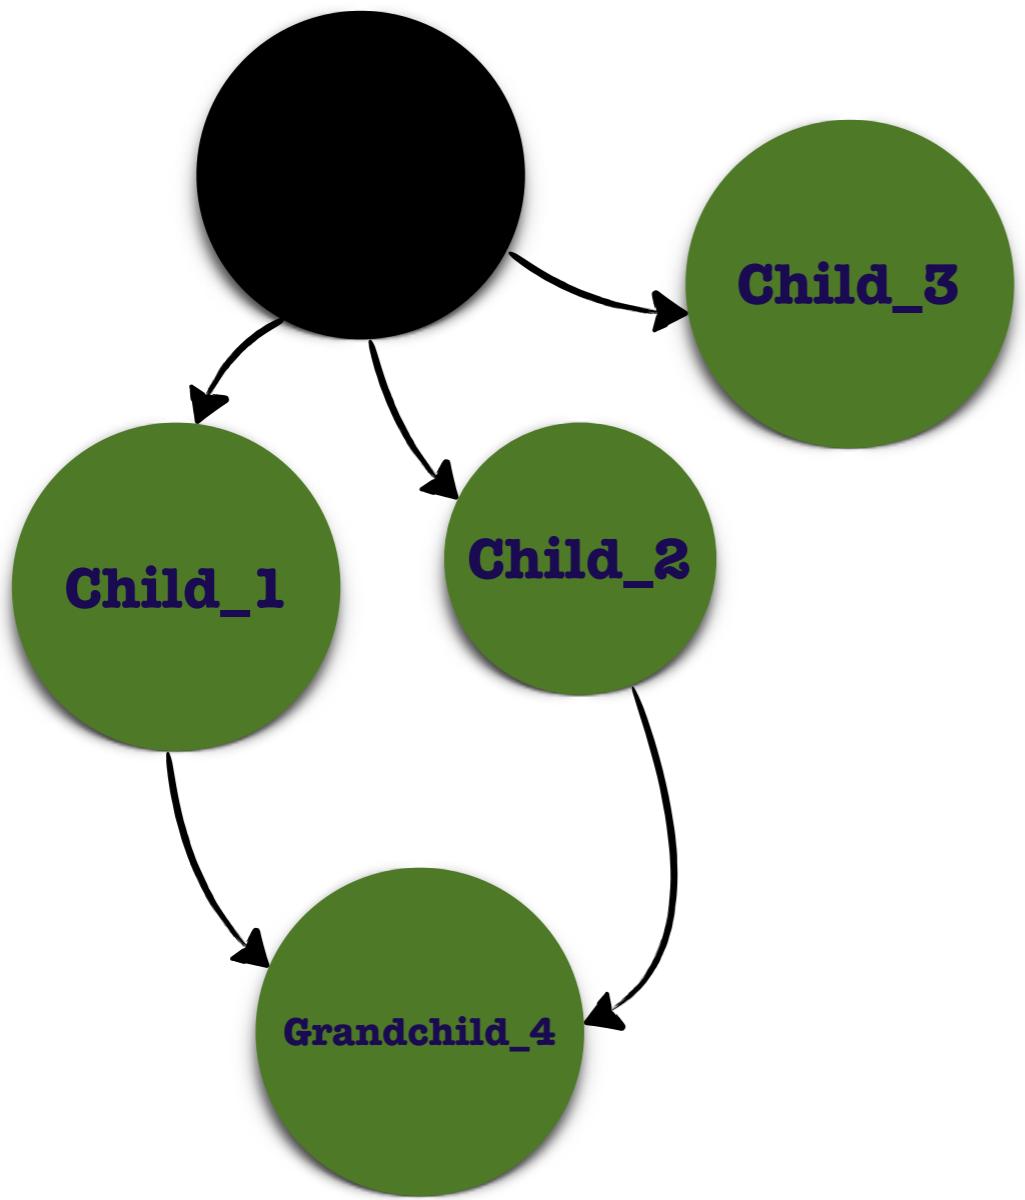
Done via **exit()** system call. Resources will be deallocated by the OS



- Parent processes can abort child processes
 - If child exceeds resources (memory /time)
 - If task is no longer required
 - If parent wants to exit
- Child process can exit using system call **exit(status)**
- Usually **status** is an int indicating whether or not the exiting process is successful
- Parent process can receive output from child process when it waits for child process to exit: **pid_t**
wait(&status)
wait() & waitpid(pid) will block the caller (the parent, basically), UNTIL any child finishes (for wait()) or until the specific child (for waitpid(pid)) finishes. Otherwise if child has finished / child w/ that id doesn't exist / no more child exists, wait() / waitpid (pid) will return on the caller can resume normally

ZOMBIE PROCESSES

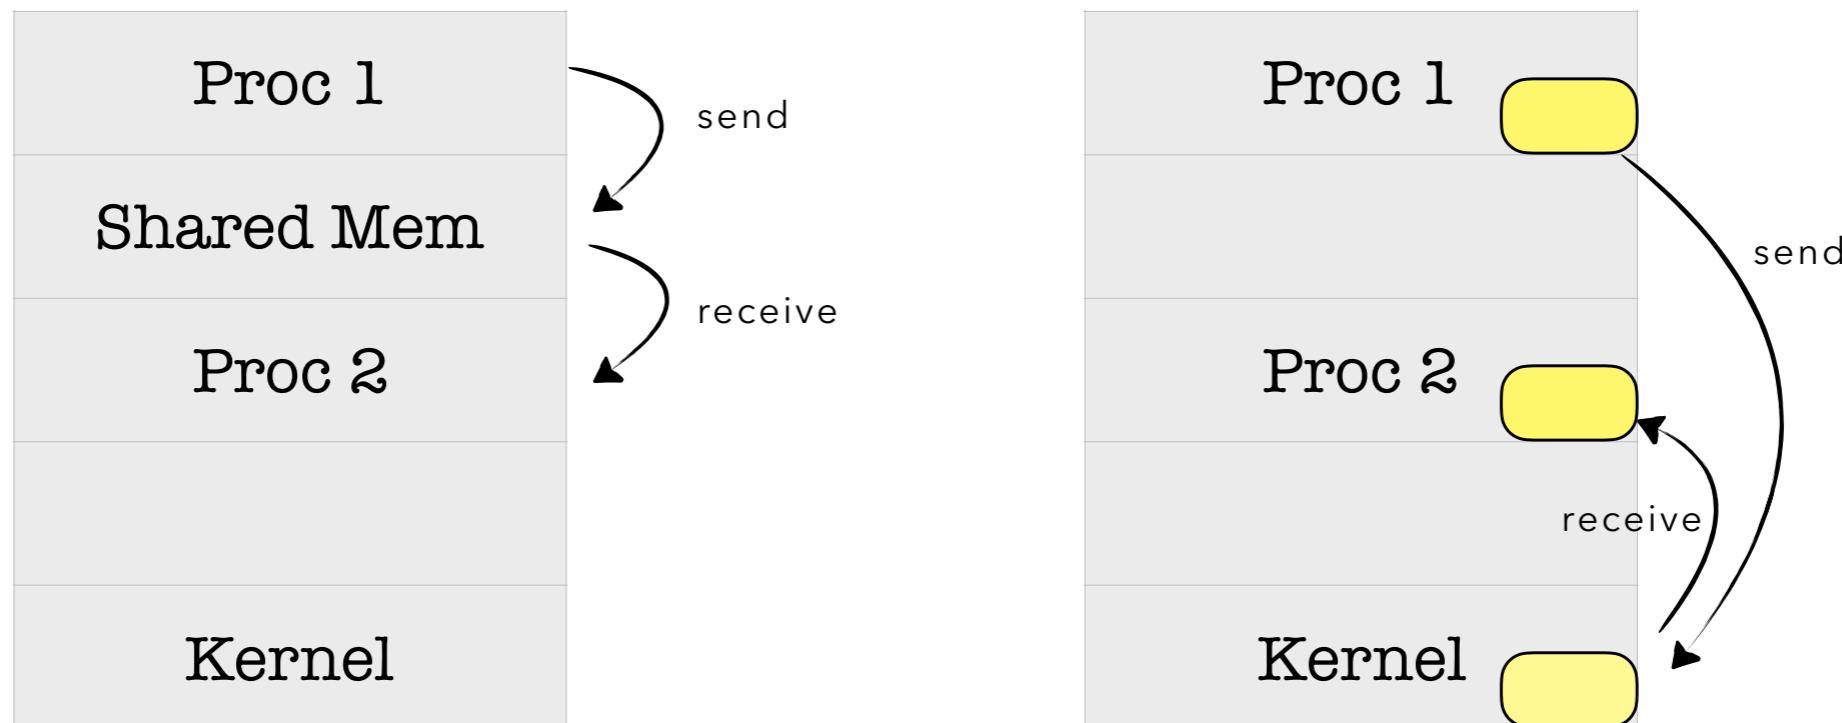
Child processes whose parents already terminated become zombie processes



- The black node: Parent_0 has terminated
- All the children processes have nobody to return to, hence turning into **zombie** processes
- OS can help abort **zombie** processes by themselves
- Otherwise, they will take up resources until computer restarts → **memory /resource leak**

① so a good practice is to always wait for child process (if any) to terminate. This wait() system call harvests the child process's exit return value and remove the child's process table entry.

INTER PROCESSES COMMUNICATION



Shared memory

- ① system call to create shared memory region
- ② Proc 1 & Proc 2 can read & write to the shared memory in user mode
- ③ Requires synchronization protocol so that processes do not overwrite each other

Sockets

- ① Requires system call to read
- ② Requires system call to write
- ③ Kernel will automatically synchronize process communication

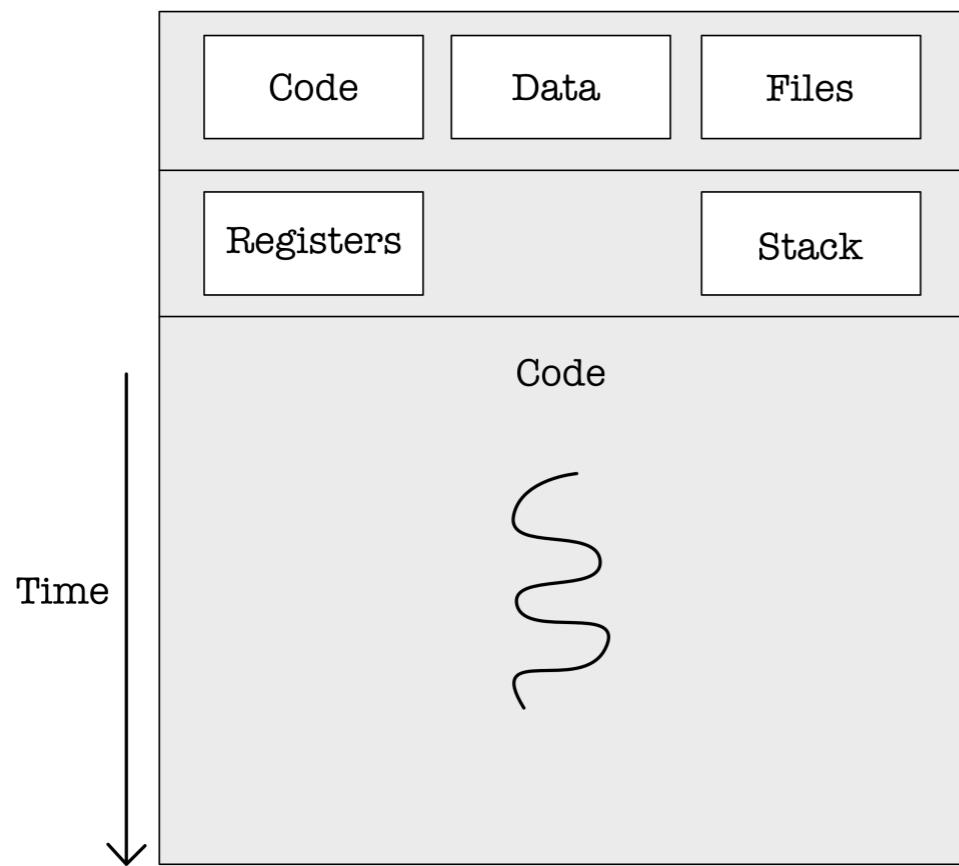
Process couples (gives) concurrency and protection

↓
interleaved execution

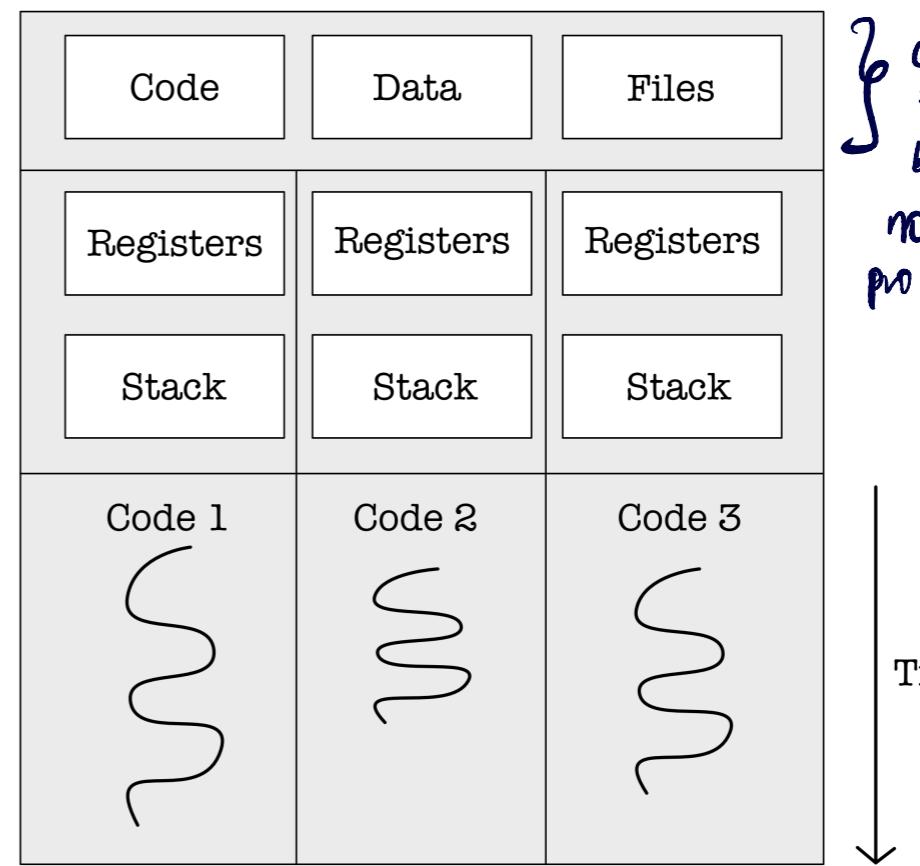
↓
address space isolated from one another

Thread gives concurrency **without protection**.

• THREADS



A single process with a single thread



A single process with three threads,
executing different entry points
of the same code. Execution
of threads of the same process
is interleaved, done by
the **thread scheduler in User mode**.

} code, data , and files are shared between threads .
No need to copy entire process space like fork()

PROCESSES

More overhead to create

because copying over the entire VM space, including code, data, & files

Requires private address space, has protection against other processes

Processes are independent of one another, no concurrency issues

No synchronization overhead, easier to program and work as intended

Can benefit from parallel execution on multiprocessor system

A process can have many threads executing different tasks

THREADS

Lesser overhead to create

shares data, code, and files

Easier to create since they share address space, only differ in program execution

Requires careful programming since threads of the same process share data structures and hence are interdependent

Can potentially suffer from synchronization overheads, harder to program and work as intended

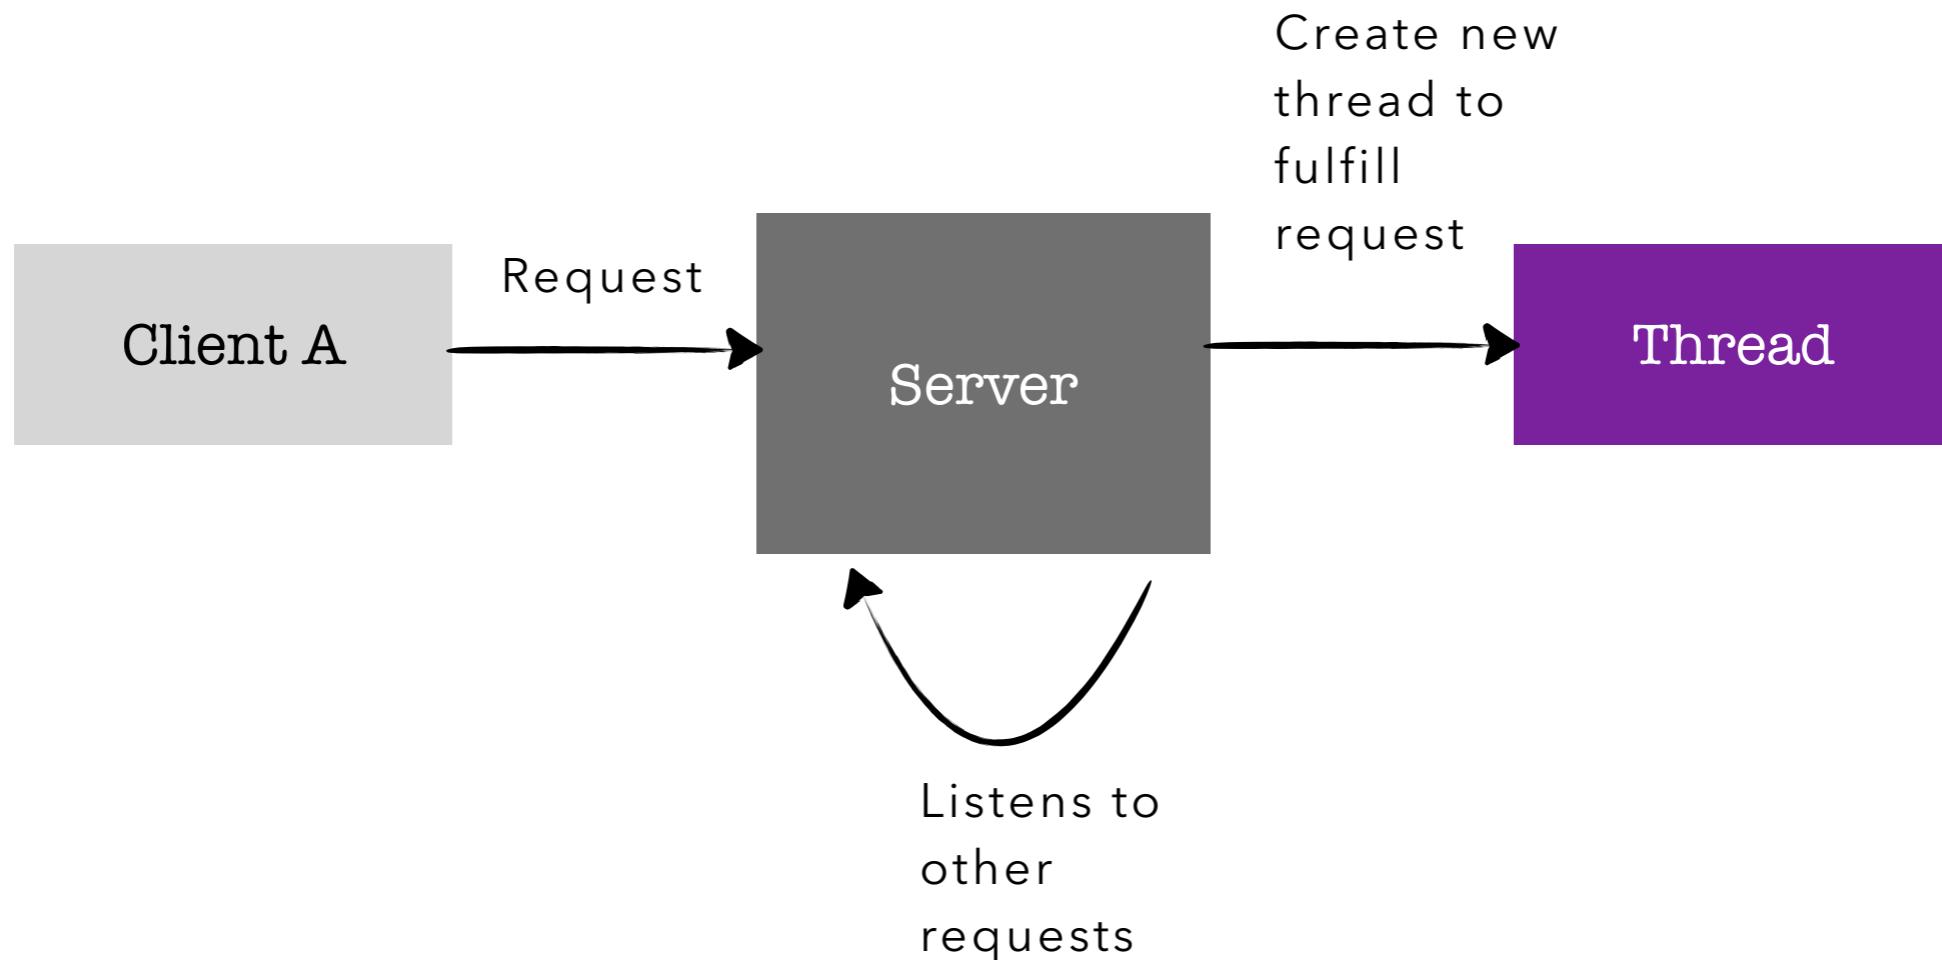
Can benefit from parallel execution on multiprocessor system

Good for responsiveness of a process

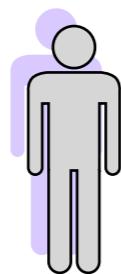
MULTITHREADED ARCHITECTURE



one of the uses of threads

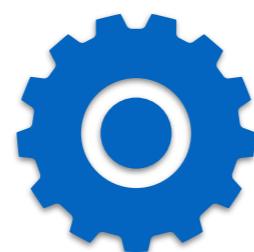


TYPES OF THREADS



User threads

- Not known to kernel
- Scheduled by thread scheduler in thread library
- Runs in user mode, cheaper to create



Kernel threads

- Way larger than user
thread data structure
- Has kernel data structure: thread control block
perform kernel thread's own
(expensive to create and context switch)
 - Scheduled and known by kernel
** runs in kernel mode*

Note: if one kernel thread is blocked, the other kernel threads that are runnable can run.

User mode cannot interrupt kernel mode.

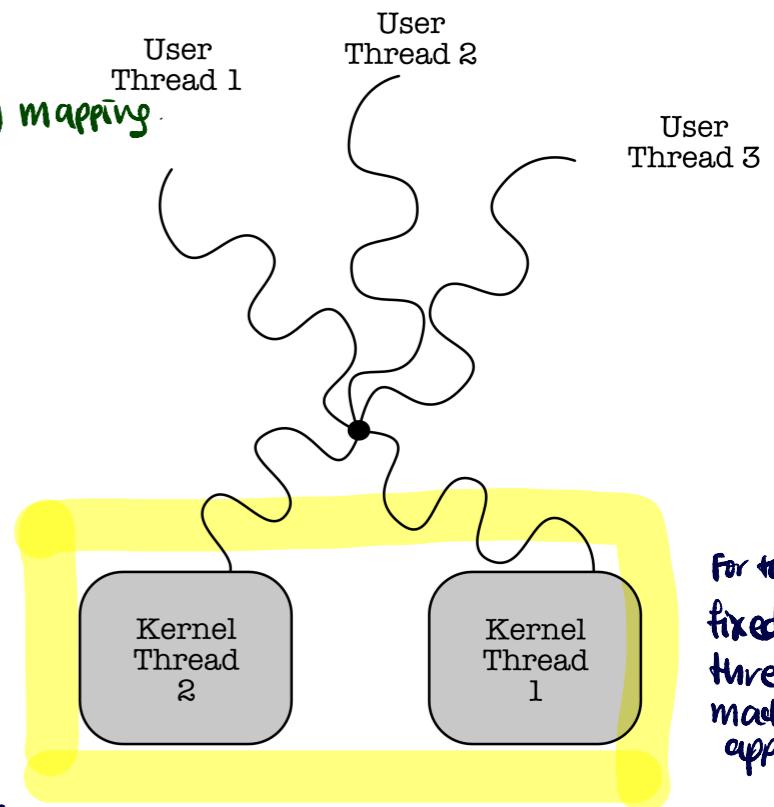
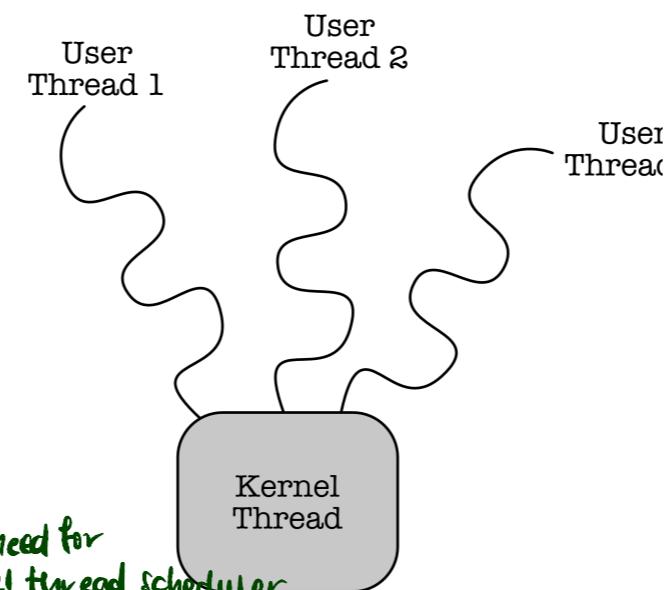
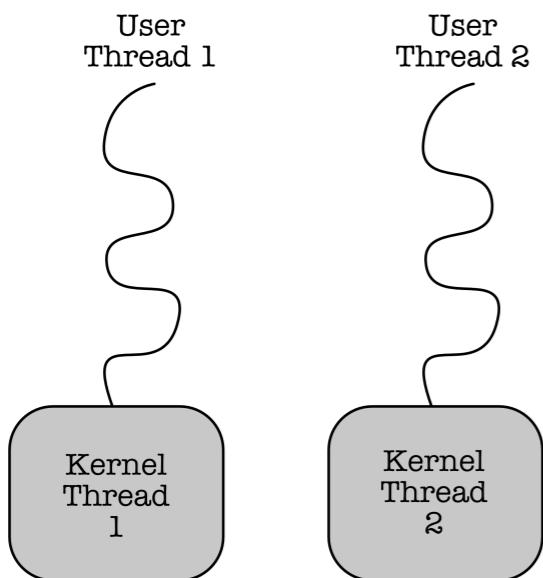
However within kernel mode itself, the kernel scheduler is free to choose which kernel thread to run / pause / kill.

Note: multiple kernel threads are required to achieve parallelism if the system has multiple cores.

↓
single
because a kernel process oversees ALL hardwares, so to use different cores the kernel process has to create kernel threads.

THREAD MAPPING

Thread mapping is SUPPORTED by the programming language's API (thread library). If you want your app to make use of your multicores then choose a library that allows either one to one mapping or many to many mapping.



For this model, fixed kernel thread per machine or per app.

one to one

- ✓ maps each user thread to kernel thread
- ✓ provides more concurrency than many to one model
- ✓ allows multiple threads to run on multiprocessor

X kernel thread creation overhead · Developer cons: May not be able to create too many user threads · (Kernel threads are large in size)

pros: no need for kernel thread scheduler

cheap, but only 1 thread can access

the kernel at a time. Multiple

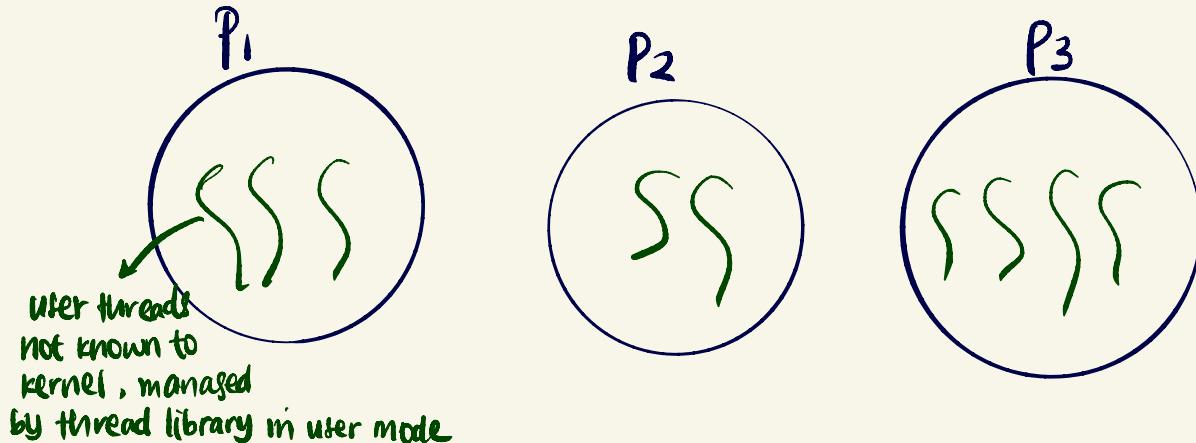
one to many threads aren't able to run on multiprocessors

X an entire process will block if a thread in that process makes a blocking system call.

many to many (BEST)

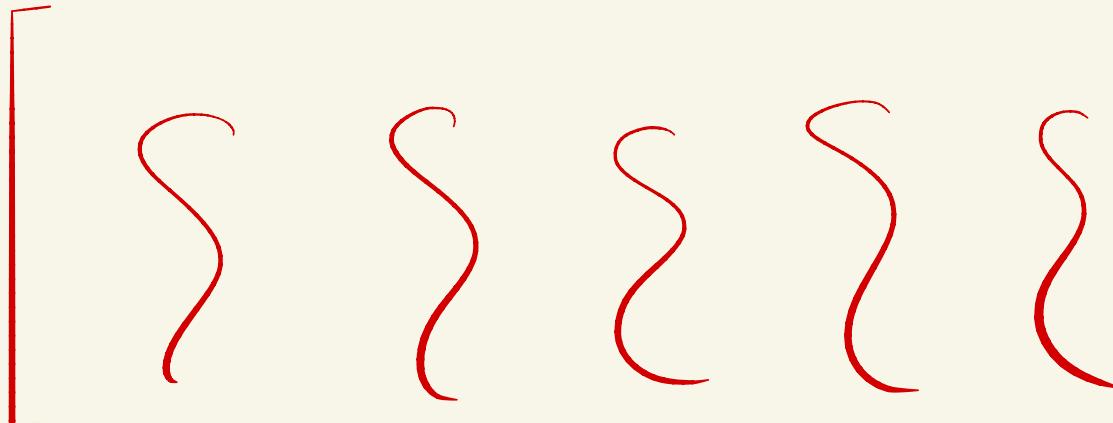
user can create as many user threads, but true parallelism (true concurrency) isn't gained because the kernel can only schedule a process → then a user thread at a time.

User mode
(in user space)

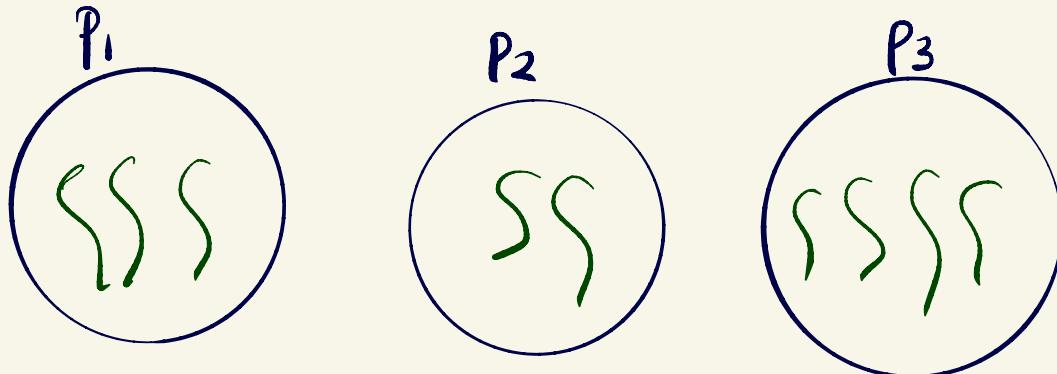


Kernel mode
(in Kernel space)

Kernel only has
a SINGLE process
but can have
multiple kernel
threads to exec
diff tasks / access
diff cores

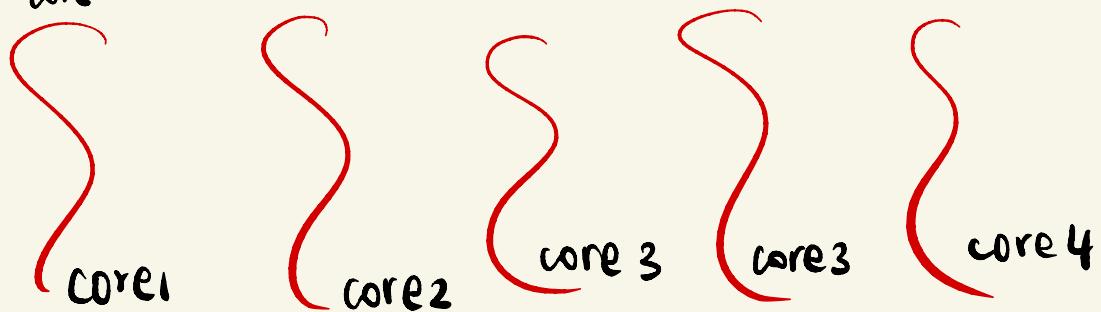


User mode
(in user space)



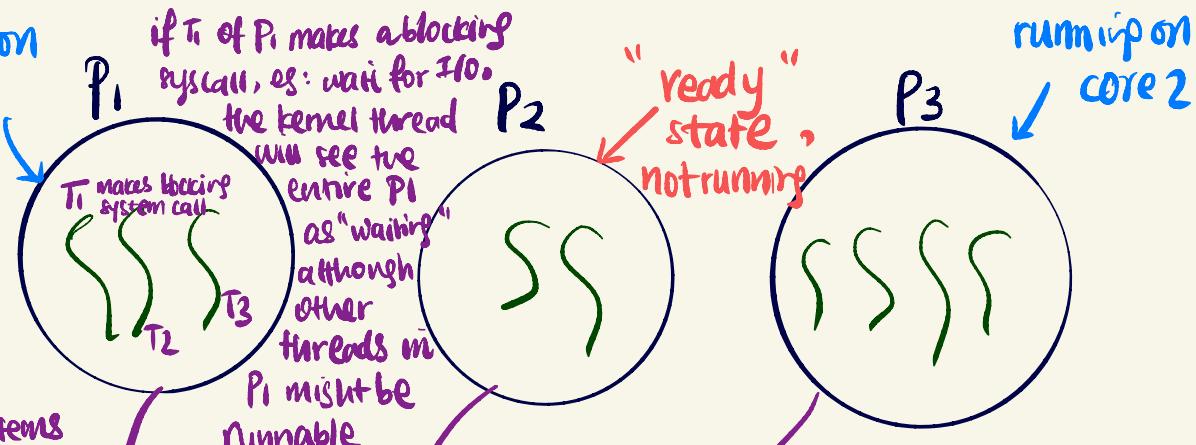
Kernel mode
(in kernel space)

The kernel can benefit from multicore hardware
iff it can create at least 1 kernel thread assigned
per core .



Eg: running on
User mode core 1
(in user space)

Many-to-1 model → each process and a bunch of its threads is assigned to ONE kernel thread. → cannot benefit from multicore systems



Kernel mode
(in Kernel space)

This thread can be scheduled to run other processes but not P_i because T_i in P_i "blocks" the entire process

core 1

core 2

core 3

core 3

core 4

User mode
(in user space)

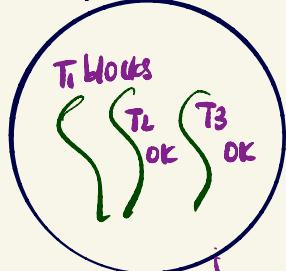
many to many model:

The user thread system library can choose to run T_2 or T_3 on kernel thread 2, so P_1 can still proceed when T_1 blocks.
also, it can benefit on multicore system,

e.g.: $P_1 T_2$ on KT_2 , $P_1 T_3$ on KT_3

Kernel mode
(in Kernel space)

Eg: running on core 1 & 2



P_2

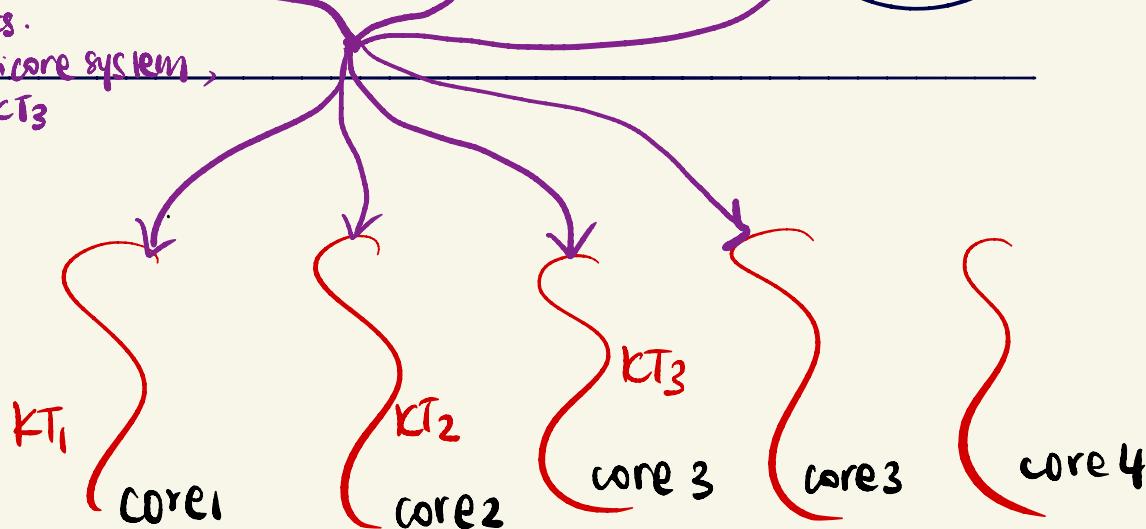
SS

ready state,
not running

P_3

SSSS

running on
core 3



User mode

(in user space)

one to one model: the earliest mapping to achieve true concurrency (parallelism).

User thread library will make system call for kernel to create a new thread for each user thread.

Kernel mode

(in Kernel space)

④ True parallelism is

limited by the # of cores. Threads, both kernel and user threads

are just enabling interleaved execution

E.g. running on core

P₁ ↕ 1, 2, 3,
true
parallelism

running on

core 3,
2 kernel
threads
interleaved

running
on
core 4,
4 kernel
threads
interleaved

SSS

SS

SSSS

core 1

core 2

core 3

core 3

core 4

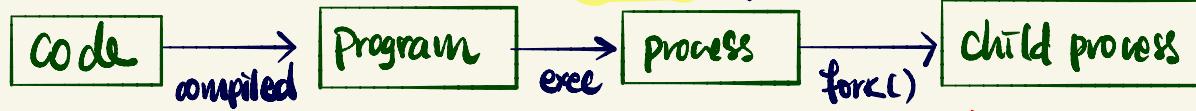
c₄

c₄

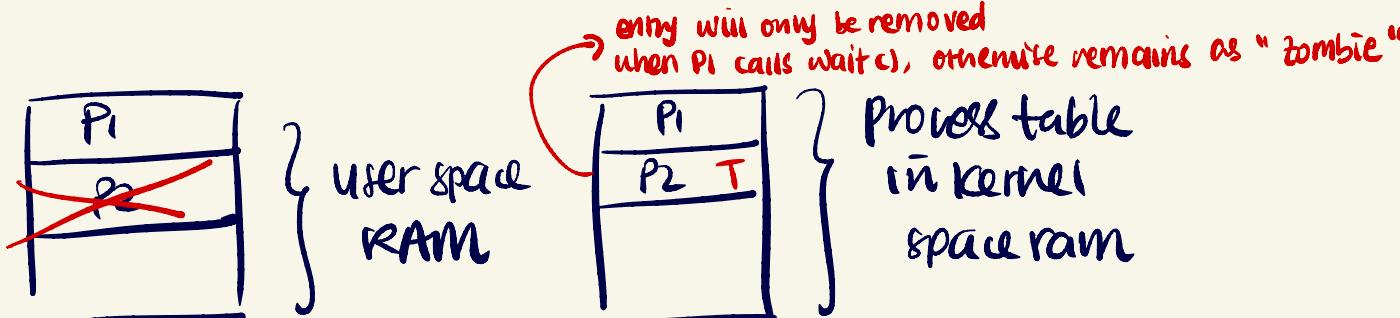
c₄

c₄

Summary :

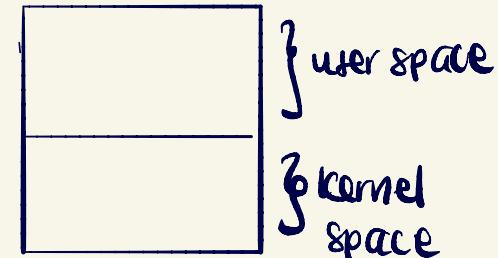


Parent process has to
wait() for child process



If $P_1 \xrightarrow{\text{fork()}} P_2$, when
 P_2 calls `exit()`, P_2 's
 user space will be deallocated,
 P_2 will be labelled as Terminated (T),
 but entry in proctable remains

Note : recall,
 RAM is divided into
 user and kernel space .



* The process table contains process control blocks for kernel to schedule processes.

* There are 3 queues : job , device , ready queue

→ to quickly determine which process is next ,
then scheduler refers to the corresponding
block on process table for further info .

→ scheduled by kernel

* Process provides:

① concurrency

② Isolation / protection (expensive) → need socket / shared mem for IPC

* Thread provides: concurrency without protection →

shares code, data,
and files

Will need to implement extra
measures to synchronize when
coding apps that creates multiple
threads within a single process .

→ scheduled by

thread library

in user mode .

→ responsible to
implement whichever
thread mapping suitable
for the user