

Deadlocks

Deadlock problem. Necessary conditions for deadlock.
Resource allocation graphs. Java deadlock examples.
Deadlock prevention. Deadlock avoidance (safe state and Banker's algorithm). Deadlock detection and recovery.

OS5


Textbook (SGG): Ch. 7.1-7.4, 7.5.1, 7.5.3, 7.6.2, 7.7

The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set (*circular wait*: simplest process P waiting for Q; Q waiting for P)
- Example 1
 - System has 2 disk drives
 - P_1 and P_2 each hold one disk drive and each needs the other one
- Example 2: (binary) semaphores A and B , initialized to 1

P_0	P_1
acquire(A);	acquire(B);
acquire(B);	acquire(A);

Bridge Crossing Example



- Traffic only in one direction
- Each section of a bridge can be viewed as a resource
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
 - Several cars may have to back up in general
- Starvation is possible (cars in one direction only keep going)
- Note: most operating systems do not prevent or resolve deadlock completely (users will deal with it when needed)

General System Model

- Resource types R_1, R_2, \dots, R_m
 - ▶ Physical interpretation of resource types: CPU, memory, I/O devices, etc
- OS has W_i instances (or units) of each resource type R_i that it can allocate to requesting processes
- Each process utilizes a resource as follows:
 - **request** (request one instance of a resource from OS; e.g., get access to a printer)
 - **use** (use the acquired resource privately to do its work; e.g., print a file to the acquired printer)
 - **release** (return an acquired instance of a resource back to the OS; e.g., give back the acquired printer so the OS can give this printer to another process)

Necessary Conditions for Deadlock

Deadlock *can* occur (but may not necessarily occur) if these four conditions hold *simultaneously*:

- **Mutual exclusion:** Only one process at a time can use a resource
- **Hold and wait:** A process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption:** A resource can be released only voluntarily by the process holding it, after process has completed its task (hence use of the resource)
- **Circular wait:** There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0

NB: These conditions are *necessary*, but *not sufficient*, for deadlock

Implication: In principle, we can solve deadlock by removing *any* of these conditions

Resource Allocation Graph

A *directed graph*, with set of vertices/nodes V and set of edges E

- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set of all the processes in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set of all the resource types in the system
- **request edge**: directed edge $P_i \rightarrow R_j$
 - P_i wants to acquire an instance of R_j
- **assignment edge**: directed edge $R_j \rightarrow P_i$
 - An instance of R_j is being held by P_i

Resource-Allocation Graph (cont'd)


- Process (first type of nodes/vertices)




- Resource type with 4 instances (second type of nodes)




- P_i requests an instance of R_j (first type of directed edge)



- P_i is holding an instance of R_j (second type of directed edge)



Example Resource Allocation Graph




Is this system deadlocked? *Hint:* Is there circular wait?

Resource Allocation Graph with a Deadlock


P_1 holds an instance of R_2 , needs an instance of R_1 to continue execution so that it may complete its task later

This graph *has* a cycle



Can any of P_1, P_2, P_3 acquire the resource it wants to complete its task?
 Hint: The processes won't release the resources they are holding until they can continue execution and complete their respective tasks.

Graph with A Cycle but No Deadlock



Multiple instances of resource: Cycle is necessary, but not sufficient, condition for deadlock

This graph has a cycle. But what is a possible completion sequence of these processes?

Summary: Graph Cycles and Deadlock

- If graph contains no cycles \Rightarrow no deadlock
- If graph has a cycle \Rightarrow
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, then *possibility* of deadlock

Java Deadlock Example

```
class A implements Runnable
{
    private Lock first, second;

    public A(Lock first, Lock second) {
        this.first = first;
        this.second = second;
    }

    public void run() {
        try {
            first.lock();
            // do something
            second.lock();
            // do something else
        }
        finally {
            first.unlock();
            second.unlock();
        }
    }
}
```

Thread A

```
class B implements Runnable
{
    private Lock first, second;

    public A(Lock first, Lock second) {
        this.first = first;
        this.second = second;
    }

    public void run() {
        try {
            second.lock();
            // do something
            first.lock();
            // do something else
        }
        finally {
            second.unlock();
            first.unlock();
        }
    }
}
```

Thread B

Note: each runnable object needs to get *two* locks to do its work

Java Deadlock Example

```
public static void main(String arg[]) {  
    Lock lockX = new ReentrantLock();  
    Lock lockY = new ReentrantLock();  
  
    Thread threadA = new Thread(new A(lockX,lockY));  
    Thread threadB = new Thread(new B(lockX,lockY));  
  
    threadA.start();  
    threadB.start();  
}
```

lockX and **lockY** are binary locks, initially available. How can deadlock occur? (Demonstrate an interleaving of the two threads' execution.)

Hint: Each runnable object in previous slide runs as a separate thread;
These threads acquire the locks **lockX** and **lockY** in *different* orders

Methods for Handling Deadlocks

■ Deadlock *avoidance*

- Before granting a resource request (even if request is valid and the requested resources are now available), check that the request will not cause the system to enter a deadlock state (not just no deadlock immediately, but *not even later*)
- Requires advance knowledge of *future* resource needs (e.g., Banker's algorithm in your OS Lab #3)

■ Deadlock *detection and recovery*

- Detect deadlock after the fact, then recover from it (e.g., preempting held resources and rolling back processes)

■ Deadlock *prevention*

- Impose conditions on resource requests to ensure that a valid request can *never* cause the system to enter a deadlock state *by design* (so no need to check the runtime system state as in deadlock avoidance)

■ Real-world OS (e.g., UNIX) may not handle deadlocks completely (i.e., possible that users will see a deadlock; then they'll deal with it)

Deadlock Avoidance

Requests can lead to deadlock (in the future); don't grant those requests (now) even if the requested resources are currently available

- Need additional a priori information: Simplest and most useful model requires each process to declare (in advance) the *maximum number* of resources of each type that it will ever need.
- At time of resource request, avoidance algorithm examines the system's resource allocation state to ensure that granting the request will *never lead to a circular-wait condition later*.
- Resource allocation state is defined by the numbers of currently available and allocated resources, and the maximum number of resources that each process may need in the future.


Safe State

- When a process requests an available resource, system must decide if granting the request will leave the system in a safe state.
- System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of all the processes in the system such that for each P_i , the resources that P_i will ever need can be satisfied by the currently available resources *plus* the resources held by all the preceding P_j , with $j < i$.
- Rationale:
 - If P_i 's resource needs are not immediately available, then P_i can wait until all the preceding P_j processes have finished.
 - When the preceding processes all finished, P_i can obtain its needed resources, do its job, *return its allocated resources*, and finish.
 - When P_i finishes, P_{i+1} can obtain its needed resources and finish, and so on.
 - I.e., $\langle P_1, P_2, \dots, P_n \rangle$ defines a feasible order for all the processes to finish.

Basic Facts

- If system is in safe state \Rightarrow no deadlocks
- If system is in unsafe state \Rightarrow *possibility* of deadlock
- Avoidance \Rightarrow ensure that system will never enter an unsafe state.
 - So we'll grant a resource request only if *after* granting the request, the system will still be in a safe state.
 - We need an algorithm to decide if a system is in a safe state (Slide 7.22)

Safe, Unsafe, Deadlock States



Avoidance algorithms

- Single instance of a resource type
 - Use a resource-allocation graph
 - We won't cover this algorithm (limited applicability)
- Multiple instances of a resource type
 - Use the Banker's algorithm (your Lab #3)
 - Subsumes the single-instance problem above

Banker's Algorithm

- Multiple instances of each resource
- Each process P must *a priori* declare its maximum needs (i.e., maximum number of instances of each resource type that P will ever need)
- When a process requests a resource, it may have to wait
- After a process P got all its resources, it must return them within a finite amount of time (when P finishes its task)

Data Structures for Banker's Algorithm

Let n = number of processes, and m = number of resources types

- **Available:** Vector (i.e., 1D array) of length m . If $Available[j] = k$, there are k instances of resource type R_j available
- **Max:** $n \times m$ matrix (i.e., 2D array). If $Max[i,j] = k$, then process P_i may request at most k instances of resource type R_j
- **Allocation:** $n \times m$ matrix. If $Allocation[i,j] = k$ then P_i is currently allocated k instances of R_j
- **Need:** $n \times m$ matrix. If $Need[i,j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:
 $Work = Available$ // NB: *Work*, *Available* are both arrays
 $Finish[i] = false$ for $i = 0, 1, \dots, n - 1$
2. Find an i such that:
 - (a) $Finish[i] = false$
 - (b) $Need[i] \leq Work$If no such i exists, go to Step 4
3. $Work = Work + Allocation[i]$ // $Allocation[i]$ is i -th row of *Allocation* matrix
 $Finish[i] = true$
go to Step 2
4. If $Finish[i] == true$ for all i , then the system is in a safe state; otherwise it's unsafe

Algorithm for Granting Request by Process P_i

$Request[i]$ = request vector for process P_i . If $Request[i, j] = k$, then process P_i wants k instances of resource type R_j

1. If $Request[i] \leq Need[i]$ go to Step 2. Otherwise, raise error, since process has exceeded its maximum claim.
2. If $Request[i] \leq Available$, go to Step 3. Otherwise P_i must wait, since the resources are not immediately available.
3. Try to allocate the requested resources to P_i by updating the resource allocation state as follows (i.e., assume we grant the new request):

$$Available = Available - Request[i];$$

$$Allocation_i = Allocation[i] + Request[i];$$

$$Need[i] = Need[i] - Request[i];$$

- If *new state is safe* \Rightarrow the resources are allocated to P_i
- If *new state is unsafe* $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored (i.e., new request not granted after all)

Activity 5.1: Banker's Algorithm

- 5 processes P_0 through P_4 ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances)

Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	

- Is the system safe? If so, give an execution sequence of the processes that demonstrates the safety. If not, why not?

Activity 5.1 (cont'd)

- Given the original system state in the previous slide, P_1 now requests (1,0,2)
- First, check that $\text{Request} \leq \text{Available}$, i.e., $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$
- Second, now assume we grant this request, which will update the data structures as follows:

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 1	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies the safety condition (i.e., system will remain safe after granting (1,0,2) to P_1)
- Can we now *further* grant (i) request (3,3,0) by P_4 ; or (ii) request (0,2,0) by P_0 ?

Deadlock Detection

- Allow system to enter deadlock state
- Detect occurrence of deadlock by *detection algorithm*
- Recover from the detected deadlock

Deadlock Detection: Multiple Instances

- **Available:** A vector (i.e., 1D array) of length m indicates the number of available resources of each type.
- **Allocation:** An $n \times m$ matrix (i.e., 2D array) defines the number of resources of each type currently allocated to each process.
- **Request:** An $n \times m$ matrix indicates the current request of each process. If $Request[i, j] = k$, then process P_i is requesting k more instances of resource type R_j .

NB: The deadlock *detection* algorithm (next slide) detects if the system is *already* deadlocked. It's similar to the safety algorithm (Slide 7.22) used by the Banker's algorithm for deadlock avoidance, but the safety algorithm detects if the system *can* become deadlocked *in the future*.

Detection Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:
Work = *Available*
For $i = 1, 2, \dots, n$, if *Allocation*[i] $\neq 0$, then
Finish[i] = false; else *Finish*[i] = true
2. Find an index i such that both:
 - (a) *Finish*[i] == false
 - (b) *Request*[i] \leq *Work*. // What i requests now is available if all its preceding processes
// finish; i isn't deadlocked alreadyIf no such i exists, go to Step 4
3. *Work* = *Work* + *Allocation*[i]. // Assume i will finish and return the resources it holds now
Finish[i] = true
go to Step 2
4. If *Finish*[i] == false for some i , then the system is (already) in deadlock state.
Moreover, if *Finish*[i] == false, then P_i is (already) deadlocked.

What is the complexity (in terms of m and n) of the algorithm?

Example of Detection Algorithm

- Five processes P_0 through P_4
- Three resource types: A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

- Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $Finish[i] = \text{true}$ for all i

Detection Example (cont'd)

- Given the original system state in the previous slide, P_2 now requests an additional instance of type C, so that we have these requests:

	<u>Request</u>		
	A	B	C
P_0	0	0	0
P_1	2	0	2
P_2	0	0	1
P_3	1	0	0
P_4	0	0	2

- What's the new state of the system?
 - Can assume that P_0 will return the resources it holds, but there'll still be insufficient resources to satisfy any other processes' requests
 - Deadlock already exists, consisting of processes P_1 , P_2 , P_3 , and P_4

Deadlock Recovery: Abort Processes

- Abort all deadlocked processes (resources held are preempted)
- Abort one process at a time until the deadlock cycle is eliminated
- Which order should we choose to abort?
 - Priority of the process.
 - How long process has computed, and how much longer until completion?
 - Resources the process has used,
 - Resources the process needs to complete.
 - How many processes will system need to abort?
 - Is process interactive or batch?
 - etc ...

Deadlock Prevention

Constrain the ways requests can be made, in order to disallow “hold-and-wait,” or “no preemption,” or “circular wait” (in prevention, valid requests won’t cause deadlock *by design*, no need to check for safety of runtime system state)

- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 - Require process to request and be allocated all its resources before it begins execution, *or* allow process to request resources only when the process has none (i.e., previously held resources must be released before new resources are requested)
 - OS must support a new system call for acquiring multiple resources *at the same time*
 - ▶ E.g., for semaphores, define a new “acquireAll([list of semaphores])” system call for acquiring all the semaphores in the argument list, in addition to normal acquire() system call
 - Disadvantages
 - ▶ Low resource utilization if a process has to acquire in the beginning all the resources it’ll ever need – why?
 - ▶ Starvation becomes more likely – Process P needs two resources R1 and R2; R1 is available from time to time, similarly R2; but R1 and R2 are never available *at the same time*

Deadlock Prevention (cont'd)

- **No Preemption** – if a process P (holding some resources already) requests another resource that cannot be immediately granted, then OS will force P to release all the resources it already holds
 - Preempted resources are added to the list of resources for which the process needs to wait for again
 - Process will be restarted only when it can obtain all the resources it needs (both the old ones preempted and the new one requested)
 - Disadvantages: preemption costs and starvation
- **Circular Wait** – impose a total ordering of all the resource types, and require that each process requests resources according to that order
 - Say: give an id for each resource type (e.g., 1 for disk, 2 for printer, etc); process can only request resource types in *increasing* id number (i.e., can request printer while holding disk, but *not* vice versa)
 - Disadvantage: Burden on programmer to ensure the order by design, without unnecessarily sacrificing utilization