

# 50.007


# Machine Learning

Lu, Wei



# Reinforcement Learning (I)


# Bayes' Ball Algorithm



Are  $X$  and  $Y$  independent?




# Bayes' Ball Algorithm




Are  $X$  and  $Y$  independent? 

# Bayes' Ball Algorithm




Are  $X$  and  $Y$  independent?

# Markov Blanket



$$p(X|\mathbf{V}_{-X})$$

# Bayesian Networks




How do we learn such probability values?

$$\begin{aligned}
 & p(x_1 = 1, x_2 = 2, x_3 = 1, x_4 = 3, x_5 = 5, x_6 = 2) \\
 &= p(x_1 = 1) \times p(x_2 = 2 | x_1 = 1) \times p(x_4 = 3 | x_1 = 1)
 \end{aligned}$$

$$\times p(x_3 = 1 | x_2 = 2, x_4 = 3) \times p(x_6 = 2 | x_2 = 2, x_3 = 1) \times p(x_5 = 5 | x_4 = 3)$$

# Bayesian Networks



What is the number of free parameters involved in this table?

$x_1$	$x_2$	1	2	...	$r_3$
1	1				
1	2				
...	...				
1	$r_2$				
2	1				
2	2				
...	...				
2	$r_2$				
...	...				
$r_1$	1				
$r_1$	2				
...	...				
$r_1$	$r_2$				



# Overview of ML

Supervised  
Learning

Unsupervised  
Learning


# Overview of ML

Supervised  
Learning

A set of  
states

Unsupervised  
Learning

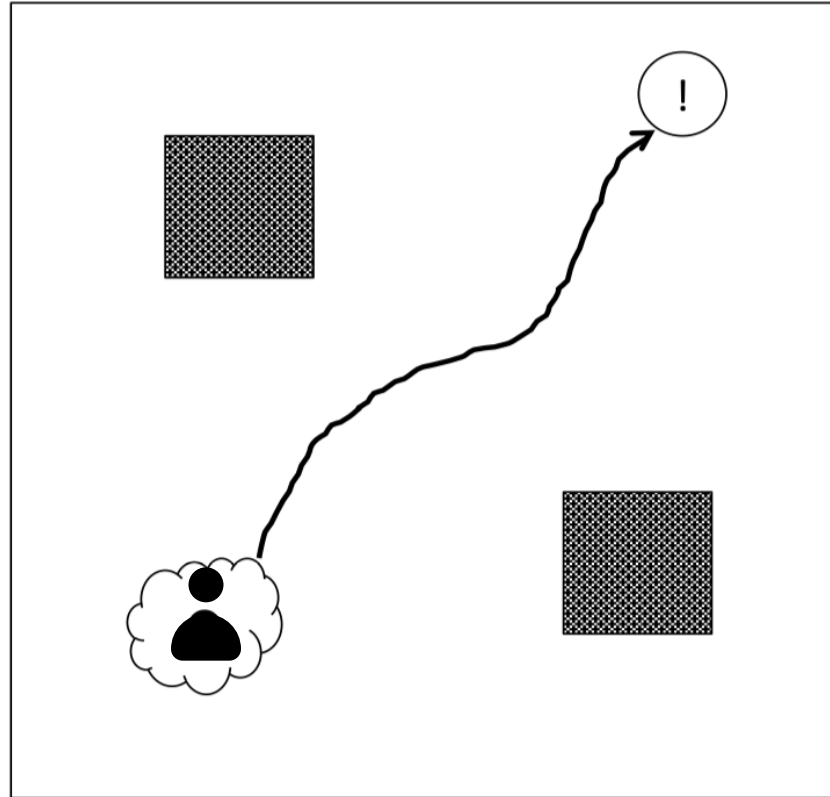
A set of  
actions



$f : S \rightarrow A$

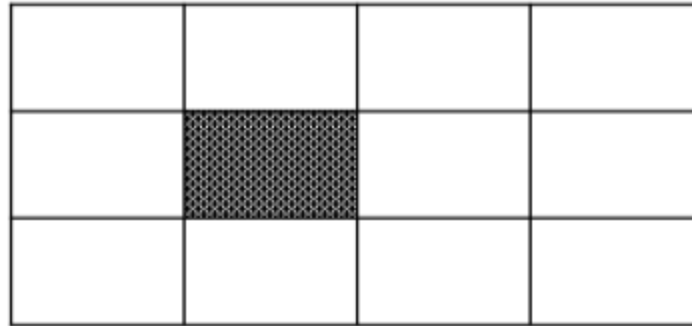
Reinforcement  
Learning

# Learn How to Act



How do we teach a robot how to act optimally in a complex environment?


# Block World Environment



The robot can be at one block at any given time.


There is a set of states  $S$

# Block World Environment




The robot can take an action from a set of predefined possible actions at each state.

There is a set of actions  $A$



# Block World

## Transition Probabilities



If the robot moves towards a particular direction, there is a 0.8 chance that it would reach the block in front, and there is a 0.1 chance to reach a state to its left (right).


A transition probability function

$$T(s, a, s') = p(s'|s, a)$$

# Block World

## Rewards



			<b>+1.0</b>
			<b>-1.0</b>



Each block is associated with a reward. The two blocks at the upper-right corner are assigned rewards +1 and -1.

# Block World

## Rewards



-0.6	+1.2	+0.1	<b>+1.0</b>
-0.1		-0.1	<b>-1.0</b>
+0.9	-0.7	-2.0	+1.3

Each block is associated with a reward. The two blocks at the upper-right corner are assigned rewards +1 and -1.

The reward is  $R(s)$  for each state.

In general it can be defined as  $R(s, a, s')$

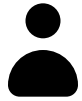
action

old state

new state



# Markov Decision Process



-0.6	+1.2	+0.1	<b>+1.0</b>
-0.1		-0.1	<b>-1.0</b>
+0.9	-0.7	-2.0	+1.3

- a set of states  $S$
- a set of actions  $A$
- a transition probability function  $T(s, a, s') = p(s'|s, a)$
- a reward function  $R(s, a, s')$  (or just  $R(s')$ )

# Block World

## Utility (Long Term Reward)



-0.6	+1.2	+0.1	<b>+1.0</b>
0.1		-0.1	<b>-1.0</b>
+0.9	-0.7	-2.0	+1.3

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots = \sum_{t=0}^{\infty} R(s_t)$$






Will this be a good way of defining long term reward?

# Block World

## Utility (Long Term Reward)




-0.6			<b>+1.0</b>
0.1		-0.1	<b>-1.0</b>
+0.9	-0.7	-2.0	+1.3

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots = \sum_{t=0}^{\infty} R(s_t)$$

Will this be a good way of defining long term reward?

# Block World

## Utility (Long Term Reward)



-0.6	+1.2	+0.1	<b>+1.0</b>
0.1		-0.1	<b>-1.0</b>
+0.9	-0.7	-2.0	+1.3



$$U([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots = \sum_{t=0}^{\infty} \gamma^t R(s_t)$$



discount  
factor

# Block World

## Utility (Long Term Reward)



-0.6	+1.2	+0.1	<b>+1.0</b>
0.1		-0.1	<b>-1.0</b>
+0.9	-0.7	-2.0	+1.3

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots = \sum_{t=0}^{\infty} \gamma^t R(s_t)$$

$$\frac{R_{min}}{1-\gamma} = \sum_{t=0}^{\infty} \gamma^t R_{min} \leq U([s_0, s_1, s_2, \dots]) \leq \sum_{t=0}^{\infty} \gamma^t R_{max} = \frac{R_{max}}{1-\gamma}$$

Lower  
Bound

↓  
Smallest Reward

↙  
Largest Reward

Upper  
Bound

# Block World


## Learning the Policy



-0.01	-0.01	-0.01	<b>+1.0</b>
-0.01		-0.01	<b>-1.0</b>
-0.01	-0.01	-0.01	-0.01



How to learn the policy?



→	→	→	<b>+1.0</b>
↓		←	<b>-1.0</b>
↑	→	↑	←

An example policy

# Value Iteration

$$\pi(s)$$

a particular *policy* that specifies the action we should take in state  $s$ .

$$V^\pi(s)$$

The *value* of state  $s$  under policy  $\pi$ .  
It is the expected long-term reward of starting in state  $s$  and act based on policy  $\pi$  thereafter.

$$Q^\pi(s, a)$$

The *Q-value* of state  $s$  and action  $a$  under policy  $\pi$ .  
It is the expected long-term reward of starting in state  $s$ , taking action  $a$  and acting based on policy  $\pi$  thereafter.

# Value Iteration

$$\pi^*(s)$$

The *optimal policy*  $\pi^*(s)$  specifies the optimal action we should take in state  $s$ .

$$V^*(s)$$

The *value* of state  $s$  under the optimal policy  $\pi^*$

$$Q^*(s, a)$$

The *Q-value* of state  $s$  and action  $a$  under the optimal policy  $\pi^*$



# Value Iteration

$$V^*(s) = ??$$

$$V^*(s)$$

The *value* of state  $s$  under the optimal policy  $\pi^*$

$$Q^*(s, a)$$

The *Q-value* of state  $s$   
and action  $a$  under the optimal policy  $\pi^*$

# Value Iteration

$$V^*(s) = \max_a Q^*(s, a)$$

$$V^*(s)$$

The *value* of state  $s$  under the optimal policy  $\pi^*$

$$Q^*(s, a)$$

The *Q-value* of state  $s$   
and action  $a$  under the optimal policy  $\pi^*$

# Value Iteration

$$V^*(s) = \max_a Q^*(s, a) = Q^*(s, \pi^*(s))$$

$$V^*(s)$$

The *value* of state  $s$  under the optimal policy  $\pi^*$

$$Q^*(s, a)$$

The *Q-value* of state  $s$   
and action  $a$  under the optimal policy  $\pi^*$

# Value Iteration

$$Q^*(s, a) = ??$$

$$V^*(s)$$

The *value* of state  $s$  under the optimal policy  $\pi^*$

$$Q^*(s, a)$$

The *Q-value* of state  $s$   
and action  $a$  under the optimal policy  $\pi^*$

# Value Iteration

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

probability of next state	immediate reward	long-term reward
------------------------------	---------------------	---------------------

$$V^*(s)$$

The *value* of state  $s$  under the optimal policy  $\pi^*$

$$Q^*(s, a)$$

The *Q-value* of state  $s$   
and action  $a$  under the optimal policy  $\pi^*$

# Value Iteration

$$V^*(s) = \max_a Q^*(s, a)$$


$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s)$$

The *value* of state  $s$  under the optimal policy  $\pi^*$

$$Q^*(s, a)$$

The *Q-value* of state  $s$   
and action  $a$  under the optimal policy  $\pi^*$

# Value Iteration

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s)$$

The *value* of state  $s$  under the optimal policy  $\pi^*$

$$Q^*(s, a)$$

The *Q-value* of state  $s$   
and action  $a$  under the optimal policy  $\pi^*$

# Value Iteration

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$



# Value Iteration

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

# Value Iteration

1. Start with  $V_0^*(s) = 0$ , for all  $s \in S$
2. Given  $V_i^*$ , calculate the values for all states  $s \in S$

$$V_{\text{💬}}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_{\text{💬}}^*(s')]$$

3. Repeat the above until convergence

# Value Iteration

1. Start with  $V_0^*(s) = 0$ , for all  $s \in S$
2. Given  $V_i^*$ , calculate the values for all states  $s \in S$   
$$V^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$
3. Repeat the above until convergence

There is a guarantee that this process will converge

# Question

How do we recover the optimal policy based on the values?

# Value Iteration

$$V^*(s) = \max_a Q^*(s, a) = Q^*(s, \pi^*(s))$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

# Value Iteration

$$V^*(s) = \max_a Q^*(s, a) = Q^*(s, \pi^*(s))$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

# Value Iteration

$$V^*(s) = \max_a Q^*(s, a) = Q^*(s, \pi^*(s))$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

# Value Iteration

Step 1  
Calculate the Q values

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Step 2  
Find the optimal action for each state

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$




# Value Iteration

$$R(s) = -0.01$$



<i>a</i>	-0.01	-0.01	-0.01	<b>+1.0</b>
<i>b</i>	-0.01		-0.01	<b>-1.0</b>
<i>c</i>	-0.01	-0.01	-0.01	-0.01
	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>




			<b>+1.0</b>
			<b>-1.0</b>

# Value Iteration

## Learned Policy



-0.01	-0.01	-0.01	<b>+1.0</b>
-0.01		-0.01	<b>-1.0</b>
-0.01	-0.01	-0.01	-0.01




→	→	→	<b>+1.0</b>
↑		←	<b>-1.0</b>
↑	←	←	↓


# Value Iteration

$$R(s) = -2.0$$



-2.0	-2.0	-2.0	<b>+1.0</b>
-2.0		-2.0	<b>-1.0</b>
-2.0	-2.0	-2.0	-2.0




			<b>+1.0</b>
			<b>-1.0</b>

# Value Iteration

## Learned Policy



-2.0	-2.0	-2.0	<b>+1.0</b>
-2.0		-2.0	<b>-1.0</b>
-2.0	-2.0	-2.0	-2.0



→	→	→	<b>+1.0</b>
↑		→	<b>-1.0</b>
→	→	→	↑

# Value Iteration

Step 1

Since the procedure relies on  $Q$ -values,  
is it possible to design an algorithm  
that directly computes these  $Q$ -values?

Step 2

Find the optimal action for each state

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

# Value Iteration

Step 1

Since the procedure relies on  $Q$ -values,  
is it possible to design an algorithm  
that directly computes these  $Q$ -values?

We will discuss  $Q$ -value  
iteration next time!

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$