


The slide features a light blue background with a pattern of thin, concentric circles and dashed lines. A large orange speech bubble is centered on the slide, containing the text "Tensorflow Overview".

Tensorflow Overview

Contents

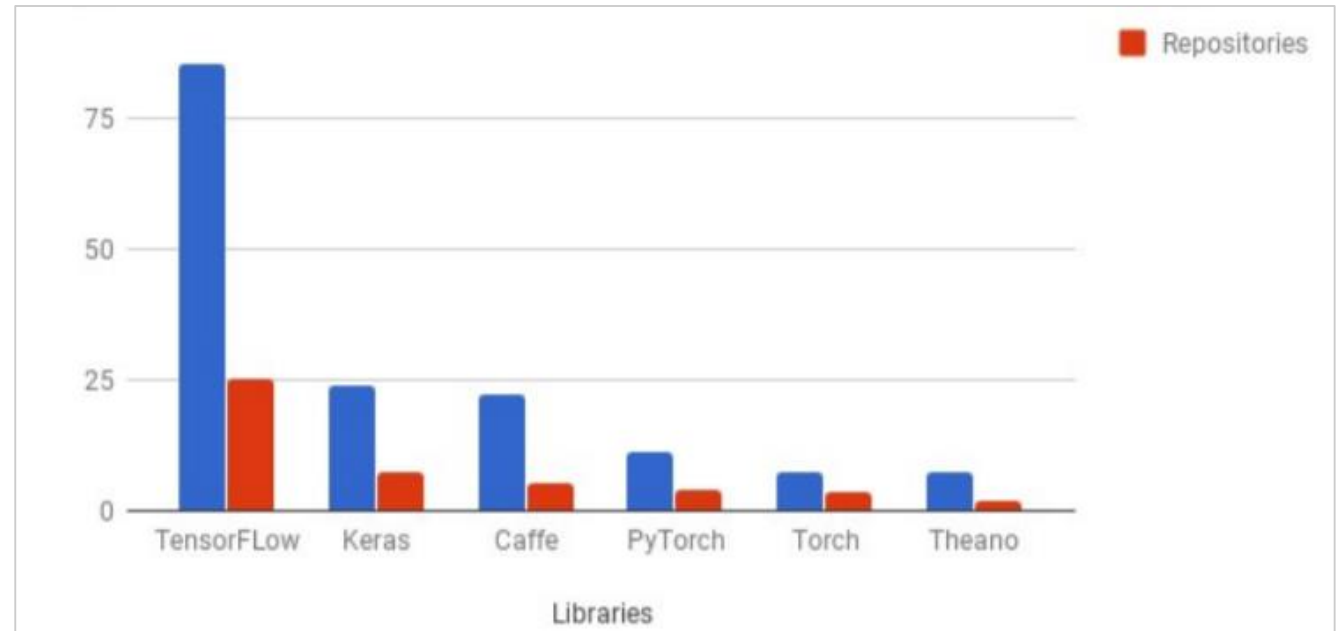
1. Introduction
2. Data Flow Graphs
3. Some tensorflow stuff
 - Tensorflow Variables
 - Tensorflow Placeholders
 - Simple Linear Regression Example
4. Tensorboard
5. Resources

What is a tensor?



- Formally, tensors are multilinear maps from vector spaces to real numbers
- A tensor can be represented as a multidimensional array of numbers. (Matrix, vector or scalar)

What is tensorflow?



- Open source software library for numerical computation using **data flow graphs**
- Other libraries eg Keras, PyTorch, Caffe etc

Tensorflow VS Numpy

- Both are N-d array libraries
- Numpy has Nddarray support, but doesn't offer methods to create tensor functions and automatically compute derivatives (+ no GPU support).

np vs tf

NUMPY

```
In [23]: import numpy as np

In [24]: a = np.zeros((2,2)); b = np.ones((2,2))

In [25]: np.sum(b, axis=1)
Out[25]: array([ 2.,  2.])

In [26]: a.shape
Out[26]: (2, 2)

In [27]: np.reshape(a, (1,4))
Out[27]: array([[ 0.,  0.,  0.,  0.]])
```

TENSORFLOW

```
In [31]: import tensorflow as tf

In [32]: tf.InteractiveSession()

In [33]: a = tf.zeros((2,2)); b = tf.ones((2,2))

In [34]: tf.reduce_sum(b, reduction_indices=1).eval()
Out[34]: array([ 2.,  2.], dtype=float32)

In [35]: a.get_shape()
Out[35]: TensorShape([Dimension(2), Dimension(2)])

In [36]: tf.reshape(a, (1, 4)).eval()
Out[36]: array([[ 0.,  0.,  0.,  0.]], dtype=float32)
```

Numpy	TensorFlow
<code>a = np.zeros((2,2)); b = np.ones((2,2))</code>	<code>a = tf.zeros((2,2)), b = tf.ones((2,2))</code>
<code>np.sum(b, axis=1)</code>	<code>tf.reduce_sum(a, reduction_indices=[1])</code>
<code>a.shape</code>	<code>a.get_shape()</code>
<code>np.reshape(a, (1,4))</code>	<code>tf.reshape(a, (1,4))</code>
<code>b * 5 + 1</code>	<code>b * 5 + 1</code>
<code>np.dot(a,b)</code>	<code>tf.matmul(a, b)</code>
<code>a[0,0], a[:,0], a[0,:]</code>	<code>a[0,0], a[:,0], a[0,:]</code>


The background features a series of concentric circles in light gray, some solid and some dashed, creating a ripple effect. In the center, there is a large orange speech bubble with a pointed bottom. Inside the bubble, the text "Data Flow Graphs" is written in white.

Data Flow Graphs

Tensorflow Computation / Data flow Graph


- “TensorFlow programs are usually structured into a construction phase, that assembles a graph, and an execution phase that uses a session to execute ops in the graph.” - TensorFlow docs
- All computations add nodes to global default graph

Simple Computation Graph



- `import tensorflow as tf`
- `a = tf.add(3, 5)`
- `print(a)`
- `>>`

Simple Computation Graph



- `import tensorflow as tf`
- `a = tf.add(3, 5)`
- `print(a)`
- `>> Tensor("Add:0", shape=(), dtype=int32) #(Not 8)`

The background features several thin, curved lines in a light gray color, some solid and some dashed, creating a sense of motion or a stylized globe. The main content is centered within an orange speech bubble shape.

HOW TO GET
VALUE OF 'a'
???

HOW TO GET
VALUE OF 'a'
???


- Create a **session**, assign it to variable `sess` so we can call it later
- Within the session, evaluate the graph to fetch the value of `a`
- `import tensorflow as tf`
`a = tf.add(3, 5)`
`with tf.Session() as sess:`
 `print(sess.run(a))`

The background of the slide features a series of thin, curved lines in a light gray color, creating a sense of motion or a stylized globe. These lines are more prominent on the left side and fade towards the right.

tf.Session()

- A Session object encapsulates the environment in which Operation objects are executed, and Tensor objects are evaluated.
- Session will also allocate memory to store the current values of variables.

Another Example



- `x = 2`
`y = 3`
`op1 = tf.add(x, y)`
`op2 = tf.multiply(x, y)`
`op3 = tf.pow(op2, op1)`
with `tf.Session()` as `sess`:
`op3 = sess.run(op3) # (x*y) ^ (x+y)`

Why graphs?

- Save computation, run subgraphs that lead to the values you want to fetch only
- Break computation into small, differential pieces to facilitate auto-differentiation
- Facilitate distributed computation, spread the work across multiple CPUs, GPUs, TPUs, or other devices

The background features a series of concentric circles in light gray, some solid and some dashed, creating a ripple effect. In the center, there is an orange speech bubble with a pointed bottom, containing the text 'tf.constants & tf.Variable' in white.

tf.constants & tf.Variable

tf.constants

```
In [95]: a = tf.constant(value=6.0, name='scalar', dtype=tf.float32, shape=[])  
print(a)
```

```
Tensor("scalar:0", shape=(), dtype=float32)
```

```
In [96]: a = tf.constant(value=6.0, name='array', dtype=tf.float32, shape=[2])  
print(a)
```

```
Tensor("array:0", shape=(2,), dtype=float32)
```

```
In [97]: a = tf.constant(value=6.0, name='matrix', dtype=tf.float32, shape=[2,2])  
print(a)
```

```
Tensor("matrix:0", shape=(2, 2), dtype=float32)
```

tf.constants

Useful constants

```
In [101]: a = tf.zeros(shape=[3,2], dtype=tf.float32, name='matrix')
          print(a)
          with tf.Session() as sess:
              a_val = sess.run(a)
              print(a_val)
```

```
Tensor("matrix_4:0", shape=(3, 2), dtype=float32)
[[ 0.  0.]
 [ 0.  0.]
 [ 0.  0.]]
```

```
In [102]: a = tf.ones(shape=[3,2], dtype=tf.float32, name='matrix')
          print(a)
          with tf.Session() as sess:
              a_val = sess.run(a)
              print(a_val)
```

```
Tensor("matrix_5:0", shape=(3, 2), dtype=float32)
[[ 1.  1.]
 [ 1.  1.]
 [ 1.  1.]]
```

tf.constants

Randomly Generated Constants

- `tf.set_random_seed(seed)`

```
tf.random_normal(shape, mean=0.0, stddev=1.0, dtype=tf.float32, seed=None, name=None)  
tf.random_uniform(shape, minval=0, maxval=None, dtype=tf.float32, seed=None, name=None)
```

Operations

Category	Examples
Element-wise mathematical operations	Add, Sub, Mul, Div, Exp, Log, Greater, Less, Equal, ...
Array operations	Concat, Slice, Split, Constant, Rank, Shape, Shuffle, ...
Matrix operations	MatMul, MatrixInverse, MatrixDeterminant, ...
Stateful operations	Variable, Assign, AssignAdd, ...
Neural network building blocks	SoftMax, Sigmoid, ReLU, Convolution2D, MaxPool, ...
Checkpointing operations	Save, Restore
Queue and synchronization operations	Enqueue, Dequeue, MutexAcquire, MutexRelease, ...
Control flow operations	Merge, Switch, Enter, Leave, NextIteration

tf.Variable

```
In [32]: W1 = tf.ones((2,2))

In [33]: W2 = tf.Variable(tf.zeros((2,2)), name="weights")

In [34]: with tf.Session() as sess:
          print(sess.run(W1))
          sess.run(tf.initialize_all_variables())
          print(sess.run(W2))

.....:
[[ 1.  1.]
 [ 1.  1.]]
[[ 0.  0.]
 [ 0.  0.]]
```

Note the initialization step `tf.initialize_all_variables()`

- “When you train a model you use variables to **hold and update parameters**. Variables are in-memory buffers containing tensors” - TensorFlow Docs
- TensorFlow variables must be initialized before they have values! Contrast with constant tensors.

Updating Variable State

```
In [63]: state = tf.Variable(0, name="counter")
```

```
In [64]: new_value = tf.add(state, tf.constant(1)) ← Roughly new_value = state + 1
```

```
In [65]: update = tf.assign(state, new_value) ← Roughly state = new_value
```

```
In [66]: with tf.Session() as sess:
....:     sess.run(tf.initialize_all_variables())
....:     print(sess.run(state))
....:     for _ in range(3):
....:         sess.run(update)
....:         print(sess.run(state))
....:
```

← Roughly
state = 0
print(state)
for _ in range(3):
state = state + 1
print(state)

```
0
1
2
3
```

The background features a series of concentric circles in light gray, some solid and some dashed, creating a ripple effect. In the center, there is a large orange speech bubble with a pointed bottom. Inside the bubble, the text 'Tensorflow Placeholders' is written in white.

Tensorflow Placeholders

Inputting Data

- Ops are functions on tensors
- `tf.constant` -> tensors with constant values
- `tf.Variable` -> tensors params which can be updated
- `tf.placeholder` -> 'dummy' tensor that holds input data.

$$\hat{y} = Wx + b$$

variable ops placeholder variable

Inputting Data

- Ops are functions on tensors
- `tf.constant` -> tensors with constant values (fixed | not trainable)
- `tf.Variable` -> tensors params which can be updated (trainable)
- `tf.placeholder` -> 'dummy' tensor that holds input data. Will need to pass actual data in `tf.Session()` using a `feed_dict`


```
# create a placeholder for a vector of 2 elements, type tf.float32
x = tf.placeholder(dtype=tf.float32, shape=[2], name='p')

y = tf.constant(value=[1, 2], dtype=tf.float32, name='c')
```

feed_dict

```
# create a placeholder for a vector of 2 elements, type tf.float32  
x = tf.placeholder(dtype=tf.float32, shape=[2], name='p')  
  
y = tf.constant(value=[1, 2], dtype=tf.float32, name='c')  
  
z = x + y  
  
with tf.Session() as sess:  
    z_val = sess.run(z, feed_dict={x: [3, 4]})  
    print(z_val)
```

[4. 6.]



The background of the slide features a series of concentric circles in a light gray color, centered around the middle of the frame. A dashed gray line also curves across the background, intersecting the circles. The overall aesthetic is clean and modern.


Linear Regression Example

1a) Create Input Data

```
import numpy as np
import seaborn
```

```
# Define input data
X_data = np.arange(100, step=.1)
y_data = X_data + 20 * np.sin(X_data/10)

# Plot input data
plt.scatter(X_data, y_data)
```



1b) Define hyper-params

```
# Define data size and batch size
n_samples = 1000
batch_size = 100
```

```
# Tensorflow is finicky about shapes, so resize
X_data = np.reshape(X_data, (n_samples,1))
y_data = np.reshape(y_data, (n_samples,1))
```

```
# Define placeholders for input
X = tf.placeholder(tf.float32, shape=(batch_size, 1))
y = tf.placeholder(tf.float32, shape=(batch_size, 1))
```

2) Define placeholders

```
# Define data size and batch size
```

```
n_samples = 1000
```

```
batch_size = 100
```

```
# Tensorflow is finicky about shapes, so resize
```

```
X_data = np.reshape(X_data, (n_samples,1))
```

```
y_data = np.reshape(y_data, (n_samples,1))
```

```
# Define placeholders for input
```

```
X = tf.placeholder(tf.float32, shape=(batch_size, 1))
```

```
y = tf.placeholder(tf.float32, shape=(batch_size, 1))
```

3) Define Variables

```
# Define variables to be learned
with tf.variable_scope("linear-regression"):
    W = tf.get_variable("weights", (1, 1),
                        initializer=tf.random_normal_initializer())
    b = tf.get_variable("bias", (1,),
                        initializer=tf.constant_initializer(0.0))
    y_pred = tf.matmul(X, W) + b
    loss = tf.reduce_sum((y - y_pred)**2/n_samples)
```

*get_variable
same as creating
a variable*

4) Define Loss Function

```
# Define variables to be learned
with tf.variable_scope("linear-regression"):
    W = tf.get_variable("weights", (1, 1),
                        initializer=tf.random_normal_initializer())
    b = tf.get_variable("bias", (1,),
                        initializer=tf.constant_initializer(0.0))
    y_pred = tf.matmul(X, W) + b
    loss = tf.reduce_sum((y - y_pred)**2/n_samples)
```

$$J(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - (Wx_i + b))^2$$

5) Define optimizer

Sample code to run full gradient descent:

Define optimizer operation

```
opt_operation = tf.train.AdamOptimizer().minimize(loss)
```

with tf.Session() as sess:

Initialize Variables in graph

sess.run(tf.initialize_all_variables())

Gradient descent loop for 500 steps

for _ in range(500):

Select random minibatch

indices = np.random.choice(n_samples, batch_size)

X_batch, y_batch = X_data[indices], y_data[indices]

Do gradient descent step

_, loss_val = sess.run([opt_operation, loss], feed_dict={X: X_batch, y: y_batch})

*Let's do a deeper.
graphical dive into
this operation*

6) Define training procedure in `tf.Session()`

```
# Sample code to run full gradient descent:  
# Define optimizer operation  
opt_operation = tf.train.AdamOptimizer().minimize(loss)
```

```
with tf.Session() as sess:  
    # Initialize Variables in graph  
    sess.run(tf.initialize_all_variables())  
    # Gradient descent loop for 500 steps  
    for _ in range(500):  
        # Select random minibatch  
        indices = np.random.choice(n_samples, batch_size)  
        X_batch, y_batch = X_data[indices], y_data[indices]  
        # Do gradient descent step  
        _, loss_val = sess.run([opt_operation, loss], feed_dict={X: X_batch, y: y_batch})
```

*Let's do a deeper.
graphical dive into
this operation*

6a) Initialize variables

```
# Sample code to run full gradient descent:
# Define optimizer operation
opt_operation = tf.train.AdamOptimizer().minimize(loss)

with tf.Session() as sess:
    # Initialize Variables in graph
    sess.run(tf.initialize_all_variables())
    # Gradient descent loop for 500 steps
    for _ in range(500):
        # Select random minibatch
        indices = np.random.choice(n_samples, batch_size)
        X_batch, y_batch = X_data[indices], y_data[indices]
        # Do gradient descent step
        _, loss_val = sess.run([opt_operation, loss], feed_dict={X: X_batch, y: y_batch})
```

*Let's do a deeper.
graphical dive into
this operation*

6b) Grad Descent Iteration loop

```
# Sample code to run full gradient descent:
# Define optimizer operation
opt_operation = tf.train.AdamOptimizer().minimize(loss)

with tf.Session() as sess:
    # Initialize Variables in graph
    sess.run(tf.initialize_all_variables())
    # Gradient descent loop for 500 steps
    for _ in range(500):
        # Select random minibatch
        indices = np.random.choice(n_samples, batch_size)
        X_batch, y_batch = X_data[indices], y_data[indices]
        # Do gradient descent step
        _, loss_val = sess.run([opt_operation, loss], feed_dict={X: X_batch, y: y_batch})
```

*Let's do a deeper.
graphical dive into
this operation*

6c) Inputs to placeholders

```
# Sample code to run full gradient descent:
# Define optimizer operation
opt_operation = tf.train.AdamOptimizer().minimize(loss)

with tf.Session() as sess:
    # Initialize Variables in graph
    sess.run(tf.initialize_all_variables())
    # Gradient descent loop for 500 steps
    for _ in range(500):
        # Select random minibatch
        indices = np.random.choice(n_samples, batch_size)
        X_batch, y_batch = X_data[indices], y_data[indices]
        # Do gradient descent step
        _, loss_val = sess.run([opt_operation, loss], feed_dict={X: X_batch, y: y_batch})
```

Let's do a deeper graphical dive into this operation

6d) Grad Descent step

```
# Sample code to run full gradient descent:  
# Define optimizer operation  
opt_operation = tf.train.AdamOptimizer().minimize(loss)
```

```
with tf.Session() as sess:  
    # Initialize Variables in graph  
    sess.run(tf.initialize_all_variables())  
    # Gradient descent loop for 500 steps  
    for _ in range(500):  
        # Select random minibatch  
        indices = np.random.choice(n_samples, batch_size)  
        X_batch, y_batch = X_data[indices], y_data[indices]  
        # Do gradient descent step  
        _, loss_val = sess.run([opt_operation, loss], feed_dict={X: X_batch, y: y_batch})
```

*Let's do a deeper.
graphical dive into
this operation*

```
feed_dict={X:X_batch, y:y_batch}
```

```
X = tf.placeholder(tf.float32, shape=(batch_size, 1))
```

```
y = tf.placeholder(tf.float32, shape=(batch_size, 1))
```

```
W = tf.get_variable(  
    "weights", (1, 1),  
    initializer=tf.random_normal_initializer())
```


```
b = tf.get_variable(  
    "bias", (1,),  
    initializer=tf.constant_initializer(0.0))
```

```
y_pred = tf.matmul(X, W) + b
```

```
loss = tf.reduce_sum((y - y_pred)**2/n_samples)
```

```
opt_operation = tf.train.AdamOptimizer().minimize(loss)
```

$$J(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - (Wx_i + b))^2$$



← *Learned model offers nice fit to data.*

Tensorboard

```
x = tf.get_variable(name='x', initializer=tf.random_normal([2]))
y = tf.get_variable(name='y', initializer=tf.random_normal([2]))
z = x+y

writer = tf.summary.FileWriter(logdir='./graphs', graph=tf.get_default_graph())

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    print(sess.run(z))
writer.close()
```

[-2.20665431 0.04801518]

- **In terminal:**
\$ tensorboard --logdir="./graphs" --port 6006
- **In web browser:**
<http://localhost:6006/>




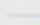

 Fit to screen Download PNGRun
(1)Session
runs (0)Upload ☐ Trace inputsColor ☒ Structure☐ Device☐ XLA Cluster☐ Compute time☐ Memory☐ TPU Compatibility

colors same substructure

☐ unique substructure

Close legend

Graph (* = expandable)

 Namespace* ? OpNode ? Unconnected series* ? Connected series* ? Constant ? Summary ? Dataflow edge ? Control dependency edge ? Reference edge ?

Resources

- <https://www.tensorflow.org/guide>
- <http://web.stanford.edu/class/cs20si/syllabus.html>
- Youtube - <https://www.youtube.com/user/hvasslabs>
- Géron, Aurélien. Hands-on machine learning with Scikit-Learn and TensorFlow: concepts, tools, and techniques to build intelligent systems. " O'Reilly Media, Inc.", 2017.