

Image, filtering, convolution

ISTD 50.035
Computer Vision

Acknowledgement: Some images are from various sources: UCF, Stanford cs231n, etc.

Image is an array of numbers

- Grayscale image

- 2D array of numbers
(pixels) / matrix

- Number indicates the intensity: [0,255] for 8-bit representation

- Image resolution /
number of pixels in an
image: 100x100,
1920x1080, etc.

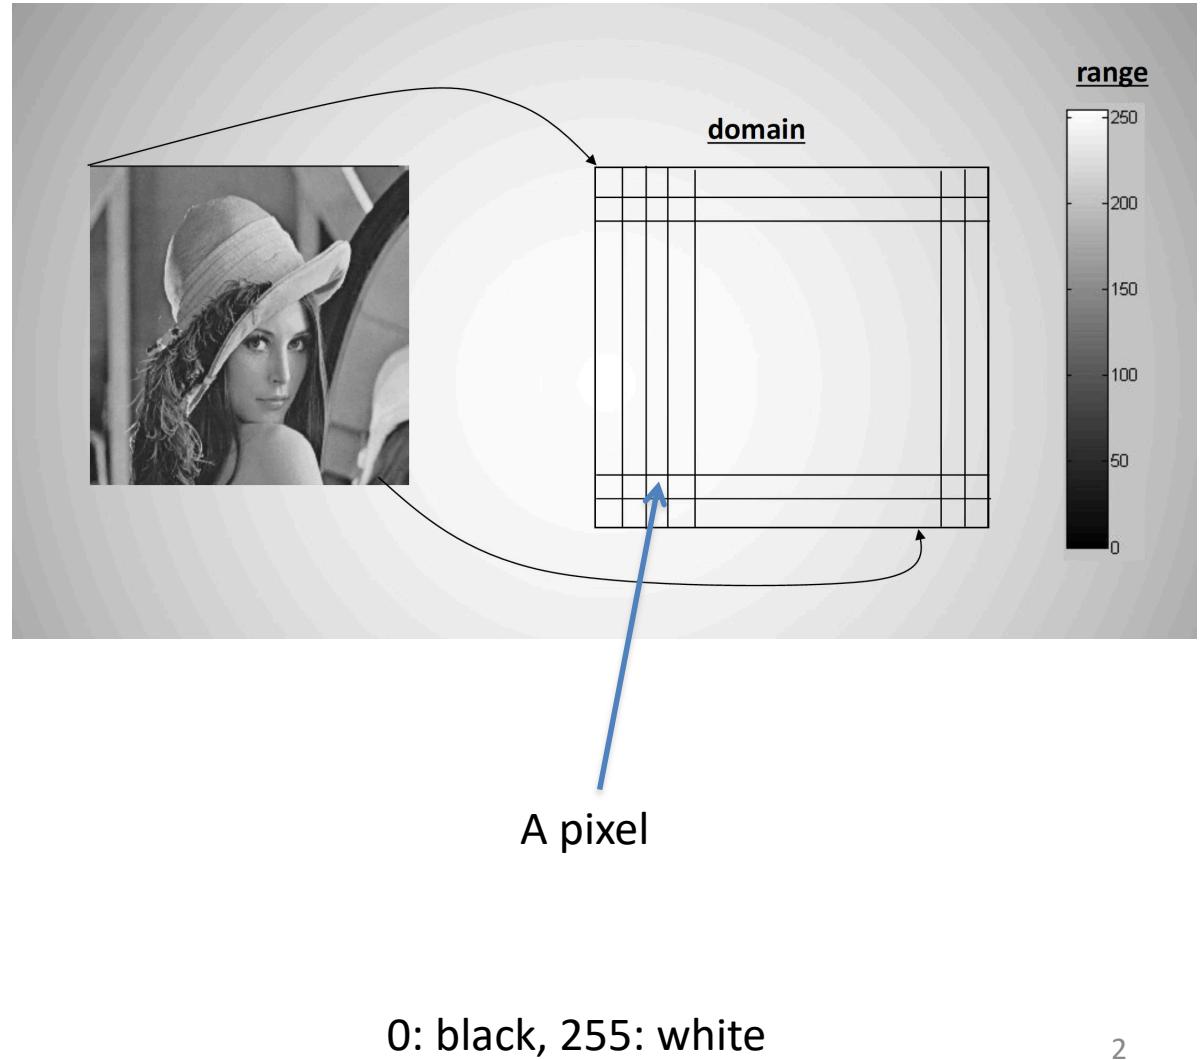
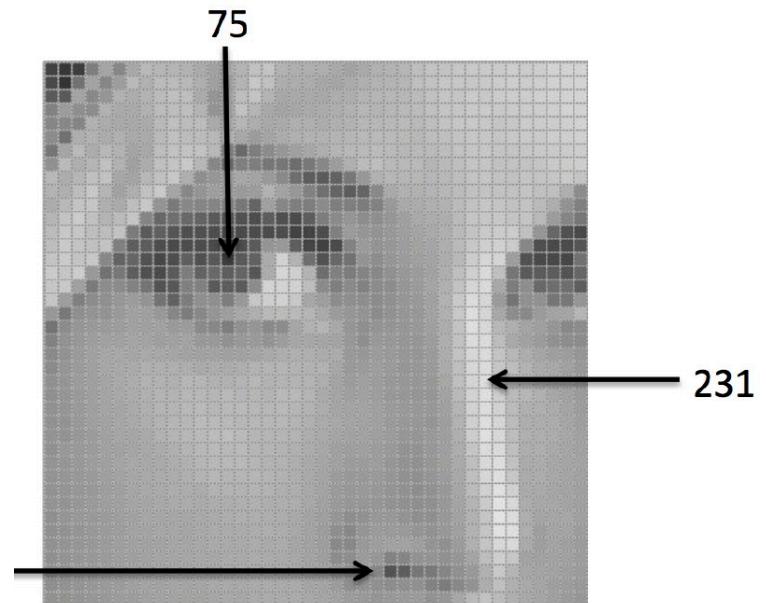


Image is a discrete array of numbers

- Samples from continuous object
- quantized to have a finite number of possible values: [0,1,2 to 255]
- Sometimes normalized from 0 to 1



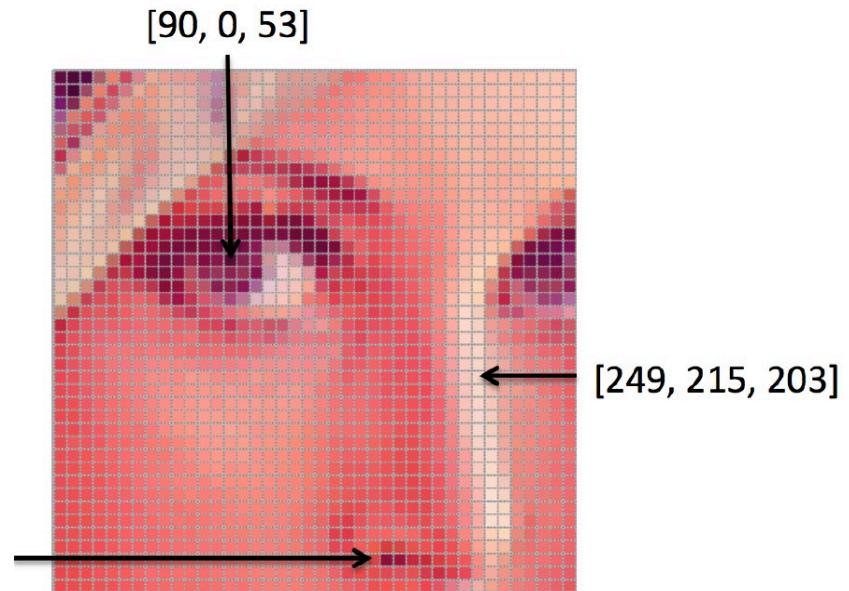
Color image

- Each pixel has three numbers to represent the red, green, blue color intensity

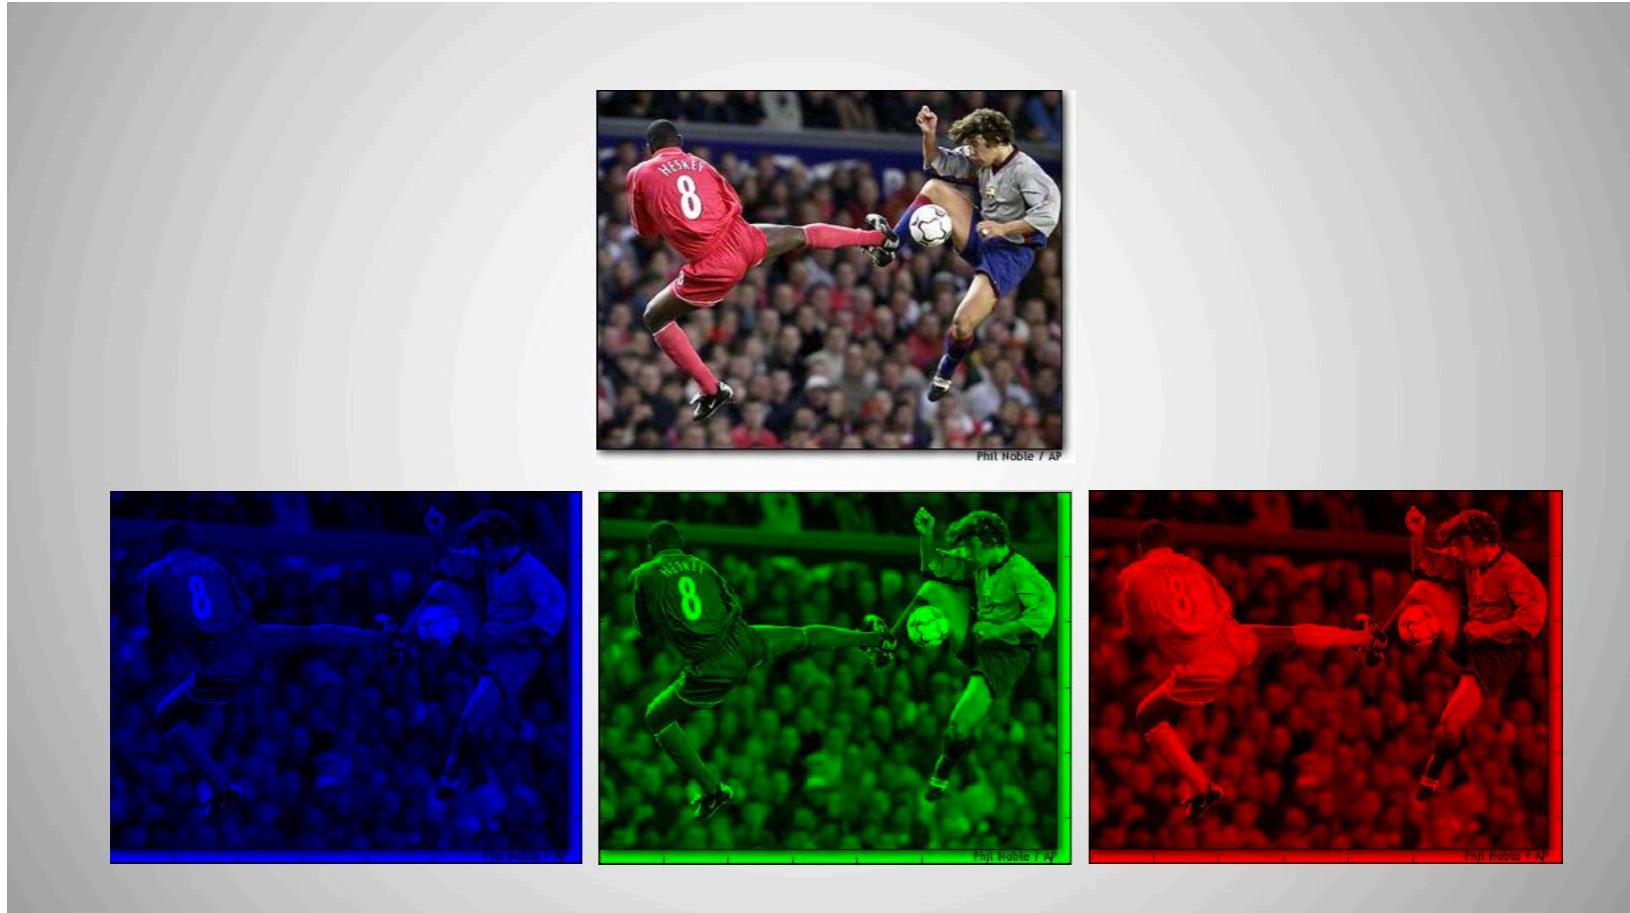
- RGB

- Other color format:
HSV, YUV

- Image can be stored in different format:
JPEG, PNG, TIF, etc.

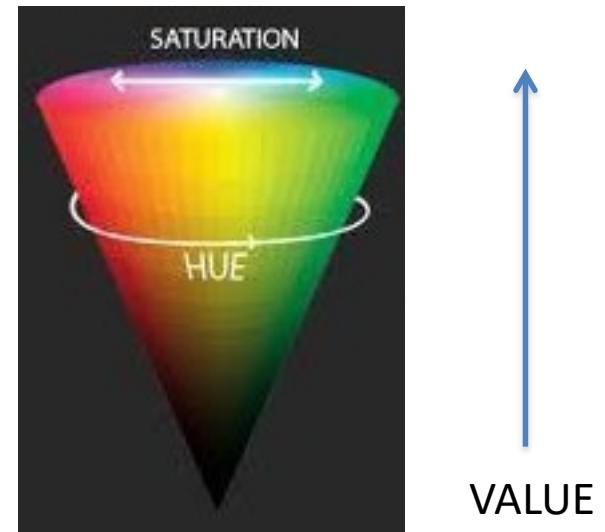


Color image with R, G, B channels



HSV color space

- **Hue** represents color in angle of HSV cone
 - 0-60: red
 - 60-120: yellow
 - 120-180: green
 - 180-240: cyan
 - 240-300: blue
 - 300-360: magenta
- **Saturation** represents the amount of ‘grey’ in the color: ‘0’ is grey, ‘1’ is pure primary color
 - Small saturation means faded color
 - It is the radius of the HSV cone
- **Value** represents intensity
 - It is the height of the HSV cone



HSV represents color in a similar way as humans perceive color, is easy to work with in some applications

Video

- A video is a sequence of images (frames)
- 30 frame per second (fps)
- Spatial dimension: x,y
- Temporal dimension: t

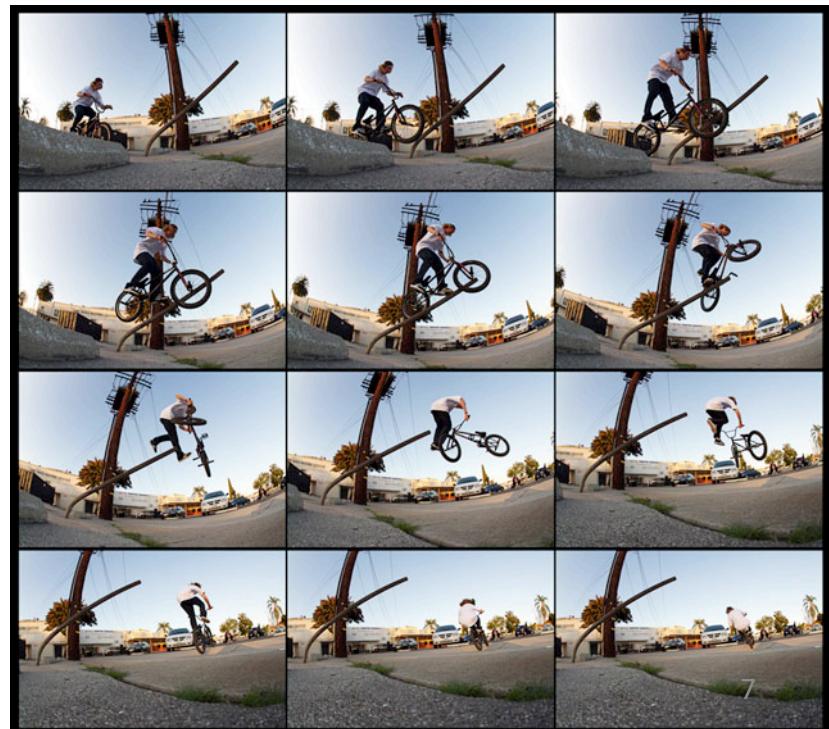


Image Filtering

- An image processing operation
- Remove some unwanted components: noise
- Extract useful information 
- Linear filtering: The output is a linear combination of pixel values in some neighborhood

Filtering is an essential operation in deep neural networks

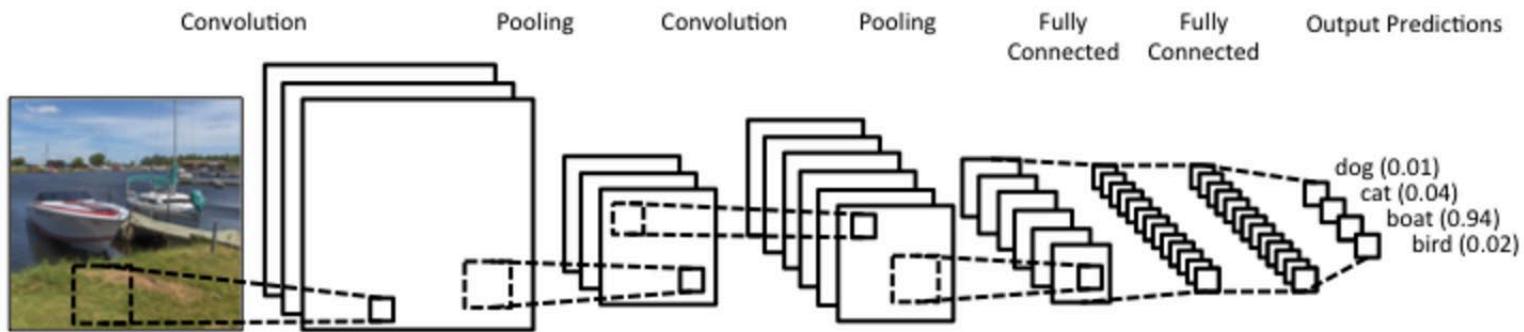


Image Filtering

- Correlation / convolution (precisely, there are subtle differences)

Correlation

f = Image

h = Kernel

Essentially weighted
sum of inputs
(over a small
neighborhood)

f

f_1	f_2	f_3
f_4	f_5	f_6
f_7	f_8	f_9

\otimes

h

h_1	h_2	h_3
h_4	h_5	h_6
h_7	h_8	h_9

$$\begin{aligned} f \otimes h = & f_1 h_1 + f_2 h_2 + f_3 h_3 \\ & + f_4 h_4 + f_5 h_5 + f_6 h_6 \\ & + f_7 h_7 + f_8 h_8 + f_9 h_9 \end{aligned}$$

Image Filtering

- Convolution (Math)
- Flipping is skipped in this course

Convolution

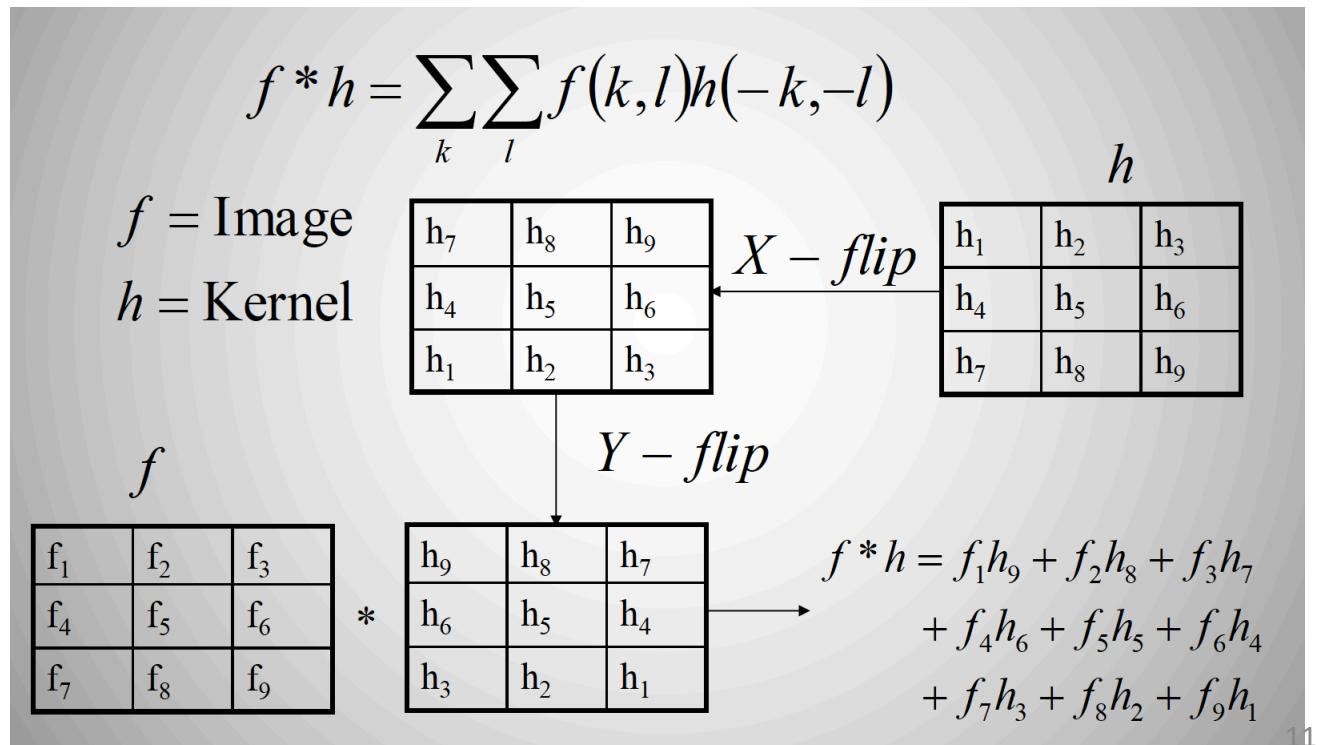


Image Filtering

- Filtering (convolution) operation

Center element of the kernel is placed over the source pixel. The source pixel is then replaced with a weighted sum of itself and nearby pixels.

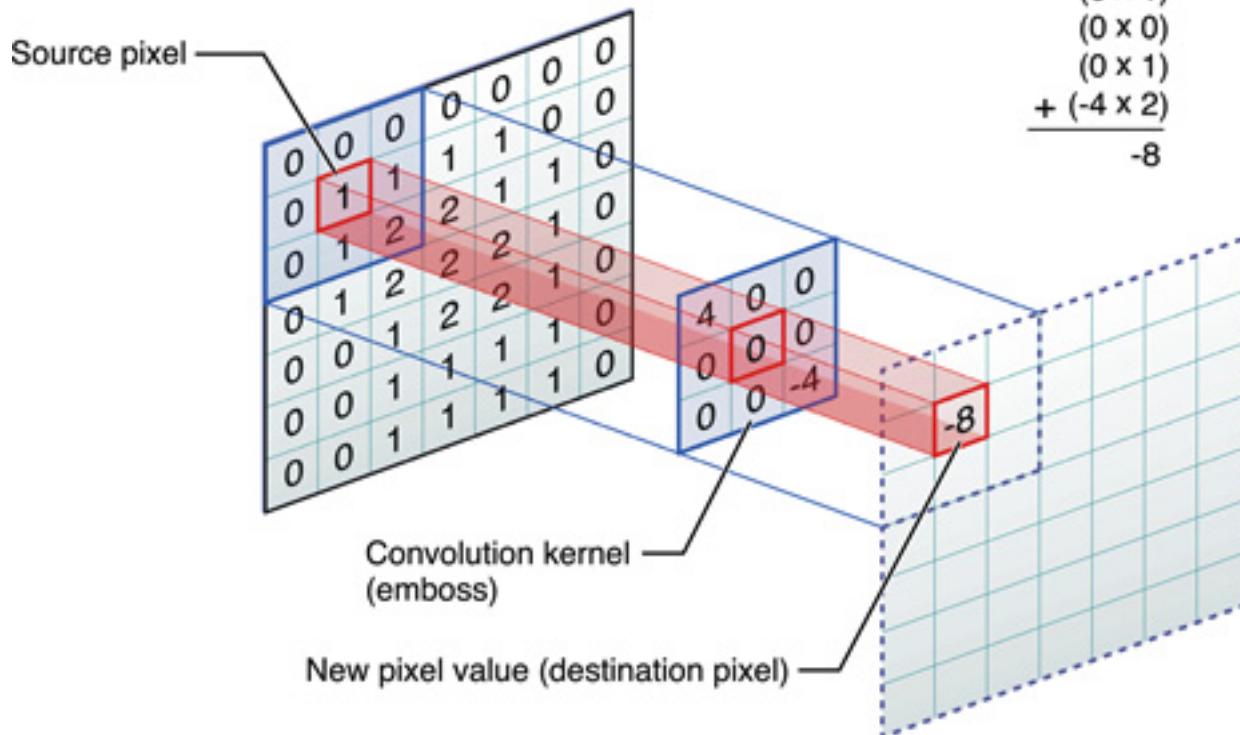


Image Filtering

- Filtering (convolution) operation
- **Slide** the filter kernel over the entire image to produce the output (image/activation)

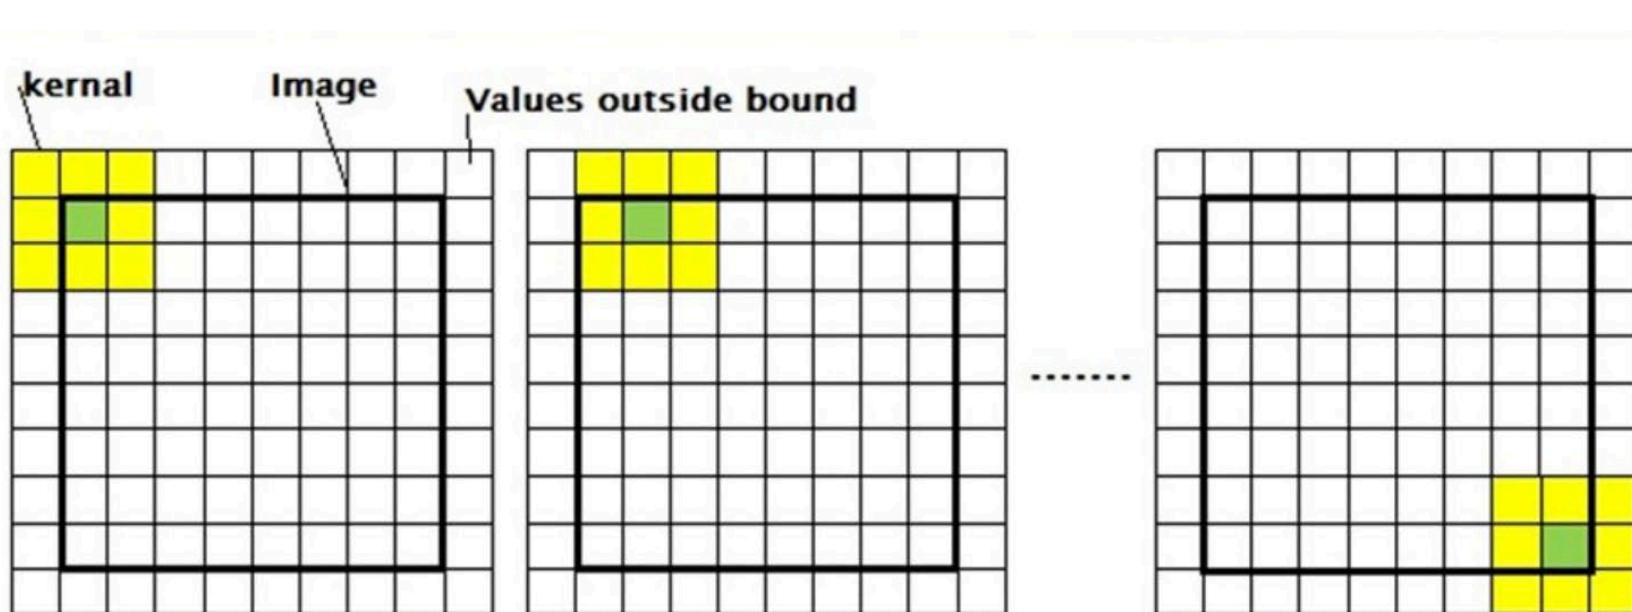


Image Filtering

Image filtering

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

$$g[\cdot, \cdot] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l] \quad k, l \text{ in } [-1,0,1]$$

Credit: S. Seitz

Image Filtering

Image filtering

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0	10										

$$g[\cdot, \cdot] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l] \quad k, l \text{ in } [-1,0,1]$$

Credit: S. Seitz

Image Filtering

Image filtering

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0	10	20								

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l] \quad k, l \text{ in } [-1,0,1]$$

Credit: S. Seitz

Image Filtering

Image filtering

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

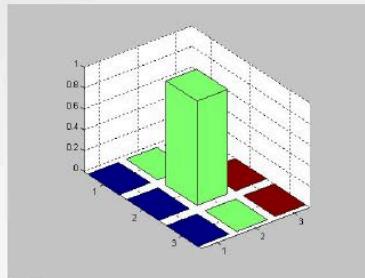
$h[.,.]$

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l] \quad k, l \text{ in } [-1, 0, 1]$$

Credit: S. Seitz

Image Filtering



*

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 1 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

=



Image Filtering

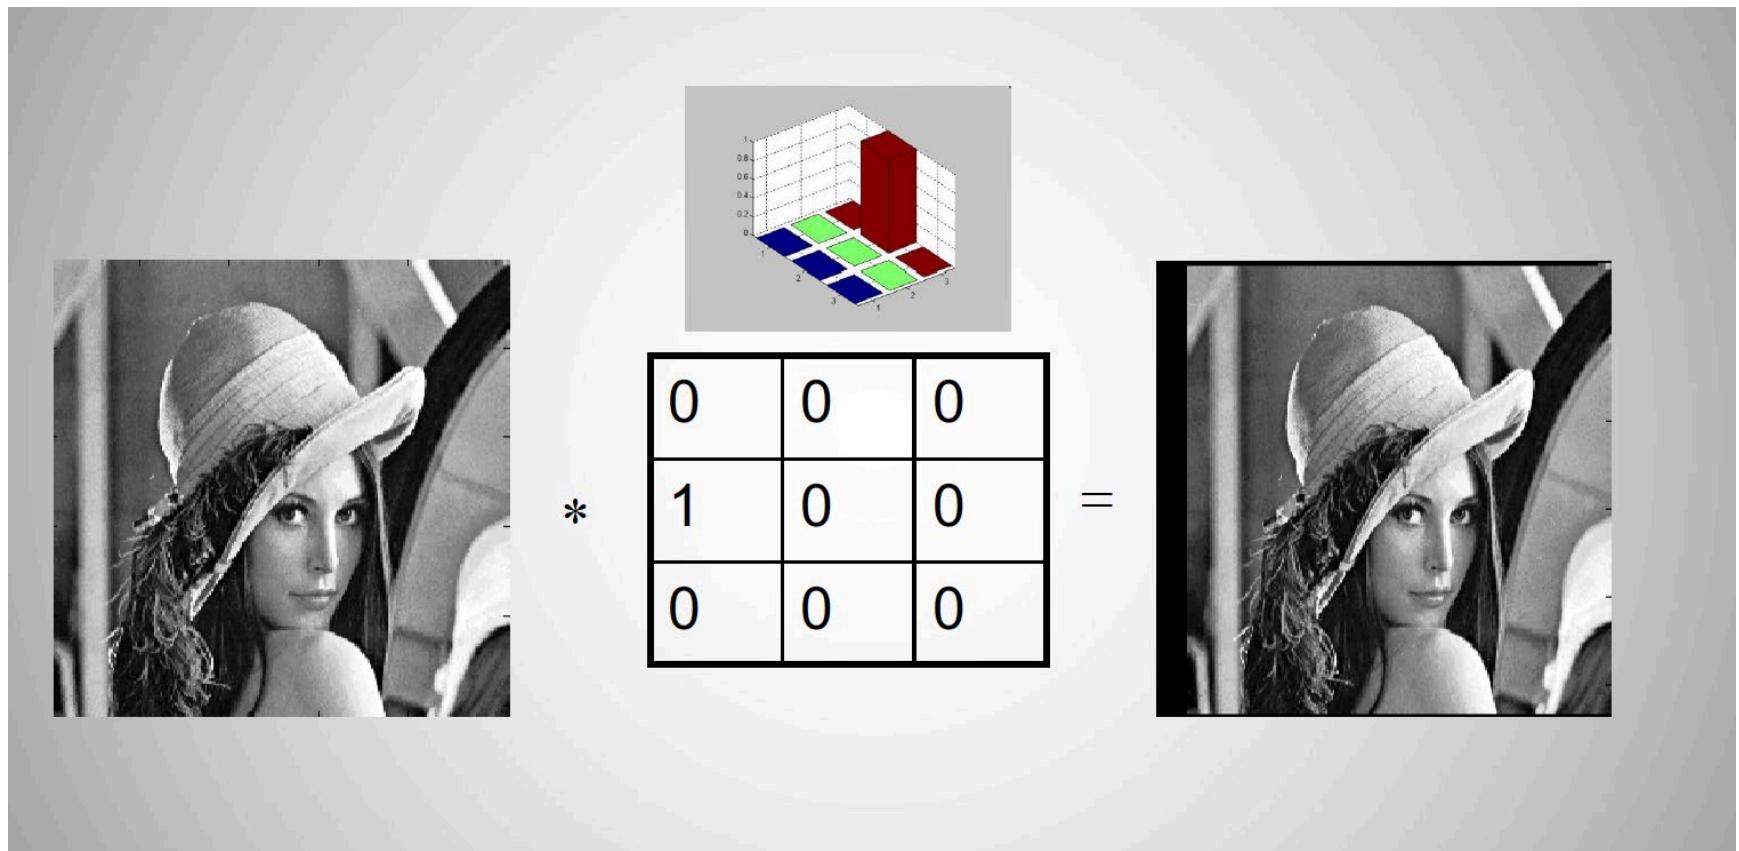


Image Filtering

- Cohort exercise: calculate the output of the following filtering (you can ignore the boundary condition)

Img =
[[30 40 20 30 40]
 [40 20 30 40 30]
 [20 30 40 30 40]
 [30 40 30 40 20]
 [40 30 40 20 30]]

kernel =

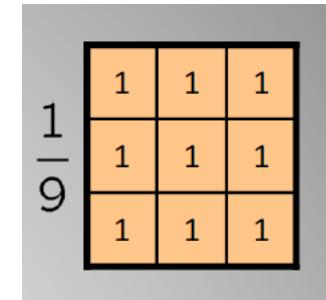
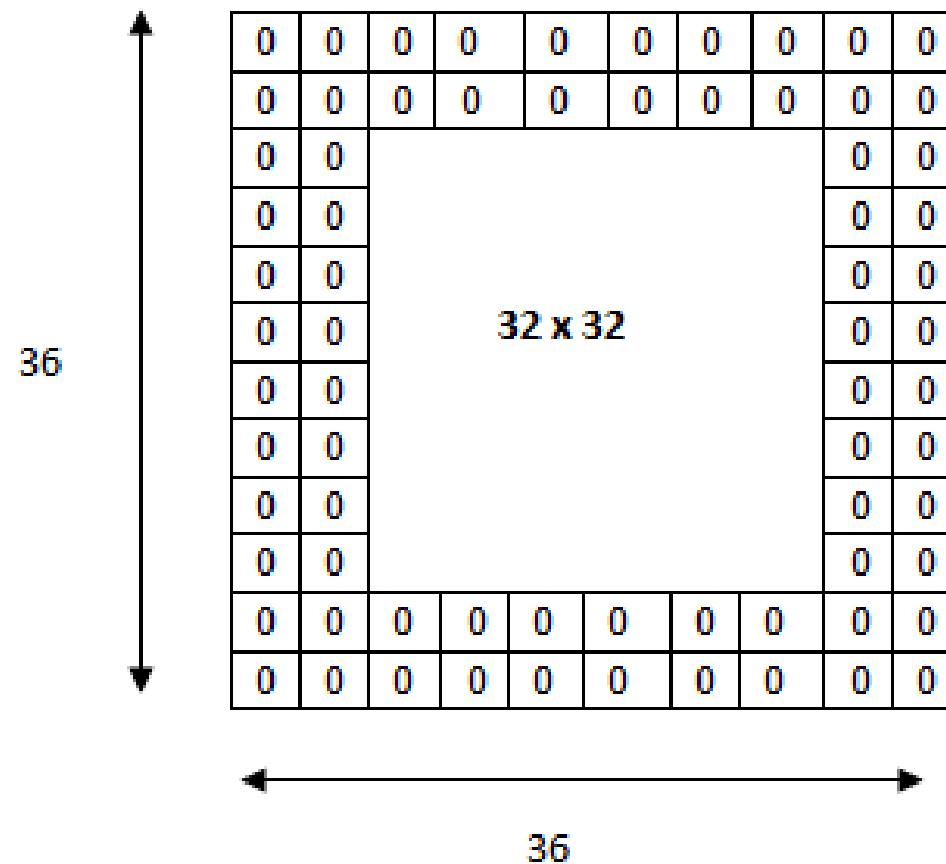


Image Filtering

- Padding to handle boundary condition
- Stride: how many pixels to shift the filter kernel in each step



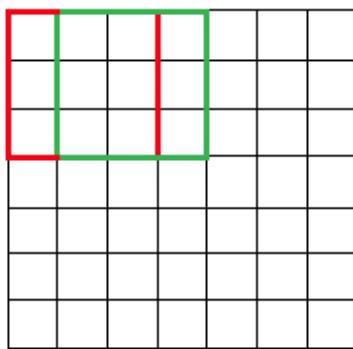
Padding



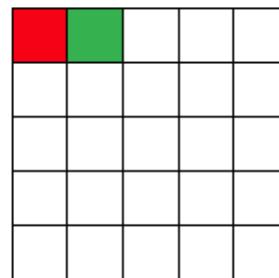
What would be the size of the kernel for an output of equal size (32x32)?

Stride

7 x 7 Input Volume

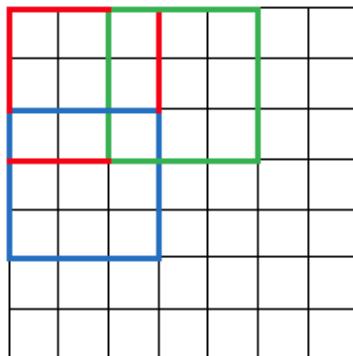


5 x 5 Output Volume

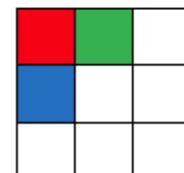


Stride=1

7 x 7 Input Volume



3 x 3 Output Volume



Stride=2

Low-pass filtering

- LPF retains low spatial frequency components,
remove high spatial frequency components
(noise, texture)

Image Frequency

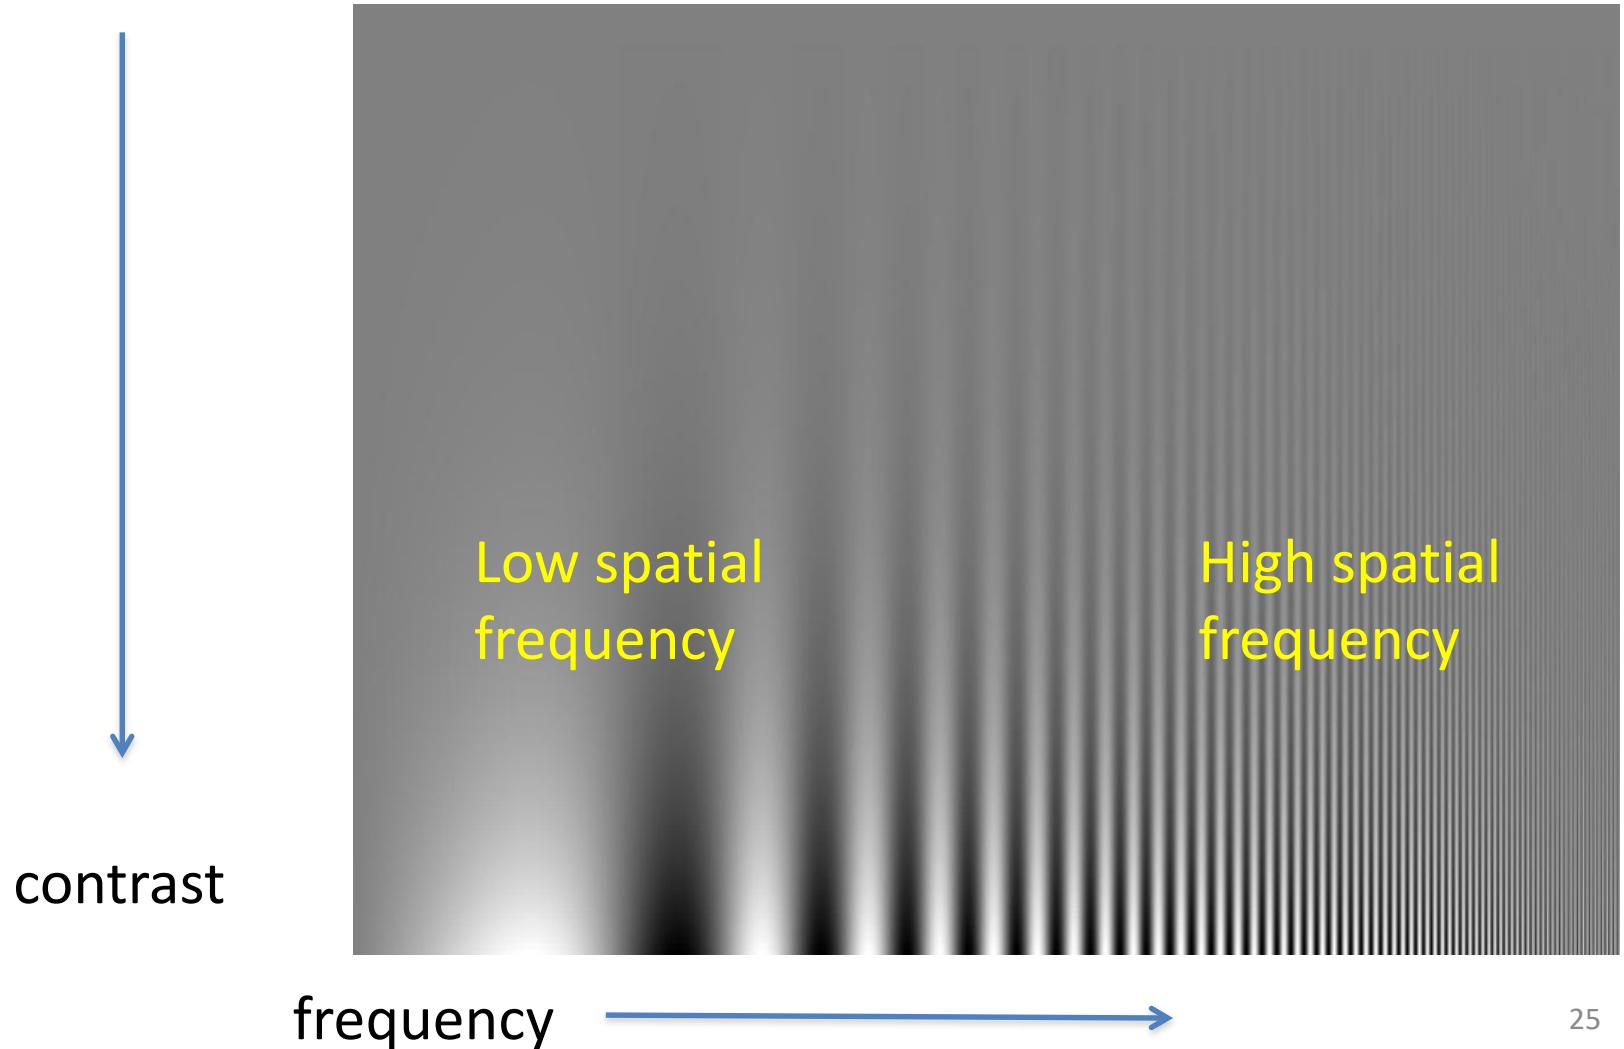


Image Frequency

- Pixel intensity at two rows

Which row has higher spatial frequency?

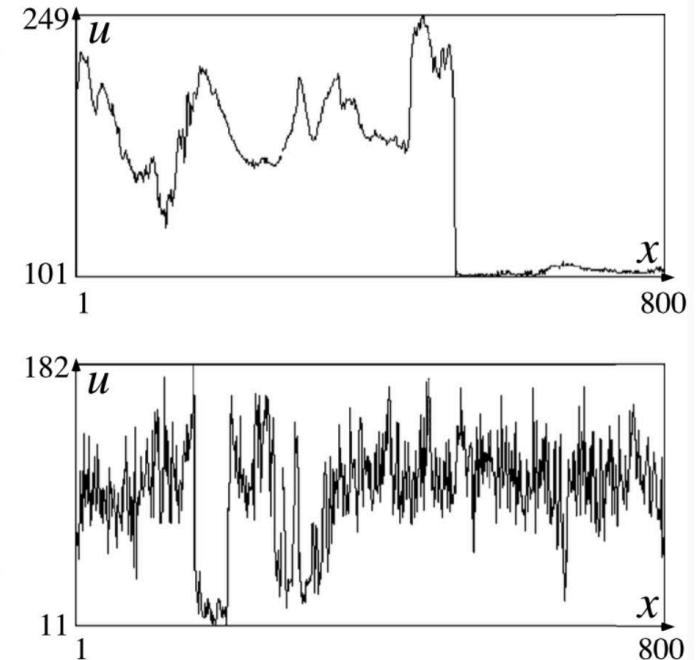
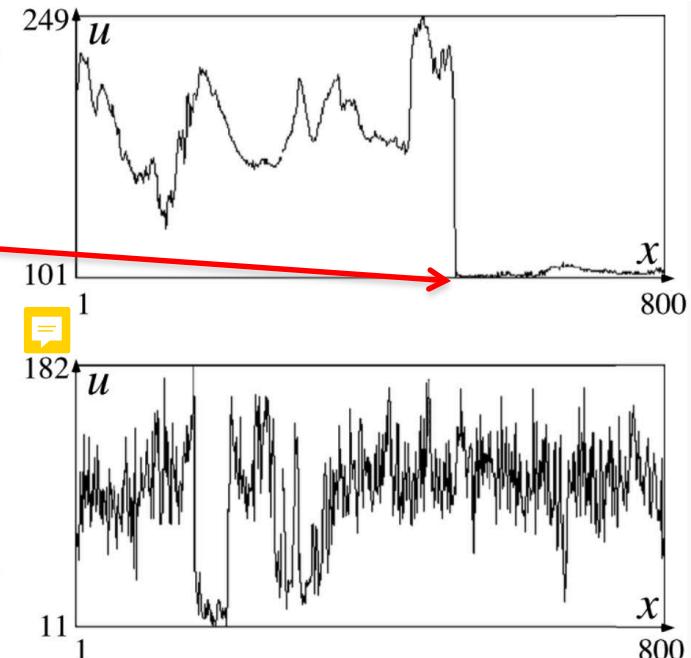


Image Frequency

- Pixel intensity at two rows



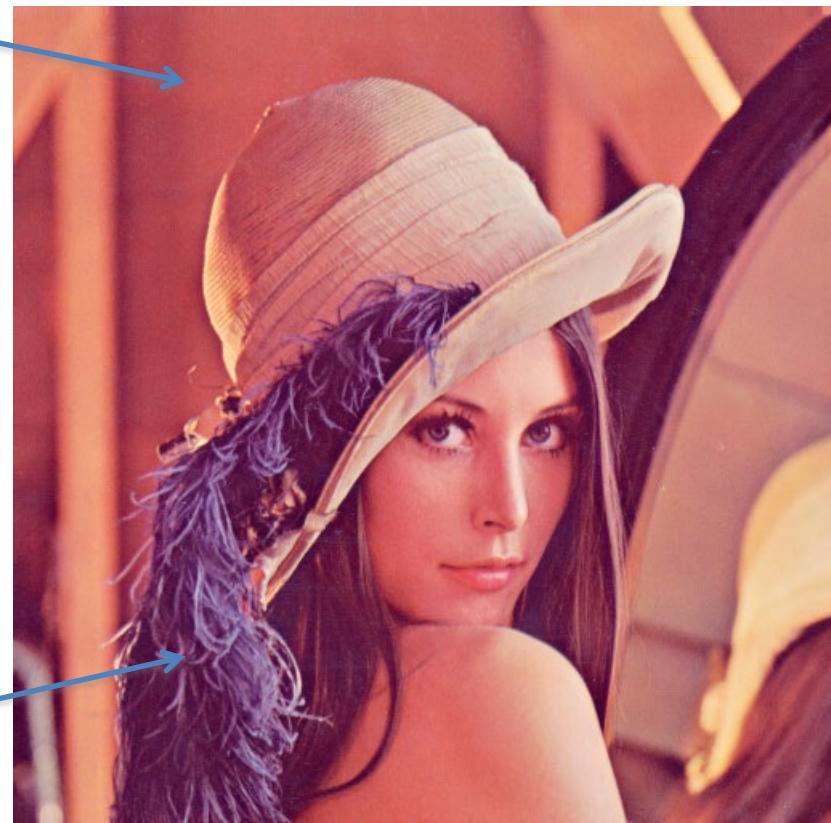
Low spatial frequency



High spatial frequency

Image Frequency

Low spatial
frequency



High spatial
frequency

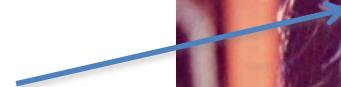




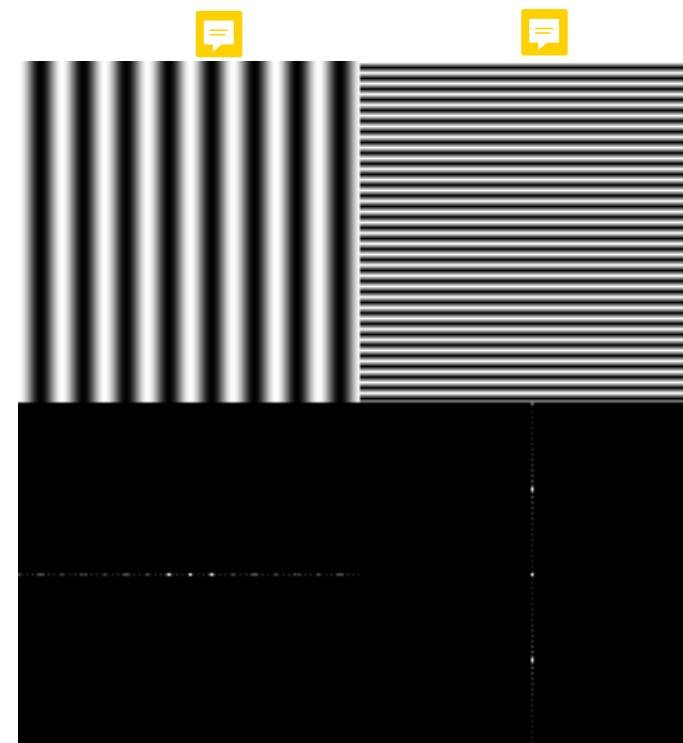
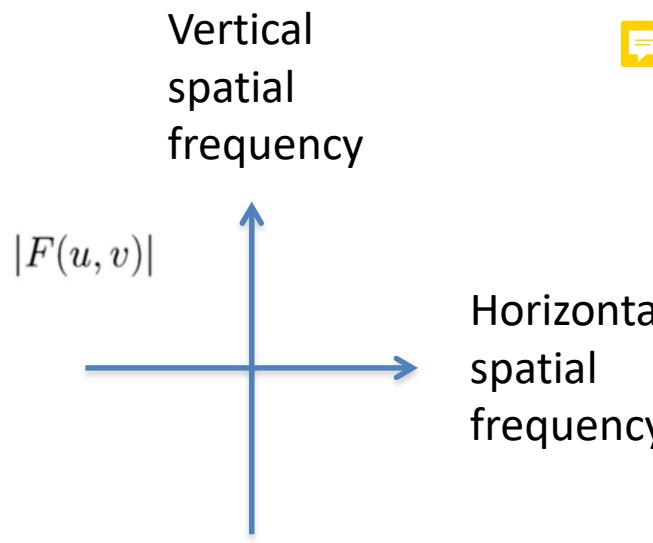
Image Frequency

- Image frequency can be obtained quantitatively using 2D Fourier Transform, which decomposes the an image into sine and cosine components

Definition

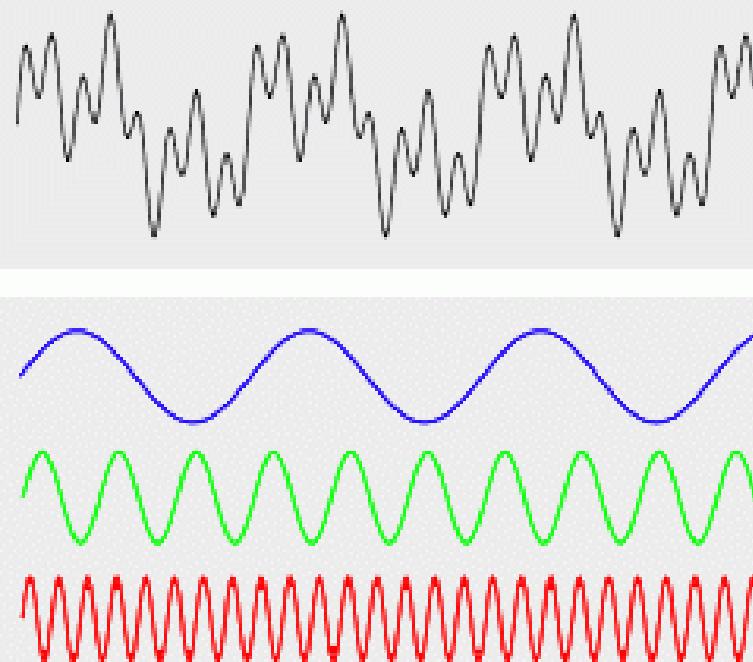
$$F(u, v) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y) e^{-j2\pi(ux+vy)} dx dy,$$

$$f(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} F(u, v) e^{j2\pi(ux+vy)} du dv$$



Fourier transform

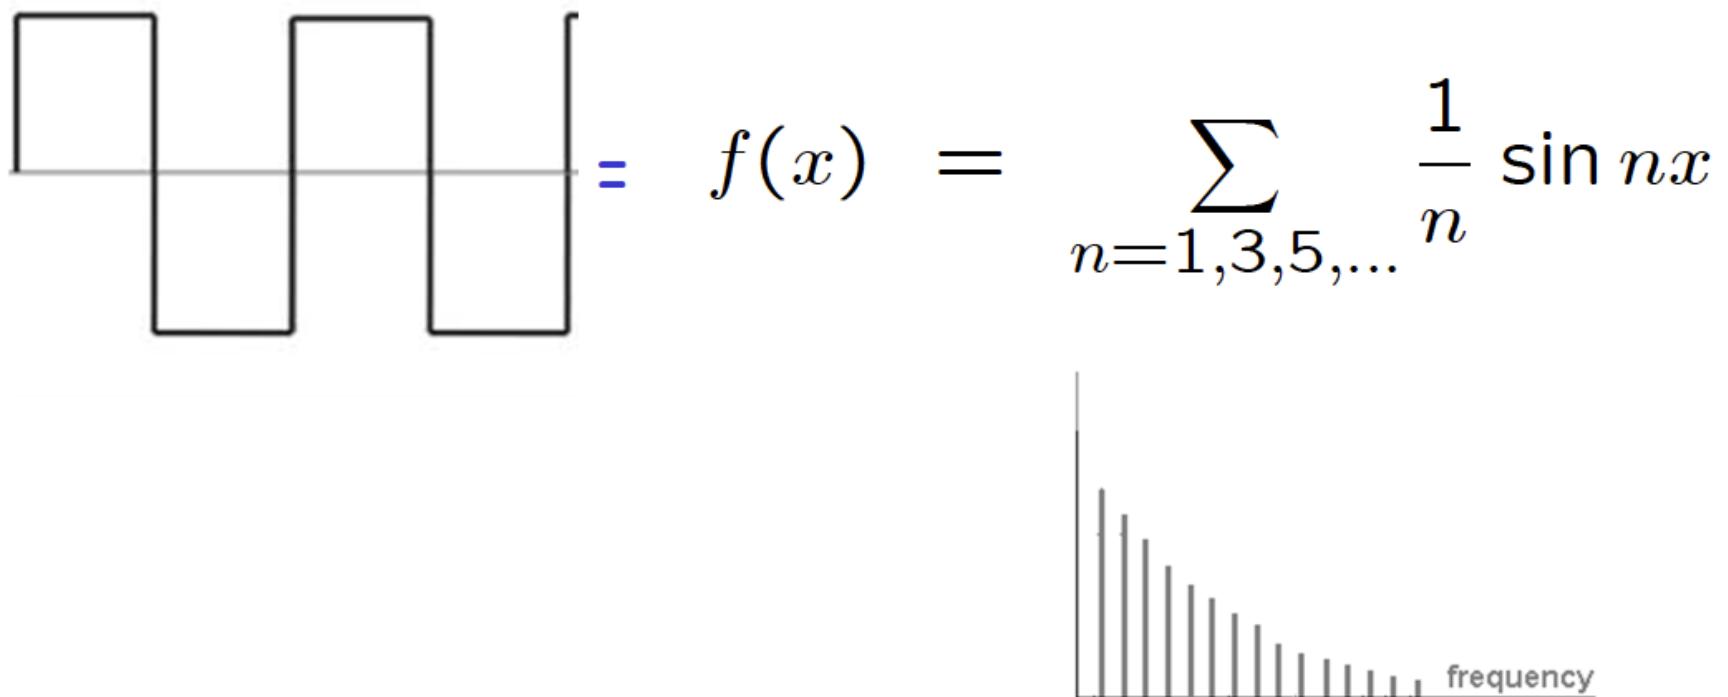
Fourier Analysis



Fourier Analysis:
The complex wave (upper) can be decomposed into the sum of the three simple sine waves (lower).



Fourier transform



[A. Zisserman]



Image Frequency

- Image frequency can be obtained quantitatively using 2D Fourier Transform, which decomposes the an image into sine and cosine components

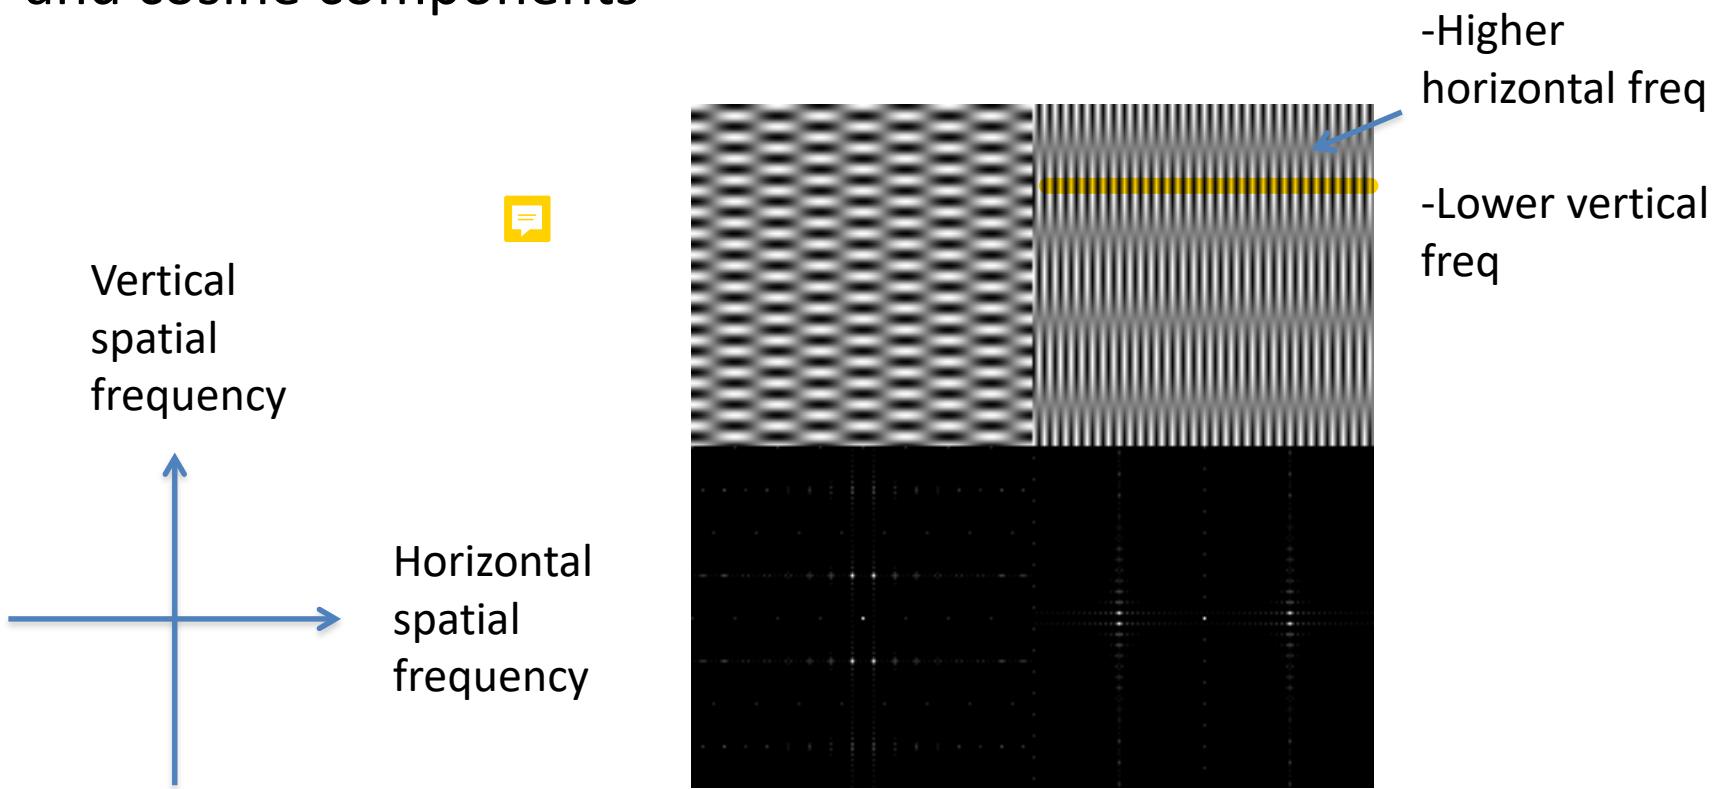
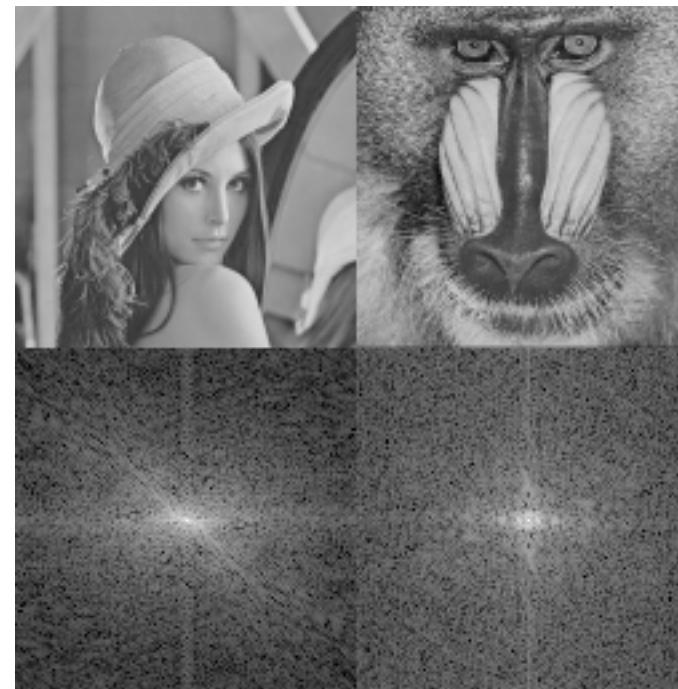
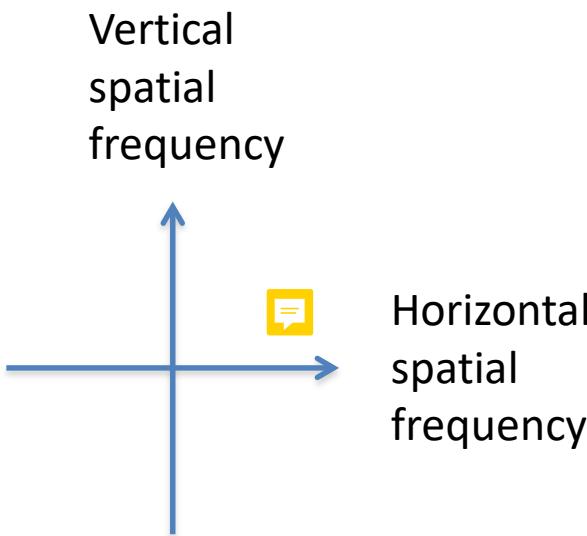




Image Frequency

- Image frequency can be obtained quantitatively using 2D Fourier Transform, which decomposes the an image into sine and cosine components



Low-pass filtering

- LPF retains low spatial frequency components, remove high spatial frequency components (noise, texture)
- An example of low-pass filter kernel

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

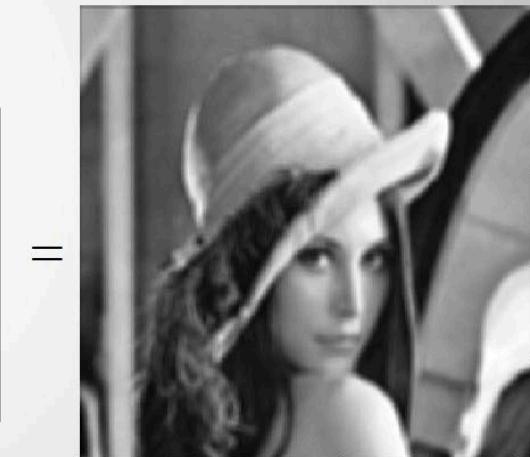
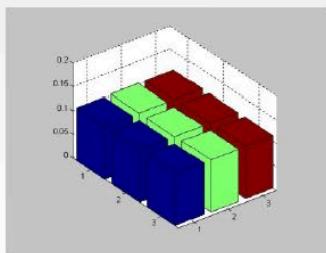
(see cohort exercise in filtering)

Low-pass filtering

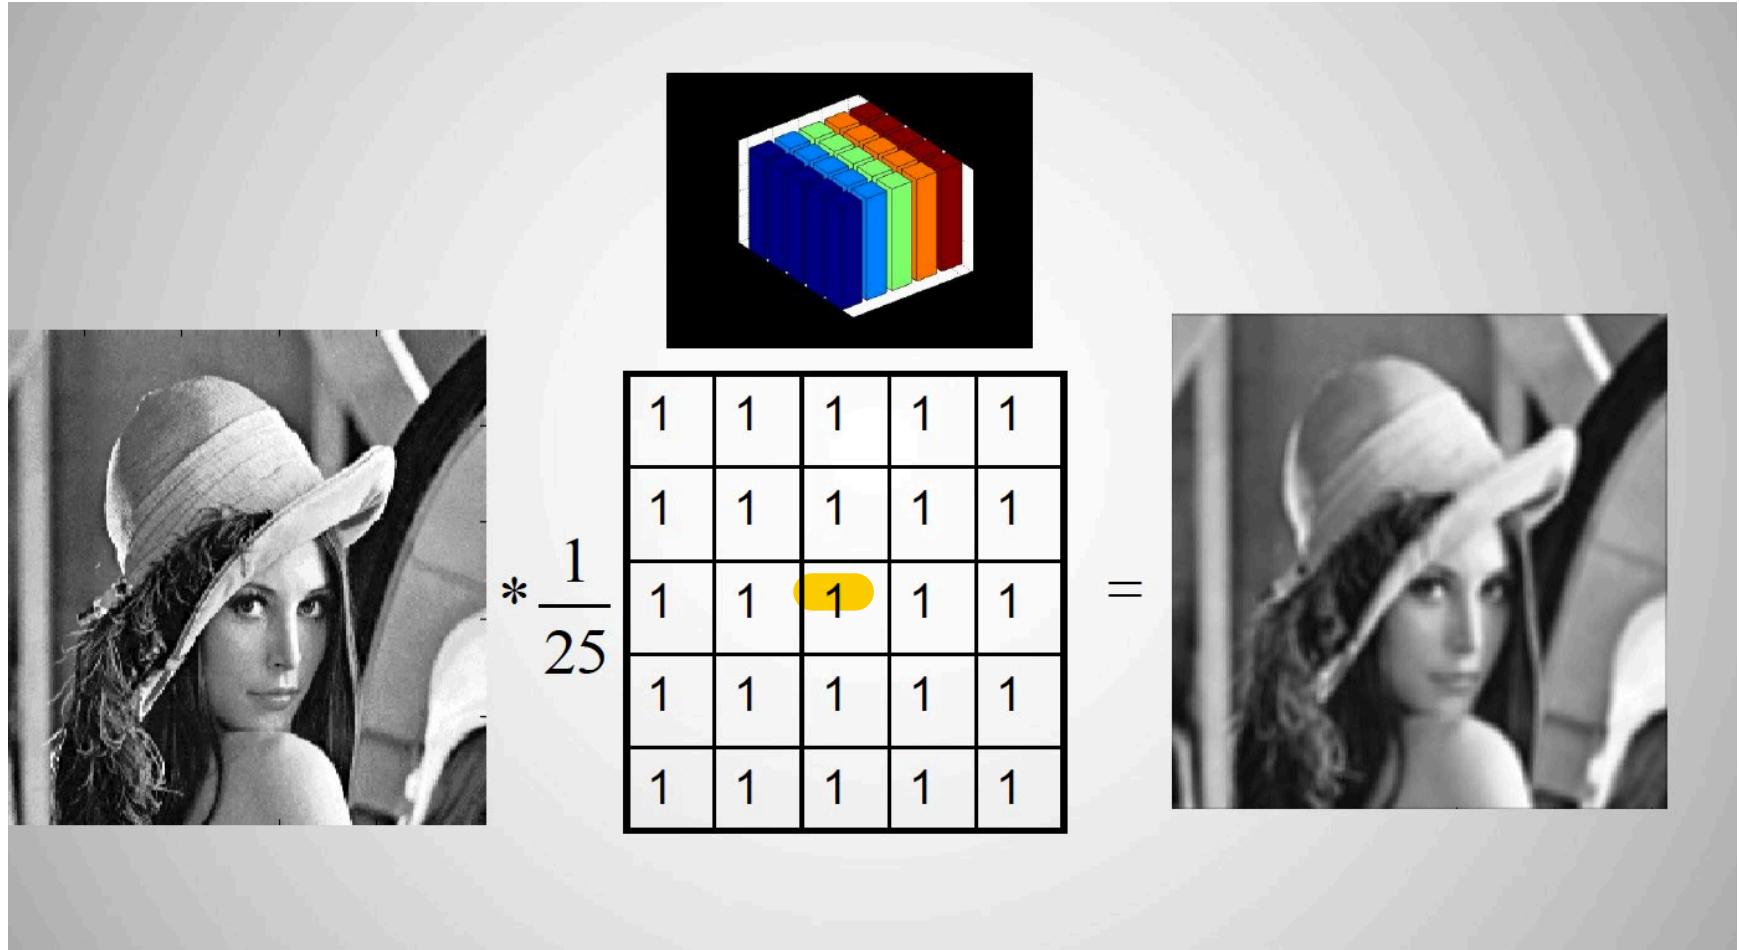


$$*\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

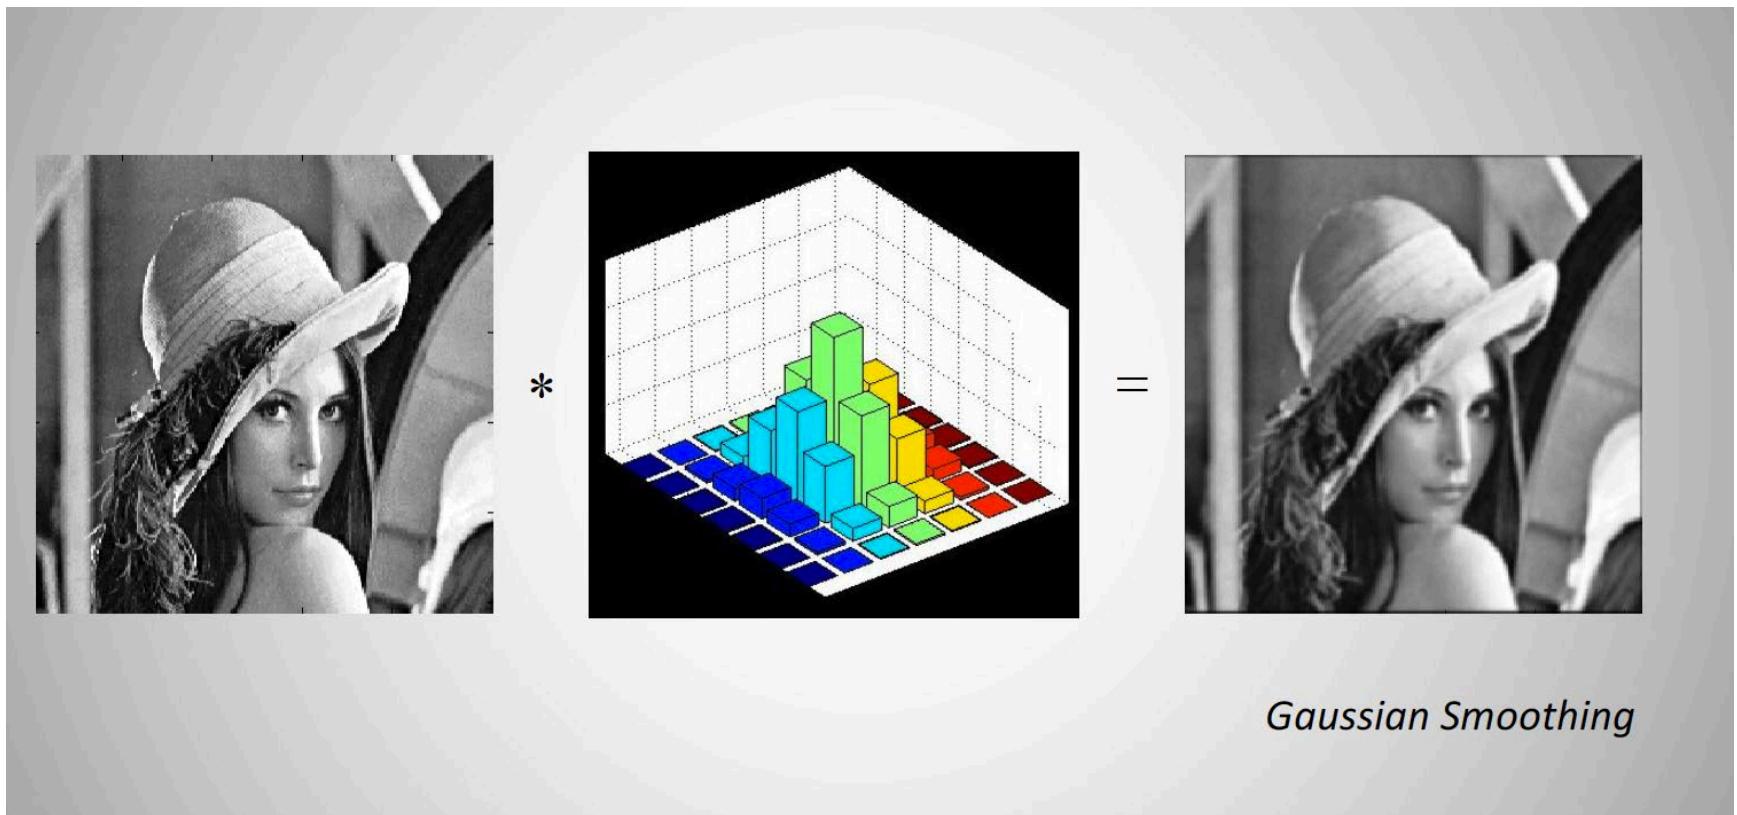


Low-pass filtering



Larger kernel -> more blurry

Low-pass filtering



Gaussian function:
$$f(x) = ae^{-\frac{(x-b)^2}{2c^2}}$$

Low-pass filtering



Gaussian Smoothing

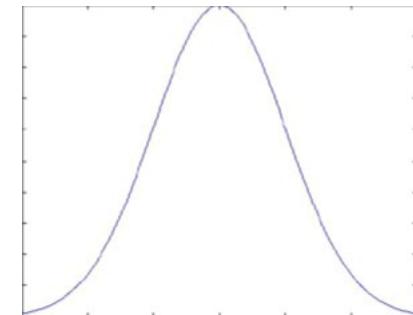


Smoothing by Averaging

Noise and denoising

Additive Gaussian Noise

$$I_{distorted} = I_{original} + n$$



n is a normally distributed

$$n \sim p(n; \mu, \sigma)$$

Noise: high spatial frequency

Denoising: low-pass filtering to remove the signal component with high spatial frequencies

Denoising



After additive
Gaussian Noise



After Averaging



After Gaussian Smoothing

Denoising by median filter

-**Non-linear filtering**: use the median as the filter output

$$\text{Median}(2,80,6) \rightarrow 6$$

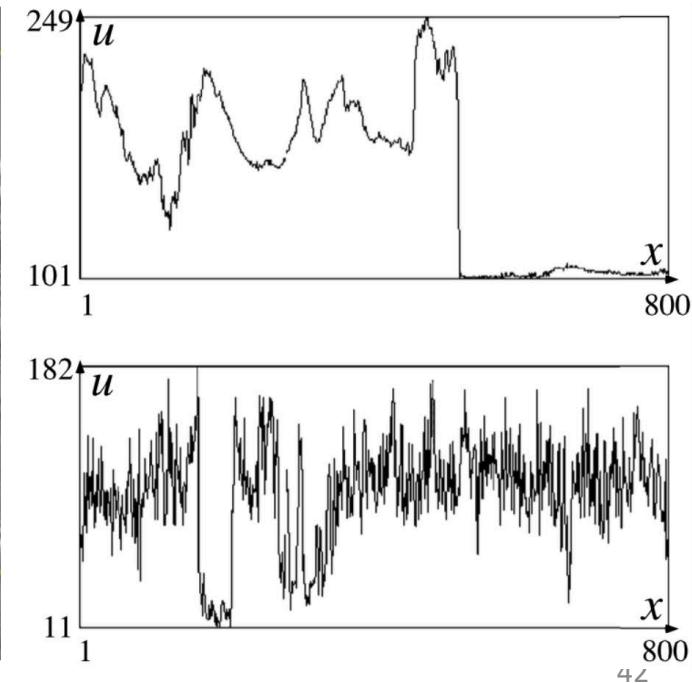
-More robust against outliers

-But more computationally-intensive



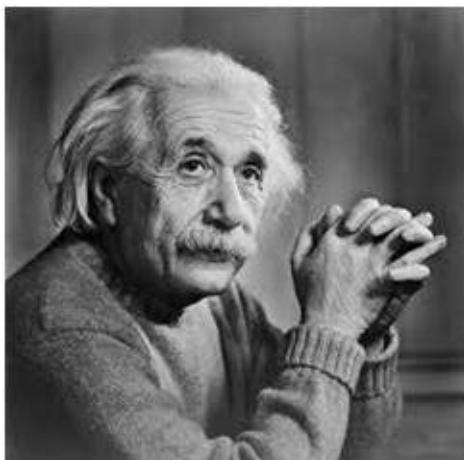
Edge detection

- Use image filtering to extract useful information: edge
- Edge: significant change in intensity values



Edge detection

- Why edge detection:
 - Edge provides important shape information: fundamental in understanding the image
 - Edge sharpen to improve visual quality



Edge detection

- Prewitt operator (as the filter kernel/mask)
 - Detect horizontal edge, vertical edge
 - Compute the difference of the neighboring pixels to indicate the likelihood of edges

$\begin{bmatrix} [-1, 0, 1], \\ [-1, 0, 1], \\ [-1, 0, 1] \end{bmatrix}$

Detect vertical edge

$\begin{bmatrix} [-1, -1, -1], \\ [0, 0, 0], \\ [1, 1, 1] \end{bmatrix}$

Detect horizontal edge

Edge detection

- Prewitt operator (as the filter kernel/mask)



$\begin{bmatrix} [-1, 0, 1], \\ [-1, 0, 1], \\ [-1, 0, 1] \end{bmatrix}$

Detect vertical edge



$\begin{bmatrix} [-1, -1, -1], \\ [0, 0, 0], \\ [1, 1, 1] \end{bmatrix}$

Detect horizontal edge

Edge detection

- Sobel operator (as the filter kernel/mask)



$\begin{bmatrix} [-1, 0, 1], \\ [-2, 0, 2], \\ [-1, 0, 1] \end{bmatrix}$

Detect vertical edge



$\begin{bmatrix} [-1, -2, -1], \\ [0, 0, 0], \\ [1, 2, 1] \end{bmatrix}$

Detect horizontal edge



Larger weight
near the
center pixel

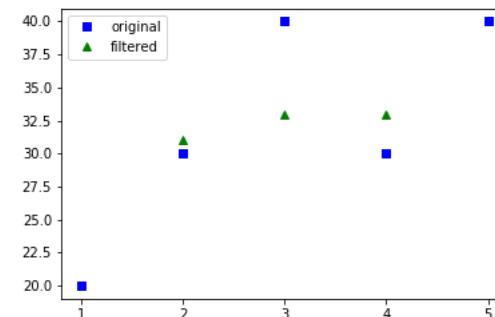
Image Filtering

- Cohort exercise: calculate the output of the following filtering (you can ignore the boundary condition)

Img = ~~[[30 40 20 30 40]
[40 20 30 40 30]
[20 30 40 30 40]
[30 40 30 40 20]
[40 30 40 20 30]]~~

kernel = $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$

out = [[30,31,33],
~~[31,33,33],~~
[33,33,32]]



After filtering,
the intensity
values are
smoother