Image classification, data-driven approach, knn

ISTD 50.035

Computer Vision

Acknowledgement: Some images are from various sources: UCF, Stanford cs231n, etc.

Image is an array of numbers

-Grayscale image

-2D array of numbers (pixels) / matrix

-Number indicates the intensity: [0,255] for 8-bit representation

-Image resolution / number of pixel in an image: 100x100, 1920x1080, etc.

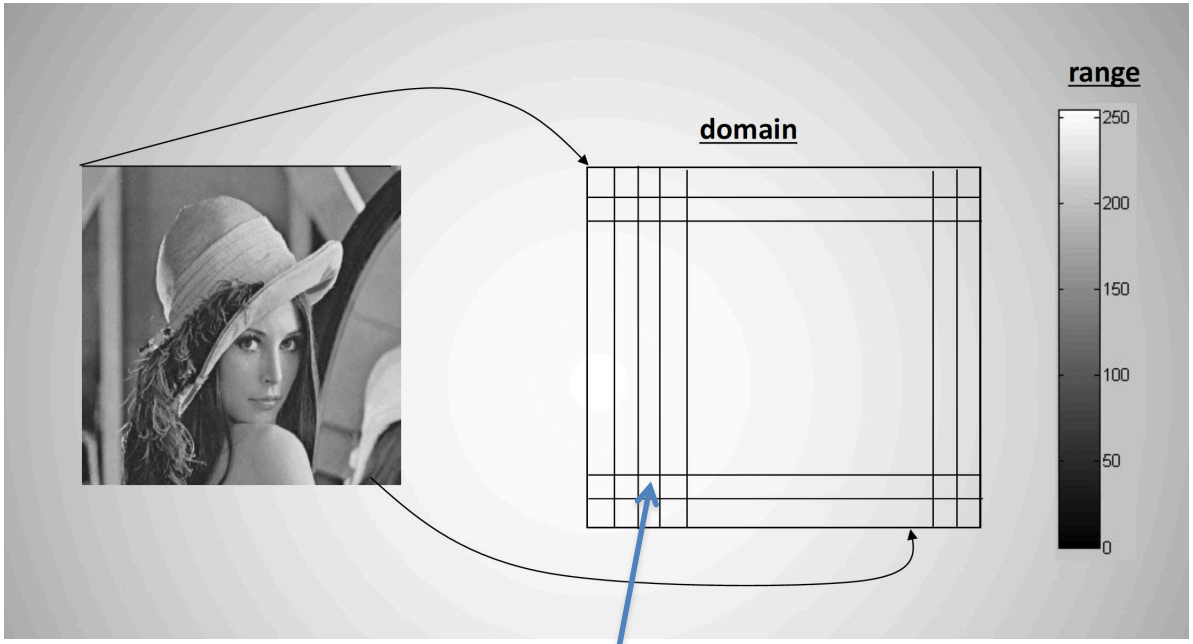

A pixel

0: black, 255: white

Can you recognize this image?

```
[[254 252 248 251 248 243 231 212 188 173 163 159 163 169 162 154]
 [255 253 251 243 227 193 158 145 159 159 154 150 153 158 159 159]
 [246 228 211 170 179 156  70  50  73 138 193 197 165 144 152 173]
 [228 218 186 149 133 130 100  48  47  61 137 192 175 168 170 169]
 [212 200 170 147 121  85 124  65  69 137 185 240 221 174 171 193]
 [222 235 228 192 162 132 150 187 218 217 200 225 228 211 214 214]
 [187 230 244 231 213 236 232 220 222 218 216 213 204 199 207 195]
 [166 173 204 230 245 247 226 193 191 225 218 209 181 132 198 202]
 [189 171 210 210 210 179 210 211 203 192 213 188 106 151 204 184]
 [205 174 219 223 219  99 121 233 231 214 214 212 195 193 151 158]
 [209 137 166 212 204 115 117 202 222 219 210 204 168 114  84 204]
 [214 174 143 155 222 137 139 182 214 193 143 157  58  78 135 198]
 [227 230 222 191 173 129 148 150 184 140 103 147 125 144 165 204]
 [241 239 238 231 216 143 143 163 203 193 175 145 143 164 198 201]
 [250 248 250 245 243 219 135  94 134 156 160 179 194 198 198 198]
 [253 251 249 249 249 246 232 223 229 232 216 213 213 203 198 197]]
```

Can you recognize this image?

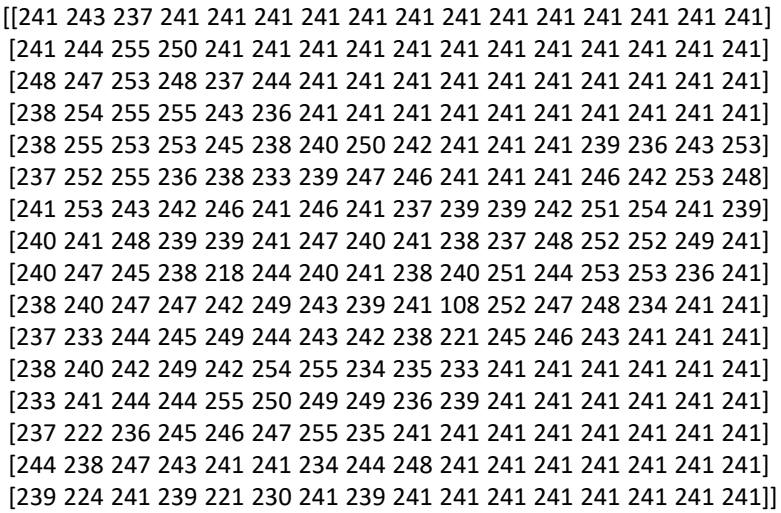

```
[[241 243 237 241 241 241 241 241 241 241 241 241 241 241 241 241]
 [241 244 255 250 241 241 241 241 241 241 241 241 241 241 241 241]
 [248 247 253 248 237 244 241 241 241 241 241 241 241 241 241 241]
 [238 254 255 255 243 236 241 241 241 241 241 241 241 241 241 241]
 [238 255 253 253 245 238 240 250 242 241 241 241 239 236 243 253]
 [237 252 255 236 238 233 239 247 246 241 241 241 246 242 253 248]
 [241 253 243 242 246 241 246 241 237 239 239 242 251 254 241 239]
 [240 241 248 239 239 241 247 240 241 238 237 248 252 252 249 241]
 [240 247 245 238 218 244 240 241 238 240 251 244 253 253 236 241]
 [238 240 247 247 242 249 243 239 241 108 252 247 248 234 241 241]
 [237 233 244 245 249 244 243 242 238 221 245 246 243 241 241 241]
 [238 240 242 249 242 254 255 234 235 233 241 241 241 241 241 241]
 [233 241 244 244 255 250 249 249 236 239 241 241 241 241 241 241]
 [237 222 236 245 246 247 255 235 241 241 241 241 241 241 241 241]
 [244 238 247 243 241 241 234 244 248 241 241 241 241 241 241 241]
 [239 224 241 239 221 230 241 239 241 241 241 241 241 241 241 241]]
```

```
[[243 241 244 243 243 243 243 243 243 243 243 243 243 243 243 243]
 [243 245 175 221 242 243 243 243 243 243 243 243 243 243 243 243]
 [239 163 164 164 194 241 243 243 243 243 243 243 243 243 243 243]
 [234 158 158 145 166 248 243 243 243 243 243 243 243 243 243 243]
 [234 156 142 144 169 215 221 244 243 243 243 243 242 246 242 238]
 [234 161 150 132 182 201 212 223 228 243 243 243 167 168 174 195]
 [226 154 153 197 210 213 218 216 211 210 187 155 156 163 164 243]
 [214 220 219 196 170 204 222 213 219 206 197 162 167 164 242 243]
 [214 216 216 183 187 203 210 215 211 215 208 164 169 161 245 243]
 [212 206 221 205 205 210 205 204 212  72 212 158 169 245 243 243]
 [209 204 216 205 208 197 193 200 201 172 239 244 242 243 243 243]
 [210 213 210 213 195 200 186 183 198 204 243 243 243 243 243 243]
 [194 215 213 202 202 171 155 192 209 241 243 243 243 243 243 243]
 [190 169 195 213 156 171 174 151 213 243 243 243 243 243 243 243]
 [197 196 240 198 213 179 148 203 242 243 243 243 243 243 243 243]
 [198 164 243 248 156 183 243 240 243 243 243 243 243 243 243 243]]
```

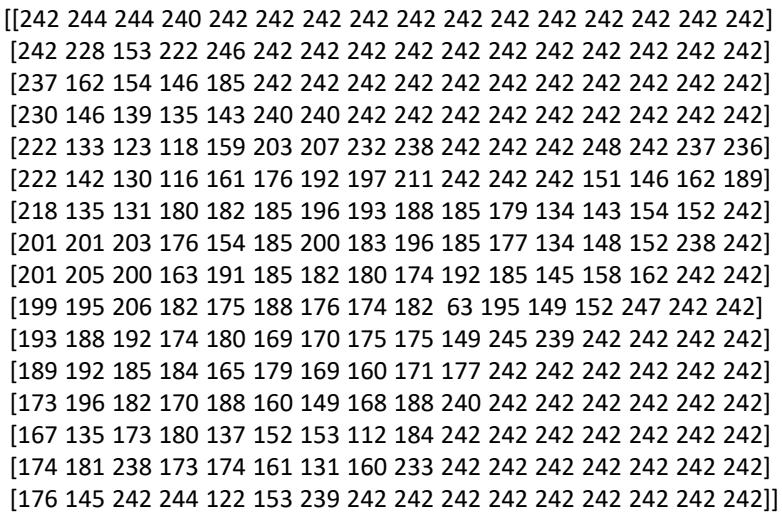

```
[[242 244 244 240 242 242 242 242 242 242 242 242 242 242 242 242]
 [242 228 153 222 246 242 242 242 242 242 242 242 242 242 242 242]
 [237 162 154 146 185 242 242 242 242 242 242 242 242 242 242 242]
 [230 146 139 135 143 240 240 242 242 242 242 242 242 242 242 242]
 [222 133 123 118 159 203 207 232 238 242 242 242 248 242 237 236]
 [222 142 130 116 161 176 192 197 211 242 242 242 151 146 162 189]
 [218 135 131 180 182 185 196 193 188 185 179 134 143 154 152 242]
 [201 201 203 176 154 185 200 183 196 185 177 134 148 152 238 242]
 [201 205 200 163 191 185 182 180 174 192 185 145 158 162 242 242]
 [199 195 206 182 175 188 176 174 182  63 195 149 152 247 242 242]
 [193 188 192 174 180 169 170 175 175 149 245 239 242 242 242 242]
 [189 192 185 184 165 179 169 160 171 177 242 242 242 242 242 242]
 [173 196 182 170 188 160 149 168 188 240 242 242 242 242 242 242]
 [167 135 173 180 137 152 153 112 184 242 242 242 242 242 242 242]
 [174 181 238 173 174 161 131 160 233 242 242 242 242 242 242 242]
 [176 145 242 244 122 153 239 242 242 242 242 242 242 242 242 242]]
```

Semantic gap

```
[[242 244 244 240 242 242 242 242 242 242 242 242 242 242 242 242]
 [242 228 153 222 246 242 242 242 242 242 242 242 242 242 242 242]
 [237 162 154 146 185 242 242 242 242 242 242 242 242 242 242 242]
 [230 146 139 135 143 240 240 242 242 242 242 242 242 242 242 242]
 [222 133 123 118 159 203 207 232 238 242 242 242 248 242 237 236]
 [222 142 130 116 161 176 192 197 211 242 242 242 151 146 162 189]
 [218 135 131 180 182 185 196 193 188 185 179 134 143 154 152 242]
 [201 201 203 176 154 185 200 183 196 185 177 134 148 152 238 242]
 [201 205 200 163 191 185 182 180 174 192 185 145 158 162 242 242]
 [199 195 206 182 175 188 176 174 182  63 195 149 152 247 242 242]
 [193 188 192 174 180 169 170 175 175 149 245 239 242 242 242 242]
 [189 192 185 184 165 179 169 160 171 177 242 242 242 242 242 242]
 [173 196 182 170 188 160 149 168 188 240 242 242 242 242 242 242]
 [167 135 173 180 137 152 153 112 184 242 242 242 242 242 242 242]
 [174 181 238 173 174 161 131 160 233 242 242 242 242 242 242 242]
 [176 145 242 244 122 153 239 242 242 242 242 242 242 242 242 242]]
```

The gap between low-level representation of an image (input to an algorithm) and high-level understanding of an image (output)

Image classification

- Given an input image, the algorithm produces one image label from a fixed set of classes (categories)

{fish, soccer ball, dog, boat}

- Image recognition (many classes)
  - 1000 categories in IMAGENET Large Scale Visual Recognition Challenge (ILSVRC): zebra, speedboat, lifeboat, …
  - 10,000+ categories in IMAGENET

zebra 

speedboat 

lifeboat 

Image classification

- **Top-n accuracy**
- The algorithm outputs k confidence for each of the k classes

Test image: 

Algorithm outputs: {cat, dog, house, mouse} = {0.1, 0.2, 0.0, 0.7}

Top-1 class: {mouse}

Top-2 class: {mouse, dog}

Incorrect for **top-1 accuracy**, correct for **top-2 accuracy** (ground truth is contained in the top-2 class)

- ILSVRC: Top-1, Top-5 accuracy

Number of correct / Number of test image

Image classification is fundamental to many computer vision tasks

- **Object localization**
- For a given image, the algorithm produces a **class label** and a **bounding box**
- Evaluation: label that best matches the ground truth label for the image, and bounding box that overlaps with the ground truth
- An error if predicted label does not match the ground truth, or the predicted bounding box has less than 50% overlap

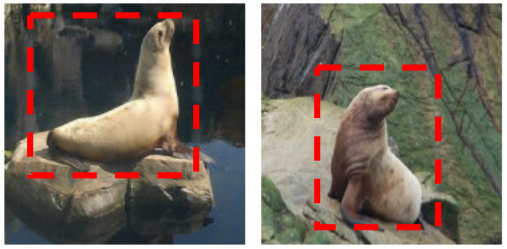

sea lion

Image classification is fundamental to many computer vision tasks

- **Object detection**
- Given an image, an algorithm produces a set of annotations (ci,si,bi): class label ci, bounding box bi and confidence score si
- Penalize: objects in the image not annotated by algorithm, more than 1 annotations for the same object in the image

- apple
- table
- bowl
- plate rack
- lamp
- chair

200 categories in ILSVRC2017

Image classification

- Challenges
  - Primitive data: Computer sees a 3d array of intensity values
  - Different variation for a certain class
    - Viewpoint variation
    - Scale variation
    - Deformation
    - Occlusion
    - Background clutter
    - Intra-class variation

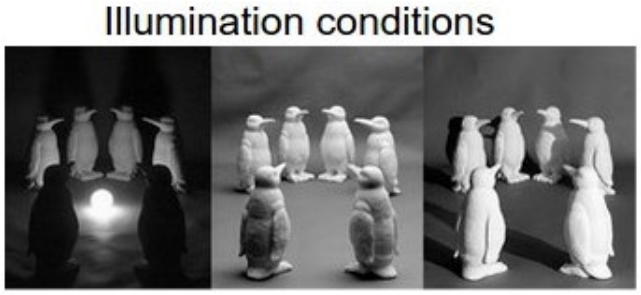

Image classification

- Challenges

Deformation of non-rigid object

Background clutter

Data driven approach

- Provide the computer with many examples of each class: training data
- Learn the visual appearance of each class: learning algorithm
- ILSVRC: 1.2 million images of 1000 categories
  - About 1k images per category

zebra

mongoose

Data driven approach

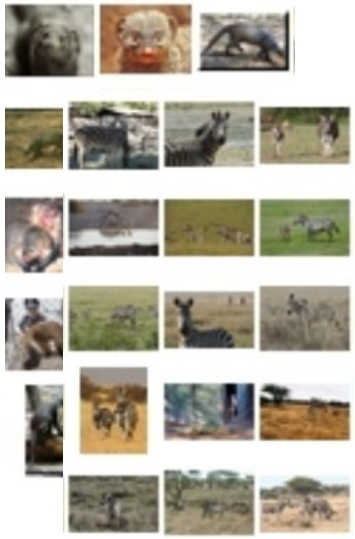

Training/Learning (usually offline)

Training set: images with known class information

→ Learn a model using some algorithm →

A model for this specific classification problem and classifier

Testing/Evaluation (usually online)

Testing set (with label during evaluation, without label in an application)

→ Classifier algorithm →

Predicted class information for this new image (compare with ground-truth during evaluation)

Nearest Neighbor Classifier

- Given a test image, compare to every one of the training images
- Use the label of the ‘closest training image’ as the predicted label

Nearest Neighbor Classifier

- Consider an image as a vector (data point) in a very high dimensional vector space
- 512x512x3 => a data point in the 786432-dim vector space
- Find the nearest neighbors of the vector representing the input test image

Nearest Neighbor Classifier

Each training image is represented by one high-dimensional vector (point)

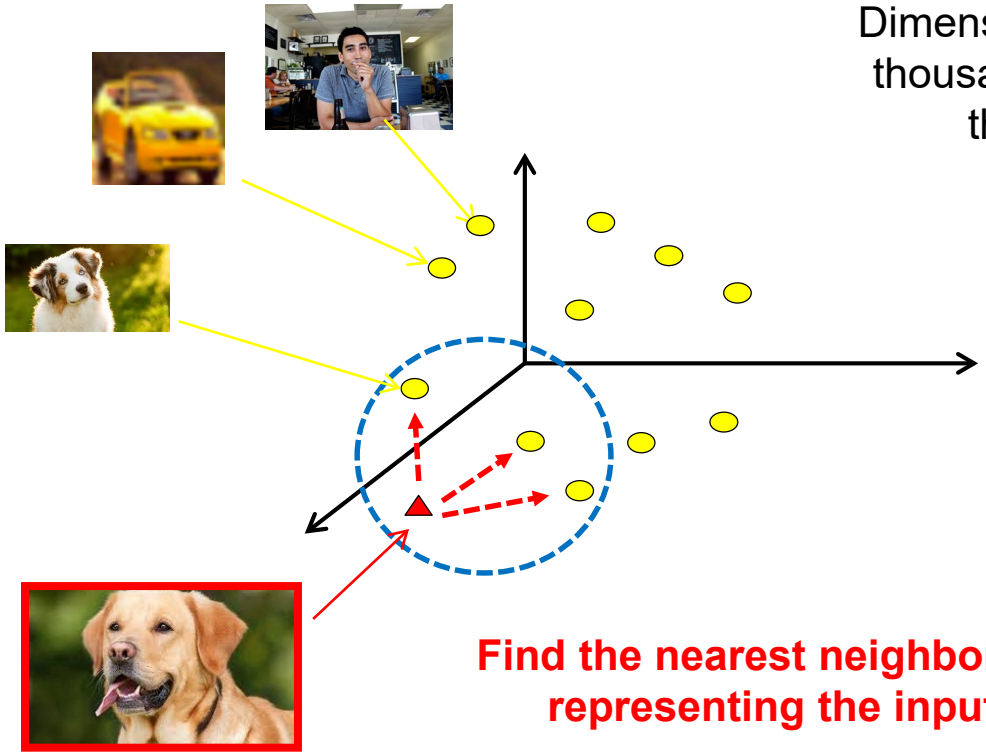

Dimensionality can be thousands or tens of thousands

Test image

Find the nearest neighbors of the vector representing the input test image

Distance

- L2 distance (Euclidean distance)

$$d_2(I_1, I_2) = \sqrt{\sum_p \left(I_1^p - I_2^p\right)^2}$$

p indicates dimension

- L1 distance (Manhattan distance)

- L2 dist Sum of abs difference

- L1 dist
  - Sum

$$d_1(I_1, I_2) = \sum_p \left|I_1^p - I_2^p\right|$$

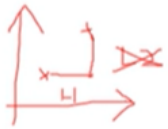

p indicates dimension

Distance

- L1/L2 circle / ball
- A circle is a set of points with a fixed distance from a point (center)

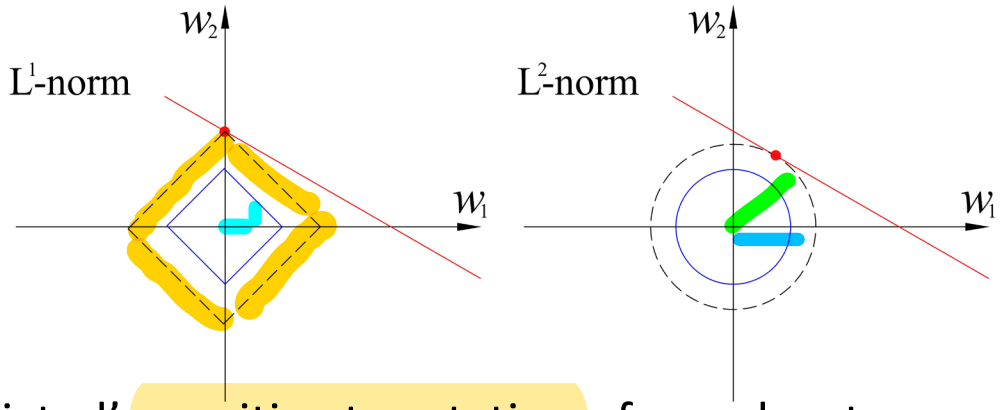

L1 is more 'restricted', sensitive to rotation of coord system
L2 emphasizes dimensions with large differences
L1: sparse model (use as regularization), robust to outliers (use as cost function)

Distance

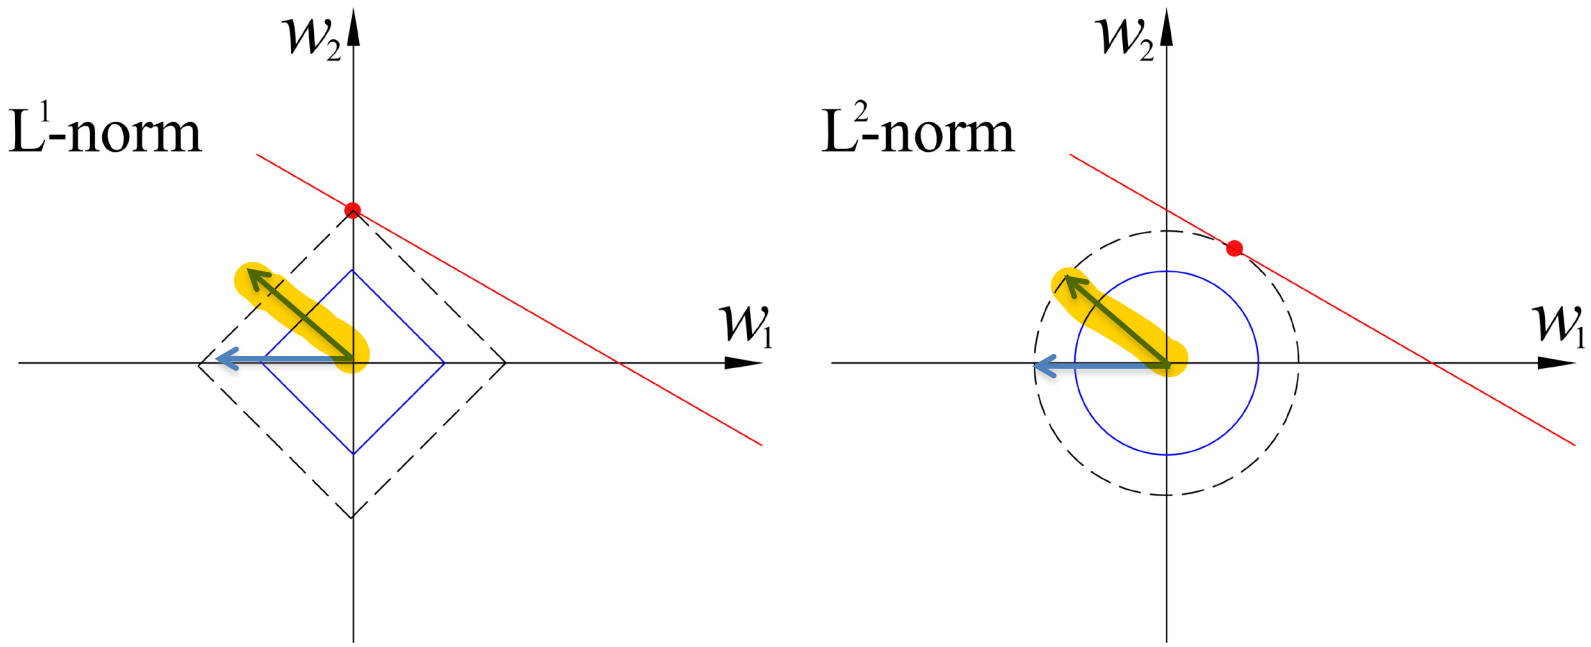

L1 is more ‘restricted’, sensitive to rotation of coord system

k-Nearest Neighbor Classifier (k-NN)

- Find the k closest images (nearest neighbors)

- Use them to vote on the label of the test image

k-Nearest Neighbor Classifier (k-NN)

- How to determine k?
- k is a hyperparameter: related to the design of the machine learning algorithm
- Another hyperparameter: L1 norm or L2 norm

Validation set for hyperparameter tuning

- Use test set to tune the hyperparameter
- Not appropriate, as your model will overfit to the test data
- Poor generalization, significant degradation during deployment / testing for other datasets

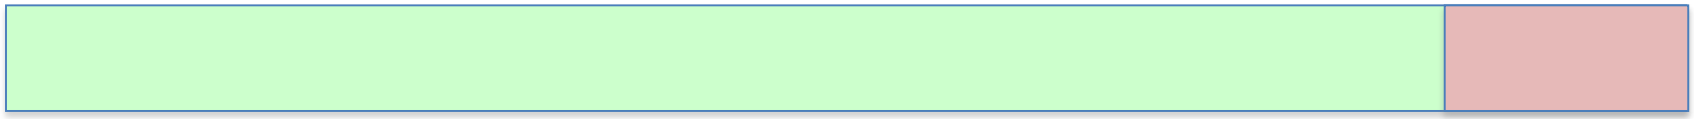

Training Test

Validation set for hyperparameter tuning

- Partition the training set into a training set and a validation set
- Use validation set to tune the hyperparameter
- Use test set to evaluate the performance

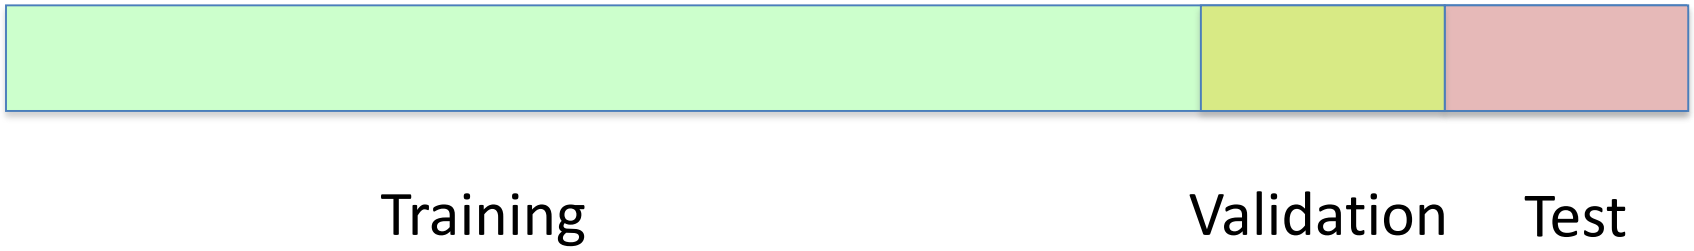

Cross validation

- If the training dataset is small, can use cross validation
- 5-fold cross validation
  - For a given k (a certain setting of hyperparameters)
  - Divide the training dataset into 5 equal folds
  - Use 4 folds for training, 1 for validation
  - Repeat using another fold as the validation set
  - Average the performance

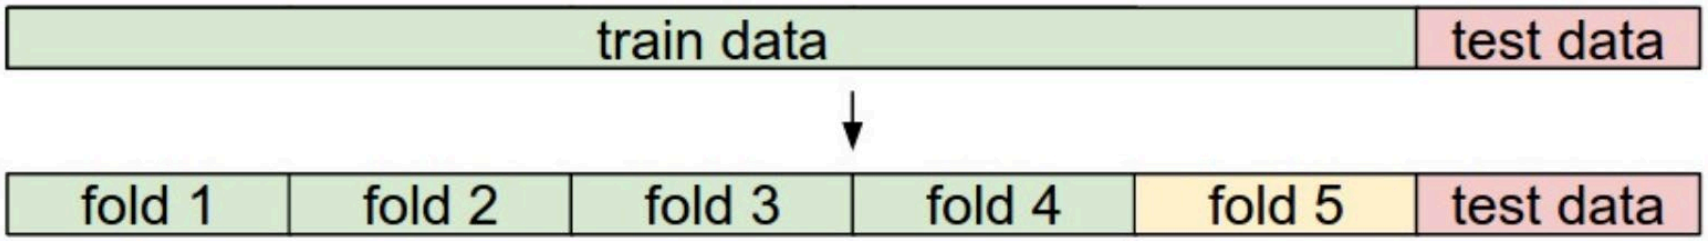

Issues of k-NN

- Memory expensive: need to remember all training data
- Computationally expensive during testing
  - Need to compare all training data
  - Not practical in an application
- Approximate nearest neighbor (ANN) algorithms accelerate the search of the nearest neighbor
- Using image intensity value for distance comparison is not robust
  - Small position or intensity shift can result in large distance

Original

Position shift

Intensity shift

Image classification: representation (feature) learning + classifier

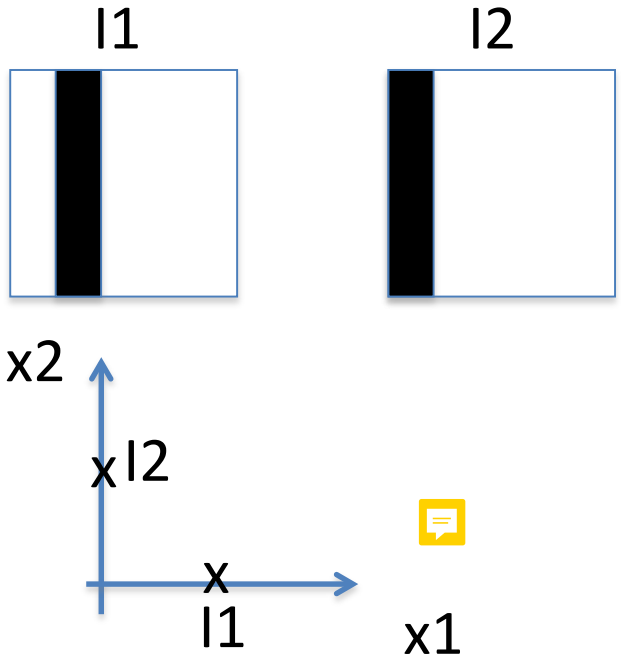

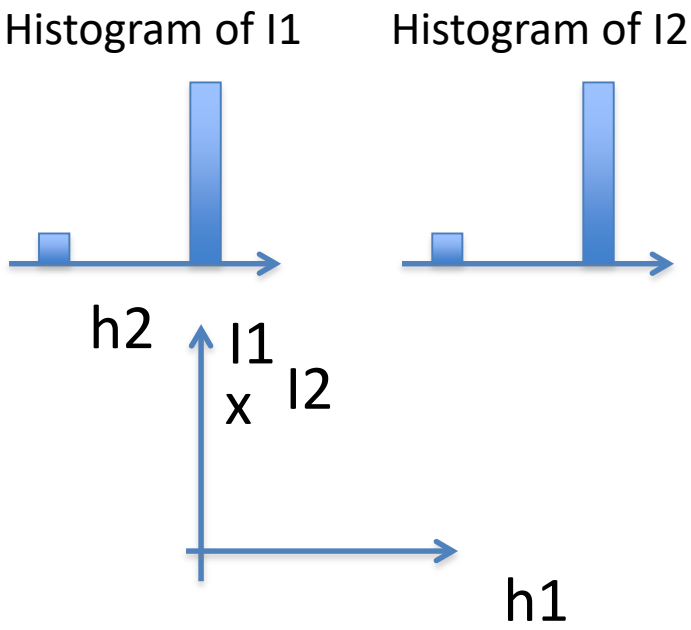

Rgb representation is sensitive to position shift (translation)

Histogram representation of an image is robust to position shift (translation)

Image classification: representation (feature) learning + classifier

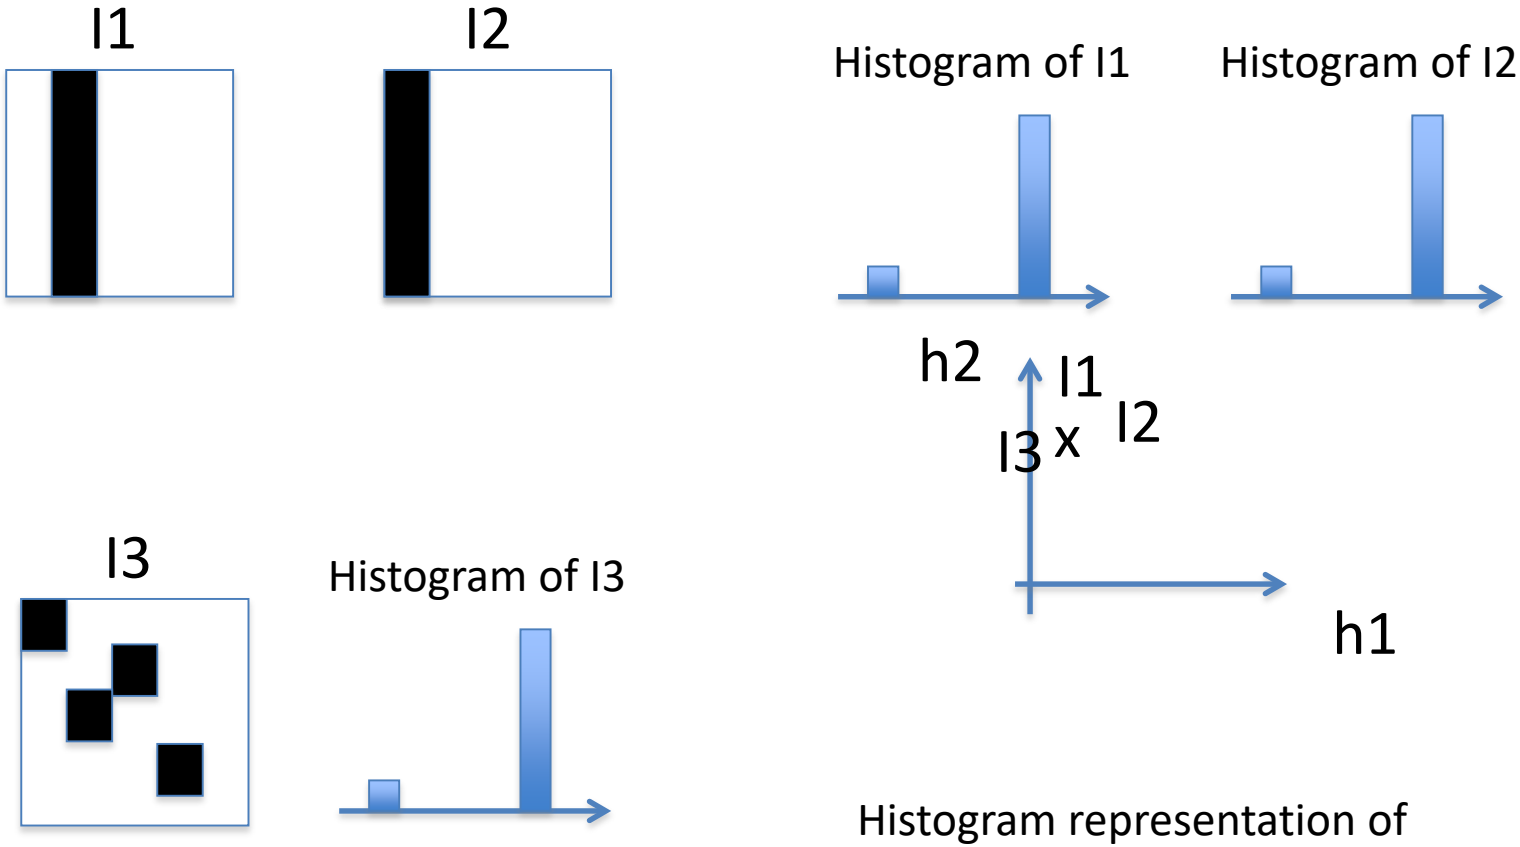

The quest for robust and discriminative representation

Histogram representation of an image is robust to position shift (translation) but not discriminative