

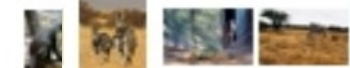
Linear classifier

ISTD 50.035

Computer Vision

Acknowledgement: Some images are from various sources: UCF, Stanford cs231n, etc.

Data driven approach



Training/Learning (usually offline)

Training set: images with known class information



Learn a model using some algorithm



A model for this specific classification problem and classifier

Testing/Evaluation (usually online)



Testing set (with label during evaluation, without label in an application)



Classifier algorithm



Predicted class information for this new image (compare with ground-truth during evaluation)

For practical use, testing time should be small

Linear Classifier

$$s = f(x; W, b) = Wx + b \quad \text{Score function}$$

- Given a test image x , produce the confidence score for each class using linear transformation (total: K classes)
- Higher confidence score for a class \rightarrow more likely to be the ground-truth class
- Test image x : flatten to a $D \times 1$ column vector, D is the image resolution times number of channel



Input x : $D \times 1$

Weight W : $K \times D$

Bias b : $K \times 1$

Score s : $K \times 1$

Linear Classifier

$$s = f(x; W, b) = Wx + b$$

- **Testing:** W, b are fixed, x is the input
- **Training:** Given N training samples (x_i, y_i) , y_i takes value in $[1, \dots, K]$, learn W and b
- The ground truth class is y_i

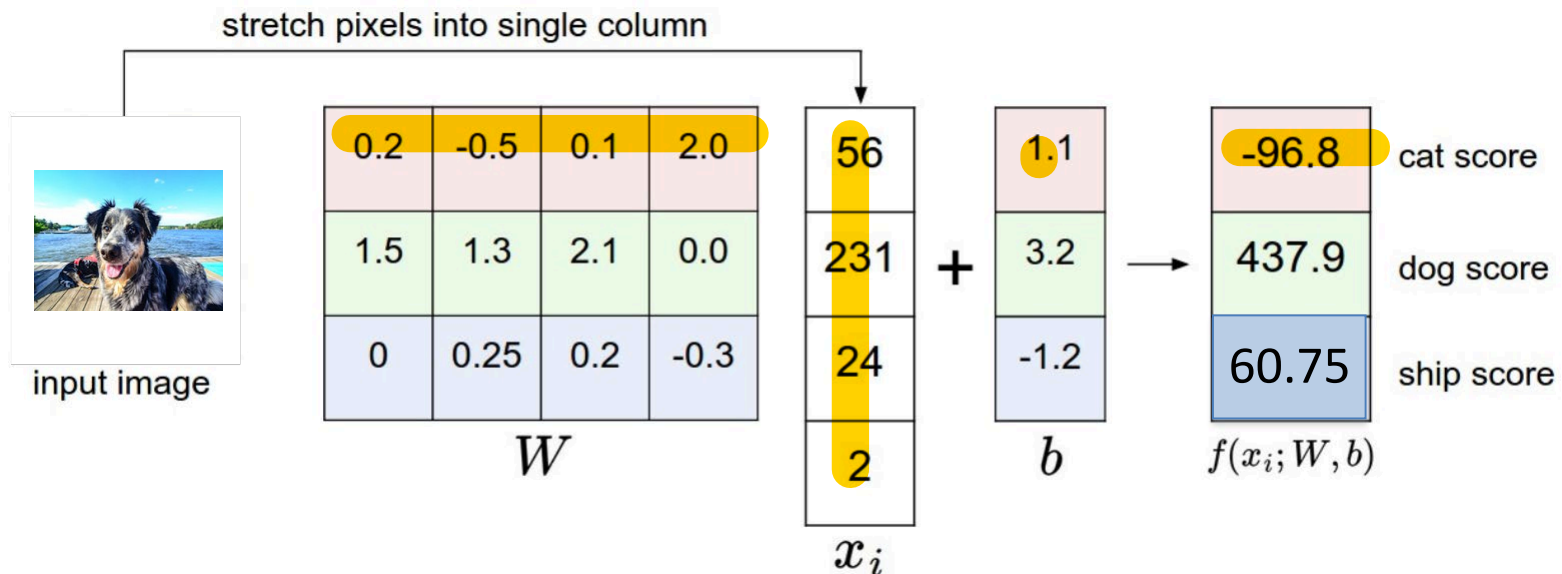
Example: $K=3$, $x_i =$  $y_i = 2$ {cat, dog, ship}

Training: (x_i, y_i) are given and fixed; W, b are the variables to be determined

Linear Classifier

$$s = f(x; W, b) = Wx + b$$

- **Testing:** W , b are fixed, x is the input



- Training: learn W , b to discriminate the classes
- Each row of W extracts the features (template) of a specific class from input

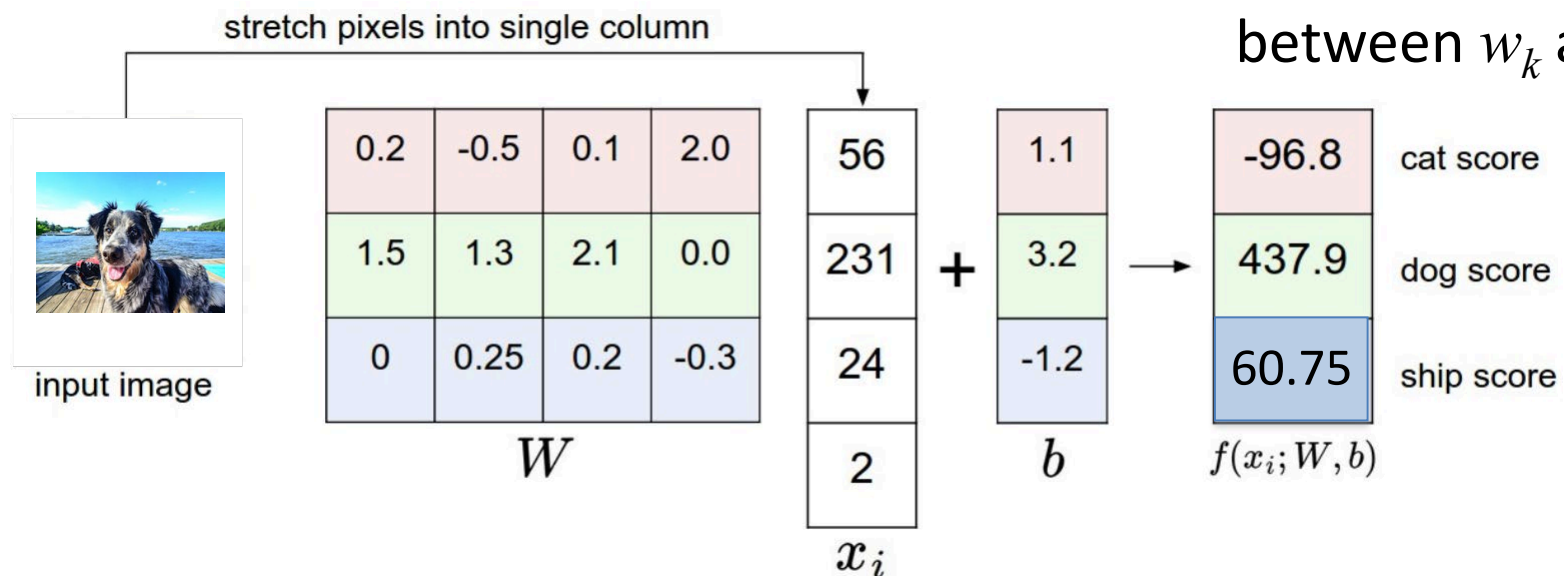
Linear Classifier

$$s = f(x; W, b) = Wx + b$$

$$w_k^T x_i = \|w_k\| \|x_i\| \cos \theta$$

- Testing: W , b are fixed, x is the input

Measure similarity
between w_k and x_i



- Training: learn W , b to discriminate the classes
- Each row of W extracts the features (template) of a specific class from input

Linear Classifier

$$s = f(x; W, b) = Wx + b$$

- **Shorthand notation**

$$s = [W \ b][x \ 1]^T$$

$$W \quad x$$

$$s = f(x; W) = Wx$$

Input x : $(D+1) \times 1$
Weight W : $K \times (D+1)$
Score s : $K \times 1$

Loss function

- **Training:** Given N training samples (x_i, y_i) , y_i takes value in $[1, \dots, K]$, learn W
- Loss function: measure how consistent are the ground-truth labels and the score function outputs, for some W
- Small loss: good W
- Softmax classifier with cross-entropy loss
- Multiclass Support Vector Machine (SVM) loss

Softmax classifier

- Regard output of the score function $f(x; W)$ as the *unnormalized log probability* of each class
- Probability of each class can be obtained by applying a **softmax function** (exp, then normalize):

$$\text{softmax}(f) = \frac{e^{f_m}}{\sum_{j=1}^K e^{f_j}}$$

For the m-th class



- **Cross-entropy loss** (apply $-\log(\cdot)$ to only the ground-truth class):

$$L_i = -\log \frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$$

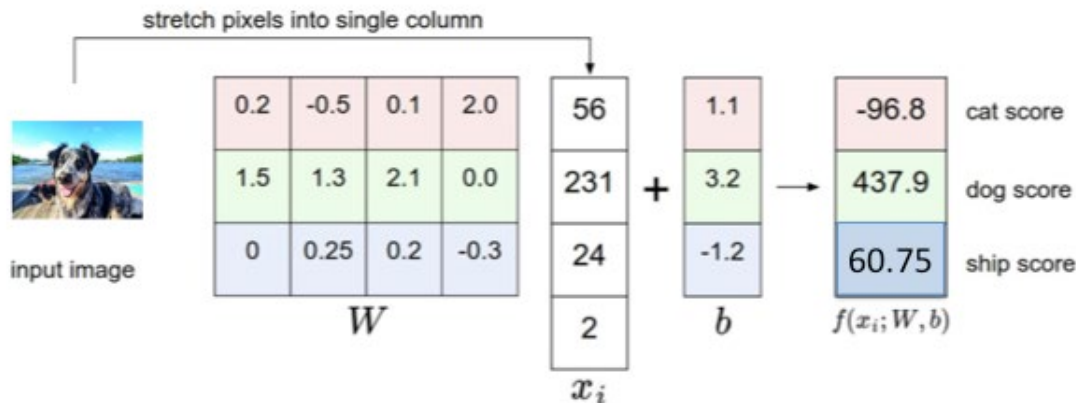
For the i-th
training
sample

Softmax classifier

- **Cross-entropy loss** (apply $-\log(\cdot)$ to only the ground-truth class):

$$L_i = -\log \frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$$

For the i -th training sample



Example: a dog (which looks like a dog)

```
print(np.exp(f))
print(np.exp(f) / sum(np.exp(f)))
```

$y_i = 2$
 $L_i = -\log(1) = 0$

```
[ 9.12628762e-043  1.50505935e+190  2.41762966e+026]
[ 6.06373937e-233  1.00000000e+000  1.60633510e-164]
```

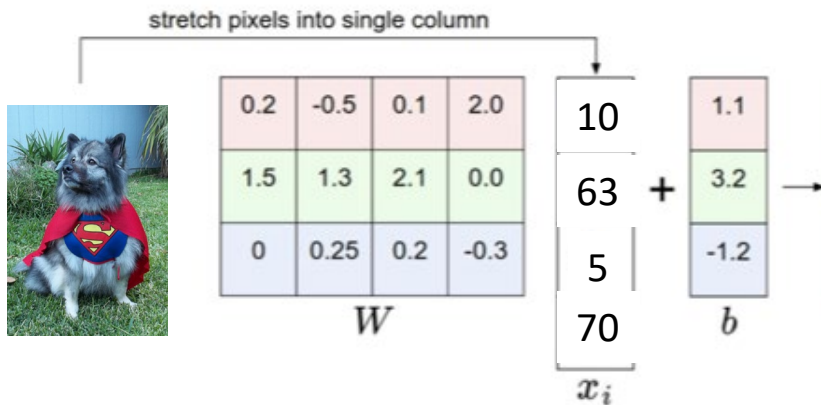
← Probability

Softmax classifier

- **Cross-entropy loss** (apply $-\log(\cdot)$ to only the ground-truth class):

$$L_i = -\log \frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$$

For the i -th training sample



Example: a dog
(which does not
look like a dog)



Softmax classifier

- The entire loss for N training samples (x_i, y_i) :

$$L = \frac{1}{N} \sum_i L_i$$

- We determine W to minimize this loss given the training dataset
- Additional regularization of W

Understand Softmax classifier

- Cross-entropy loss
 - First apply softmax function
 - Then apply $-\log(\cdot)$ to only the ground-truth class

$$L_i = -\log \frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}} \quad \text{For the } i\text{-th training sample}$$

- The term $\frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$ is the probability of the correct class (i.e., y_i)
- Therefore, want this to be large, i.e., $\max_w \log(\cdot)$
- Thus, want this to be small $\min_w -\log(\cdot)$

Understand Softmax classifier

$$L_i = -\log \frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$$

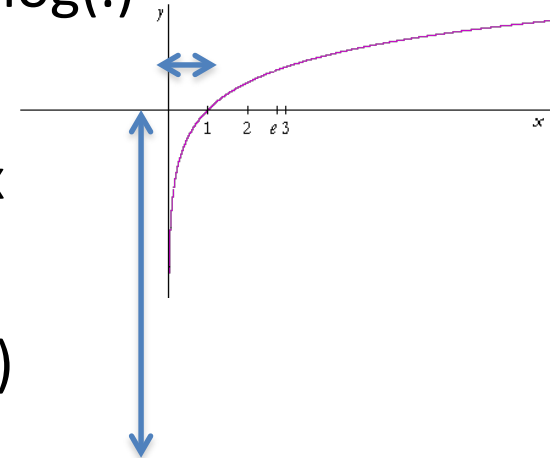
- Why min $-\log(\cdot)$ or max $\log(\cdot)$



- The term $\frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$ is the probability, between $[0,1]$


$\log(\cdot)$

- Stretch the numerical range during min/max
- Often used when working with probability
- $p_1 p_2$ is small: $\log(p_1 p_2) = \log(p_1) + \log(p_2)$
- Maximum Likelihood Estimation (MLE)
 - Minimize the negative log likelihood of the correct class



Understand Softmax classifier

$$L_i = -\log \frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$$

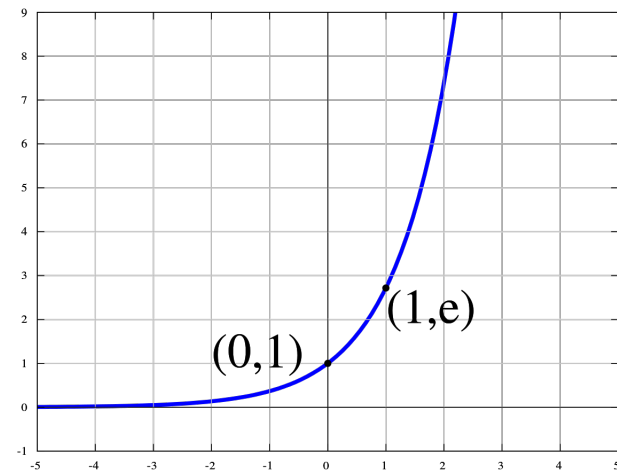
- Why exp before normalization? 
- Much higher confidence if the activation is large (clear images)

```
s = np.array([1,2])  
print(np.exp(s) / sum(np.exp(s)))
```

```
[ 0.26894142  0.73105858]
```

```
s = np.array([10,20])  
print(np.exp(s) / sum(np.exp(s)))
```

```
[ 4.53978687e-05  9.99954602e-01]
```



Understand Softmax classifier

$$L_i = -\log \frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$$

- Another interpretation of the cross entropy loss 
- Difference between:

- The model (estimated) prob Q: $\text{softmax}(f) = \frac{e^{f_m}}{\sum_{j=1}^K e^{f_j}}$
- The data (true) prob P: $[0, 0, \dots, 1, \dots, 0]$ (1 at the y_i -th position)
- Measured by Kullback-Leibler (KL) divergence ($D_{\text{KL}}=0$ when P, Q are “the same”)

$$D_{\text{KL}}(P\|Q) = -\sum_i P(i) \log \frac{Q(i)}{P(i)} \quad \text{Want Q to be close to P}$$

Understand Softmax classifier

$$L_i = -\log \frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$$

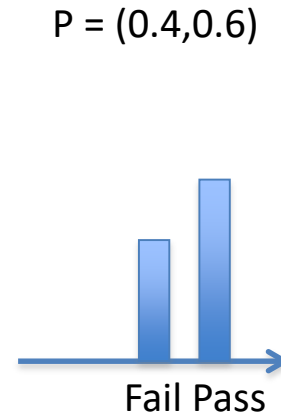
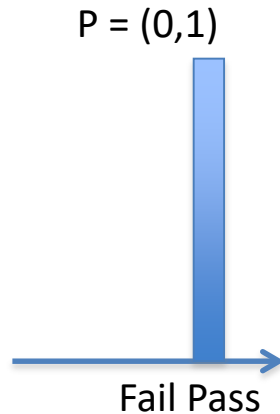
- Another interpretation of the cross entropy loss
- Difference between P and Q as measured by Kullback-Leibler (KL) divergence

$$\begin{aligned} D_{\text{KL}}(P\|Q) &= -\sum_x p(x) \log q(x) + \sum_x p(x) \log p(x) \\ &= H(P, Q) - H(P) \end{aligned}$$

Cross
entropy of P
and Q

Entropy of P
 $H(P) = 0$ in
this case

Probability, self information, entropy



Self information =

$$\log \frac{1}{P(i)}$$

Information theoretic entropy =
average amount of information an
observer would gain when sampling
a random variable

Rare event has more
information

Information \sim "surprise"

$$H(P) = \sum_i P(i) \log \frac{1}{P(i)}$$

bit if base 2; nat for base e

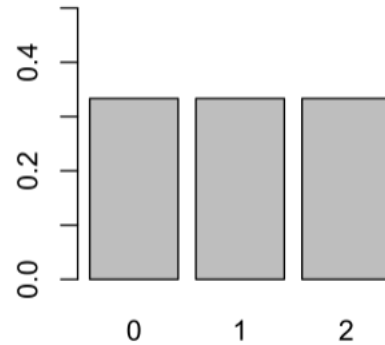
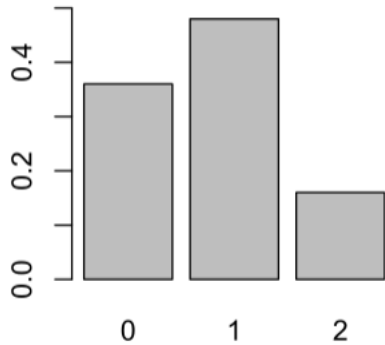
KL Divergence

$$D_{\text{KL}}(P||Q) = - \sum_i P(i) \log \frac{Q(i)}{P(i)} = \sum_i P(i) (\log \frac{1}{Q(i)} - \log \frac{1}{P(i)})$$

- Measure how one probability distribution is different from the other
- $D_{KL}(P||Q) \geq 0$; Equality holds iff $P=Q$ almost everywhere



KL Divergence



x	0	1	2
Distribution $P(x)$	0.36	0.48	0.16
Distribution $Q(x)$	0.333	0.333	0.333

$$\begin{aligned} D_{\text{KL}}(P \parallel Q) &= \sum_{x \in \mathcal{X}} P(x) \ln \left(\frac{P(x)}{Q(x)} \right) \\ &= 0.36 \ln \left(\frac{0.36}{0.333} \right) + 0.48 \ln \left(\frac{0.48}{0.333} \right) + 0.16 \ln \left(\frac{0.16}{0.333} \right) \\ &= 0.0852996 \end{aligned}$$

Is it true: $D_{KL}(P||Q) = D_{KL}(Q||P)$



Softmax classifier

- Additional regularization of W (L2 or L1 norm of weights)

$$R(W) = \sum_k \sum_l W_{k,l}^2 \quad R(W) = \sum_k \sum_l |W_{k,l}|$$

- Prefer small $W_{k,l}$, less likely to overfit the training dataset
- Regularize only W , not the bias b
- The entire loss for N training samples (x_i, y_i) : data loss and regularization loss

$$L = \frac{1}{N} \sum_i L_i + \lambda R(W)$$

- We determine W to minimize this loss given the training dataset

Multiclass SVM loss

- SVM loss: The correct class has a score higher than the incorrect class by some fixed margin d

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + d)$$

- For this test image i , **zero score** contributed by class j iff

$$0 \geq s_j - s_{y_i} + d$$

$$s_{y_i} \geq s_j + d$$

Correct class score s_{y_i} is higher than class j score s_j by at least d (otherwise, +ve. contribution of loss from class j)

Multiclass SVM loss

- $S = [13, -7, 11]$
- $y_i = 1$ For test image i
- $d = 10$

$$L_i = \max(0, -7 - 13 + 10) + \max(0, 11 - 13 + 10)$$
$$= 0 + 8$$

- Ground-truth score 13 is higher than -7 by more than the margin $d = 10$
- Ground-truth score 13 is not higher than 11 by $d = 10$

Train W so that the correct class y_i has a score higher than the incorrect classes by at least d

Linear Classifier

$$s = f(x; W, b) = Wx + b$$

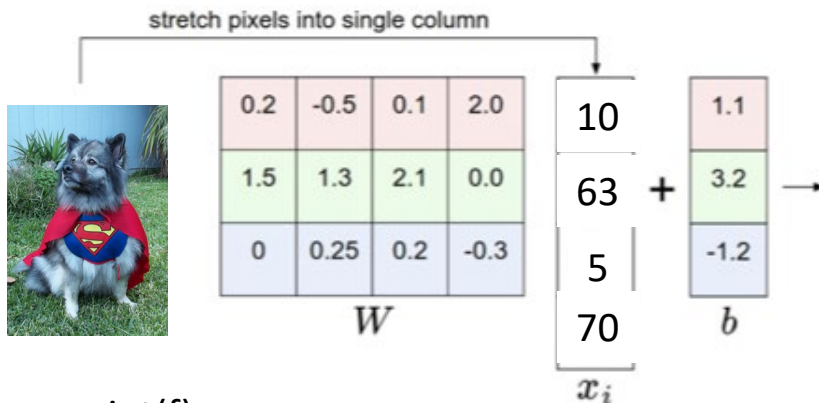
- After learning of the parameter W , do not need the training data in deployment
- Fast in deployment
- How to learn W ?

Softmax classifier

- **Cross-entropy loss** (apply $-\log(\cdot)$ to only the ground-truth class):

$$L_i = -\log \frac{e^{f_{y_i}}}{\sum_{j=1}^K e^{f_j}}$$

For the i -th training sample



Example: a dog (which does not look like a dog)

```
print(f)
print(np.exp(f))
print(np.exp(f) / sum(np.exp(f)))
[ 112.1   110.6   -5.45]
[ 4.83516636e+48  1.07887144e+48  4.29630469e-03]
[ 8.17574476e-01  1.82425524e-01  7.26458780e-52]
```

$y_i=2$

$L_i = -\log(0.182) = 1.7$

← Probability