

Sol. for selected questions

[20 points] Question 2: Inception module and GoogLeNet

What is the purpose of using 1x1 convolutions before 5x5 and 3x3 convolutions in the Inception module?

see "class_09" notes

Calculate the number of trainable parameters in the following inception module. There is no need to include the bias of the filter. **Show your calculation.**

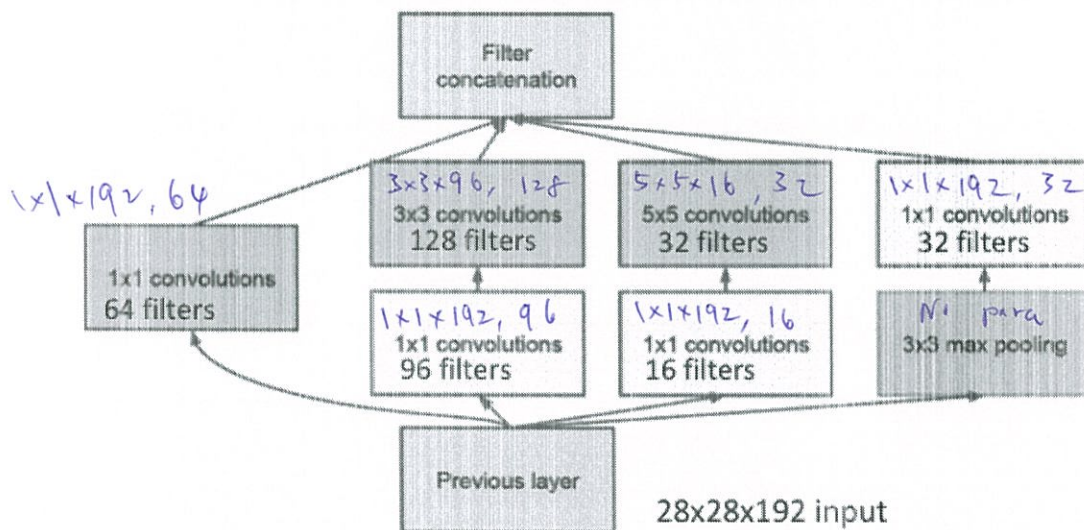
What is the number of channels after "Filter concatenation"?

Note:

1K = 1024

1M = 1024 * 1024

Recall that in an Inception module, all convolution and max pooling outputs have the same spatial dimension as their inputs. Moreover, max pooling uses stride=1.




$$\begin{aligned}
 \# \text{ trainable parameters} &= 1 \times 1 \times 192 \times 64 + 1 \times 1 \times 192 \times 96 + 1 \times 1 \times 192 \times 16 \\
 &\quad + 3 \times 3 \times 96 \times 128 + 5 \times 5 \times 16 \times 32 + 1 \times 1 \times 192 \times 32 \\
 &= 159.5k
 \end{aligned}$$

$$\begin{aligned}
 \# \text{ channels after filter concatenation} &= 64 + 128 + 32 + 32 \\
 &= 256
 \end{aligned}$$

[20 points]

c) The figure below shows the last several layers of a deep neural network trained with cross-entropy loss L .

Note that the first class is the ground-truth class for this training sample.



$$W = \begin{bmatrix} 5 & 4 & 3 \\ 3 & 4 & 2 \end{bmatrix}$$

Given the 3-by-1 activation vector: $h^{(l-2)} = [0.3, 0.2, 0.2]^T$ as input to the fully-connected layer, compute: (show your steps)


(i) L

(ii) $\frac{\partial L}{\partial W}$

(iii) $\frac{\partial L}{\partial h^{(l-2)}}$

(iv) The new W updated with gradient descent with a step size:

$$\gamma = 2$$



$$W = \begin{bmatrix} 5 & 4 & 3 \\ 3 & 4 & 2 \end{bmatrix}$$

(i) L :

$$h^{(l-1)} = W h^{(l-2)} = \begin{bmatrix} 5 & 4 & 3 \\ 3 & 4 & 2 \end{bmatrix} \begin{bmatrix} 0.3 \\ 0.2 \\ 0.2 \end{bmatrix} = \begin{bmatrix} 2.9 \\ 2.1 \end{bmatrix}$$

$$h^{(l)} = \text{softmax}(h^{(l-1)}) = \begin{bmatrix} \frac{e^{2.9}}{e^{2.9} + e^{2.1}} \\ \frac{e^{2.1}}{e^{2.9} + e^{2.1}} \end{bmatrix} = \begin{bmatrix} 0.68997 \\ 0.31003 \end{bmatrix}$$

$$L = -\log(0.68997) = 0.37110$$

$$(11) \frac{\partial L}{\partial w}: \quad \frac{\partial L}{\partial h^{(1)}} = \begin{bmatrix} -1 \\ 0.68997 \end{bmatrix}$$

$$\frac{\partial L}{\partial h^{(2)}} = \frac{\partial L}{\partial h^{(1)}} \frac{\partial h^{(1)}}{\partial h^{(2)}}$$

$$h^{(1)} = \begin{bmatrix} 0.68997 \\ 0.31003 \end{bmatrix} \quad (A)$$

From (A):

$$= \begin{bmatrix} -1 \\ 0.68997 \end{bmatrix} \begin{bmatrix} 0.68997(1-0.68997) & -0.68997(0.31003) \\ -0.31003(0.68997) & 0.31003(1-0.31003) \end{bmatrix}$$

$$= \begin{bmatrix} -(1-0.68997) & 0.31003 \end{bmatrix}$$

$$\frac{\partial L}{\partial h^{(2)}} = \begin{bmatrix} -0.31003 & 0.31003 \end{bmatrix}$$

$$\frac{\partial L}{\partial w_i} = \frac{\partial L}{\partial h_i^{(2)}} [h^{(2)}]^T \quad (B)$$

$$h^{(2)} = \begin{bmatrix} 0.3 \\ 0.2 \\ 0.2 \end{bmatrix}$$

From (B):

$$\frac{\partial L}{\partial w} = \begin{bmatrix} (-0.31003) [0.3 \quad 0.2 \quad 0.2] \\ (0.31003) [0.3 \quad 0.2 \quad 0.2] \end{bmatrix}$$

$$= \begin{bmatrix} -0.09301 & -0.06201 & -0.06201 \\ 0.09301 & 0.06201 & 0.06201 \end{bmatrix}$$

(iii)

$$\frac{\partial L}{\partial h^{(l-2)}} = \frac{\partial L}{\partial h^{(l-1)}} \cdot W$$

$$= [-0.31003, 0.31003] \begin{bmatrix} 5 & 4 & 3 \\ 3 & 4 & 2 \end{bmatrix}$$

$$= [-0.62005 \quad 0 \quad -0.31003]$$

(iv)

$$W' = W - \alpha \frac{\partial L}{\partial W}$$

$$= \begin{bmatrix} 5 & 4 & 3 \\ 3 & 4 & 2 \end{bmatrix} - 2 \begin{bmatrix} -0.09301 & -0.06201 & -0.06201 \\ 0.09301 & 0.06201 & 0.06201 \end{bmatrix}$$

$$= \begin{bmatrix} 5.186 & 4.124 & 3.124 \\ 2.814 & 3.576 & 1.876 \end{bmatrix}$$

[15 points]

Question 4: OpenCV and numpy

Given a 3 channel color square image:

```
img = cv2.imread('sample.jpg', cv2.IMREAD_COLOR)
img = cv2.resize(img, (100, 100))
```

Choose the best matches between OpenCV code and numpy code that perform the same functions. Briefly discuss the meaning of the corresponding codes.

OpenCV:

```
a/ cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
b/ cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
c/ cv2.warpAffine(img,
cv2.getRotationMatrix2D((W/2,H/2),90,1), (W,H)) # W, H
are width and height of the image
d/ cv2.getGaussianKernel(5, 1)
e/ cv2.threshold(img[:, :, 0], 40, 255,
cv2.THRESH_BINARY)
```

Numpy:

```
1/ img[:, :, np.array([0, 1, 2])] = img[:, :,
np.array([2, 1, 0])]

2/ img_new = np.zeros((img.shape[0], img.shape[1]),
dtype=np.uint8)
img_new[img[:, :, 0] > 40] = 255

3/ np.around(img[:, :, 0]*0.114 + img[:, :, 1]*0.587 +
img[:, :, 2]*0.299).astype(np.uint8)

4/ np.transpose(img, [1, 0, 2])

5/ x = np.array([-2, -1, 0, 1, 2])
scale = np.sum(np.exp(-np.power(x,2)/2))
np.exp(-np.power(x,2)/2)/scale
```

Answer:

OpenCV	Numpy	Meaning
a	3	Convert from color to grayscale
b	1	Convert BGR order to RGB
c	4	Rotate by 90 degrees
d	5	Obtain a 5x1 Gaussian kernel
e	2	If the first channel of img is larger than 40, it will become 255