


50.040

Natural Language Processing

Lu Wei



Instructor




Lu Wei
1.302.10
luwei@sutd.edu.sg
www.statnlp.org

NLP Research

 StatNLP

About People Publications Tools



Our Vision

Conduct fine research & Nurture world class researchers



NLP Research



Cheng Liying
PhD Student



Research Interests:

Text Generation



Guo ZhiJiang
PhD Student



Research Interests:

Syntactic Parsing
Semantic Parsing



Jie Zhanming
PhD Student



Research Interests:

Semantic Parsing
Named Entity Recognition



Nan Guoshun
Postdoc



Research Interests:

Named Entity Recognition
Relation Extraction
Question & Answering



Qiao Rui
Research Asst.



Research Interests:

Machine Learning
Structured Prediction



Zhuang Boyuan
Research Asst.



Research Interests:

Discourse Analysis
Implicit Discourse Analysis
Neural Network Interpretability



Joel Ong
PhD Student



Research Interests:

Natural Language Parsing
Machine Learning in Finance
Reinforcement Learning



Li Hao
PhD Student



Research Interests:

Machine Learning
Artificial Intelligence
Natural Language Processing



Li Haoran
PhD Student



Research Interests:

Machine Learning
Graph Neural Networks
Natural Language Processing



Liao Jing
Visiting Student



Research Interests:

Sentiment Analysis
Intelligent Design
Human Centered Computing



Ivan Sekulic
Research Intern



Research Interests:

Machine Learning
Natural Language Processing



Wang Huiming
Research Intern



Research Interests:

Named Entity Recognition
Semantic Parsing
Machine Learning



Perry Lam
PhD Student



Research Interests:

Speech Generation
Conversational Information Extraction
Computational Psycholinguistics



Sun Xiaobing
PhD Student



Research Interests:

Target Sentiment Analysis
Deep Learning Interpretability



Thilini
PhD Student



Research Interests:

Visual Language Tasks
Visual Question & Answering



Zhou Yuxiang
Visiting Student



Research Interests:

Text Classification
Sentiment Analysis
Information Extraction



Cheng Juan
Postdoc



Research Interests:

Sentiment Analysis
Relation Extraction
Deep Learning Interpretability



Xu Lu
PhD Student



Research Interests:

Graph Neural Network
Low Resource Task
Relation & Information Extraction



Zhang Yan
PhD Student



Research Interests:

Graph Neural Networks
Knowledge Representation & Reasoning



Zou Yanyan
PhD Student



Research Interests:

Natural Language Understanding
Semantic Parsing
Text Summarization



Lu Wei
Associate Professor



Research Interests:

Machine Learning
Artificial Intelligence
Natural Language Processing

Major NLP Conferences



Course Plan



**50.040 Natural Language Processing
(Undergraduate)**
Summer 2020

Instructor	Lu Wei	luwei@sutd.edu.sg	1.302.10
Teaching Assistant	Haoran Li	haoran2_li@mymail.sutd.edu.sg	1.417
Office Hours	Lu Wei Haoran Li	Thursday 15:30 - 16:30pm, or by appointment Friday 14:00 - 15:00pm, or by appointment	1.302.10 1.417
Lessons	Every Week	Mon 10:30 – 12:30 Thu 13:30 – 15:30	
Classroom	Think Tank 24 (2.503)		
Prerequisites	<ol style="list-style-type: none">01.112 Machine Learning or 01.113 Statistical and Machine Learning andA good foundation in: 1) programming, 2) design and analysis of algorithms, 3) mathematics including linear algebra, calculus, optimization, probability, and statistics.		
Assessment	Homework (coding) (30%), Homework (written) (20%), Mini Project (20%), Final Project (30%)		

Prerequisites

Machine Learning

Algorithms

Deep Learning

Linear Algebra

Prob & Stats

Programming

Optimization

Assessment




Item	Weightage	Remark
Homework	30%	Biweekly
Mini Project	20%	Individual
Written Homework	20%	6-weekly
Final Project	30%	Group (2-3)

Learning Objectives

By the end of the course, students will be able to:

1. Explain the fundamental tasks within NLP
2. Explain possible algorithms as solutions to NLP tasks
3. Implement the algorithms used for various NLP tasks
4. Design novel algorithms for solving new NLP tasks,
and use existing NLP technologies for solving real
problems

Textbooks




Honour Code

You are encouraged to form study groups to work on the homework together. However, please write out the solutions on your own, without referring to notes from other students, online solutions or answers from past courses. Reproducing the solution from scratch will help you to find out if you fully understand the material. Please also list your group members on the first page of your submission.


What is Natural Language Processing

Tasks in NLP




POS Tagging
Chunking
Document Classification
Information Extraction
Syntactic Parsing
Semantic Parsing
Natural Language Generation
Machine Translation
Sentiment Analysis
Coreference Resolution
Question Answering

Tasks in NLP



Supervised



Unsupervised

Part-of-Speech Tagging

A	N	V	D	N
<i>Fruit</i>	<i>flies</i>	<i>like</i>	<i>a</i>	<i>banana</i>

Noun-Phrase Chunking

NP

Fruit

flies


like

a


banana

NP

Constituency Parsing



Dependency Parsing



Semantic Parsing

LIKE(F102, B87)



Fruit


flies

like

a

banana

Semantic Parsing



Fruit

flies

like

a

banana

Sentiment Analysis



(**neutral**)

Fruit

flies


like

a

banana

(**positive**)


Tasks in NLP



Supervised

Many (Supervised)
Problems in NLP are
Structured Prediction
Problems!

Tasks in NLP



Supervised

Question

How to Build a Document Classification Model?

Word Representation

cat	eat	dog
$\begin{pmatrix} 0 \\ \textcolor{red}{1} \\ 0 \\ 0 \\ 0 \\ \cdot \\ \cdot \\ \cdot \\ 0 \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ 0 \\ \textcolor{red}{1} \\ 0 \\ 0 \\ \cdot \\ \cdot \\ \cdot \\ 0 \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ 0 \\ 0 \\ \textcolor{red}{1} \\ 0 \\ \cdot \\ \cdot \\ \cdot \\ 0 \\ 0 \end{pmatrix}$



1-hot representations are sparse, which cannot capture
semantic relations between words!



Need a way to capture *semantics*!


WordNet

<https://wordnet.princeton.edu>



WordNet

A Lexical Database for English



WordNet

Limitations:
Only for English
Incomplete
Subjective
Expensive to Maintain




Need a way to capture *semantics automatically!*

Question

How to Automatically Group Similar Data Points?



Tasks in NLP



Unsupervised

Word Embeddings

cat

$$\begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ \cdot \\ \cdot \\ \cdot \\ 0 \\ 0 \end{pmatrix}$$

eat


$$\begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ \cdot \\ \cdot \\ \cdot \\ 0 \\ 0 \end{pmatrix}$$

dog


$$\begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ \cdot \\ \cdot \\ \cdot \\ 0 \\ 0 \end{pmatrix}$$

Word Embeddings

$$\begin{array}{lll} \text{cat} & \text{eat} & \text{dog} \\ \left(\begin{array}{c} 0 \\ 1 \\ 0.232 \\ -0.903 \\ 0.213 \\ 0.239 \\ -0.679 \\ 0.923 \\ 0 \\ 0 \end{array} \right) & \left(\begin{array}{c} 0 \\ 0 \\ -0.932 \\ -0.903 \\ 0.213 \\ 0.590 \\ 0.609 \\ 0.388 \\ 0 \\ 0 \end{array} \right) & \left(\begin{array}{c} 0 \\ 0 \\ 0.190 \\ -0.887 \\ 0.193 \\ 0.257 \\ -0.708 \\ 0.906 \\ 0 \\ 0 \end{array} \right) \end{array}$$



Tasks in NLP




Supervised




Unsupervised

Tasks in NLP



Word Embeddings



Maps words in a vocabulary to vectors of continuous real numbers, and these vectors are in a lower-dimensional space


Question

How to Learn Word Embeddings?

Distributional Semantics

You shall know a word by the company it keeps


- Firth, J. R. 1957



Distributional Semantics

You shall know a word by the company it keeps

- Firth, J. R. 1957



Two Major Approaches to Learning Word Embeddings

Count-based

Prediction-based

Co-Occurrence Matrix

	cat		eat		dog		
cat	0	2	1	0	0	0	0
eat	2	0	0	1	0	1	0
dog	1	0	0	0	0	0	1
	0	1	0	0	1	0	0
eat	0	0	0	1	0	0	1
	0	1	0	0	0	0	1
dog	0	0	1	0	0	0	1
	0	0	0	0	1	1	0



Co-Occurrence Matrix

	cat			eat		dog	
cat	0	2	1	0	0	0	0
	2	0	0	1	0	1	0
	1	0	0	0	0	0	1
	0	1	0	0	1	0	0
eat	0	0	0	1	0	0	1
	0	1	0	0	0	0	1
dog	0	0	1	0	0	0	1
	0	0	0	0	1	1	0

Co-Occurrence Matrix

	cat			eat		dog	
cat	0	2	1	0	0	0	0
	2	0	0	1	0	1	0
	1	0	0	0	0	0	1
	0	1	0	0	1	0	0
eat	0	0	0	1	0	0	1
	0	1	0	0	0	0	1
dog	0	0	1	0	0	0	1
	0	0	0	0	1	1	0

Still very **high dimensional!** Still very sparse!

Dimensionality reduction: for example - SVD!

Co-Occurrence Matrix

$$\begin{matrix} & \text{cat} & & \text{eat} & & \text{dog} \\ \text{cat} & \left[\begin{array}{ccccc} 0 & 2 & 1 & 0 & 0 \\ 2 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 0.232 & 0 & -0.932 & 0 \\ 0 & -0.903 & 0 & -0.903 & 0 \\ 0 & 0.213 & 0 & 0.213 & 0 \\ 0 & 0.239 & 0 & 0.590 & 0 \\ 0 & -0.679 & 0 & 0.609 & 0 \\ 0 & 0.923 & 0 & 0.388 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & 1 \end{array} \right] \end{matrix}$$

Matrix decomposition / dimension reduction methods
will project them into a lower-dimensional space

Alternative Matrices


	cat		eat		dog		
cat	0	2	1	0	0	0	0
	2	0	0	1	0	1	0
	1	0	0	0	0	0	1
	0	1	0	0	1	0	0
eat	0	0	0	1	0	0	0
	0	0	0	1	0	0	1
	0	1	0	0	0	0	0
dog	0	0	1	0	0	0	1
	0	0	0	0	1	1	1

Binary Coccurrence Matrix, TF-IDF Matrix, PMI Matrix, Shifted PMI Matrix ...

Distributional Semantics

You shall know a word by the company it keeps

- Firth, J. R. 1957



Two Major Approaches to Learning Word Embeddings

Count-based



Prediction-based

Distributional Semantics

Count
based

Prediction
based

- Compute the statistics of how often each word co-occurs/interacts with its neighboring words in a large text corpus.
- Learn to map these count-statistics down to a small, dense vector for each word.

Two Approaches


Count
based

Prediction
based

- Directly predict a word from its neighboring words.
- This process will yield dense word embeddings

Word2Vec


Tomas Mikolov et.(2013)



CBOW

Skip-gram

Word2Vec




1. Project each context word from $|V|$ dimensional space to its (current) embedding space
2. Sum these vectors
3. Project it to another space of dimension $|V|$ (we expect this vector to be similar to the 1-hot vector of the target word)

Where are the embeddings?



Word2Vec




Word2Vec




INPUT PROJECTION OUTPUT

Values on these edges form
a low-dimensional vector




Word2Vec


Some early interesting properties found in word embeddings



Male-Female




Verb tense



Country-Capital

Tasks in NLP



Supervised

Unsupervised