

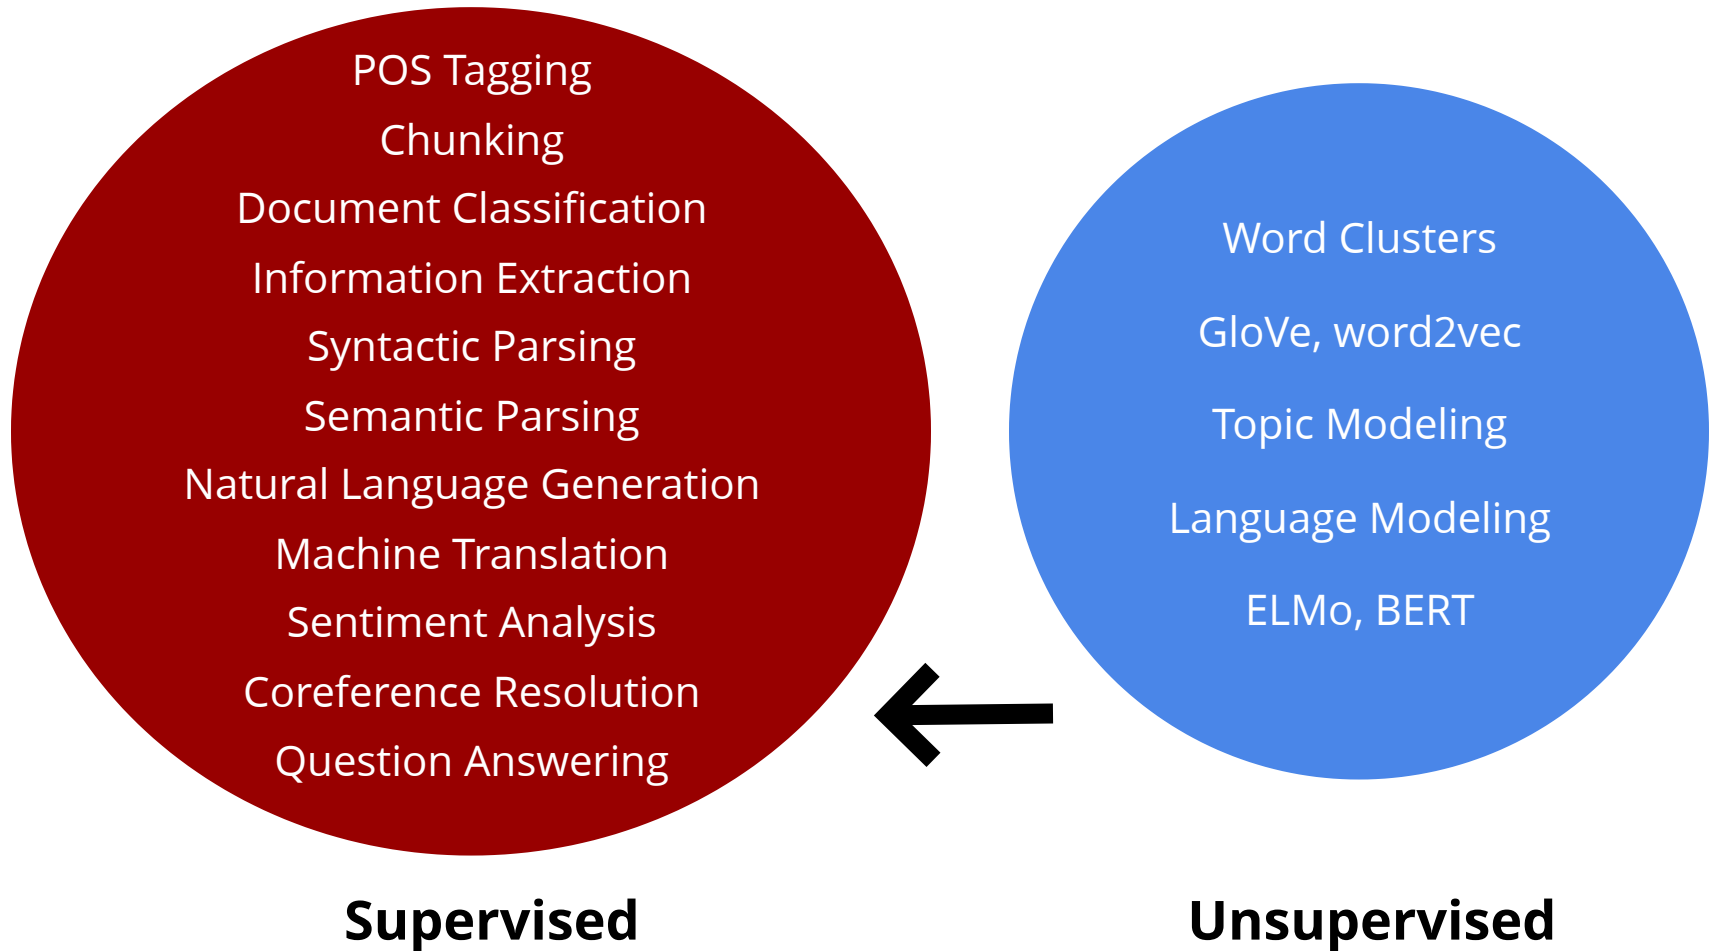
50.040

Natural Language Processing

Lu, Wei



Tasks in NLP




Unsupervised Learning




Unsupervised

Word Embeddings




CBOW

How about Sentences

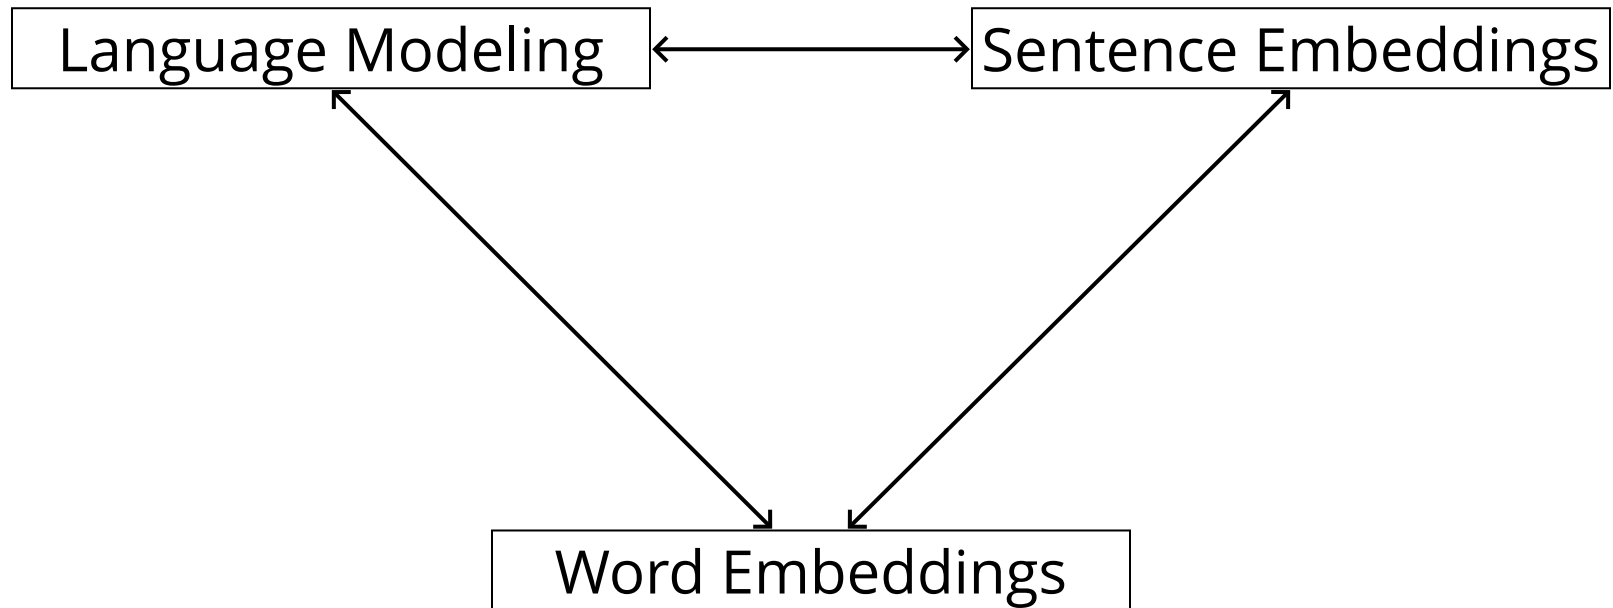


How likely can we see this sentence?



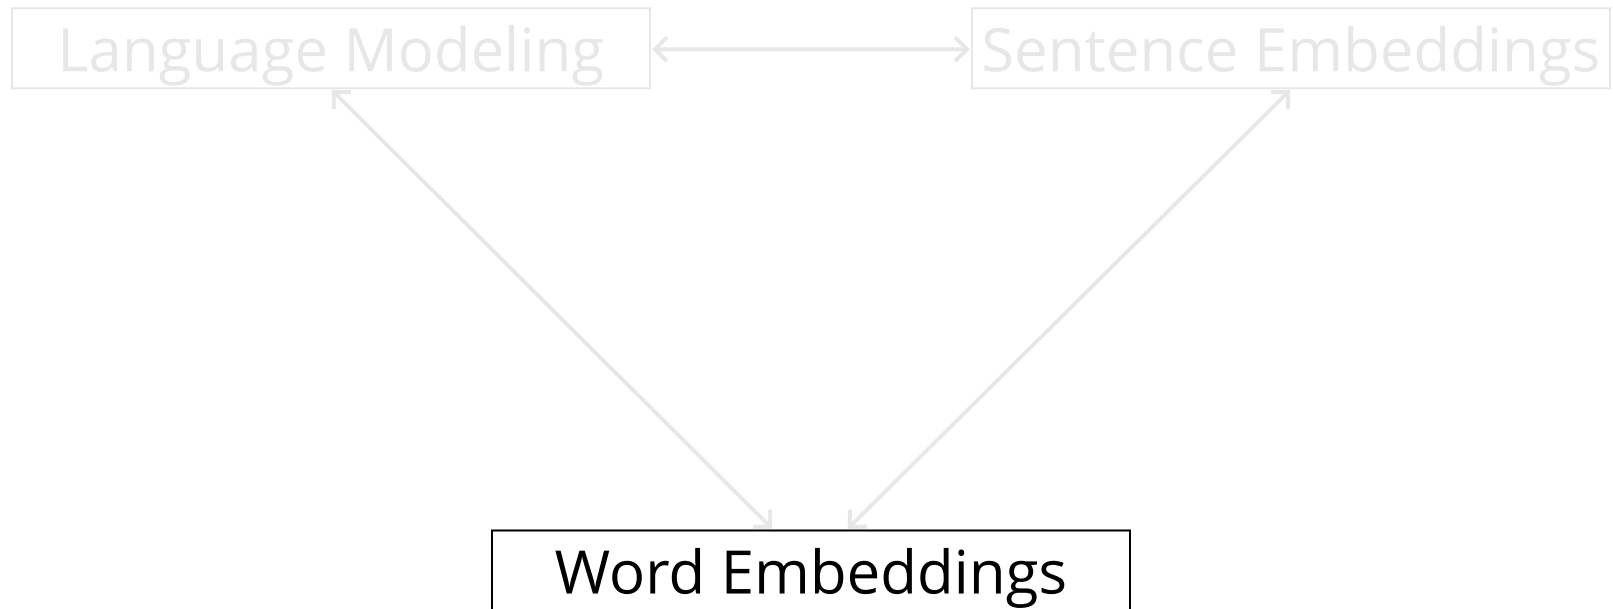
How to represent this sentence?

Three Tasks



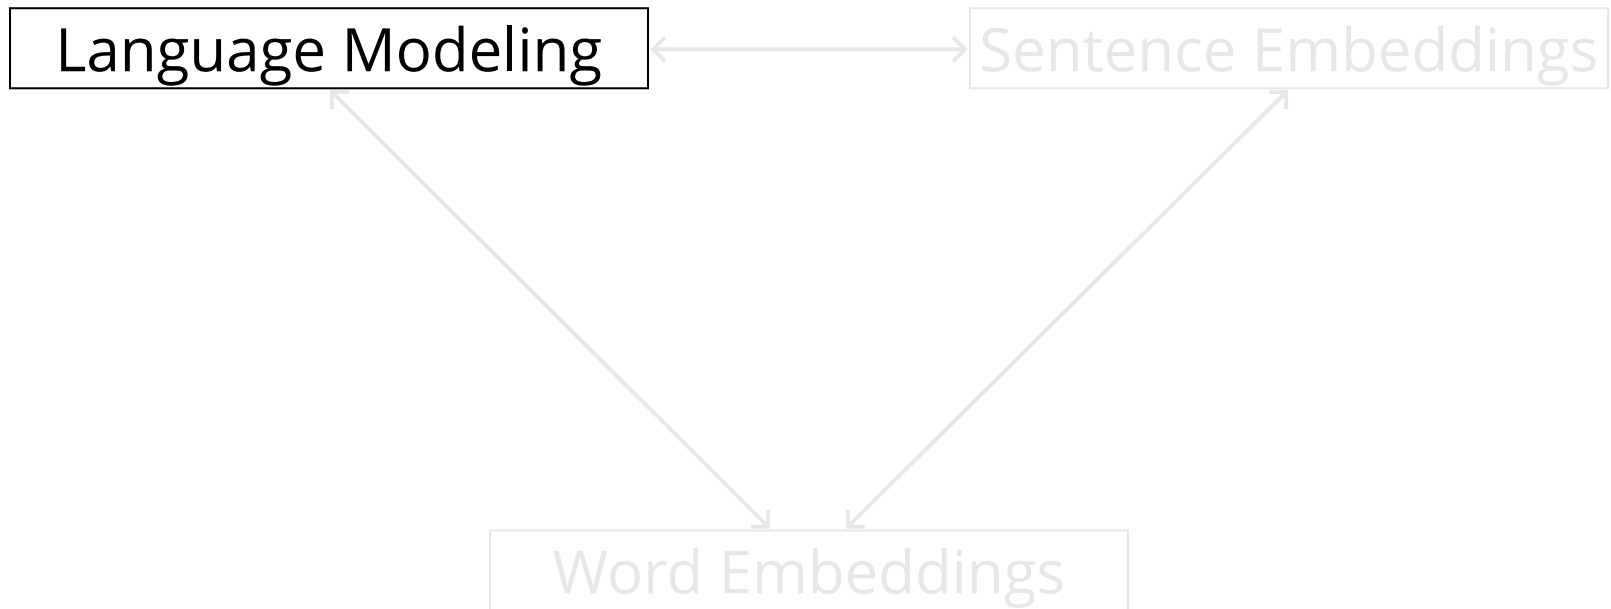
These three tasks are closely related!

Three Tasks



We have already looked at word embeddings

Three Tasks



Now let us look at language modeling

Language Modeling

Language Modeling

How likely can we see this sentence?

Fruit flies like a banana → 0.00078

I love NLP → 0.00428

Fruit flies → ????

Language Modeling

First of all, similar to HMM, we assume each sentence is attached with a special symbol at its end: **STOP**

Fruit flies like a banana STOP

I love NLP STOP

Fruit flies STOP

Language Modeling

A language model consists of a finite set V , and a function $p(x_1, x_2, \dots, x_m)$ such that:

1. For any sequence $\langle x_1, \dots, x_m \rangle \in V^+$

$$p(x_1, \dots, x_m) \geq 0$$

2. In addition,


$$\sum_{\langle x_1, \dots, x_m \rangle \in V^+} p(x_1, x_2, \dots, x_m) = 1$$

Hence, $p(x_1, \dots, x_m)$ is a probability distribution over the sentences defined by V^+ .



Question

How to Learn a Language
Model Based on a Corpus?



Language Model

A simple Maximum Likelihood Estimator gives:

$$p(x_1, \dots, x_m) = \frac{\text{count}(x_1, \dots, x_m)}{\sum_{s \in \mathcal{D}} \text{count}(s)}$$

Is this feasible?

Language Model

A simple Maximum Likelihood Estimator gives:

$$p(x_1, \dots, x_m) = \frac{\text{count}(x_1, \dots, x_m)}{\sum_{s \in \mathcal{D}} \text{count}(s)}$$

Is this feasible?

NO! The training set does not contain all possible sentences! It does not generalize to new sentences!

Alternative Approach

Recall what we did in a Generative Model?

$$p(x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} p(x_i | x_1, \dots, x_{i-1})$$

Is this feasible?

Alternative Approach

Recall what we did in a Generative Model?

$$p(x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} p(x_i | x_1, \dots, x_{i-1})$$

Is this feasible?

NO! The sequence that we condition on may only appear a few times in the training set. Poor generalization again!

Markov Assumption

Recall what we did in Naive Bayes / HMM?

$$p(x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} p(x_i | x_1, \dots, x_{i-1})$$

In other words, we only consider the previous $(n - 1)$ words.

n-Gram Language Model

$$p(x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} p(x_i | x_{i-n+1}, \dots, x_{i-1})$$

Bigram Language Model

$$n = 2$$

$$p(x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} p(x_i | x_{i-1})$$

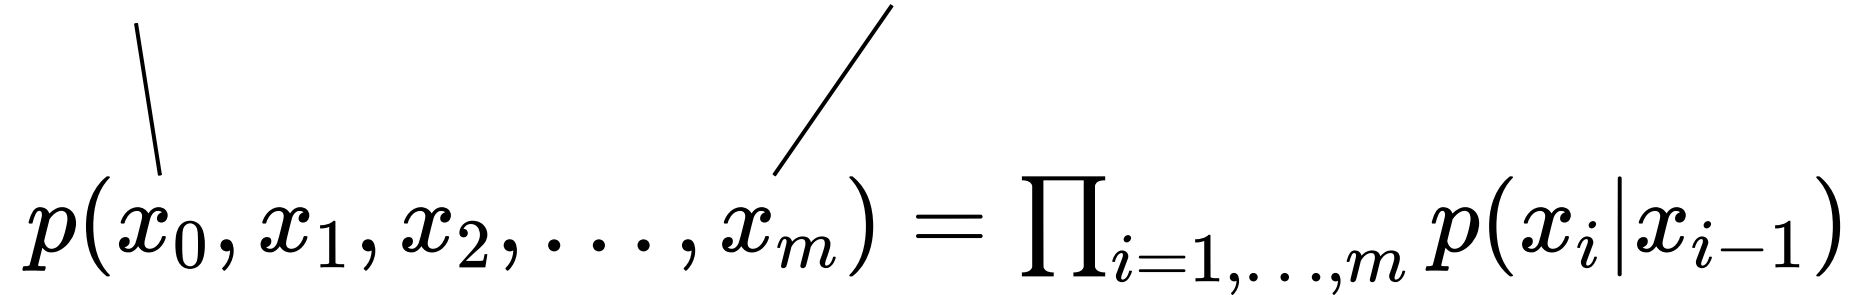
Similar to what we did for HMM, we may introduce

$$x_0 = \text{START}$$

Bigram Language Model

START

STOP



The diagram shows the words "START" and "STOP" in a serif font. A thin black line extends from the bottom of "START" down to the x_0 term in the formula. Another thin black line extends from the bottom of "STOP" down to the x_m term in the formula.

$$p(x_0, x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} p(x_i | x_{i-1})$$

Question

How to Learn a Bigram
Language Model?

$$p(x_0, x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} p(x_i | x_{i-1})$$

Bigram Model

How to do parameter estimation?

$$p(x_0, x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} \underbrace{p(x_i | x_{i-1})}_{\text{These are the Model Parameters}}$$

These are the Model
Parameters

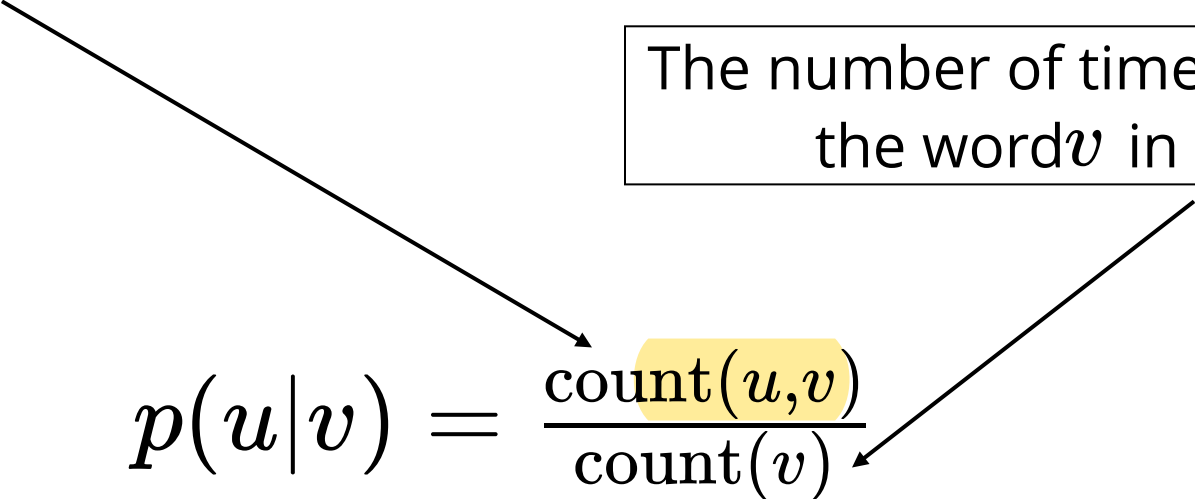
Objective:

$$\mathcal{L} = \prod_{\langle x_1, \dots, x_m \rangle \in \mathcal{D}} p(x_0, x_1, \dots, x_m) \quad \text{💬}$$

Bigram Model

The number of times we see the word u is followed by the word v in \mathcal{D} .


The number of times we see the word v in \mathcal{D} .


$$p(u|v) = \frac{\text{count}(u,v)}{\text{count}(v)}$$

for all $u \in V, v \in V \cup \{\mathbf{START}\}$

Trigram Model


START STOP


$$p(x_{-1}, x_0, x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} p(x_i | x_{i-2}, x_{i-1})$$

$$p(u | w, v) = \frac{\text{count}(w, v, u)}{\text{count}(w, v)}, \text{ for all } u \in V, w, v \in V \cup \{\text{START}\}$$

Unigram Model

STOP


$$p(x_1, x_2, \dots, x_m) = \prod_{i=1, \dots, m} p(x_i)$$

$$p(u) = \frac{\text{count}(u)}{c}, \text{ for all } u \in V.$$



The total number of words in the corpus \mathcal{D} .


***n*-Grams**

Are they useful?

n -Gram Model


“It must be recognized that the notion ‘probability of a sentence’ is an entirely useless one, under any known interpretation of this term.”

- Noam Chomsky (1967)



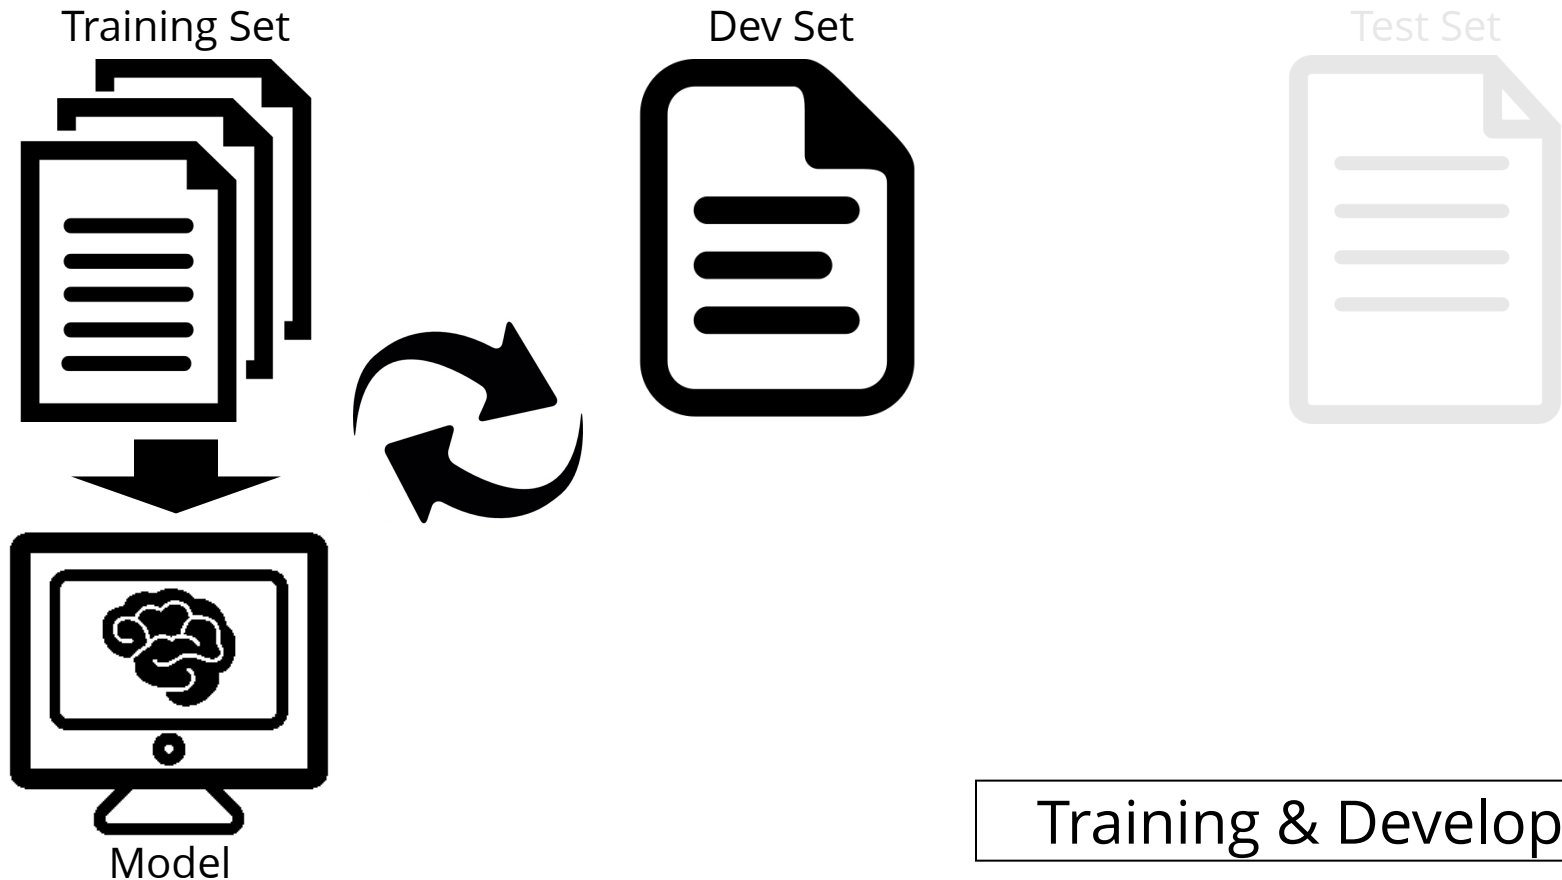
n -Gram Model

“Every time I fire a linguist, the performance
of the recognition system goes up”
- Fred Jelinek (1988)



Language Model Evaluation

How do we evaluate a model?



Language Model Evaluation

How do we evaluate a model?



Evaluation Metric

One idea: the likelihood of the entire test/dev set \mathcal{D}' !

$$\prod_{j=1}^{|\mathcal{D}'|} p(\mathbf{x}^{(j)})$$



$$\log_2 \prod_{j=1}^{|\mathcal{D}'|} p(\mathbf{x}^{(j)})$$

$$\sum_{j=1}^{|\mathcal{D}'|} \log_2 p(\mathbf{x}^{(j)})$$

Perplexity

$$\sum_{j=1}^{|\mathcal{D}'|} \log_2 p(\mathbf{x}^{(j)})$$

$$\frac{1}{c'} \sum_{j=1}^{|\mathcal{D}'|} \log_2 p(\mathbf{x}^{(j)})$$

Total number of words in the set \mathcal{D}' .



$$2^{-\ell} \text{ where } \ell = \frac{1}{c'} \sum_{j=1}^{|\mathcal{D}'|} \log_2 p(\mathbf{x}^{(j)})$$

Perplexity

Perplexity

$$2^{-\ell} \text{ where } \ell = \frac{1}{c'} \sum_{j=1}^{|\mathcal{D}'|} \log_2 p(\mathbf{x}^{(j)})$$



$$1 / \sqrt[c']{\prod_{j=1}^{|\mathcal{D}'|} p(\mathbf{x}^{(j)})}$$



c' terms after
expansions

Multiplicative inverse of the geometric mean of the terms $p(x_k^{(j)} | x_{k-2}^{(j)}, x_{k-1}^{(j)})$.

Trigram Model

$$p(u|w, v) = \frac{\text{count}(w, v, u)}{\text{count}(w, v)}$$

What are the potential problems with such a model (or in general, an n -gram model)?

Trigram Model

$$p(u|w, v) = \frac{\text{count}(w, v, u)}{\text{count}(w, v)}$$

What are the potential problems with such a model (or in general, an n -gram model)?

Many counts could be zeros!

Trigram Model

$$p(u|w, v) = \frac{\text{count}(w, v, u)}{\text{count}(w, v)}$$

What are the potential problems with such a model (or in general, an n -gram model)?

Many counts could be zeros!

One solution:

Smoothing

Smoothing: Interpolation

$$p(u|w, v) = \frac{\text{count}(w, v, u)}{\text{count}(w, v)}$$

$$p(u|v) = \frac{\text{count}(v, u)}{\text{count}(v)}$$



Smoothed probability

$$p(u) = \frac{\text{count}(u)}{c}$$

$$\boxed{q(u|w, v)} = \lambda_1 p(u) + \lambda_2 p(u|v) + \lambda_3 p(u|w, v)$$

$$\boxed{\lambda_1 + \lambda_2 + \lambda_3} = 1$$

hyperparameters

Smoothing: Interpolation

$$q(u|w, v) = \lambda_1 p(u) + \lambda_2 p(u|v) + \lambda_3 p(u|w, v)$$

$$\lambda_1 + \lambda_2 + \lambda_3 = 1$$

Hyperparameters will be tuned on the development set.

Each is indicating the **significance** / contribution / confidence on the corresponding model.

Laplace Smoothing

"Add-one" Smoothing

$$p(u|w, v) = \frac{\text{count}(w, v, u)}{\text{count}(w, v)}$$

$$p(u|w, v) = \frac{\text{count}(w, v, u) + 1}{\text{count}(w, v) + U}$$

What is U ? The U is chosen such that the sum of these p terms will be 1!

In other words, $U = |V|$



Smoothing

Other smoothing techniques exist:

- Good-turing smoothing
- Kneser-Ney smoothing
- Witten-Bell smoothing
- Katz smoothing
- Church and Gale smoothing



Read relevant book chapters to learn more.

Trigram Model

$$p(u|w, v) = \frac{\text{count}(w, v, u)}{\text{count}(w, v)}$$

What are the potential problems with such a model (or in general, an n -gram model)?

Many counts could be zeros!

One solution:

Smoothing



Question

After smoothing, how many
model parameters do we
have to store?

Curse of Dimensionality

in the order of $|V|^n$

Too many model parameters!


One Solution:

Learn language model with word embeddings!

Neural Language Model

"The model learns simultaneously (1) a distributed representation for each word along with (2) the probability function for word sequences, expressed in terms of these representations."

- Yoshua Bengio et al. (2003)




Neural Language Model


Bengio et al. (2003)

1. associate with each word in the vocabulary a distributed *word embedding*,
2. express the joint probability function of word sequences in terms of the embeddings of these words in the sequence, and
3. learn simultaneously the *word embeddings* and the parameters of that *probability function*.

Neural Language Model

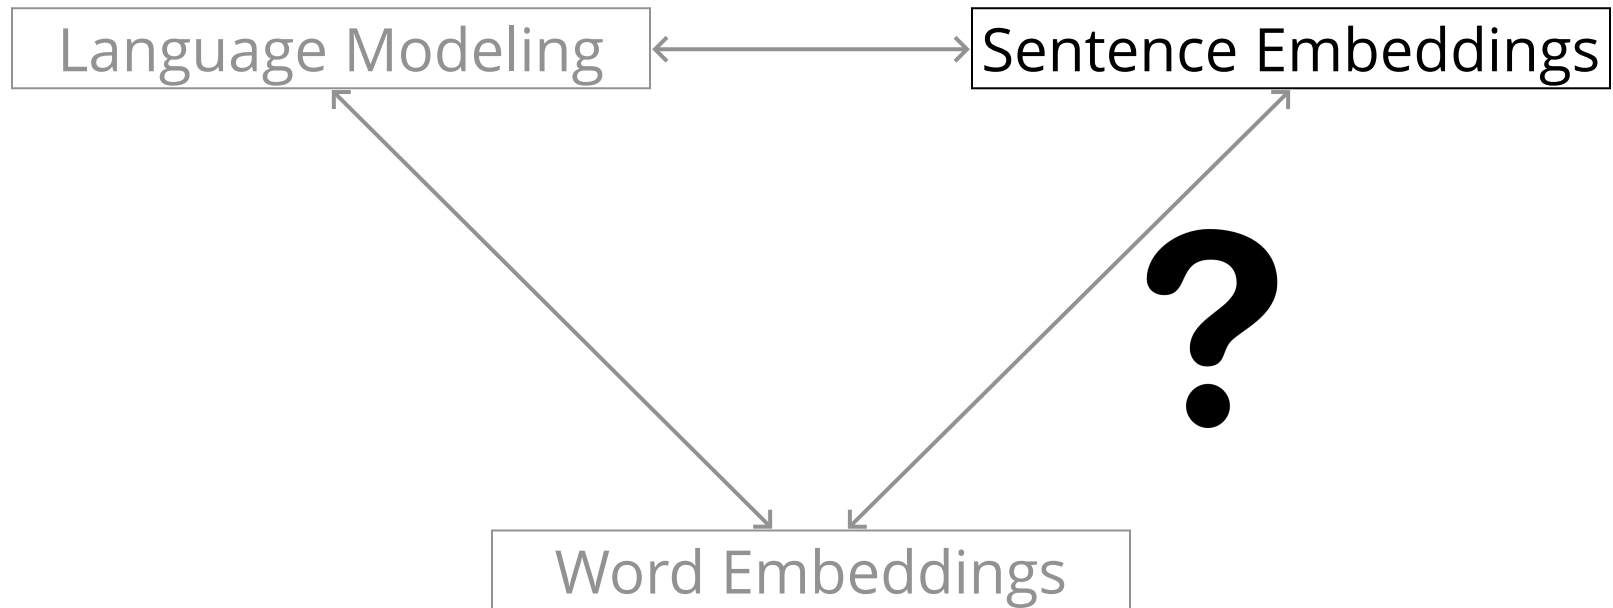


Three Tasks



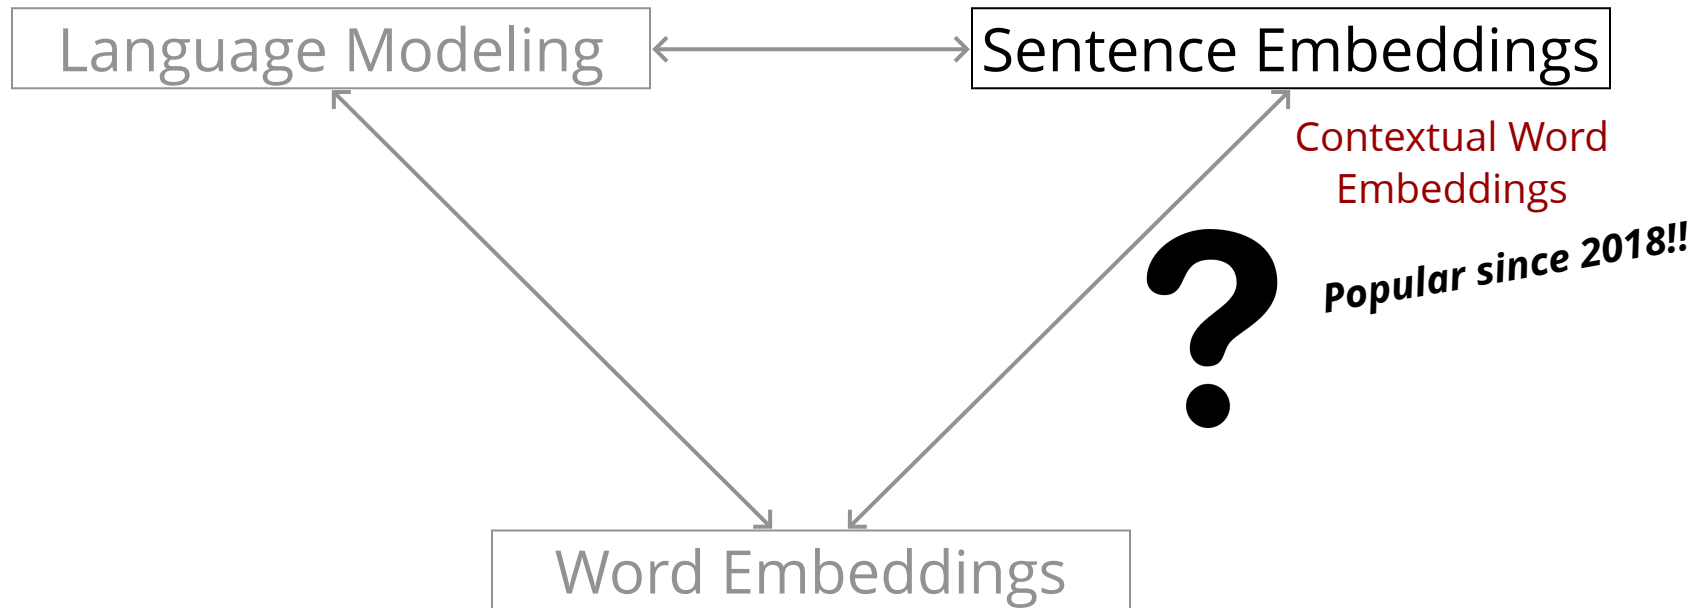
These three tasks are closely related!

Three Tasks



These three tasks are closely related!

Three Tasks



These three tasks are closely related!