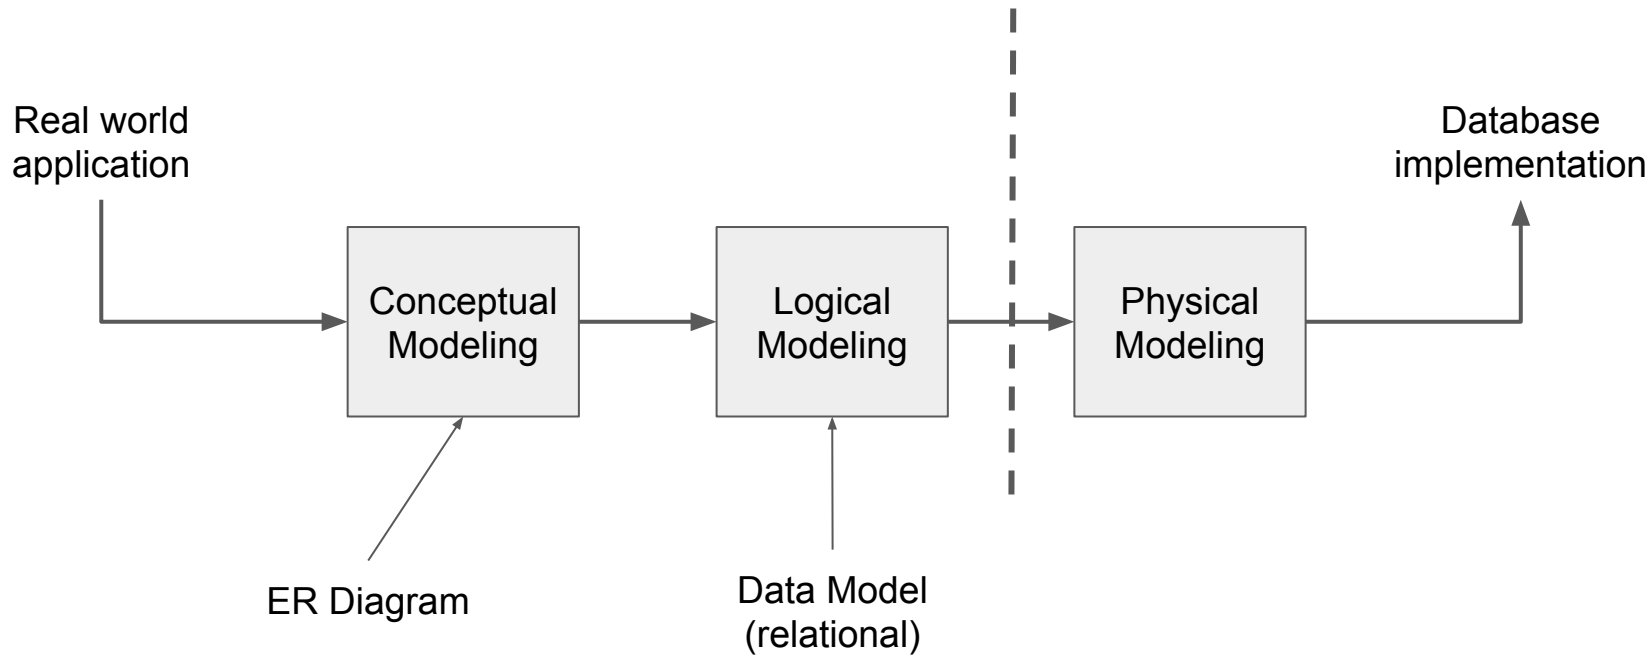



Databases and Big Data

SQL

Recap



Data Model



Relational Algebra

What data
I want

SQL

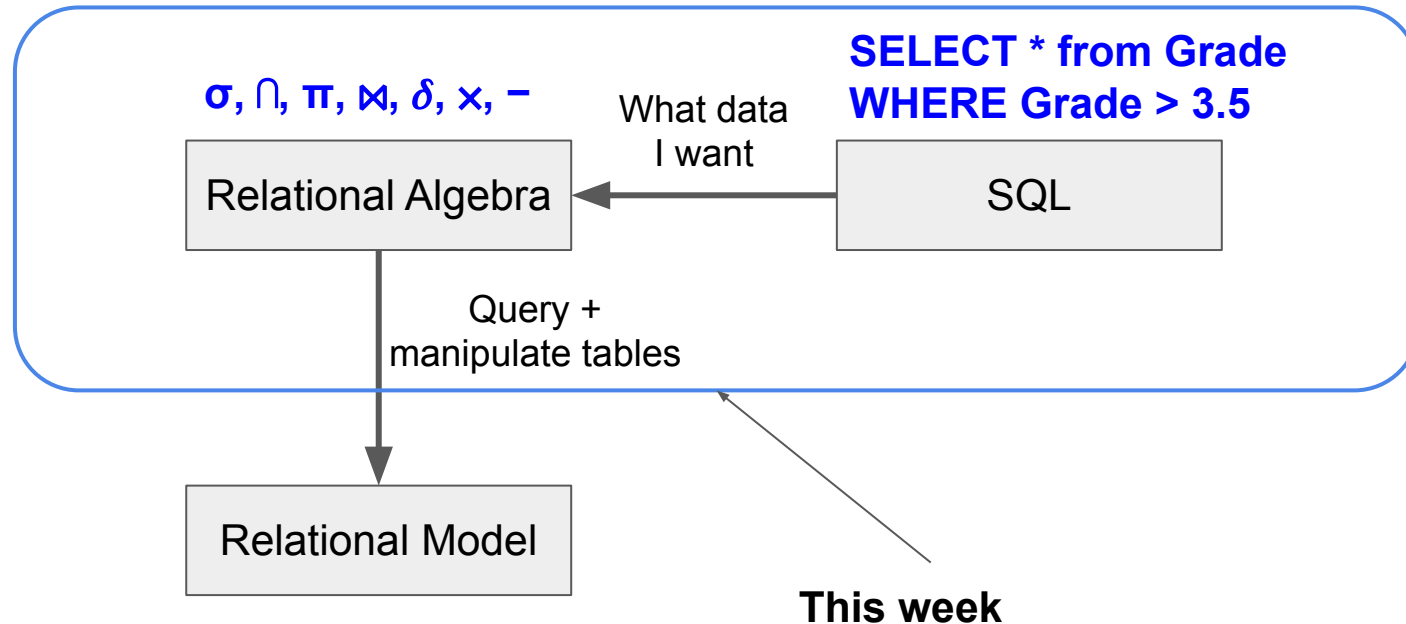
Query +
manipulate tables

Relational Model

Relational Model of Uni-DB

Professor				Student				Lecture			
Person-ID	Name	Level	Room	Person-ID	Name	Semester	Year	No	Title	CP	People
2125	Mayer	PP	225	20002	Gierke	18		5001	Databases	4	2127
2126	Kossmann	PP	232	25403	Zallinger	12		5041	Networks	4	2125
2127	Roscoe	AP	310	26120	Frey	10		5043	Operating Systems	3	2126
2133	Petrig	AP	52	24820	Kling	8		5049	Programming	2	2125
2134	Spatene	AP	309	27150	Petri	6		4052	Architecture	4	2125
2136	Wahl	PP	30	28106	Luttenberger	3		5012	Theory	3	2126
2137	Natter	PP	7	29120	Schweiser	2		5216	Graphics	2	2126
				29555	Meyer	2		5259	Distributed Systems	2	2133
Requires				Attends				Assistant			
Professor	Follow-up			Leg	No			AcadId	Name	Room	Class
5001	5041			26120	5001			3002	Harris	Databases	2125
5001	5041			27550	5001			2100	Heller	Theory	2125
5041	5049			28106	5041			3004	Kossmann	Networks	2126
5041	5052			28106	5052			3005	Frey	Graphics	2127
5041	5052			28106	5216			2106	Petri	Operating Systems	2127
5052	5259			28106	5259			2007	Kroska	Formal Methods	2126
Lectures				Lectures				Lectures			
Leg	No	People	Grade	Leg	No	People	Grade	Leg	No	People	Grade
28106	5001	2120	1	26120	5041			26120	5049		
25403	5041	2120	2	29555	5022			25403	5022		
27550	4630	2127	2	25403	5022						

Data Model



Example

Payroll

UserID	Name	Job	Salary
123	Alice	Asst. Prof	100
456	Bob	TA	80
789	Carol	Asst. Prof	120
101	David	Prof	150


Give me
salaries of all
Asst. Prof

Here

UserID	Salary
123	100
789	120



Example



SQL query

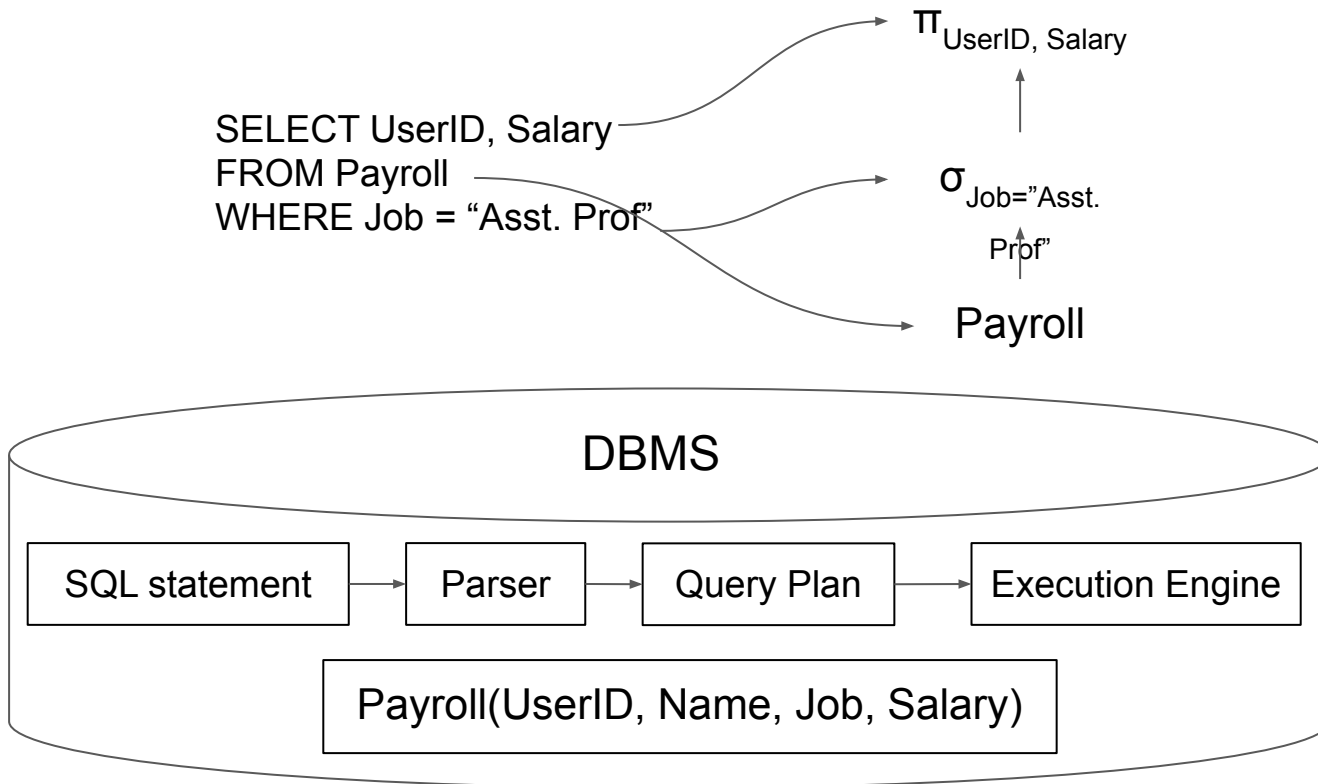
```
SELECT UserID, Salary  
FROM Payroll  
WHERE Job = "Asst. Prof"
```



Example

Query Plan



Acts as an algorithm for the DBMS to carry out the operation.




Example

*Goes to each row in Payroll table, if the column Job satisfies some condition, then you output the values you want.
(This algorithm is a simplified version, real execution engines are much faster)*

UserID	Name	Job	Salary
123	Alice	Asst. Prof	100
456	Bob	TA	80
789	Carol	Asst. Prof	120
101	David	Prof	150



UserID	Salary
123	100
789	120



Execution Engine

```
foreach row in Payroll:
    if (row.Job == "Asst. Prof")
        output (row.UserID, row.Salary)
```


Relational Algebra

- Algebra:

- Study of symbols
 - Their meanings
 - Their relationships

Algebra (from [Arabic](#) "*al-jabr*", literally meaning "reunion of broken parts"^[1]) is one of the [broad parts](#) of [mathematics](#), together with [number theory](#), [geometry](#) and [analysis](#). In its most general form, algebra is the study of [mathematical symbols](#) and the rules for manipulating these symbols;^[2] it is a unifying thread of almost all of mathematics.^[3] It includes everything from elementary equation solving to the study of abstractions such as [groups](#), [rings](#), and

$$x^2 - 2x - 4 = 0$$

- Relational Algebra

- Study of symbols that manipulates relations
- Symbols = **operators**

$$\sigma_{A > 3}(\pi_{X,Y}(R) \bowtie_X \pi_{X,Z}(T)) =$$

Relational Algebra

- Fundamental operators to:
 - Retrieve
 - Manipulate relations
- Each operator:
 - Take one or more relations as input
 - Output a new relation

σ	Selection
π	Projection
$-$	Difference
\cup	Union
\cap	Intersection
\bowtie	Join
\times	Product

Relational Algebra

- We focus on 7 operators
 - They can be chained (or composed) to create more complex operators
- Based on **set** algebra
 - Because relation is a set

BRACE YOURSELVES

THE DISCRETE MATH IS
COMING

σ	Selection
π	Projection
$-$	Difference
\cup	Union
\cap	Intersection
\bowtie	Join
\times	Product

All tuples are unique, proper relation :)

Selection

- Syntax: $\sigma_{\text{predicate}}(R)$
- Choose a subset of tuples from R that satisfy the predicate
 - Like a filter
 - Predicate can be complex, with conjunction (AND) and disjunction (OR)

R(A,B)

A	B
a1	101
a2	102
a2	103
a3	104

$\sigma_{A=="a2"}(R)$

A	B
a2	102
a2	103

$\sigma_{A=="a2" \text{ AND } B>102}(R)$

A	B
a2	103

```
select * from R
where A='a2' and B > 102;
```

Projection

- Syntax: $\pi_{A_1, A_2, \dots}(R)$
- Generate a relation with tuples containing only the specified attributes
 - Can rearrange attribute order
 - Can transform values

R(A,B)

A	B
a1	101
a2	102
a2	103
a3	104

```
select B-100, A
from R where A='a2';
```

$\pi_{B-100, A}(\sigma_{A='a2'}(R))$

B-100	A
2	a2
3	a2

Union

Result will be a set as well

- Syntax: **(R U S)**
- Generate a relation with tuples appearing in any of the two relations
 - Exactly like set union

```
(select * from R)
UNION
(select * from S);
```

R(A,B)

A	B
a1	101
a2	102
a3	103

S(A,B)

A	B
a3	103
a4	104
a5	105

R U S

A	B
a1	101
a2	102
a3	103
a4	104
a5	105

Intersection

- Syntax: $(R \cap S)$
- Generate a relation with tuples appearing in both relations
 - Exactly like set intersection

```
(select * from R)
INTERSECT
(select * from S);
```

R(A,B)

A	B
a1	101
a2	102
a3	103

S(A,B)

A	B
a3	103
a4	104
a5	105

$R \cap S$

A	B
a3	103

Difference

- Useful for selecting 2 criteria

select cat='A' - select cat='B' EQUAL select cat='A' AND cat='B'

- Useful for selecting 'only' etc

get all runners who only run 100m category: $R - \sigma(100m)(R)$

- Syntax: **(R – S)**
- Generate a relation with tuples appearing in R but not in S
 - Exactly like set difference

R(A,B)

A	B
a1	101
a2	102
a3	103

S(A,B)

A	B
a3	103
a4	104
a5	105

R – S

A	B
a1	101
a2	102

```
(select * from R)
EXCEPT
(select * from S);
```


Product

- Syntax: **(R × S)**
- Generate a relation with all possible combination of tuples from R and S
 - Exactly like set Cartesian product
 - R, S can have different schema

```
select * from R cross join S;
```

R(A,B)

A	B
a1	101
a2	102

S(C,D)

C	D
a3	103
a4	104

A × B

R.A	R.B	S.C	S.D
a1	101	a3	103
a1	101	a4	104
a2	102	a3	103
a2	102	a4	104

Join

- Bread and butter! Very important
- Already seen cross-join: ✕
- We focus on three variants:
 - Inner Join (Equi-Join)
 - Natural Join
 - Left/Right/Full Outer Join




After cross product, there will be 6 columns, 9 tuples.

After applying selection, only 2 tuples are returned

Inner Join

- Syntax: $(R \bowtie_{R.A = S.B, R.B = S.D, \dots} S)$
- Generate a relation with tuples appearing in $R \times S$ and satisfying condition
 - Product followed by Selection
 - Can join on multiple columns



INNER JOIN

```
select * from R, S
where R.A = S.D;
```

R(A,B,C)

A	B	C
a1	101	0
a2	102	1
a3	103	0

S(D,E,F)

D	E	F
a3	103	'a'
a1	107	'b'
a5	105	'c'

$R \bowtie_{R.A = S.D} S$

R.A	R.B	R.C	S.D	S.E	S.F
a1	101	0	a1	107	'b'
a3	103	0	a3	103	'a'

Natural Join

- Syntax: $(R \bowtie S)$
- Like Inner Join, but:
 - Automatically detect **all common attributes** (by names), and use them as join condition
 - Automatically **remove duplicate column**

```
select * from R natural join S;
```

R(A,B,C)

A	B	C
a1	101	0
a2	102	1
a3	103	0

S(A,E,F)

A	E	F
a3	103	'a'
a1	107	'b'
a5	105	'c'

$R \bowtie S$

R.A	R.B	R.C	S.E	S.F
a1	101	0	107	'b'
a3	103	0	103	'a'

Left Outer Join

- Syntax: $(R \bowtie_{R.A=S.B} S)$
- Same as Inner Join, except:
 - All tuples of R appear in the result


R(A,B,C)

A	B	C
a1	101	0
a2	102	1
a3	103	0

S(D,E,F)

D	E	F
a3	103	'a'
a1	107	'b'
a5	105	'c'

```
select * from R left outer join S
where R.A = S.D;
```



LEFT OUTER JOIN

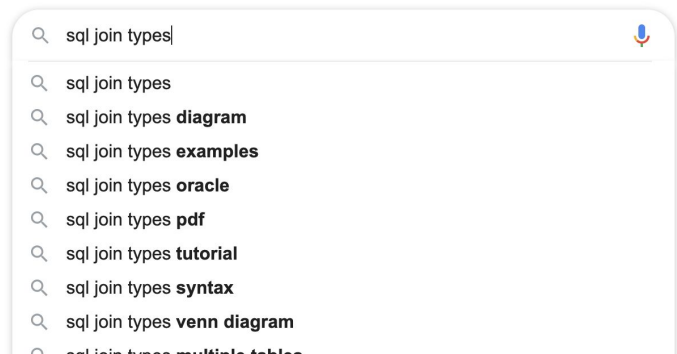
*NULL if they don't match with any values in S.
(e.g. a2 does not match with any values in S, it will still appear in the results)*

$R \bowtie_{R.A=S.D} S$

R.A	R.B	R.C	S.D	S.E	S.F
a1	101	0	a1	107	'b'
a3	103	0	a3	103	'a'
a2	102	1	NULL	NULL	NULL

Other Operators

- Many other operators
- Check them out yourself



ρ	Rename
δ	Duplicate Elimination
γ	Aggregation
τ	Sorting
\bowtie	Right Outer Join
\Join	Full Outer Join


Relational Algebra

- Observation:

$$\sigma_{B=102}(R \bowtie S) = R \bowtie (\sigma_{B=102}(S))$$

- Given input relations, there are **> 1** ways to get the desired output
- Good design = only specify the output
 - Let the machine select the best way

Glimpse Into Database Internal



In Class Exercise

R x S ---- 5 columns, 25 tuples

R natural join S -- 4 columns, 5 tuples

1x03, 2y21, 2y33, 2y20, 3x03

R inner join S (where A=D) -- 5 columns, 3 tuples

1xy21, 3xx03, 3xy33

difference of B in R and B of (C<3) in S

1 column, 2 tuples (z, a)

R(A,B)

A	B
1	x
2	y
2	z
3	x
9	a

S(B,C,D)

B	C	D
x	0	3
y	2	1
y	3	3
w	3	0
y	2	0

Expression	Size of results
$R \times S$	25
$R \bowtie S$	5
$R \bowtie_{A=D} S$	3
$\pi_B(R) - \pi_B(\sigma_{C<3}(S))$	2

In Class Exercise

Reader (ReaderID, FirstName, LastName)

Book (ISBN, Title, Author, PublicationDate, PublisherName)

Publisher (PublisherName, PublisherCity)

Loan (ReaderID, ISBN, Copy, ReturnDate)

Who are the readers who borrow more than 10 copies of a book at a time.

$\Pi_{\text{FirstName, LastName}}(\text{Reader} \bowtie (\sigma_{\text{Copy} > 10}(\text{Loan})))$

Which books (Author, Title) are from publishers in New York or London?

Which books did “Anh Dinh” borrowed?

In Class Exercise

Reader (ReaderID, FirstName, LastName)

Book (ISBN, Title, Author, PublicationDate, PublisherName)

Publisher (PublisherName, PublisherCity)

Loan (ReaderID, ISBN, Copy, ReturnDate)

Who are the readers who borrow more than 10 copies of a book at a time.

Which books (Author, Title) are from publishers in New York **or** London?

$\Pi_{\text{Title, Author}} (\sigma_{\text{PublisherCity} == \text{'NewYork'} \text{ OR } \text{PublisherCity} == \text{'London'}} (\text{Publisher}) \bowtie \text{Book})$

Which books did “Anh Dinh” borrowed?

In Class Exercise

Reader (ReaderID, FirstName, LastName)

Book (ISBN, Title, Author, PublicationDate, PublisherName)

Publisher (PublisherName, PublisherCity)

Loan (ReaderID, ISBN, Copy, ReturnDate)

Who are the readers who borrow more than 10 copies of a book at a time.

Which books (Author, Title) are from publishers in New York or London?

Which books did “Anh Dinh” borrowed?

$\Pi_{\text{Title, Author}}((\text{Loan} \bowtie \sigma_{\text{LastName} == \text{'Dinh'} \text{ AND } \text{FirstName} == \text{'Anh'}}(\text{Reader})) \bowtie \text{Book})$