

Databases and Big Data

Query Optimization

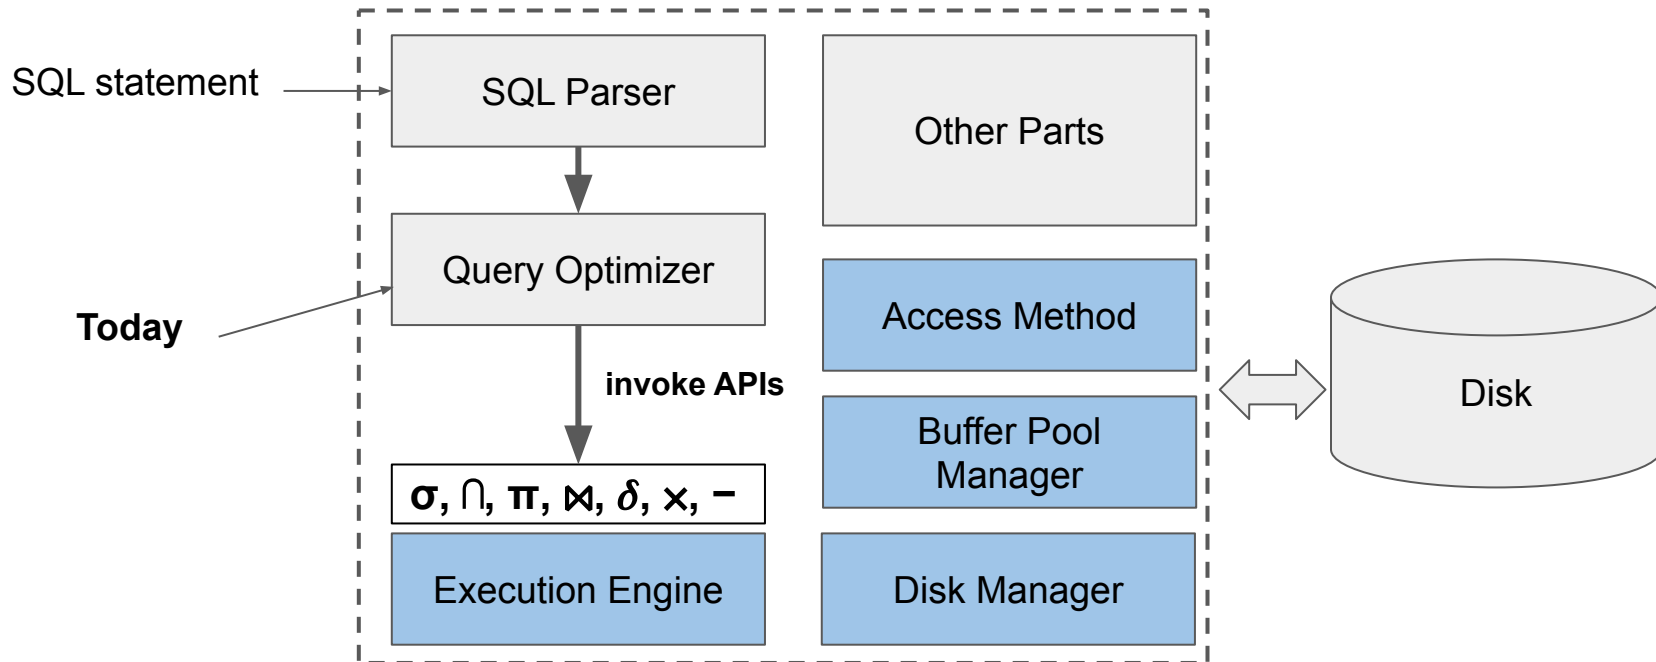
Schedule

Date	Topics	Notes
Week 1	Intro Data Model	No lab
Week 2	SQL	
Week 3	No SQL MongoDB	Group project launch
Week 4	Functional Dependencies Normal Form	
Week 5	Storage Index	Project checkpoint 1
Week 6	Sort, Join Query optimization	Homework
Week 7	Recess	

Schedule

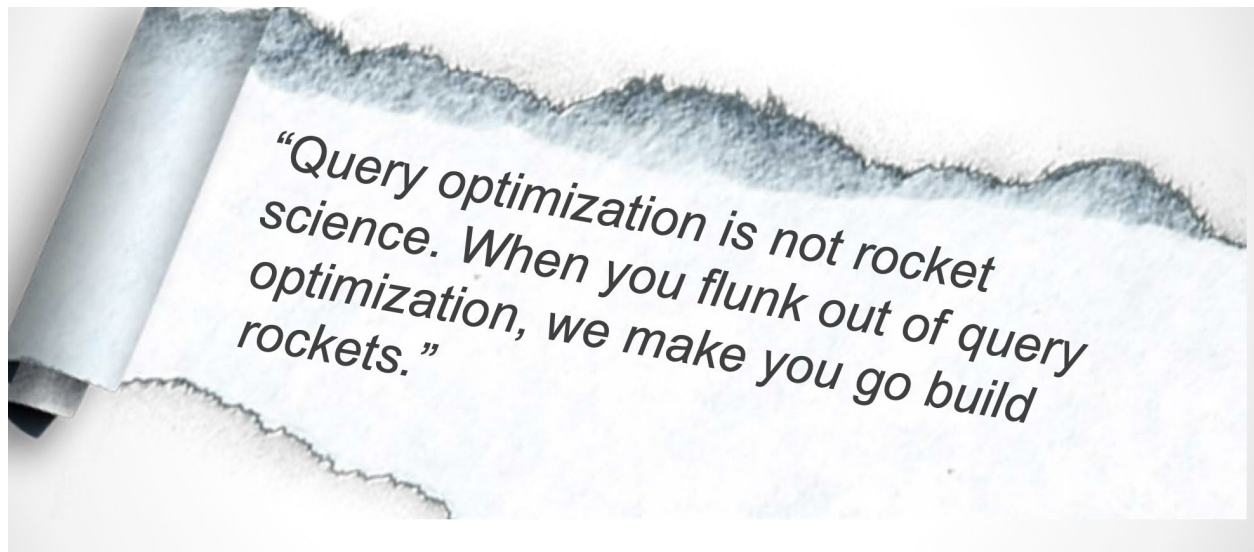
Date	Topics	Notes
Week 7	Recess	
Week 8	Transactions	In-class quiz
Week 9	Introduction to Big Data Cloud Computing	Project checkpoint 2
Week 10	Hadoop	
Week 11	Spark	
Week 12	Spark Ecosystem	
Week 13	Guest speakers Revision	In-class quiz Project due
Week 14	Revision Exam	

So Far



Today

- How hard it is to get it right



Query Optimization

- Recall

Product (pid, name, price)

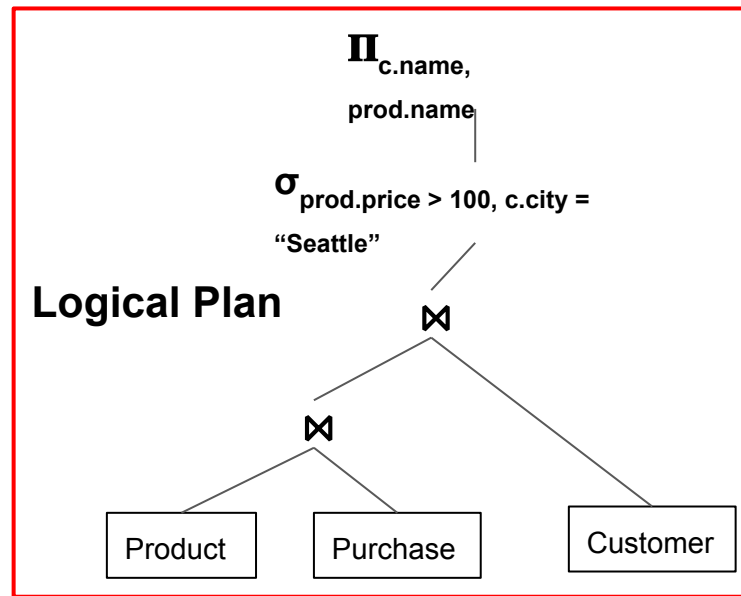

Purchase (pid, cid, store)

Customer (cid, name, city)

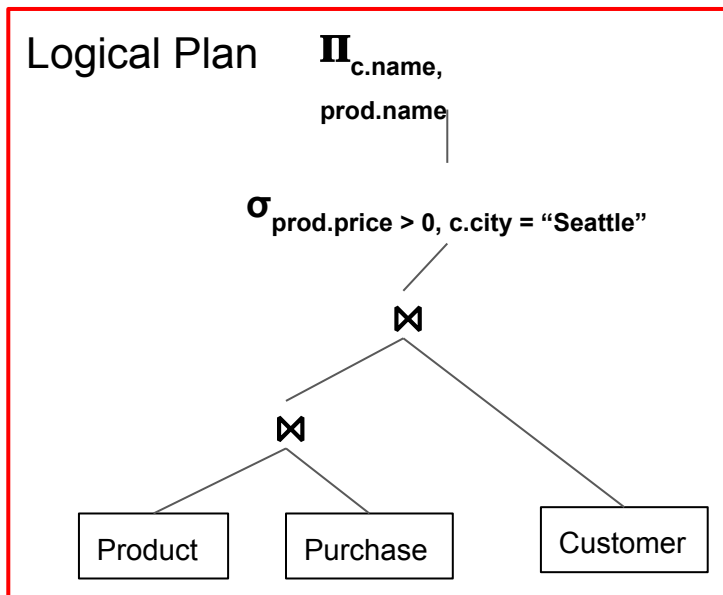


```
select c.name, prod.name
from customer c, product prod, purchase pur
where prod.pid = pur.pid and c.cid = pur.cid
      and prod.price > 100 and c.city = "Seattle";
```

Find name of the customer in Seattle who buys anything over \$100, and name of the product he buys




Query Optimization



Week 2



The final piece



Grace Hash Join

Nested Loop Join

Sort Merge Join

Index Scan

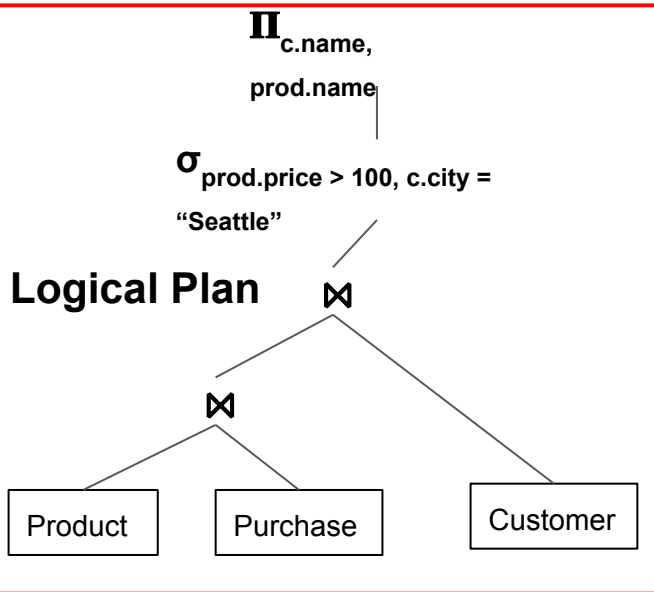
Heap File Scan

External Sort

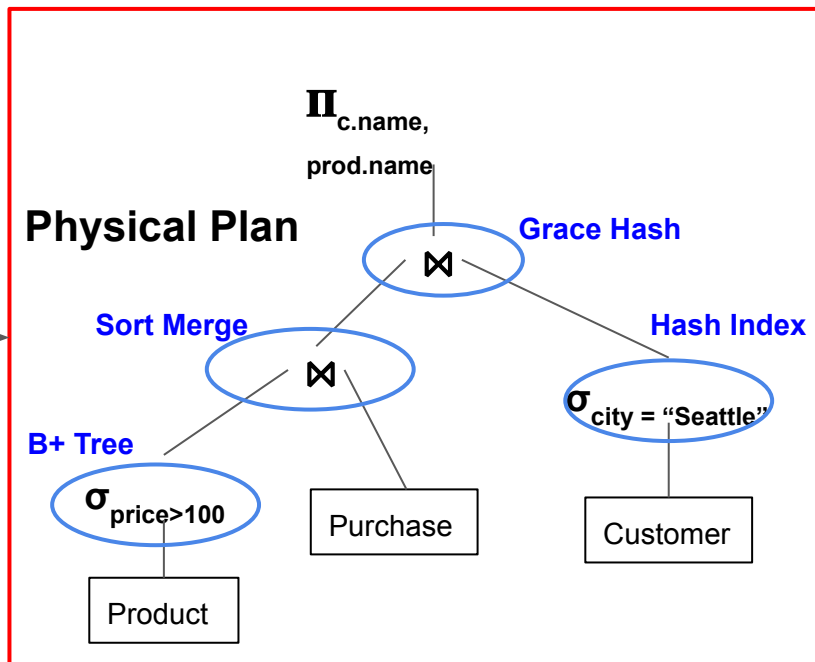
...

Week 5,6

Query Optimization



Query Optimizer



That's why DBMS is declarative

Query Optimizer

- Turn a logical plan into physical plan
- Questions:
 - How to execute a physical plan?
 - How to enumerate equivalent plans?
 - How to estimate cost of each plan?
 - How to search for the best plan?




Execution Model

Execution Model

- Given a physical query plan, how does DBMS execute it?
 - Iterator Model
 - Materialization Model
 - Vector Model



```
select A.id, B.value  
from A, B  
where A.id = B.id and B.value > 100
```



Execution Model

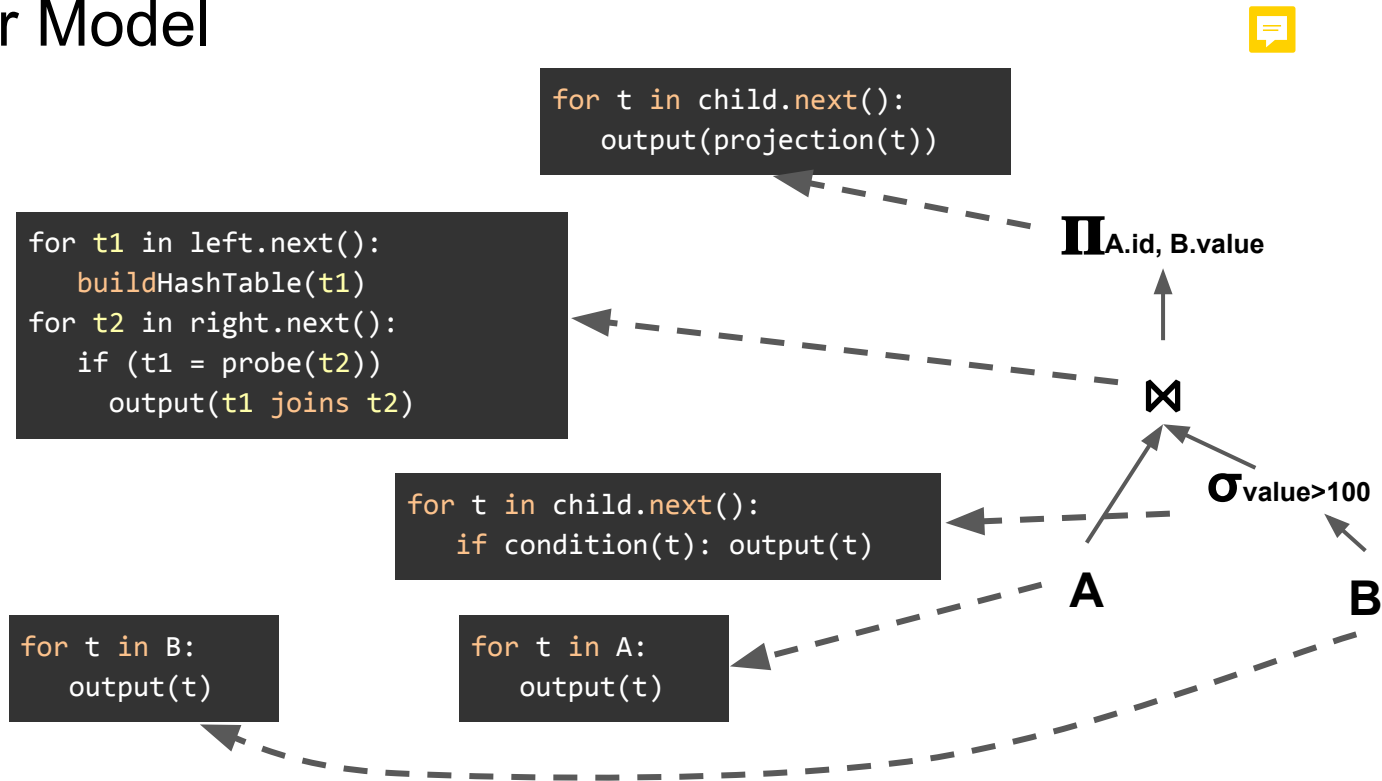
- Iterator Model

- Everything in the plan is an Iterator
- Each operator implement a **next()** method
- Upstream operator calls next(.) of its children



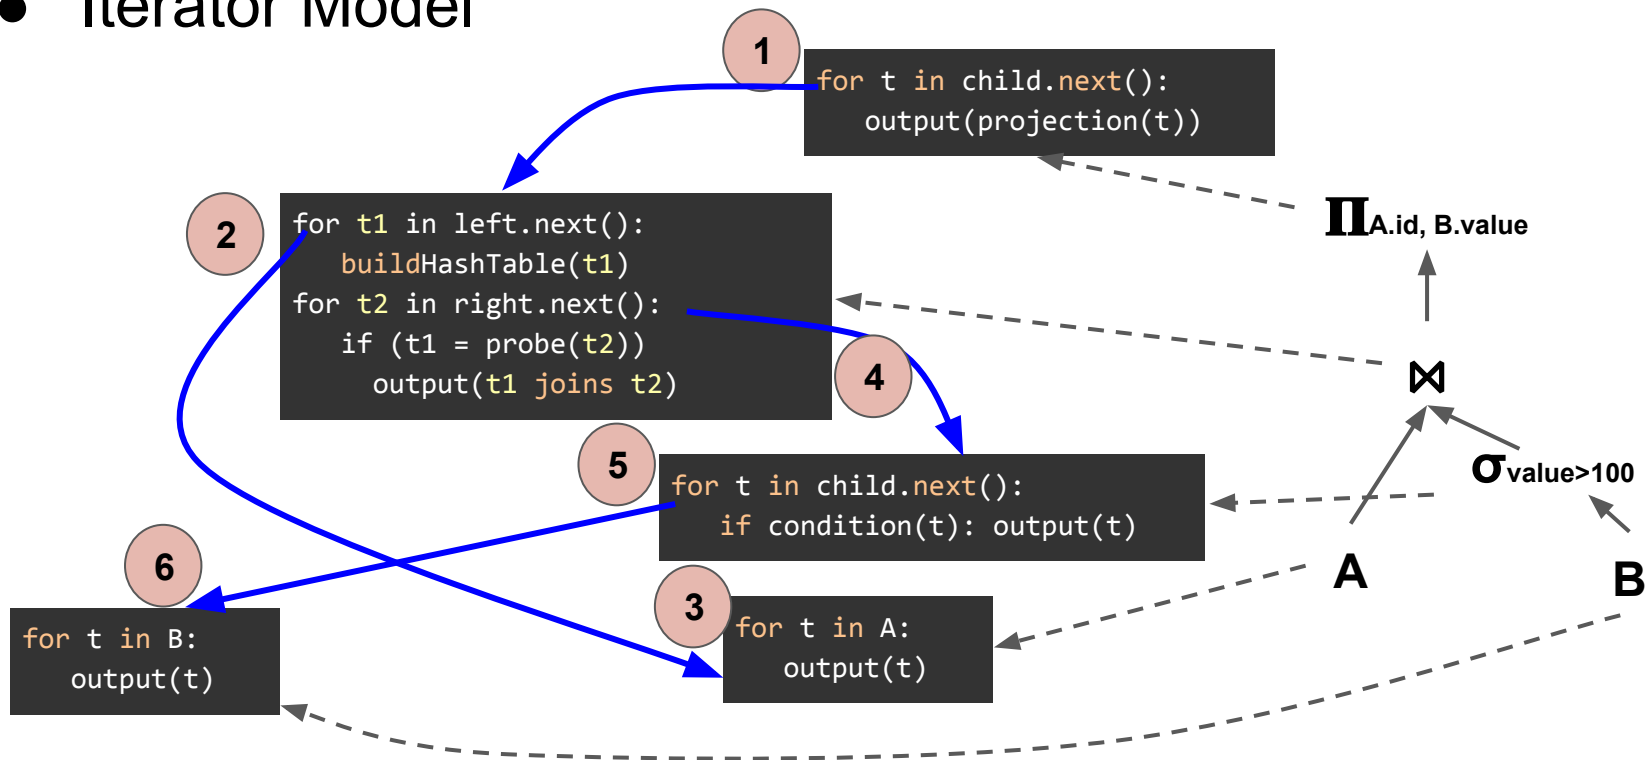
Execution Model

- Iterator Model



Execution Model

- Iterator Model



Execution Model

- Iterator Model

- Data flows bottom up
- Control (call to next(.)) from top down

- Almost all DBMS use this

- Simple to implement
- Allows pipelining
- Great if only subset of results consumed: LIMIT operator



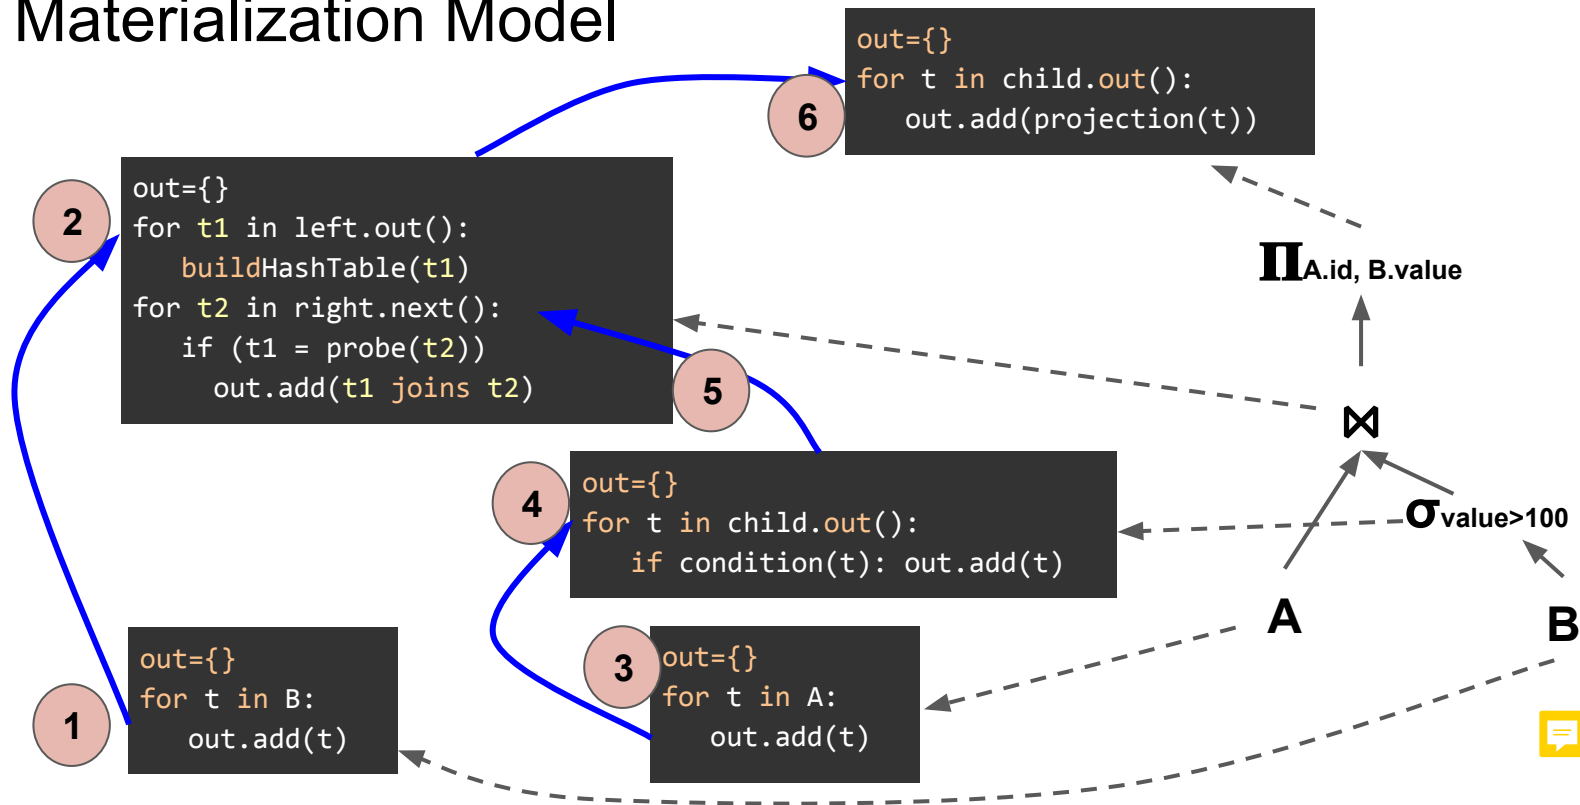
Execution Model

- Iterator Model
 - But high overhead from method calls
- Materialization Model
 - Process all input, and emit all output **at once**
 - Better than Iterator when immediate results **not too much larger** than final result



Execution Model

- Materialization Model



Execution Model

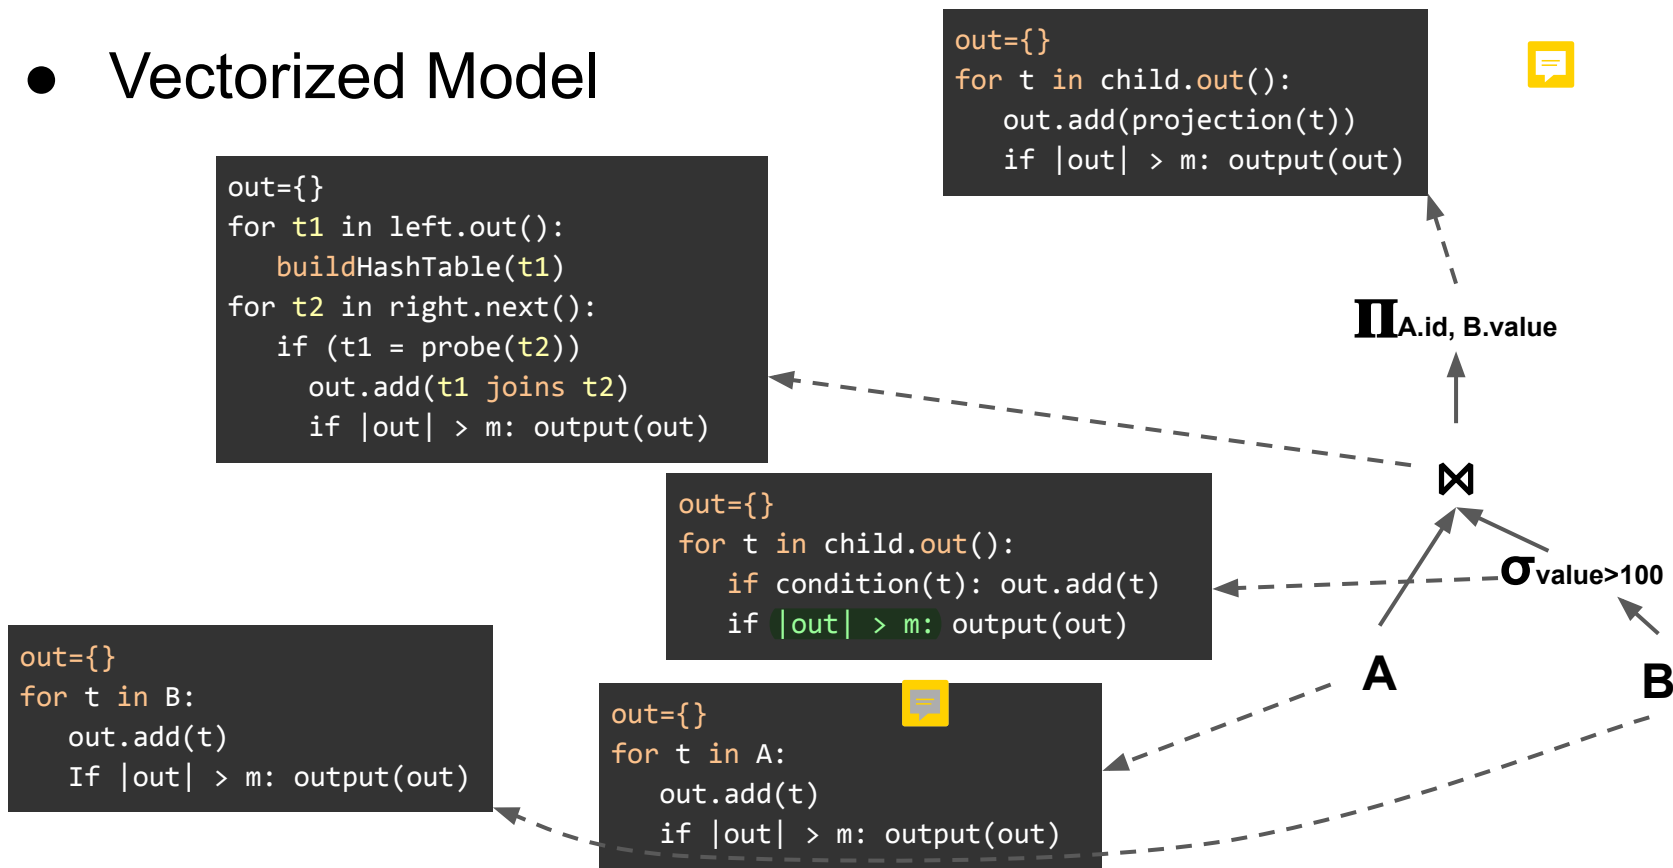
- Vectorized Model

- Combining Iterator + Materialization
- Every operator implement next(.)
- next(.) processes in **batch**
 - Reduce number of invocations in Iterator model



Execution Model

- Vectorized Model



Execution Model

Iterator Model:

- + Direction: top down
- + One tuple at a time
- + General purpose

Vectorized:

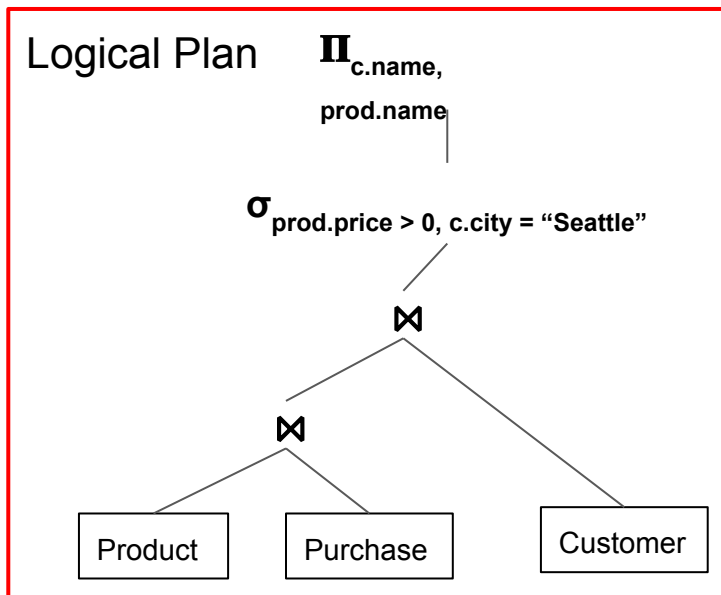
- + Direction: top down
- + Batch of tuple at a time
- + Read-heavy workloads

Materialization:

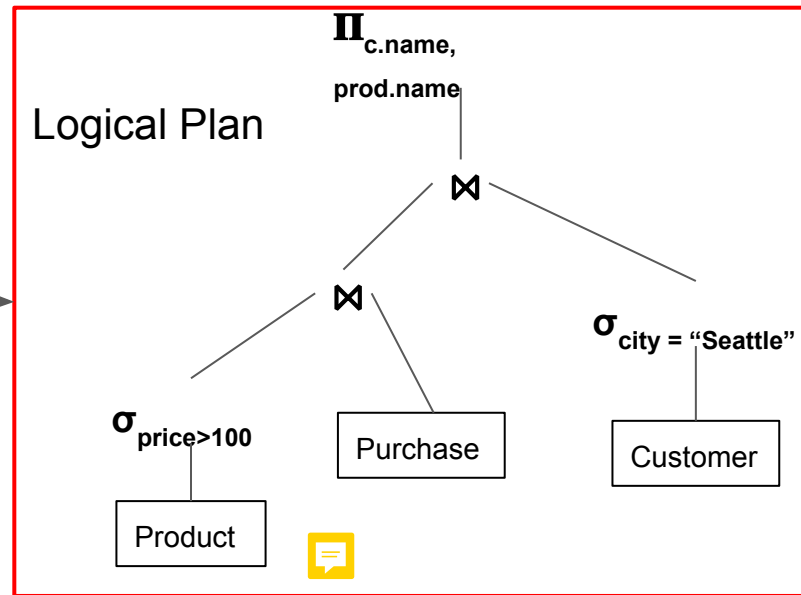
- + Direction: bottom up
- + Entire output set at a time
- + Write-heavy workloads

Query Rewriting

Query Rewriting



Equivalent Plan




$\Pi(\sigma(\text{Product} \Join \text{Purchase}) \Join \text{Customer})$

$\Pi(\sigma(\text{Product}) \Join \text{Purchase}) \Join \sigma(\text{Customer}))$

Query Rewriting

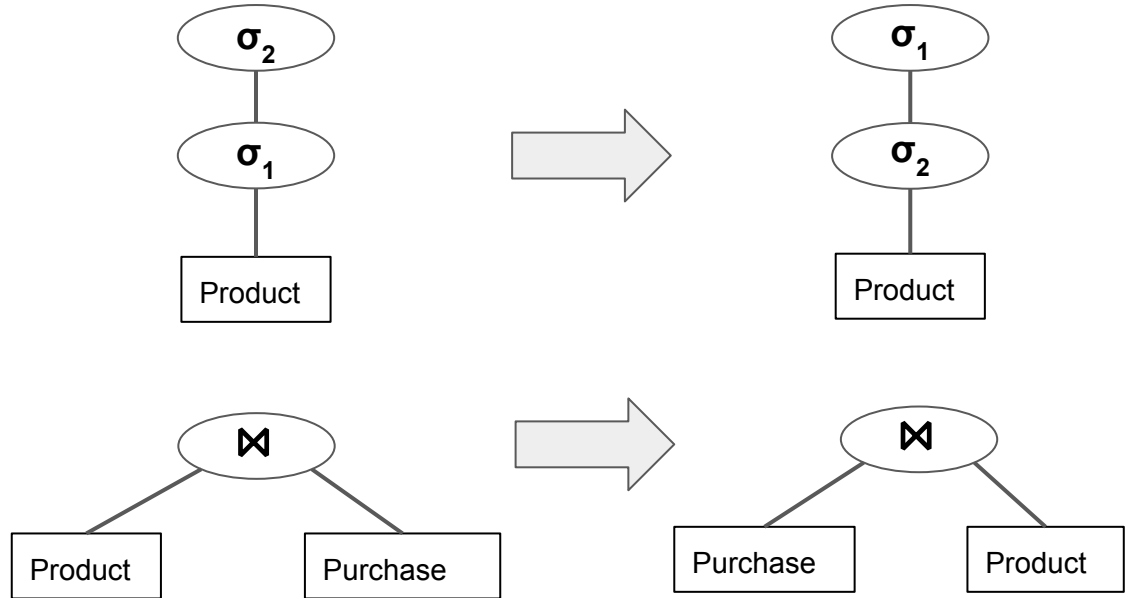
- E, E' are relational expression
 - e.g. $\Pi(\sigma(\text{Product} \bowtie \text{Purchase}) \bowtie \text{Customer})$
- E, E' are equivalent:
 - Let \mathcal{I} be all possible database instances

$$\forall I \in \mathcal{I}. E(I) = E'(I)$$



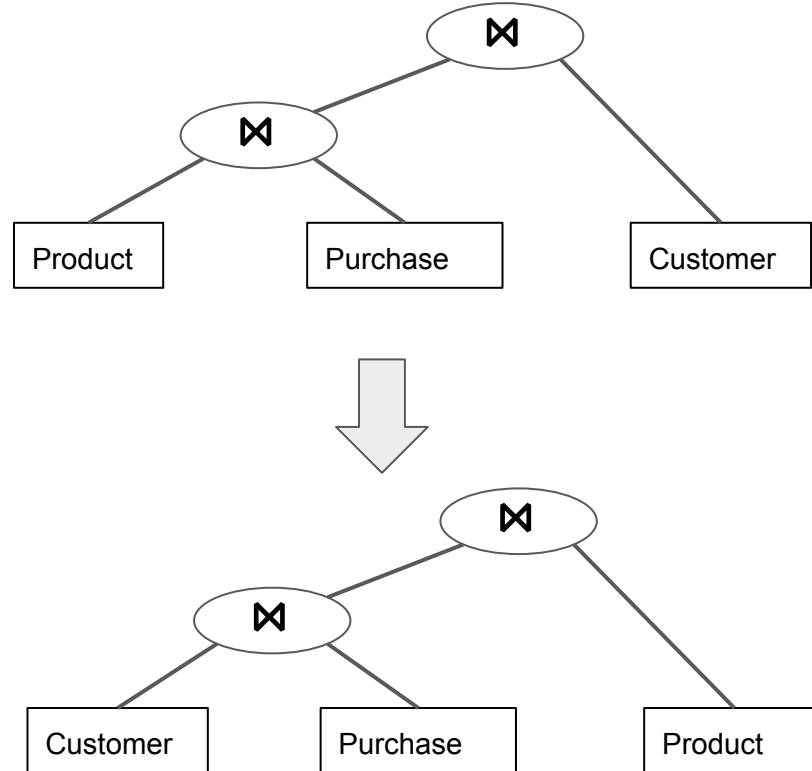
Equivalence Rules

Select and Join operators
commute with each other




Equivalence Rules

Join operators are
associative




Equivalence Rules

Select operator
distributes over Join



Project operator
cascades




Example

Plan 0

$\Pi_{c.name, prod.name}$


$\sigma_{prod.price > 100, c.city = \text{"Seattle"}}$



Plan 1

$\Pi_{c.name, prod.name}$

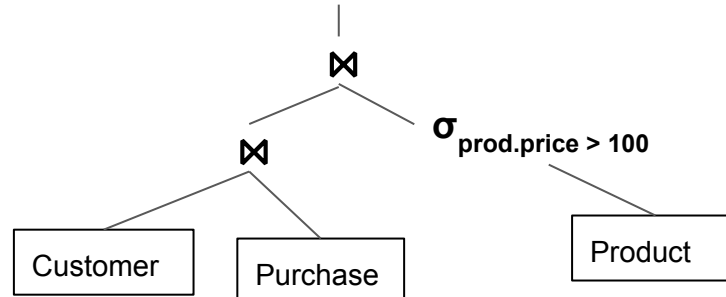
$\sigma_{prod.price > 100, c.city = \text{"Seattle"}}$



Plan 2

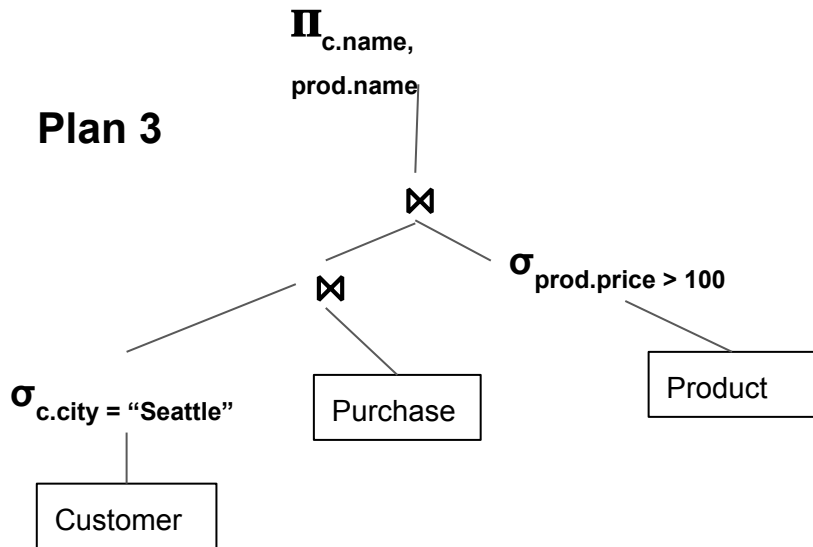
$\Pi_{c.name, prod.name}$

$\sigma_{c.city = \text{"Seattle"}}$

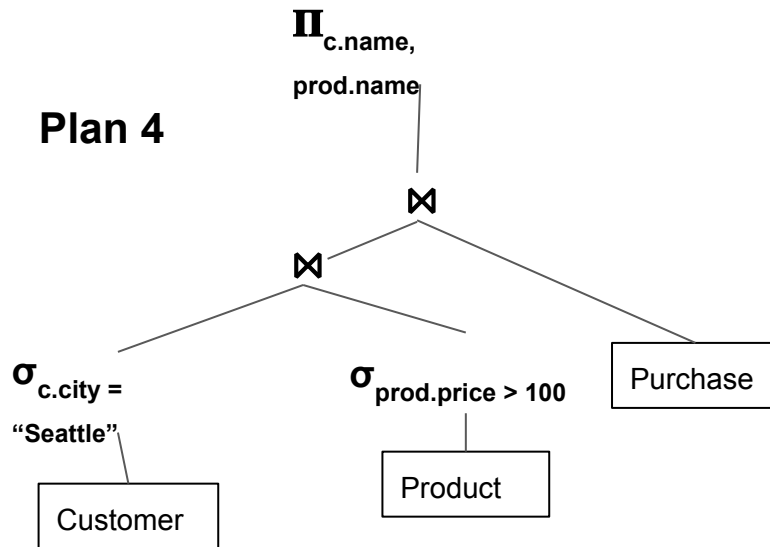


Example

Plan 3



Plan 4



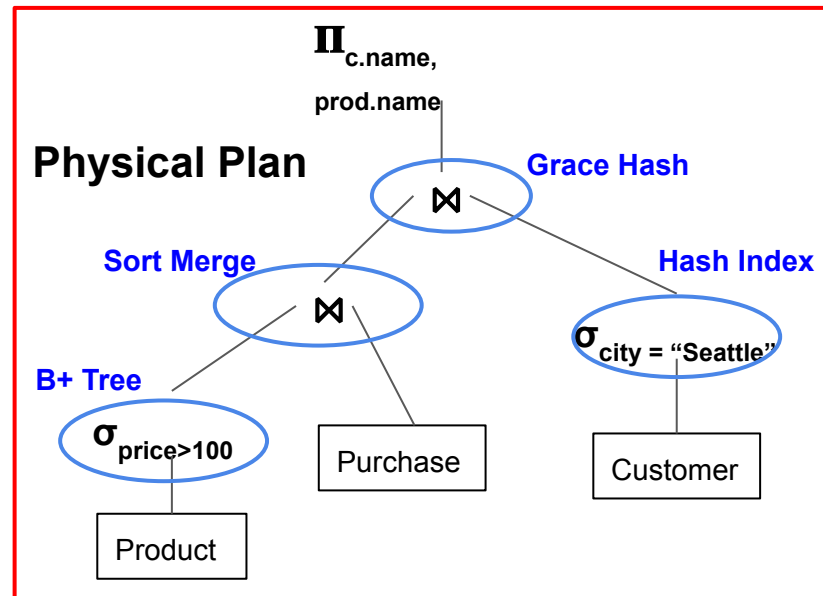
Any a few more logical plans

For each logical plan, there are many physical plans

Cost Estimation

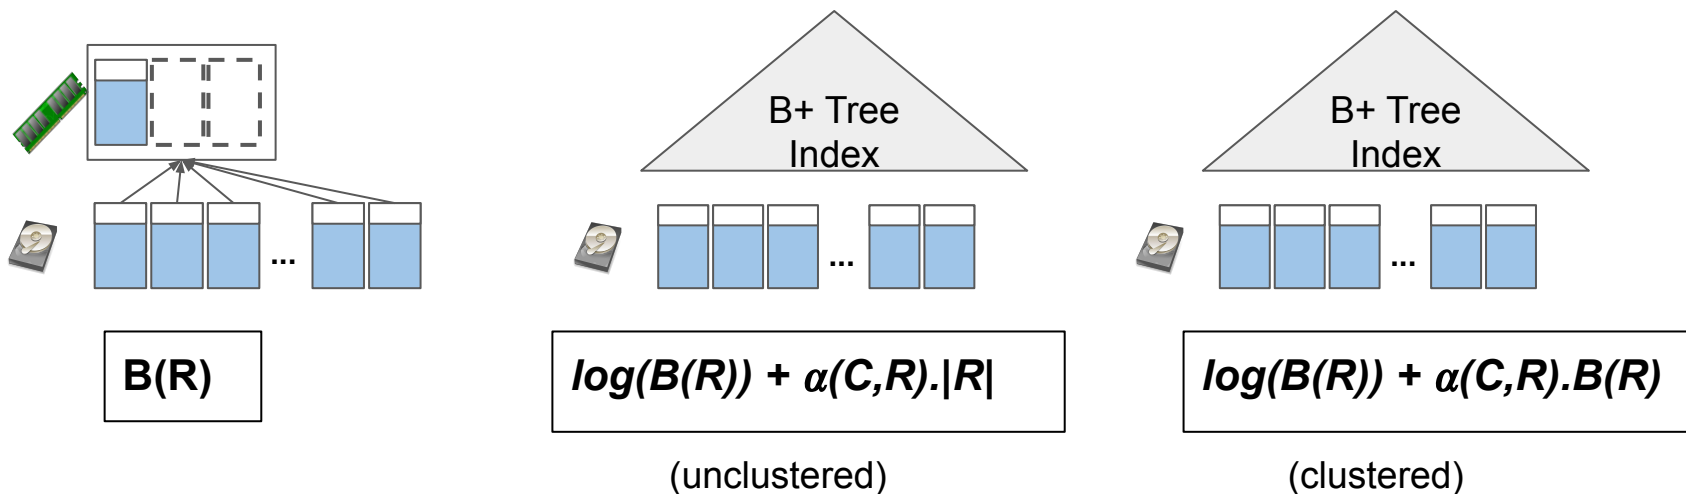
Cost Estimation

- Question: given a physical plan
 - What's the estimated cost of running it
 - Without running it



Cost Estimation

- Recall
 - Cost of select **depends on selectivity**



Cost Estimation

- Cost of Join:
 - Depends on the size of the results

$$(R \bowtie S) \bowtie T$$

$$3.(B(R) + B(S)) + 3.(B(R \bowtie S)) + B(T)$$

$$R \bowtie (S \bowtie T)$$

$$3.(B(S) + B(T)) + 3.(B(S \bowtie T)) + B(R)$$

Cost Estimation

- Unsolved problems
- Still active area of research
- Two main approaches:
 - Build histogram over attributes
 - Sampling



Query Plan Search

Query Plan Search

- So far, we have:
 - Generated lots of physical plans
 - Estimated their costs
- Now: how to pick the best one
 - With lowest cost

Query Search

- If your queries are like these:

```
select * from R,S
where R.id = S.id
```

- But in practice...

```
-- start query 14 in stream 0 using template query14.tpl
WITH cross_items
  AS (SELECT i_item_sk ss_item_sk
        FROM
          (SELECT iss.i_brand_id brand_id,
                  iss.i_class_id class_id,
                  iss.i_category_id category_id
           FROM store_sales,
                item iss,
                date_dim d1
           WHERE ss_item_sk = iss.i_item_sk
                 AND ss_sold_date_sk = d1.d_date_sk
                 AND d1.d_year BETWEEN 1999 AND 1999 + 2)
        INTERSECT
        SELECT ics.i_brand_id,
                ics.i_class_id,
                ics.i_category_id
        FROM catalog_sales,
             item ics,
             date_dim d2
        WHERE cs_item_sk = ics.i_item_sk
              AND cs_sold_date_sk = d2.d_date_sk
              AND d2.d_year BETWEEN 1999 AND 1999 + 2)
        INTERSECT
        SELECT iws.i_brand_id,
                iws.i_class_id,
                iws.i_category_id
        FROM web_sales,
             item iws,
             date_dim d3
        WHERE ws_item_sk = iws.i_item_sk
              AND ws_sold_date_sk = d3.d_date_sk
              AND d3.d_year BETWEEN 1999 AND 1999 + 2)
  WHERE i_brand_id = brand_id
        AND i_class_id = class_id
        AND i_category_id = category_id),
avg_sales
  AS (SELECT Avg(quantity * list_price) average_sales
        FROM (SELECT ss_quantity quantity,
                      ss_list_price list_price
                FROM store_sales,
                      date_dim
                WHERE ss_sold_date_sk = d1.d_date_sk
                      AND d1.d_year BETWEEN 1999 AND 1999 + 2)
        CROSS JOIN cross_items)
```



Query Search

- But in practice...
- In fact:
 - Consider: $R_1 \bowtie R_2 \dots \bowtie R_N$
 - # plans \sim Catalan number

The first Catalan numbers for $n = 0, 1, 2, 3, \dots$ are


1, 1, 2, 5, 14, 42, 132, 429, 1430, 4862, 16796, 58786, 208012, 742900, 2674440, 9694845, 35357670, 129644790, 477638700, 1767263190, 6564120420, 24466267020, 91482563640, 343059613650, 1289904147324, 4861946401452, ...

```
-- start query 14 in stream 0 using template query14.tpl
WITH cross_items
AS (SELECT i_item_sk ss_item_sk
      FROM item,
      (SELECT iss.i_brand_id brand_id,
              iss.i_class_id class_id,
              iss.i_category_id category_id
      FROM store_sales,
           item iss,
           date dim d1
      WHERE ss_item_sk = iss.i_item_sk
            AND ss_sold_date_sk = d1.d_date_sk
            AND d1.d_year BETWEEN 1999 AND 1999 + 2
      INTERSECT
      SELECT ics.i_brand_id,
              ics.i_class_id,
              ics.i_category_id
      FROM catalog_sales,
           item ics,
           date dim d2
      WHERE cs_item_sk = ics.i_item_sk
            AND cs_sold_date_sk = d2.d_date_sk
            AND d2.d_year BETWEEN 1999 AND 1999 + 2
      INTERSECT
      SELECT iws.i_brand_id,
              iws.i_class_id,
              iws.i_category_id
      FROM web_sales,
           item iws,
           date dim d3
      WHERE ws_item_sk = iws.i_item_sk
            AND ws_sold_date_sk = d3.d_date_sk
            AND d3.d_year BETWEEN 1999 AND 1999 + 2)
WHERE i_brand_id = brand_id
      AND i_class_id = class_id
      AND i_category_id = category_id),
avg_sales
AS (SELECT Avg(quantity * list_price) average_sales
      FROM (SELECT ss_quantity quantity,
                    ss_list_price list_price
      FROM store_sales,
           date dim
      WHERE ss_sold_date_sk = d.date_sk
            AND d_year BETWEEN 1999 AND 1999 + 2
      UNION ALL
      SELECT cs_quantity quantity,
              cs_list_price list_price
      FROM catalog_sales,
           date dim
      WHERE cs_sold_date_sk = d.date_sk
            AND d_year BETWEEN 1999 AND 1999 + 2
      UNION ALL
      SELECT ws_quantity quantity,
              ws_list_price list_price
      FROM web_sales,
           date dim
      WHERE ws_sold_date_sk = d.date_sk
            AND d_year BETWEEN 1999 AND 1999 + 2)
```

Query Search

- Another unsolved problem
- Most common approach: Dynamic Programming
 - Pass 1: find best 1-relation plans
 - Pass 2: find best 2-relation plans by joining results from Pass 1
 - ...
 - Pass N: find best N-relation plans by joining results from Pass N-1
- Still exponential in search space!

Query Search




Legend:

SS – sequential scan

IS – index scan

5 – cost

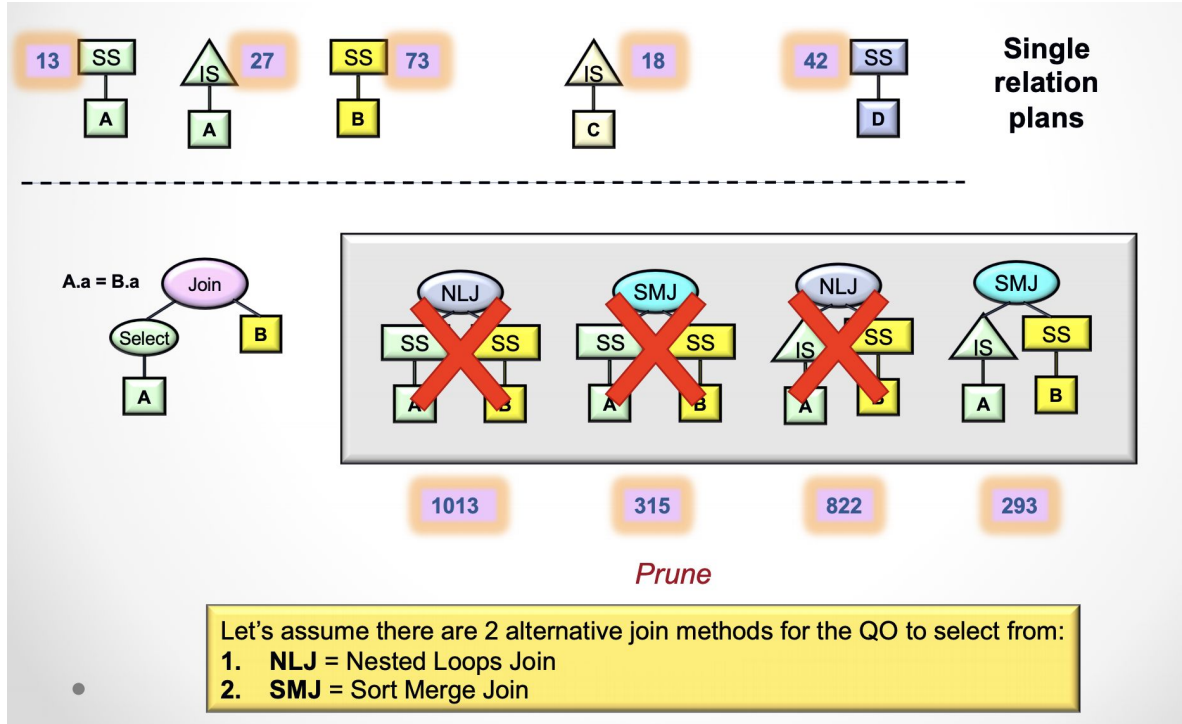
First, generate all single relation plans:




All tables

All single
relation plans


Query Search



Query Search




Query Search



Query Search

- DBMS often picks bad plans
 - Search is heuristic, space is too big
 - Too sensitive to errors in cost estimation
 - And cost estimation is very hard!
 - Cost models go out of date quickly
 - Hardware, software upgrade



Summary

- Query Optimization far from solved problem
 - ML for help!

Learning to Optimize Join Queries With Deep Reinforcement Learning

Sanjay Krishnan^{1,2}, Zongheng Yang¹, Ken Goldberg¹, Joseph M. Hellerstein¹, Ion Stoica¹

¹RISELab, UC Berkeley ²Computer Science, University of Chicago

skr@cs.uchicago.edu {zongheng, goldberg, hellerstein, istoica}@berkeley.edu

DOI: 10.1145/nnnnnnnn.nnnnnnn

ABSTRACT

Exhaustive enumeration of all possible join orders is often avoided, and most optimizers leverage heuristics to prune the search space. The design and implementation of heuristics are well-understood when the cost model is roughly linear, and we find that these heuristics can be significantly suboptimal when there are non-linearities in cost. Ideally, instead of a fixed heuristic, we would want a strategy to guide the search space in a more data-driven way—tailoring the search to a specific dataset and query workload. Recognizing the link between classical Dynamic Programming enumeration methods and recent results in Reinforcement Learning (RL), we propose a new method for learning optimized join search strategies. We present our RL-based DO optimizer, which

1 INTRODUCTION

Join optimization has been studied for more than four decades [44] and continues to be an active area of research [33, 40, 49]. The problem's combinatorial complexity leads to the ubiquitous use of *heuristics*. For example, classical System R-style dynamic programs often restrict their search space to certain shapes (e.g., “left-deep” plans). Query optimizers sometimes apply further heuristics to large join queries using genetic [4] or randomized [40] algorithms. In edge cases, these heuristics can break down (by definition), which results in poor plans [29].

Summary

